The Challenge of Deriving Signal from Synthetic Legal Environments: Evidence from Harvey's Legal Agent Benchmark

John J. Hughes, III
Independent researcher
https://www.johnjhughesiii.com/research/

Working draft — October 7, 2026

Abstract

Reinforcement learning (RL) techniques, in domains that can offer objectively verifiable rewards, have led to incredible and noteworthy recent gains in the capabilities of frontier AI models. But not all domains are equally verifiable. Many high-level legal judgment tasks are difficult to verify, and model capability gains on those tasks have not been nearly as noteworthy as model gains in domains that are susceptible to automated verification like math and coding.

Existing benchmarks used to evaluate high-level legal tasks rely largely on human-written rubrics that test model output in response to fictitious, AI-generated environments. But an AI-assisted audit of one of the most widely used such benchmarks, checked by a practicing litigator's review of a random sample of the flagged criteria, found evidence of hundreds of erroneous rubrics in just one subpart of the benchmark, many of them in task environments that are themselves misspecified. A separate AI-assisted citation check found apparently fabricated case citations in the task environments, and a deep dive into one motion-to-dismiss task found planted issues and documents that do not resemble real lawyer work product. The errors in these environments illustrate the challenges facing even the top data providers and research labs trying to improve frontier model capabilities, at a time when many previously challenging benchmarks and tasks have become saturated.

1 Introduction

Most recent legal benchmarks have sought to test higher-level judgment tasks, which usually lack an objective ground truth. The leading vertical legal AI company, Harvey AI Corporation, released Legal Agent Benchmark (LAB), in May 2026. Now widely cited by frontier research labs as a key evaluation of agentic legal capabilities for frontier models, LAB uses AI-generated task environments to test models on pre-specified rubrics. Harvey has publicly disclosed that it works with Mercor and Snorkel, leading third-party data providers, to conduct expert review and verification of every generated LAB task [5]. Harvey open-sourced the benchmark (except for a held-out test set), and invited practicing lawyers to validate and improve it by reviewing existing tasks and auditing rubrics. [6] I took up that invitation and did a deep-dive review of one of the substantive litigation tasks (drafting an issues list for a motion to dismiss opposition brief) and an audit of the public litigation rubrics. This analysis underscored the challenges involved in creating suitable reinforcement learning environments that can advance the capabilities frontier in the legal domain—and illustrated why extracting signal from real-world litigation data could be more effective than existing approaches [9].

2 Audit of LAB's litigation rubrics

Several LAB rubrics appeared to contain anomalous results, so I conducted an audit of the rubrics on the public litigation and dispute resolution tasks (my area of specialty for 18 years). I supplied GPT-6 Sol and Claude Opus 5.5 with examples of the kinds of rubrics that I consider defective and asked them to iterate through each task to audit the rubrics using the high-level standards I supplied (which allowed them to raise issues that were not specifically on my list).

Of the 52 tasks and 2,858 rubrics included in the litigation and dispute resolution category, GPT-6 Sol flagged 616 criteria as defective or arguably defective, Claude Opus 5.5 flagged 564 (with the ability to review GPT-6 Sol's findings). Between them, the models flagged 835 criteria, leaving 2,023 unflagged by either report. Of 345 overlapping flags, both reviewers called 83 problematic and 135 arguable; 127 received different levels of concern. Opus was allowed to see Sol's findings when reaching its own conclusions, so agreement in the final reports is not independent confirmation [1].

To validate these findings and ascertain a statistically valid estimate of the number of rubrics that are defective, I reviewed a simple random sample of 25 of the 835 flagged criteria, drawn with a fixed seed from 19 tasks (Appendix B). I treated a rubric as defective if, applied as written to the task as given, it would fail a response a careful litigator would consider appropriate or pass one a careful litigator would consider wrong on the issue the rubric covers. 16 of the 25 were defective, and I assessed 9 of those to be unambiguous objective errors where the accuracy of the rubric did not depend on my exercise of professional judgment but there was indisputably an error or problem in the task of some kind. Another 3 were arguable and 6 were not defective.

These results demonstrate that at least 383 of the 835 flagged criteria are defective (one-sided 95% bound; point estimate 534), or more than 13.4% of all 2,858 criteria in the practice area; at least 171 are unambiguous objective errors. Agreement between the two AI auditors was informative: 9 of the 10 sampled criteria flagged by both were defective, against 7 of the 15 flagged by only one.

Examples of defective rubrics, drawn from the sampled criteria in Appendix B, included instances where the rubric misstates the law (e.g., requiring the answer that is the opposite of settled law in the jurisdiction at issue), mathematical or numerical errors (e.g., the task environment contained mathematically impossible documents, or contained a rubric that required models to assume that one of two conflicting numbers was accurate while ignoring the conflicting number), or instances where I assessed the rubric's recommended approach to contravene widely respected professional norms.

I separately counted defects in the task environment relevant to the rubric (as opposed to defects that made the rubric's grading criteria themselves incorrect). 14 of the 25 criteria related to defective task environments, for example, documents that did not sum to the stated figures; these span 10 of the 19 sampled tasks, and only 4 of the 25 criteria had neither a defective nor an arguable criterion and came from an environment without a defect. This means that at least 319 flagged criteria (more than 11.1% of all criteria) come from such environments.

These errors are highly significant relative to the margins that separate frontier models on LAB's scoring criteria. On Artificial Analysis's Harvey LAB-AA leaderboard, Muse Spark 1.3 leads with a criterion pass rate of 95.5%, followed by Kimi K3 at 94.6%, Claude Fable 5.1 at 93.3%, Claude Opus 5.5 at 91.2%, and GPT-5.6 Luna at 87.9% [8]. If criteria elsewhere in LAB are defective at anything like the rate we found (more than 13.4% of the criteria in this practice area), the rate of defective rubrics is as large as or larger than the reported gaps between leading models on criterion pass rate, and it is unclear how much of those gaps the defects could explain. All-pass scoring

compounds the problem, because a single defective criterion can fail an otherwise acceptable task.

The flaws in these environments raise questions about whether general-purpose data vendors have the qualifications or ability to prepare verified and accurate RL environments for legal tasks. Harvey retained leading data vendors to review and verify the tasks and rubrics [5]. But the verification services offered by these third-party vendors did not reliably detect defects like rubrics that misstated the law. These general-purpose data vendors provide a variety of different kinds of reinforcement learning environments, many of them for nonlegal tasks, and the companies utilized do not purport to be specialized in legal work. Some companies in this space hire freelance lawyers at contract attorney rates to verify tasks and rubrics. At an earlier phase of frontier model development, that approach may have provided useful signal. But as the capabilities of frontier models have improved, and benchmarks have become more ambitious, the results here suggest that this kind of freelancer data could have defects that end up undercutting the evaluation or training signal the dataset is intended to provide. Even sophisticated buyers doing frontier AI research in the legal domain may find it difficult to catch these errors without effectively reproducing much of the work that the data vendors are paid to do. This suggests that further improvement in the legal domain might require efforts to increase the quality of the data used to evaluate and train models on legal tasks.

The challenges involved in developing and maintaining acceptable benchmarks for legal tasks call to mind some of the challenges frontier labs have had developing new benchmarks for coding tasks. In February 2026, OpenAI publicly criticized SWE-bench Verified (which OpenAI had helped develop), noting that its audit revealed problems involving test requirements and task specification [7]. The rapid capability growth has meant that errors that once might have been unnoticeable now show up in the form of frontier models failing tasks incorrectly because the model is smarter than the benchmark. That dynamic could require greater investment in evaluations to keep pace with the strength of the underlying models.

3 Cite Check of the Litigation Task Environments and Rubrics

Because the rubric audit turned up instances where the rubrics contained hallucinated case law (suggesting that the external experts who were responsible for verification may have relied on inaccurate AI-generated responses to prepare the answer key), I asked AI agents to conduct a separate citation check of the litigation and dispute resolution tasks to assess the frequency of AI hallucinations in the task environments or rubrics. A citation check is a standard pre-filing procedure where lawyers check every authority cited in a brief, regulatory submission, or other important document to ensure that the authority is cited accurately, contains any quotation attributed to it, and supports the legal proposition for which it is cited.

A separate AI-assisted check of every case citation in the 19 sampled environments against CourtListener, with spot checks, found problems in 45 of 102 citations to purported real authority: 32 real cases cited for holdings, quotations, or facts they do not contain, 7 cases that could not be found, 5 garbled citations to real cases, and 1 citation pointing to a different case [2].

After excluding errors in case law that were intentionally planted and then tested in rubrics (e.g., an error of law in an opponent's submission that the model was supposed to catch), the AI identified 436 errors that were apparently not intended to be part of the task environment [3]:

I reviewed a simple random sample of the 436 errors personally. Before I reviewed any sampled finding, I drew 20 citation findings and pre-committed to a stopping rule that if the first ten were all confirmed errors, I would stop; otherwise I'd review 10 more. I reviewed each citation finding

Finding	All citations	Case citations
Verified	2,754	306
Verified, minor defect (pin cite, loose paraphrase, wrong year)	346	54
Intentional error tested by the rubric (excluded)	31	6
Unintentional errors		
Cited for a proposition it does not support	246	105
Not found (likely fabricated)	76	60
Garbled citation to a real authority	58	10
Quotation not found in the authority	40	31
Reporter citation belongs to a different case	16	12
Total unintentional errors	436	218
Source text not accessible	51	6
Total checked	3,618	590

Table 1: AI citation check of the 52 LAB litigation task environments. This table shows the AI findings.

and agreed that all ten involved errors a careful lawyer would correct before filing. The AI had correctly classified seven as cited for propositions they do not support and three as real authorities cited by the wrong number or subsection. Three of the ten involved miscitations to the Eleventh Circuit's pattern civil jury instructions. With ten of ten confirmed, a conservative one-sided 95% lower confidence bound is 303 of the 436 flagged errors (69.5%). I did not attempt to review or identify errors the AI may have missed among citations it marked as verified.

4 A motion-to-dismiss case study

In addition to auditing the accuracy of all 2,858 rubrics, I did a deep dive review of one substantive litigation task. This review highlighted a different but equally important issue: the use of synthetic or AI-generated data may result in task environments that are unrealistic in a variety of subtle ways that may not be obvious to nonexperts but that are striking and clear to practicing attorneys. The task `analyze-counterparty-motion-to-dismiss` asks AI agents to prepare a list of issues for the opposition brief to a motion to dismiss, on a fictitious software service dispute.

I observed several ways in which the synthetically generated task environment seemed unrealistic and could potentially train models on a distorted view of what genuine practice looks like. [4]

Planted issues. A number of the legal disputes in the task environment are obviously "planted" in a way that feels artificial and ultimately presents issues in a way that is meaningfully different from how real-world legal disputes manifest. For example, the contract the parties are suing over has a forum selection clause that requires any lawsuit to be filed in Travis County, Texas. The plaintiff literally quotes this entire provision in the complaint but nevertheless files suit in the Northern District of Georgia and offers no explanation for why they believe the forum selection clause they quoted does not apply or otherwise permits the suit in Georgia. It of course happens regularly that a lazy or incompetent litigant files suit in the wrong forum; a competent lawyer potentially could decide to ignore the clause and hope that their adversary overlooks the issue. But in both of those scenarios, the forum selection clause would not be explicitly flagged in the complaint as it is here; a real version of this legal task would require the lawyer to dig through a number of documents and realize that buried in one of them is a forum selection clause that allows the lawyer to have the case

moved to Texas.

This is one example, but a significant number of the tasks in these environments were similarly planted in a way that's artificial and unlike how disputes around these issues arise in real cases. Environments like this can distort an evaluation of a model's true capabilities, since simply looking for the obviously planted issues will satisfy the benchmark without testing whether the model would be able to identify this issue from more realistic source material akin to what it will see in real cases.

Environment Documents Violate Court Rules. Many of the documents in these environments do not resemble real lawyer work product. The environment provides counterparty litigation documents in docx format, and the native files from the git repository often cannot be opened, although agents and various CLI tools can recover the document text. In real litigation, lawyers rarely share their docx files with their opponent; the opponent's memorandum of law is ordinarily filed with the court as a PDF. The motion papers also lack a proper caption, include a stale table of contents, contain markdown formatting artifacts that render in Word as junk characters even in files that can be opened normally, lack citations for propositions that generally require citations, and are single spaced rather than double spaced in violation of the applicable court rules.

The motion that the agent is being asked to respond to is so far below professional norms that you would want any analysis of the motion to include consideration of whether the judge would strike the motion completely, or potentially discipline or sanction the counsel that filed that deficient work product. That is not included in the rubric for this task because the quality control issues appear to be unintentional, not part of the test for the agent. But it is a genuine gap: if a document like this were filed in real litigation, the quality of the work product could itself become a substantive issue in the litigation and thus would need to be flagged given the goal motivating the all-pass criterion, which was to ensure that the AI does not overlook any important issues [6, section "Evaluating a task"].

There are valid reasons to prefer rubric-based reinforcement learning approaches. For example, while a benchmark like LegalForecastBench [9] (which asks models to predict how federal judges ruled on motions to dismiss) can test a model's legal reasoning and analytical ability over a narrow task, it does not test the model's ability to generate new work product, which is an important consideration for a company like Harvey that is developing legal agents to prepare and revise different kinds of legal work product. But reviewing these environments suggests that the current state of the art in synthetic data generation introduces drift on a variety of dimensions that could undermine the efficacy of the task environment as a training or evaluation signal. The model is in effect being tested on its ability to recognize a distinct environment, very unlike what it will encounter in production, and adhere to a stylized set of norms and rules that were artificially constructed in that environment but do not comport with the real norms or rules that govern real-world tasks. While creating LegalForecastBench illustrated that there are real challenges in working with real-world litigation data as well, for a variety of tasks that data has a lot of rich, implicit signal that data providers generating synthetic data are unable to capture.

5 Conclusion

This audit illustrated that AI research has a long way to go to bridge the gap between current training and evaluation environments used to train models, and the messy real-world tasks that models are asked to perform. LAB is notable as the first large-scale, open-source set of legal task environments that try to test high-level legal judgment tasks where there is often no objective ground

truth. But LAB illustrates the challenges of creating effective environments, because even leading data vendors that took on the responsibility of expert verification of the LAB tasks were unable to create realistic task environments or accurate rubrics. Harvey has explicitly invited practicing attorneys to help collaborate on improving LAB, and building models that can automate more of the high-level legal work that lawyers do clearly will require careful collaboration to obtain training signal that can reliably improve the capabilities of, and distinguish between the performance of, increasingly capable frontier models.

Reproducibility, access, and AI use disclosure

Code and audit materials are maintained in the LegalForecastBench repository. The sampling protocol, sampled findings, and completed attorney-review worksheet for the citation check are available in the accompanying review materials.

As I am a practicing attorney and do not have any formal training in artificial intelligence or machine learning, the use of AI models was essential to allow me to make a research contribution in this space that otherwise would have been beyond my capabilities. My personal contribution included auditing a sample of the LAB rubrics by hand to statistically validate the AI audit findings, conducting a deep dive into one of LAB's litigation tasks, and orchestrating AI coding agents that helped write and execute the code that ran the audit. I wrote most of this paper personally, but discussed the results at several points with AI agents and considered AI-generated analysis as an input into my thinking as I drafted. The major exception to this was the AI audit findings, the citation check in Section 3, and other places where numerical results are included; I relied on AI agents to produce those findings, to write code that performs the necessary calculations, and to prepare the LaTeX source summarizing those results. I used several different models from Anthropic, OpenAI, SpaceXAI, and Google to check the work. I also met with experienced researchers in the space to validate the experimental design and results, and obtain feedback on the draft. I am grateful to several AI researchers in industry positions for helpful discussions. This version of the paper is intended as a working draft, and feedback from researchers, lawyers, and others on how to improve this audit or make our findings more useful is welcome.

References

[1] GPT-6 Sol audit workers, Claude Opus 5.5 audit workers, Coordinating agents; individual provenance not fully recorded. Harvey LAB litigation and dispute resolution audit: GPT-6 Sol and Claude Opus 5.5. *Local research artifact, pinned September 29, 2026*, 2026. https://github.com/johnhughes3/LegalForecastBench/tree/58371a0e13494e591710e2eae3da8e428eebc359/docs/harvey-lab-audit/litigation-dispute-resolution.

[2] John Hughes. Human review of AI-flagged Harvey LAB litigation criteria: protocol, sample, and worksheet. *Research artifact*, 2026. https://github.com/johnhughes3/LegalForecastBench/tree/main/docs/harvey-lab-audit/litigation-dispute-resolution/human-review.

[3] Claude Sonnet 5.5 and Claude Opus 5.5 citation-check agents, coordinated by the author. Citation check of the Harvey LAB litigation and dispute resolution task environments. *AI-generated research artifact, not attorney-reviewed*, 2026. https://github.com/johnhughes3/LegalForecastBench/tree/main/docs/harvey-lab-audit/litigation-dispute-resolution/citation-check.

[4] Harvey AI task authors; upstream authorship not individually listed. Pinned LAB task.json for analyze-counterparty-motion-to-dismiss. *Primary online source*, 2026. https://github.com/harveyai/harvey-labs/blob/1dd81403b2fbb60596f7aea3fcecafad7bf73143/tasks/litigation-dispute-resolution/analyze-counterparty-motion-to-dismiss/task.json.

[5] Gabe Pereyra. How Harvey Built a Research Lab on a Budget. *Sequoia Capital, Own Your Intelligence*, August 11, 2026. Slide at 4:11: "Scale with expert data providers." https://www.youtube.com/watch?v=MGouk8W51v0&t=251s.

[6] Niko Grupen, Gabe Pereyra, Julio Pereyra. Introducing Harvey's Legal Agent Benchmark. *Primary online source*, 2026. https://www.harvey.ai/blog/introducing-harveys-legal-agent-benchmark.

[7] OpenAI. Why SWE-bench Verified no longer measures frontier coding capabilities. *Primary online source*, 2026. https://openai.com/index/why-we-no-longer-evaluate-swe-bench-verified/.

[8] Artificial Analysis. Harvey LAB-AA evaluation results. *Leaderboard, accessed October 5, 2026*, 2026. https://artificialanalysis.ai/evaluations/harvey-lab-aa.

[9] John J. Hughes, III. Law as a Verifiable Domain: Using Real-World Judicial Outcomes to Evaluate AI Legal Reasoning. *Working paper*, 2026. https://www.legalforecastbench.org/paper/.

A Model outputs and native LAB judging

To understand how the rubrics evaluated actual AI-generated work product, we ran LAB using GPT-6 Luna (xhigh thinking) and Claude Opus 5.5 (low thinking) on the litigation tasks (Harvey commit `1dd81403`). LAB's standard judge pair, Claude Sonnet 4.6 and GPT-5.5, graded those outputs. The examples in Appendix B that I reviewed include these outputs, so I could consider how actual model outputs are being graded in evaluating the rubrics.

Solver	Reasoning effort	Task outputs	Mean criterion pass rate
GPT-6 Luna	xhigh	52	83.80%
Claude Opus 5.5	low	52	90.81%

Table 2: Native LAB AI-judge grades on the 52 public litigation tasks. For each task, we calculate each judge's fraction of criteria passed, average the two judges' fractions, and then average across tasks with equal task weights. These are criterion pass rates, not the fraction of tasks passing every criterion.

These grades reflect how LAB's standard pair of AI judges (Claude Sonnet 4.6 and GPT-5.5) applied the original rubrics. Since the models that prepared these outputs used different reasoning settings, the difference in their scores here is not a controlled comparison of model capability; instead, it was merely used as a diagnostic.

B Human review of AI-flagged LAB criteria

I asked two frontier models, GPT-6 Sol and Claude Opus 5.5, to audit the 52 tasks in LAB's litigation tasks (specifically Harvey commit `1dd81403`). Table 3 reports how many criteria each pass flagged as defective or arguably defective. The combined row is Opus's later review of rubrics already flagged by either pass. A criterion "flagged by either model" was flagged at least once; that is 835 of the 2,858 criteria (29.2%) [1].

Pass	What the model saw	Criteria flagged
GPT-6 Sol	The tasks, without Opus's results	616
Claude Opus 5.5, first pass	The tasks, without Sol's results	650
Claude Opus 5.5, combined	The rubrics flagged by either pass	564
Flagged by either model	Share of all 2,858 criteria	835 (29.2%)

Table 3: AI audit of the 52 litigation tasks at Harvey commit `1dd81403`.

In order to validate the AI findings and determine how many of the AI-flagged rubrics were genuinely defective, I reviewed a simple random sample of 25 of the flagged rubrics, drawn with a fixed seed from 19 tasks. I committed the protocol for that review before it began [2]. For each rubric, I reviewed the grading rubric itself, the findings from both AI auditors, how LAB's AI judges graded two real model runs (GPT-6 Luna and Opus 5.5) on that rubric, and the relevant documents from the task environment. In a few cases I also did research into the relevant governing law. I marked a rubric as defective when it could produce an inaccurate grade, meaning: either the rubric could fail a response that a thoughtful litigator would consider appropriate, or the rubric could pass a response that a thoughtful litigator would consider wrong on the issue the rubric is testing. A rubric that forces one of several permissible answers is therefore defective, because it would fail the others.

During the review, I decided to distinguish between defects in the rubrics themselves and defects in the task environment. The line between the two was based on whether the defect could change the grade a model is given on the task. For example, one rubric wanted the work product to contain a specific metric sourced from the task environment, but the AI auditors pointed out that the task environment actually contained two contradictory numbers for that particular metric. Accordingly, an AI agent that relied on one of the numbers (which was not the one the rubric was expecting) would be failed incorrectly because of the underlying defect in the task environment. I treated that as a defect in both the rubric and the task environment. By contrast, a different task environment contained a notice of an OSHA regulatory violation from a state-level regulatory agency in Ohio that does not actually have jurisdiction over the violations charged. The rubric wanted work product to simply flag the existence of the charge in the work product. I coded that as a defect in the task environment but not as a defect in the rubric (because an AI agent handed that task could, and likely still would, pass the rubric by simply mentioning the charged violation).

Of the 25 criteria, I graded 16 rubrics defective in ways that affect grading, including 9 that I considered to be unambiguous objective errors in which the record, the law, the cited authority, or the arithmetic contradicts the rubric. 3 were arguable defects, and 6 were not defective. 14 of the 25 come from task environments that contain defects (10 of the 19 tasks). Only 4 had neither a defective nor an arguable criterion and came from an environment without a defect (Table 4).

Since the sample was a draw without replacement from the 835 criteria that AI agents had flagged as defective, the one-sided 95% lower confidence bounds are: at least 383 flagged criteria are defective

in ways that affect grading (point estimate 534; the bound is 13.4% of all criteria in the practice area); at least 171 are unambiguous objective errors (point estimate 301); and at least 319 come from task environments that contain defects (point estimate 468). At least 524 involve a defective criterion, a defective environment, or both (point estimate 668). In the sample, 9 of the 10 criteria flagged by both auditors were defective, against 7 of the 15 flagged by only one.

I did not attempt to estimate defects in criteria that were not flagged by the AI agents; it is possible there could be some additional defects that the agents did not flag. These estimates count criteria/rubrics, not the individual mistakes that give rise to the defects. In some cases, the defect was specific to the particular rubric at issue, but in other cases, a single error in a task environment could make several criteria defective. After I prepared an initial set of verdicts, I asked AI models (GPT-6.1 Sol and Opus 5.5) to stress test my answers and to draft additional arguments both for and against each call and a summary of those they would change. I changed a handful of verdicts after reflecting on these points, and did a second pass through the full set of my verdicts to ensure that they applied the principles I adopted during the review consistently to each call.

The audit of rubrics revealed instances where the task environment, and sometimes grading rubrics, appeared to rely on AI-generated hallucinated case law citations. As a result, I had AI agents (Opus 5.5) conduct a separate AI-assisted check of every case citation in the 19 sampled task environments against CourtListener. One agent checked each task. A synthesis agent re-verified a sample of the findings, and I independently re-checked seven of them. This process found problems in 45 of 102 citations to real authority, including 7 cases that could not be found and 32 real cases cited for propositions they do not contain [2].

Criterion verdict	Environment defective	Environment correct	Total
Defective	10	6	16
Arguable	2	1	3
Not defective	2	4	6
Total	14	11	25

Table 4: My verdicts on the 25 sampled criteria, by whether the task environment itself contains a defect.

The sampled criteria

Each entry has four parts: (1) the criterion quoted verbatim from Harvey's task file (that text is from Harvey's data vendors, not me, and in many cases the text appears to be AI-generated), (2) my verdict on the criterion and task environment, (3) results from a sample run of LAB using GPT-6 Luna and Claude Opus 5.5, to assess how LAB's AI graders (Claude Sonnet 4.6 and GPT-5.5) graded outputs from Luna and Opus on the specific criteria at issue, (4) my explanation of the verdict. Harvey task files are cited at commit `1dd81403`.

1. Assess Reasonableness of Staffing Levels on Litigation Invoice (C-049)

Quoted from Harvey LAB

> **Calculates budget overage percentage as approximately 17%.** PASS if the memo calculates the percentage by which the attorney fee subtotal ($193,372.50) exceeds the approved budget ceiling ($165,000) as approximately 17% (exact figure is 17.2%). FAIL if no percentage overage calculation is provided or if the percentage stated is materially

incorrect (off by more than 2 percentage points).

— Harvey LAB, criterion C-049, task file at commit 1dd81403

My assessment	
Criterion	Defective — unambiguous objective error, deficiency in task environment
Task environment	Defective: the invoice's 117 time entries add up to $215,680.00, while its summary and subtotal show $193,372.50.
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric is testing the AI on a specific percentage calculation from an invoice, but the document in the environment is internally inconsistent. The invoice's summary page and its subtotal row show fees of $193,372.50 (17.2% over the $165,000 budget), but the 117 detailed time entries add up to $215,680.00, which puts the overage at $50,680, or 30.7% (13.5 points from the rubric's required 17.2% answer and far outside the 2-point tolerance).

2. Categorize Document Production Set by Relevance and Privilege — Privilege Log and Relevance Classification Report (C-012)

Quoted from Harvey LAB

ISSUE_003: DOC_012 identified as Partially Privileged. PASS if DOC_012 (chandrasekaran-meeting-notes-aug28.docx, Chandrasekaran's handwritten meeting notes from August 28, 2023) is classified as Partially Privileged. FAIL if classified as fully Privileged or Not Privileged.

— Harvey LAB, criterion C-012, task file at commit 1dd81403

My assessment	
Criterion	Defective — low-quality professional judgment, rubric is incorrect
Task environment	Correct
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric is wrong because it fails fully Privileged, which is the better answer here (even if Partially Privileged is defensible).

- The documents are handwritten lawyer notes from a fictitious meeting discussing problems with the R&D Department. The headings on individual topics seem intended to create the impression that there is a discussion of business topics (which would not be privileged) followed by a separate discussion of legal topics (privileged), so from a superficial glance, the rubric does seem to match the topic headings. However, the bullets under each heading make clear that the entire meeting is just dedicated to the topic of the company's desire to find a way to fire the head of R&D without getting sued. The business topics (which look to be nonprivileged) are discussions of how the R&D department is a bit of a disaster, such that management

wants to fire the head. The legal sections then involve discussion of the fact that the R&D head has some sort of environmental whistleblower status, so termination is likely to trigger a lawsuit/whistleblower retaliation claim.

- The meeting was specifically convened because outside counsel was concerned about the retaliation claim, and even some of the 'ops' sections that are supposed to be nonprivileged reflect counsel's advice, so they cannot be cleanly treated as nonlegal business discussions.
- While partially privileged might be defensible, my preferred answer would be fully privileged, which the rubric incorrectly fails.

3. Compare Document Production Against Discovery Requests — Discovery Gap Analysis Memorandum (C-029)

Quoted from Harvey LAB

> **ISSUE_011: Notes only 5 documents actually relate to counterclaim damages.** PASS if the memorandum notes that only 5 of the 14 documents coded to RFP 27 (MLC-004056 through MLC-004102) actually relate to Pacific Corridor's volume shortfall and the counterclaim damages calculation, meaning the counterclaim damages support is thin. FAIL if this is not noted.
>
> — Harvey LAB, criterion C-029, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Defective: the production index codes 15 rows to RFP 27, but its summary tab gives a count of 14 for a range that contains 13.
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. This task requires models to make the point that "only" five documents in the bundle "actually relate to Pacific Corridor's volume shortfall and the counterclaim damages calculation, meaning the counterclaim damages support is thin." This rubric is flawed on multiple levels.

- First, it overlooks MLC-003781, a sixth document that plainly relates to the volume shortfall, so the count is objectively wrong (and the agents pointed to several other examples, at least some of which also seem to fit within the stated criteria). There is a document in the bundle that claims that only five documents concern the volume shortfall, but that assertion is undercut by other evidence supplied in the environment.
- Second, this rubric is someone's argumentative characterization of the documents and probably would not hold up under adversarial testing. For example, there is case law that the term "relate to" is fairly broad and flexible, so quite plausibly a number of other documents would apply.
- Third, the premise underlying this rubric–that "only" five documents is such "thin" evidence that it must be called out–is wrong. Even one document could constitute conclusive proof of damages (the general rule is that damages only need to be reasonably estimated; courts do not require great precision); to my mind, this is actually quite a lot of evidence at the fact

stage compared to what one would normally see.

4. Draft Case Assessment Memorandum — Litigation Risk Analysis for Distribution Agreement Dispute (C-025)

Quoted from Harvey LAB

> **ISSUE_007: Recommends motion to dismiss on unjust enrichment.** PASS if the memo recommends filing a motion to dismiss (or motion for judgment on the pleadings, or similar dispositive motion) targeting the unjust enrichment claim (Count IV) on the ground that an express contract governs the parties' relationship. FAIL if no such motion recommendation is made regarding the unjust enrichment claim.
>
> — Harvey LAB, criterion C-025, task file at commit `1dd81403`

My assessment	
Criterion	Defective — low-quality professional judgment, rubric is incorrect
Task environment	Correct
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. A different rubric on this task concludes that plaintiff has a strong breach of contract claim. (See id: C-006.) The prompt instructs models to review case documents and "prepare a comprehensive case assessment memo," and models are dinged on this rubric if they do not "recommend[] filing a motion to dismiss" on the unjust enrichment claim, which is expressly pled in the alternative to the (strong) breach of contract claim.

- The rubric is wrong because a careful litigator could reasonably conclude that a motion to dismiss that claim is not a good idea: Count IV is pleaded in the alternative (Complaint ¶127) and seeks restitution of "no less than $22,521,312.50" (¶135), which is exactly the Count I contract damages ($17,721,312.50, ¶94) plus the Count III trade-secret damages ($4,800,000, ¶120). Counts I and III would survive the motion and double recovery is barred, so dismissing Count IV would not reduce the defendant's net exposure; and the factual underpinnings of both claims are the same so a successful motion would not reduce the scope of discovery. Filing a dispositive motion is costly and, on these facts, there seems to be no benefit to doing so. So the required recommendation would fail many thoughtful analyses that could discuss the pros and cons and recommend against filing such a motion.

5. Draft Conflict Check Memorandum for Litigation Engagement Clearance (C-030)

Quoted from Harvey LAB

> **ISSUE_005: Recommends screening or removal of Caleb Strand.** PASS if the memo recommends that Caleb Strand be screened from the engagement or removed from the proposed staffing. FAIL if no such recommendation is made.
>
> — Harvey LAB, criterion C-030, task file at commit `1dd81403`

My assessment	
Criterion	Defective — low-quality professional judgment, rubric is incorrect
Task environment	Correct
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric requires the model to recommend that an associate who shares an apartment with his sister (an in-house IP paralegal at TriPoint, the adverse party) be screened from the matter. Screening is generally used in certain specific situations where required by the professional conduct rules, but no rule requires screening in this scenario.

- I agree that a prudent memo would flag the fact that a junior team member is sharing an apartment with an employee of the adverse party, but that is covered by other rubrics (C-026 through C-029). This rubric specifically fails the work product (and results in a 0 for all work on this task) even if the model recommended, say, obtaining the client's informed consent, which would be perfectly permissible and appropriate here.
- Other rubrics in this task resolve similar kinds of conflicts with informed consent (C-024 accepts "clearance with conditions (such as obtaining informed consent)" for Lisa Chow's spousal conflict, and C-016 recommends informed consent for the Ridgeline conflict). Accordingly, informed consent plus a home confidentiality protocol could suffice as an alternative to screening.

6. Draft Conflict Check Memorandum for Litigation Engagement Clearance (C-049)

Quoted from Harvey LAB

> **ISSUE_011: Recommends integrating annual disclosure questionnaires with ConflictTracker.** PASS if the memo recommends that annual conflict-disclosure questionnaires be cross-referenced with or integrated into the ConflictTracker system. FAIL if no recommendation about integrating disclosure questionnaires with the conflicts database is made.
>
> — Harvey LAB, criterion C-049, task file at commit `1dd81403`

My assessment	
Criterion	Defective — requires specific outputs not expressly or impliedly requested in the prompt
Task environment	Correct
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The assignment in the prompt is to draft a "comprehensive conflict check memorandum for the proposed engagement." The rubric requires that the memo contain suggested firm-wide process improvements (memo must "recommend[] that annual conflict-disclosure questionnaires be cross-referenced with or integrated into the ConflictTracker system"). That's simply not the

assignment, and there's many reasons why it's undesirable to jam broad process suggestions into a conflicts memo for one specific case. This rubric actually failed work product from Luna that did everything a clearance memo needs on this point. The two judges split on that same Luna memo (Sonnet 4.6 failed it; GPT-5.5 passed it), which shows the criterion does not grade reliably.

7. Draft Counterclaim Against Plaintiff for Breach of Joint Development Agreement (C-045)

Quoted from Harvey LAB

> **Count: Patent Infringement (ISSUE_005).** PASS if the counterclaim includes a separate count for patent infringement under 35 U.S.C. § 271, alleging that the LumiSense 400 infringes U.S. Patent No. 11,234,567 covering the SensorCore™ architecture. FAIL if no patent infringement counterclaim is asserted.
>
> — Harvey LAB, criterion C-045, task file at commit `1dd81403`

My assessment	
Criterion	Arguable
Task environment	Defective: the record lacks the patent the counterclaim must plead, and the two expert reports give different titles for it.
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The task is to draft a counterclaim based on a handful of documents (including an expert report, some emails, and a counterclaim strategy memo that another lawyer prepared). The rubric dings the AI if the counterclaims do not include "a separate count for patent infringement under 35 U.S.C. § 271, alleging that the LumiSense 400 infringes U.S. Patent No. 11,234,567."

- The environment does not contain the allegedly infringed patent, which one would need to review in order to have a sufficient basis for drafting that counterclaim, so this environment appears to be incomplete. That said, a defensible approach might be for the agent to draft the counterclaim but flag in brackets that it requires verification.

8. Draft Counterclaim Against Plaintiff for Breach of Joint Development Agreement (C-046)

Quoted from Harvey LAB

> **Patent Infringement Count: Patent Number Cited.** PASS if the patent infringement count specifically identifies U.S. Patent No. 11,234,567. FAIL if the patent number is omitted or incorrect in the patent infringement count.
>
> — Harvey LAB, criterion C-046, task file at commit `1dd81403`

My assessment	
Criterion	Arguable
Task environment	Defective: same task environment as item 7.
AI auditors' reasoning	Partly correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. Same issue as prior rubric. This one is essentially derivative because it depends on pleading the patent count that was required in the prior rubric.

9. Draft Case Assessment Memorandum for Defective Industrial Equipment Product Liability Claim (C-003)

Quoted from Harvey LAB

> **Executive Summary mentions death of Marco Reyes.** PASS if the Executive Summary mentions the death of Marco Reyes. FAIL if the fatality is not mentioned in the Executive Summary.
>
> — Harvey LAB, criterion C-003, task file at commit `1dd81403`

My assessment	
Criterion	Defective — requires specific outputs not expressly or impliedly requested in the prompt
Task environment	Defective: same task environment as item 10 (Ohio state agency issuing OSHA citations).
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. The case involved a (fictitious) safety incident at a manufacturing plant that killed one employee and seriously injured two others. The rubric requires the fact that one employee died be discussed specifically in the "executive summary." When we tested Opus 5.5 and GPT-6 Luna on this benchmark, both discussed the death in several places in their memo, especially the factual background section, but the executive summary consisted of an overview of the legal analysis that did not restate the fact of the death. They were failed by one or more judges.

- This fails perfectly competent and appropriate work product. The company CEO who retained counsel to advise on this incident is the same person who delivered news of the (fictitious) tragic death to the decedent's widow and surviving family members, so the most senior management of the company is obviously very familiar with the incident and the agents' approach of focusing the executive summary on the legal exposure is appropriate.
- The two judges reached opposite results on Luna's memo: GPT-5.5 failed it, while Sonnet 4.6 passed it only by quoting a sentence from the memo's facts section as if it appeared in the executive summary.

10. Draft Case Assessment Memorandum for Defective Industrial Equipment Product Liability Claim (C-026)

Quoted from Harvey LAB

> **OSHA LOTO citation identified with specific violation and penalty.** PASS if the memo identifies the OSHA serious citation for LOTO deficiency under 29 CFR 1910.147(c)(4)(i) with a penalty of $18,900. FAIL if this specific citation, regulation, or penalty amount is not referenced.
>
> — Harvey LAB, criterion C-026, task file at commit `1dd81403`

My assessment	
Criterion	Not defective — defective environment, but rubric does not fail correct work product
Task environment	Defective: Ohio's Division of Safety & Hygiene cannot issue OSHA citations to this private employer. The criterion itself does not misgrade work (environment-only defect).
AI auditors' reasoning	Wrong
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. Although the specific rubric being tested does not fail work product incorrectly, the environment is defective in an unintended way. The environment contains an OSHA citation from Ohio's Division of Safety & Hygiene. But Ohio does not have an OSHA-approved State Plan, so only federal regulators have jurisdiction to issue OSHA citations against this company. The environment does not test that so it is not being presented as a "mistake" by one's adversary that the agents should pick up on.

- The rubric itself is not wrong because it simply wants a memo to address the purported citation, so an agent could write the memo while noticing that the task is obviously a fictitious evaluation with an incoherent environment.

11. Draft Discovery Plan Memorandum for Breach of Contract and Fraud Defense (C-036)

Quoted from Harvey LAB

> **Provides custodian priority rankings.** PASS if the memo provides priority rankings or tiers for custodians (e.g., high/medium/low priority, or Tier 1/Tier 2, or similar prioritization). FAIL if custodians are listed without any indication of relative priority.
>
> — Harvey LAB, criterion C-036, task file at commit `1dd81403`

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The AI agent is right that the rubric is really enforcing a stylistic choice, but the choice here is well-grounded in convention and an appropriate expectation for quality work product. The case file contains a court order where the judge encouraged the parties to discuss whether phased discovery would be appropriate at the Rule 26(f) conference, and there is some existing prioritization work in the environment. The rubric requires that the agent's discovery plan include some kind of prioritization or phasing of custodians, which is a common expectation in these cases. I think it's fair to judge anything that lacks any form of prioritization as deficient given the context.

12. Draft Federal Complaint for Breach of Contract and Fiduciary Duty — Placement Agent Dispute (C-025)

Quoted from Harvey LAB

> **ISSUE_012: Complaint requests injunctive relief.** PASS if the prayer for relief includes a request for preliminary and/or permanent injunctive relief — specifically to enforce the non-solicitation provision and prevent Graydon from continuing to divert Fund III investor prospects to competing funds. FAIL if no injunctive relief is requested.
>
> — Harvey LAB, criterion C-025, task file at commit `1dd81403`

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong

LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The client intake memo specifically describes an interest in seeking injunctive relief ("We may also wish to seek preliminary and permanent injunctive relief to enforce the twelve-month non-solicitation tail"), and the client's demand letter threatens it. While it's true that the intake memo flags the topic for discussion, the correct approach is to include the proposed prayer for relief in the draft, so the case team can see how it reads and discuss any downsides. As a practical matter, there's unlikely to be much downside to including this prayer for relief in the complaint. (Whether it's worth the cost of actually pursuing a preliminary injunction on this is debatable, so the rubric's approach makes sense.)

13. Draft First Set of Interrogatories to Defendant Veridian Health Systems in Trade Secret Misappropriation Case (C-008)

Quoted from Harvey LAB

> **Definitions section defines 'Document' or 'Documents'.** PASS if the Definitions section includes a definition of 'Document' or 'Documents' that is broad enough to encompass electronically stored information (ESI), emails, and physical documents. FAIL if there is no definition of 'Document' or it is so narrow as to exclude electronic records.
>
> — Harvey LAB, criterion C-008, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Defective: the supplied documents mischaracterize or miscite case law (citation sweep; peripheral to this criterion).
AI auditors' reasoning	Partly correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric fails any draft that does not supply a definition of "document" encompassing electronically stored information (ESI). There are two problems:

- First, these are interrogatories, not document requests. Well-drafted interrogatories do not necessarily need to use the word "document." The rubric does not really make sense as a requirement for this kind of work product.
- Second, even if we were dealing with document requests here (which we are not), the local rules of many courts, including the Western District of Texas (where this fictitious case is being litigated), define the term "document" for all discovery requests served by litigants in those cases. There is thus no requirement to define the term "documents," and in fact, supplying a definition could just trigger objections from the other side that the local rule controls. See Local Rule CV-26, available at https://www.txwd.uscourts.gov/court-information/lcr-civil-rules/ ("The term 'document' means any document or electronically stored information as described in Federal Rule of Civil Procedure 34(a). A draft of a nonidentical copy is a separate document within the meaning of this term.").
- The rubric does demonstrably fail rules-compliant work product. For example, Luna was failed for not including a definition of the term "Document" but that's actually what the local rules encourage in this Court.

14. Draft First Set of Interrogatories to Defendant Veridian Health Systems in Trade Secret Misappropriation Case (C-026)

Quoted from Harvey LAB

> **Interrogatory re: USB drive or CMT files on Veridian systems.** PASS if at least one interrogatory asks whether Oshiro's USB drive (referencing serial number CX-8827491 or otherwise describing the external storage device) was ever connected to, accessed from, or present on any Veridian system, device, or network, or asks whether any CMT files or materials were ever possessed, accessed, viewed, or copied by Veridian or its employees. FAIL if no interrogatory addresses the USB drive or CMT files on Veridian systems.
>
> — Harvey LAB, criterion C-026, task file at commit `1dd81403`

My assessment	
Criterion	Not defective
Task environment	Defective: same task environment as item 13 (environment-only defect).
AI auditors' reasoning	Wrong
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. This one does not really belong on the list, as GPT-6 Sol was complaining about a different rubric (C-048) which it argues is inconsistent with this one. It does not identify any issue in this rubric, standing alone, so this counts as an instance where Sol added something to the "flagged" list that does not belong.

15. Draft Discovery Responses and Objections to RFAs and RFPs in Breach of Supply Agreement Litigation (C-028)

Quoted from Harvey LAB

> **ISSUE_011: Identifies RFA No. 20 as calling for a legal conclusion.** PASS if the plan identifies that RFA No. 20 (asking Prismavale to admit the products were 'not merchantable as defined by the Uniform Commercial Code') calls for a legal conclusion regarding UCC § 2-314 merchantability and recommends objecting on that basis. FAIL if the plan does not identify the legal conclusion problem with RFA No. 20.
>
> — Harvey LAB, criterion C-028, task file at commit `1dd81403`

My assessment	
Criterion	Defective — low-quality professional judgment, rubric is incorrect
Task environment	Defective: the supplied Answer cites complaint paragraph ranges that do not match the complaint (e.g. Count III as ¶¶ 88–94; the complaint's Count III is ¶¶ 86–96).
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. Plaintiff served a request for admission on the defendant that says, "Admit that the SB-102 and PS-302 products shipped by Prismavale to Greenleaf were not merchantable as defined by the Uniform Commercial Code." The prompt asks the agent to "prepare a discovery strategy memo and formal RFA/RFP responses." The rubric requires "the plan" (i.e., the strategy memo) to recommend objecting to the request on the grounds that it calls for a legal conclusion. The rubric is wrong for two main reasons:

- First, the modern view is requests for admissions can call for application of law to facts. While it is within acceptable professional norms for a lawyer to object on the ground that such a request calls for a legal conclusion, and then proceed to answer, the better reading of the law is that this request is not objectionable such that the party should simply answer. On this legal conclusion in particular, there's really no reason why one would prefer to object rather than simply deny the requested admission. So this rubric fails competent work product (and

probably the better answer than the one the rubric calls for).

- Second, even assuming one should include the objection at all, the task fails the model if the issue is not expressly addressed in the strategy memo. Luna made the "legal conclusion" objection in the draft responses and then proceeded to a substantive denial (an acceptable approach), but did not discuss the "legal conclusion" objection in the memo. That was judged a failure. That's really requiring duplicative responses on a very minor point that probably shouldn't be in the drafts at all.

16. Draft Motion in Limine to Exclude Expert Testimony and Prejudicial Evidence in Commercial Breach of Contract Case (C-009)

Quoted from Harvey LAB

> **ISSUE_004: Identifies the $11,000 arithmetic discrepancy.** PASS if the motion identifies that the sum of Marchetti's five itemized damages categories totals $14,211,000, but her report states total damages as $14,200,000, creating an $11,000 discrepancy. FAIL if this arithmetic error is not identified.
>
> — Harvey LAB, criterion C-009, task file at commit `1dd81403`

My assessment	
Criterion	Defective — low-quality professional judgment, rubric is incorrect (a judgment call, not an objective error); defective environment (directs advocacy outside professional norms); companion rubric C-010 also incorrect (not sampled)
Task environment	Defective: the task directs advocacy outside professional norms (an $11,000 rounding point offered as a *Daubert* ground) and labels a rounding difference an "arithmetic error."
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The damages expert includes a summary table in her report listing various items of damages, which collectively add up to $14,211,000; the Total Damages line lists the total as $14,200,000. Opposing counsel questioned the expert about it at her deposition, and the expert calls it "a rounding discrepancy." An internal strategy memo highlights the issue as a ground for seeking to exclude the expert's testimony under *Daubert*, arguing "the arithmetic error supports a cumulative argument that the report was prepared without the rigor required of expert testimony in federal court."

- This task is poorly specified, and the recommended strategy of raising this with the judge is a poor one:
- First, the task environment repeatedly states that this is an "arithmetic error." That is false. The expert has rounded the total in the table to the nearest $100,000. While that is somewhat unorthodox in an expert report (particularly since it is written out as $14,200,000 not $14.2 million), numbers often are rounded in court filings so the judge is unlikely to agree that this is an error.
- Second, even if there were an error (which there is not), that is generally not a basis for

excluding the expert's testimony under *Daubert*. At most one could move to exclude the erroneous version and ask for more accurate testimony, but that does nothing to help the hypothetical client here; it just raises their damages exposure by the $11,000 that the expert was willing to round off.

- Third, the proposed framing comes close to a breach of decorum and carries a real risk of drawing a rebuke from the judge. The strategy memo wants the motion to argue the expert's 30 prior engagements were "in state courts and arbitrations that do not impose the same methodological gatekeeping requirements" as federal courts. While some state courts use the *Frye* standard instead of *Daubert*, many federal judges would take offense at an argument that state courts do not observe adequate methodological rigor. And the reality is that many busy judges would be annoyed at spending time on an argument that literally amounts to a rounding error. It's so trivial that raising it as a substantive point risks making the judge think you had no better argument to put forward, so your motion must be weak.
- This rubric is scored as Defective, as a professional-judgment call rather than an objective error. Our view is that the better approach is to drop this issue entirely, and the sampled rubric (C-009) fails a motion that does so, even though a motion that merely mentions the discrepancy in a footnote would pass. The companion rubric, C-010, is in my view objectively wrong because it requires that this issue be presented as a basis for excluding the expert's testimony under *Daubert*. Luna was failed incorrectly on C-010 (by the Sonnet 4.6 judge) for downplaying this issue; in fact, Luna was right to do so, as it should not be raised at all. But since that's not the sampled criterion, we do not mark it here.

17. Draft Rule 12(b)(6) Motion to Dismiss Brief — Commercial Software Licensing Dispute (C-037)

Quoted from Harvey LAB

> **ISSUE_008 — Cites Delaware parol evidence authority.** PASS if the motion brief cites relevant Delaware authority on the parol evidence rule and integration clauses, such as SIGA Technologies v. PharmAthene or Eagle Industries v. DeVilbiss Health Care. FAIL if no Delaware parol evidence authority is cited.
>
> — Harvey LAB, criterion C-037, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error (including hallucinations), rubric is incorrect
Task environment	Defective: the supplied research memo cites fabricated or mischaracterized authority (citation sweep: 28 of 53 citations in this task have problems).
AI auditors' reasoning	Correct

LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. The rubric here involves a legal error, and the task environment contains hallucinations about case law. This one is extremely defective for multiple reasons:

- First, the rubric is demanding case law on the wrong legal doctrine. Under Delaware law,

the parol evidence rule and an integration clause do not preclude fraud claims. The issue here is whether the plaintiff can reasonably allege that they relied on their counterparty's pre-signing representations despite their agreement in the contract that they were not relying on any representations other than those written in the contract itself.

- Second, the cases in the environment appear to contain multiple hallucinations. The rubric refers to SIGA v. PharmAthene, but Opus noted that case does not address integration or parol evidence. The task environment contains a legal research memo that cites *SIGA Technologies, Inc. v. PharmAthene, Inc.*, 67 A.3d 330 (Del. 2013), and claims that it states (at page 344), "an integration clause is the clearest declaration that the parties intend the writing to be the complete and final expression of their agreement." The quoted language does not appear in the cited case. Page 344 of that case addresses the duty to negotiate in good faith, not integration clauses. Similarly, the memo cites *Eagle Industries, Inc. v. DeVilbiss Health Care, Inc.*, 702 A.2d 1228 (Del. 1997) and says: "The Delaware Supreme Court held that 'a party to a contract cannot promise, in a clear integration clause of a negotiated agreement, that it is not relying on promises or representations made by the other party outside of the agreement, and then assert a claim for fraud based on those very representations.' *Id.* at 1232." The quoted language does not appear in the cited case.
- Third, the rubric failed responses that stated the correct law on this point. In our tests, Opus reviewed SIGA and Eagle, correctly realized they were inapposite and declined to cite them, cited more on-point and accurate authority, and was then failed for doing a better job than the rubric's author did at this task.

18. Draft Rule 12(b)(6) Motion to Dismiss Brief — Commercial Software Licensing Dispute (C-045)

Quoted from Harvey LAB

> **ISSUE_010 — References Arcadia's own delays contributing to Go-Live delay.** PASS if the motion brief identifies Arcadia's own contributing failures, including the 45-day Project Manager gap (Kevin Liu's departure) and/or the 75-day late API specifications, as causes of the Go-Live delay. FAIL if Arcadia's own contributory delays are not mentioned.
>
> — Harvey LAB, criterion C-045, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Defective: same task environment as item 17.
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. The rubric makes a rookie mistake, demanding that the agent argue a factual dispute on a motion to dismiss, where the court is required to assume the plaintiff's allegations are true and cannot consider factual dispute. The 45-day and 75-day delays complained of come from the defendant's internal report, which the plaintiff's complaint disputes. Moreover, courts generally cannot resolve messy fact disputes over "causation" on a motion to dismiss.

- The criterion failed agents that correctly recognized that this argument was outside the scope of what can be argued on a motion to dismiss and thus decided to exclude it.

19. Draft Verified Responses and Objections to Plaintiff's First Set of Interrogatories in Commercial Breach of Contract and Fraud Action (C-028)

Quoted from Harvey LAB

> **Correct financial figures for Tri-Basin purchase history.** PASS if any purchase history figures cited in the responses are consistent with the canonical data: 2015: $9.2M; 2016: $10.1M; 2017: $10.8M; 2018: $12.4M; 2019: $13.6M; 2020: $9.7M; 2021: $15.3M; 2022: $16.8M; 2023 (Jan–Aug): $11.9M. Minor rounding differences are acceptable. Also PASS if the response invokes Rule 33(d) rather than citing specific figures. FAIL if specific purchase figures are cited but are materially incorrect.
>
> — Harvey LAB, criterion C-028, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error, deficiency in task environment; rubric is incorrect
Task environment	Defective: the purchase-history spreadsheet is internally inconsistent (three 2020 series; line items that do not sum to subtotals), and the record has no contract-year data.
AI auditors' reasoning	Partly correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. There are two issues:

- First, Interrogatory No. 4(b) asks for Tri-Basin's actual purchases for the 2020 contract year (the contract year runs from January 15 through January 14 of the following year; Interrogatory No. 14 separately asks for calendar-year totals), but the canonical data referred to in the rubric is calendar year data, so the rubric's numbers are not correct for that request. The correct value cannot be discerned from the environment: the monthly data cannot be split at January 15, and nothing in the record supports assuming purchases are spread evenly within January (seasonal variation would make that unlikely). The rubric therefore grades wrong in the opposite direction from most of the other items: rather than failing good work product, it passes responses that supply the calendar-year figure for a contract-year request without acknowledging that the record does not contain the requested data. That is approving fundamentally deficient work product.
- Second, the data that is provided has some internal inconsistencies. For example, an "annual summary" tab lists 2020 purchases at $9.7 million, but the monthly detail lists 2020 three times, with totals of $8,700,000, $8,549,400 and $9,700,000. The units sold do not match up between those two sheets for 2021, 2022, 2023, and series 3 of 2020. One of the monthly rows does not add up (gross $1,210,000 + credits −$17,400 = $1,192,600, but the net shown is $1,192,500). The annual summary minimums do not add up. In the Product Line Breakdown, line-item net revenue never sums to the subtotal. So a response that relied on the line items would get graded as a failure.

- The environment contains law firm discovery guidelines that recommend that interrogatory responses match the client's Answer, but the Answer here contains only approximate figures, which is consistent with the Answer using calendar year data as a rough approximation instead of providing the contract year data requested.

20. Extract Key Obligations from Litigation Hold and Document Preservation Notice — Obligation Summary Memorandum (C-002)

Quoted from Harvey LAB

> **ISSUE_001: Flags financial records date discrepancy as requiring DOJ clarification.** PASS if the memo recommends seeking clarification from the DOJ regarding the conflicting date ranges (Jan 1, 2017 vs. Jan 1, 2019) for financial records. FAIL if no recommendation to seek clarification is made.
>
> — Harvey LAB, criterion C-002, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Correct
AI auditors' reasoning	Correct

LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: fail

My analysis. The task involves a document preservation notice from federal criminal authorities that defines the relevant period (for which the recipient must preserve documents) as January 1, 2019, through March 3, 2025 "unless otherwise specified herein." For one specific category of documents, the notice states that, "Notwithstanding the Relevant Period defined in Paragraph 12 of this Notice, Ridgeline shall preserve [certain kinds of] records . . . from January 1, 2017 to the present date."

- The rubric suggests that the agent should seek clarification of the "conflicting" date ranges but there is no conflict. The notice quite clearly imposed a longer period for one specific category of documents, while all the other categories are subject to a shorter period. This is a fairly common request.
- All the models were failed on this criterion, but the criterion makes no sense.

21. Identify Excessive or Duplicative Research Charges in Litigation Invoice (C-013)

Quoted from Harvey LAB

> **Issue 3: Identifies duplicative MTCA research across Takahashi and Wendt.** PASS if the memo identifies that multiple attorneys (Kenji Takahashi and Tyler Wendt) billed for overlapping MTCA liability/contractor liability/summary judgment research during July 2024, citing Section 6.4 (Duplicative Research Prohibition). FAIL if the MTCA duplication across timekeepers is not identified.
>
> — Harvey LAB, criterion C-013, task file at commit `1dd81403`

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. While Sol is right that the facts here are insufficient to establish that the client's billing guidelines were violated, for purposes of the task (a client-side memo identifying changes to dispute), it seems reasonable to flag that multiple attorneys were doing research into closely related issues. The task requested "a memo identifying all non-compliant or excessive charges with recommended adjustments," so flagging potential violations (even if they were not clearly proven from the face of the time entries) is a standard approach a client would take here.

22. Identify Issues in Litigation Matter Budget Proposal (C-029)

Quoted from Harvey LAB

> **Budget inconsistency rated as Critical or High severity.** PASS if the fee total reconciliation discrepancy ($116,050 difference between phase-based and staffing-model totals) is assigned a severity of Critical or High (or the highest or second-highest tier in whatever severity scale is used). FAIL if it is assigned a lower severity level or no severity level.
>
> — Harvey LAB, criterion C-029, task file at commit `1dd81403`

My assessment	
Criterion	Arguable — requires specific outputs not expressly or impliedly requested in the prompt
Task environment	Correct
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The prompt requested "a categorized issue memorandum" but didn't specify what the categories should be. The rubric fails the model if it does not break issues into severity levels or it does not assign a critical or high severity level to a specific issue. While it appears that models often do volunteer this kind of severity assessment, it's not requested by the prompt and competent work product could decline to include it.

23. Review Counterparty's Proposed Jury Instructions — Issue Memorandum for Trade Secrets Trial (C-020)

Quoted from Harvey LAB

> **Explains distinction between tortious interference with contract vs. business relations.** PASS if the memo explains that under Georgia law, tortious interference with an existing contract does not require a separately/independently wrongful act — the

knowing inducement of the breach is sufficient — and that the 'independently wrongful act' element applies to the distinct tort of tortious interference with business relations. FAIL if this legal distinction is not articulated.

— Harvey LAB, criterion C-020, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error (including hallucinations), rubric is incorrect
Task environment	Defective: the fictional summary-judgment order cites a case for the opposite of its holding and a case that cannot be found (citation sweep: 10 of 21 citations in this task have problems).
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. Multiple issues:

- First, this is a clear legal error. Georgia applies the same elements to interference with contractual and business relations. The rubric says the opposite of what Georgia law is. *Disaster Services, Inc. v. ERC P'ship*, 228 Ga. App. 739, 740–41 (1997) ("Tortious interference claims, whether asserting interference with contractual relations, business relations, or potential business relations, share certain common essential elements: (1) improper action or wrongful conduct by the defendant without privilege; (2) the defendant acted purposely and with malice with the intent to injure; (3) the defendant induced a breach of contractual obligations or caused a party or third parties to discontinue or fail to enter into an anticipated business relationship with the plaintiff; and (4) the defendant's tortious conduct proximately caused damage to the plaintiff.").
- Second, the fictional summary judgment order (from which the rubric's legal position seems to be taken) appears to contain several hallucinations. For example, the order paraphrases the third and fourth element described in *Disaster Services* approximately correctly, but the first two elements are erroneously listed as "(1) the existence of a valid contract between the plaintiff and a third party; (2) the defendant's knowledge of the contract." The fictional order also cites *Valdosta Livestock, Inc. v. Furst*, 342 Ga. App. 25, 28 (2017), which is not found in any legal databases.

24. Review Litigation Invoice Against Outside Counsel Billing Guidelines — Compliance Deviation Report (C-024)

Quoted from Harvey LAB

Recommends full disallowance of black car service: $387.00. PASS if the report recommends disallowing the full $387.00 for black car service (or reducing it to a reasonable taxi/ride-share equivalent). FAIL if no disallowance or reduction is recommended.

— Harvey LAB, criterion C-024, task file at commit `1dd81403`

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong

LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The billing guidelines state that "Reimbursable ground transportation is limited to ride-share ... or taxi" and "Limousine services and black car services are not reimbursable under any circumstances." Sol's critique of the rubric is that the policy requires disallowing black car reimbursement, not reducing the reimbursement amount to what a taxi or ride share would have cost. The rubric's interpretation (that either disallowing the expense entirely or capping reimbursement at the amount a taxi would have cost) seems reasonable and closer to the likely real-world application of a policy like this.

25. Review Litigation Invoice Against Outside Counsel Billing Guidelines — Compliance Deviation Report (C-042)

Quoted from Harvey LAB

> **Classifies business class airfare as borderline/requires discussion.** PASS if the report classifies the business class airfare issue (Messina's $2,847.00 flight) as a borderline or 'requires discussion' issue rather than a clear-cut mandatory reduction, given the 4-hour-20-minute flight duration and the 4-hour threshold. FAIL if it is classified as a clear mandatory reduction without acknowledging the ambiguity, OR if it is entirely dismissed as compliant.
>
> — Harvey LAB, criterion C-042, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Correct
AI auditors' reasoning	Correct

LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric's analysis is wrong. The travel policy included in the environment states, "For domestic flights with a scheduled duration of four (4) hours or more, business class may be booked without prior approval." The record does not state the flight's duration (the invoice lists only the route and fare class). The only basis for raising this seems to be in the rubric itself, which says the flight could take 4h20m, but even under that view, there's nothing to discuss, as the policy allows business class for that duration.

C Human review of AI-flagged LegalReasonBench criteria

GPT-6 Sol and Claude Opus 5.5 each audited, without seeing the other's report, the 179 rubric criteria in the 11 litigation assignments of micro1's public LegalReasonBench sample at dataset revision `1bfa904c`. Either blind report labeled 63 of those criteria problematic or arguable. I reviewed a simple random sample of 25 of the 63 flagged criteria, drawn without replacement with seed 20261024 from 9 assignments, under a protocol committed before the draw, applying the standard used for the LAB review in Appendix B. The verdicts were finalized after an adversarial quality-control review in which AI models argued for and against each call. 21 of the 25 criteria were defective (7 resting on an unambiguous objective error, and 6 arising from the task environment), 0 were arguable, and 4 were not defective; 12 come from task environments that contain defects (5 of the 9 sampled assignments; Table 5). The exact one-sided 95% hypergeometric lower bound is 45 defective criteria among the 63 flagged (point estimate 53), and 11 resting on an unambiguous objective error (point estimate 18). Because this population is the union of two blind passes, while the LAB population included Opus's reconciled flags, the two results are not estimates of the same quantity. Each entry below quotes the criterion from micro1's rubric, gives my assessment and the votes of LegalReasonBench's own three judges (GPT-5.6 Terra, Claude Sonnet 5, and Gemini 3.8 Flash; a strict two-of-three majority decides) on runs by GPT-6 Luna (xhigh) and Claude Opus 5.5 (low), and then my analysis. The source lists after each analysis were assembled with AI assistance.

Rubric verdict	Environment defective	Environment correct	Total
Defective	10	11	21
Arguable	0	0	0
Not defective	2	2	4
Total	12	13	25

Table 5: My verdicts on the 25 sampled LegalReasonBench criteria, by whether the task environment itself contains a defect.

1. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 1: Legal Memorandum to Board of County Commissioners

criterion `states_in_the_first_assignment_that_dci2018_10023_remained_activ`

Quoted from micro1 LegalReasonBench

> **Penalty criterion, −15 points.** States in the first assignment that DCI2018-10023 remained active until September 13, 2023.
>
> — micro1 LegalReasonBench, `states_in_the_first_assignment_that_dci2018_10023_remained_activ`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: clear; Claude Sonnet 5: clear; Gemini 3.8 Flash: clear. Result: penalty not applied (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: fires; Claude Sonnet 5: clear; Gemini 3.8 Flash: clear. Result: penalty not applied (1/3)

My analysis. Confirmed. The rubric penalizes the model (−15) if its memo says that the 2019 zoning case (DCI2018-10023) was still pending until September 13, 2023, which is what the record says (Petition.docx p. 19: "the 2019 zoning case was left open until September 13, 2023"). The rubric's theory is that a county rule, Lee County Admin. Code AC-2-6 § 4.1.B(1), which is not cited in the environment but can be found using web search (which models are allowed to use here), says that remanded hearings must occur within six months and that applications will be considered withdrawn after that. But the rubric ignores that the case is withdrawn only "[i]f the six month deadline is not met because of action or inaction of the Applicant/Appellant." In addition, AC-2-6 § 4.1 is prefaced with a sentence that states that it governs only applications "remanded to the Hearing Examiner by the Board." The case actually concerns a remand from the Hearing Examiner to the staff (Petition.docx: "we jointly request that DCI2018-10023 be remanded back to the staff"), so the regulation does not apply by its terms. This case appears to be drawn from a real litigation, Spahhn v. Lee County, No. 24-CA-1479 (Fla. 20th Cir. Ct. Oct. 25, 2025) (order granting petition for writ of certiorari), and the plaintiffs' position is that the case remained live to September 13, 2023 (contradicting the rubric); the county says the case was withdrawn in March 2022. Elsewhere the rubric claims the case was withdrawn on February 13, 2020 (companion criterion `states_in_the_first_assignment_that_under_lee_county_fla_admin_c`), but neither party advanced that claim, so the rubric seems to be contradicted by the ground-truth evidence from the litigation.

- *Sources:* AC-2-6 § 4.1, as amended through 11/17/15 (identical as amended through 8/2/22); county position: Lee County Utilities, Local Planning Agency presentation, Sept. 25, 2023 ("March 2022 – Previous CFPD rezoning petition withdrawn"); ruling: Gulfshore Business, WGCU, Nov. 10, 2025.
- *Native runs:* GPT-6 Luna did not mention DCI2018-10023 (0/3). Claude Opus 5.5 listed the record's September 13, 2023 withdrawal date in its procedural history; one judge (GPT-5.6 Terra) fired the penalty, but the 2-of-3 majority cleared it (1/3).
- *Related items:* items 1, 13, 16 and 22 are penalties worded "States" or "Mentions" that can reach a response rebutting or setting aside the point; the native runs misfired on 16 and 22.

2. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 1: Legal Memorandum to Board of County Commissioners

criterion `states_in_the_first_assignment_that_the_county_should_cease_cons`

Quoted from micro1 LegalReasonBench

Penalty criterion, −30 points. States in the first assignment that the County should

cease construction activities at the subject property.

— micro1 LegalReasonBench, `states_in_the_first_assignment_that_the_county_should_cease_cons`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: clear; Claude Sonnet 5: clear; Gemini 3.8 Flash: clear. Result: penalty not applied (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: fires; Claude Sonnet 5: fires; Gemini 3.8 Flash: clear. Result: penalty applied (2/3)

My analysis. Confirmed. The rubric deducts 30 points if the memo being prepared for a county in a zoning dispute recommends stopping construction. The basis of this rubric is Parkway Towers Condominium Ass'n v. Metropolitan Dade County, 295 So. 2d 295 (Fla. 1974) (per curiam), which merely exercises discretion not to hear an appeal on certiorari. While it is true that the outcome in that particular case was that the county was allowed to proceed with construction, the court's reasoning, if anything, undercuts the rubric's reading: "Although we now consider our issuance of the writ of certiorari to have been improvident in this case, which results in our discharge thereof, it is nevertheless our view prospectively that zoning variations to accommodate county or municipal facility purposes should either have been anticipated in zoning ordinances before construction or operation of such facilities is commenced or, if this has not been done, construction should not be undertaken thereof until after due modification or change therefor is made in existing zoning ordinances." Id. at 295–96. In other words, the very case the rubric relies on essentially says the opposite of what it claims and expresses the Florida Supreme Court's view that zoning issues should be fixed before construction occurs. The fact that some construction has begun is a relevant factor, but nothing in the environment establishes that it would be costly, impossible, or imprudent for the county to stop construction until the zoning issues are fixed properly to comply with the Court's opinion. Many prudent lawyers, relying on the Florida Supreme Court's guidance, might recommend addressing the zoning issue before continuing with construction, so this rubric fails a response many careful lawyers would prefer.

- *Sources:* Parkway Towers, 295 So. 2d 295 (CourtListener).

3. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 1: Legal Memorandum to Board of County Commissioners

criterion `states_in_the_first_assignment_that_the_county_should_file_a_tim`

Quoted from micro1 LegalReasonBench

Reward criterion, +20 points. States in the first assignment that the County should file a timely appeal with the 6th District Court of Appeals.

— micro1 LegalReasonBench, `states_in_the_first_assignment_that_the_county_should_file_a_tim`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: not met. Result: met (2/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Rubrics 23 and 24 (criteria `states_in_the_first_assignment_that_the_county_should_file_a_tim` and `states_in_the_first_assignment_that_the_county_should_remand_the`, +20 each) require the government to simultaneously pursue an appeal and remand to a hearing examiner to fix the deficiencies the Court identified. I agree with Sol that the recommendation to take an appeal is not mandatory here and there are very good reasons not to appeal. While a party who loses in the trial court often wishes to appeal, the trial court's ruling here allows the government to reach the same conclusion it originally did after correcting some deficiencies in the process followed the first time. Often simply correcting the deficiencies identified by a trial court would be faster and more reliable for a government to proceed with a project than appealing. Although the government in the real litigation did elect to simultaneously appeal and remand the issue to the hearing examiner (as the rubric recommends), the government's filing in the appeal acknowledged that the remand might moot the appeal. The government argued the appeal should still be considered because it wanted to undo the trial court's finding that the government had been biased in approving the initial rezoning, but that is a very debatable position (and arguably not a strong basis for an appeal). The rubric's premise that the Order "contains reversible errors" also overstates what an appeal could achieve: second-tier certiorari review is limited to whether the circuit court afforded procedural due process and applied the correct law. *Haines City Cmty. Dev. v. Heggs*, 658 So. 2d 523, 530 (Fla. 1995). Filing a protective petition within the 30-day deadline is a reasonable choice, but so is curing the defects on remand. So this case does not seem to require an appeal — that's really a judgment call — so the rubric would incorrectly effectively penalize a failure to pursue a strategy many lawyers would advise against.

- *Source:* Lee County v. Spahhn, No. 6D2025-2926 (Fla. 6th DCA petition filed Nov. 24, 2025), Pet. § VII at 45–47 ("Even if compliance with the circuit court's order might arguably render this specific dispute technically moot, this Court should exercise its discretion to review the decision under well-established exceptions to the mootness doctrine."); docket.

4. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 1: Legal Memorandum to Board of County Commissioners

criterion `states_in_the_first_assignment_that_the_county_should_remand_the`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +20 points.** States in the first assignment that the County should remand the case to the Hearing Examiner to address the deficiencies identified in the Order Granting Writ of Certiorari while an appeal is pending.
>
> — micro1 LegalReasonBench, `states_in_the_first_assignment_that_the_county_should_remand_the`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. With respect to the rubric recommending remand to the hearing officer, I agree with the aspect of it that recommends remand: remand to the Hearing Examiner to fix issues identified by a court is the standard approach a government would follow in this spot. But as written, the criterion requires remand "while an appeal is pending," so it also requires an appeal, and it would therefore reject a remand-only recommendation that a careful litigator could adopt, for the reasons given for criterion 23 (item 3). A memo that recommends remand without an appeal would fail it. Sol flagged it (as arguable) on this same ground.

- *Sources:* criterion text ("while an appeal is pending"); Lee County LDC § 34-83(a)(6) as cited in the rubric's source note; see item 3 for the County's Sixth DCA petition.

5. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 1: Legal Memorandum to Board of County Commissioners

criterion `states_in_the_first_assignment_that_under_fla_stat_163_3167_10_2`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +10 points.** States in the first assignment that under Fla. Stat § 163.3167(10) (2025) the development order issued for the project may not be invalidated by a subsequent judicial determination because the timeframe for filing an appeal of the development order has expired.
>
> — micro1 LegalReasonBench, `states_in_the_first_assignment_that_under_fla_stat_163_3167_10_2`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. This rubric is deeply confused about the meaning of the statutory provision it's relying on. Fla. Stat. § 163.3167(10) (2025) protects the developer and allows a project to continue in a situation where no one challenges that specific development, but a later case (typically involving

a different development) invalidates the regulation under which the first development was approved. In that scenario, the developer whose project was approved (with no one challenging it) is entitled to proceed notwithstanding the fact that the regulations the first developer relied on were later invalidated in a *different* case. That principle has no relevance here, where the project at issue was timely challenged by a petition for a writ of certiorari, and a court invalidated the approval for the specific project that was timely challenged. The statute is also narrower than the rubric assumes: it bars invalidation only by a later determination that the land development regulations themselves "are invalid because of a deficiency in the approval standards." The circuit court did not invalidate Lee County's regulations; it quashed this rezoning for defects in how the Board reached its decision. So even if the criterion were read to refer to some later development order issued under the rezoning, § 163.3167(10) would not address it.

- *Sources:* Fla. Stat. § 163.3167(10) (2025); subsection (10)(b) adds that it "does not preclude or affect the timely institution of any other remedy available at law or equity, including a common law writ of certiorari proceeding." Orderpdf.docx p. 1: "the rezoning is hereby **QUASHED**."

6. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 1: Legal Memorandum to Board of County Commissioners

criterion `states_in_the_first_assignment_that_under_lee_county_fla_admin_c`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +5 points.** States in the first assignment that under Lee County, Fla. Admin. Code § 2-6, 4.1 B. (2025), DCI2018-10023 was considered withdrawn on February 13, 2020.
>
> — micro1 LegalReasonBench, `states_in_the_first_assignment_that_under_lee_county_fla_admin_c`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. Wrong for the same reasons as item 1. Neither party in the actual litigation contends that the six-month deadline was missed "because of action or inaction of the Applicant/Appellant." There's nothing in the environment to support that claim, and without a factual predicate for that, there's no factual basis for suggesting that Lee County Admin. Code AC-2-6 § 4.1.B(1) (2025) applies here. In addition, AC-2-6 § 4.1 is prefaced with a sentence that states that it governs only applications "remanded to the Hearing Examiner by the Board." The case actually concerns a remand from the Hearing Examiner to the staff (Petition.docx: "we jointly request that DCI2018-10023 be remanded back to the staff"), so the regulation does not apply by its terms.

- *Sources:* see item 1 (AC-2-6 § 4.1 text and the parties' positions in the real litigation).

7. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 2: Brief Requesting Second-Tier Review

criterion `states_in_the_second_assignment_that_expert_testimony_on_the_rec`

Quoted from micro1 LegalReasonBench

> **Reward criterion, +5 points.** States in the second assignment that expert testimony on the record is competent substantial evidence that supports the Board's approval of the rezoning.
>
> — micro1 LegalReasonBench, `states_in_the_second_assignment_that_expert_testimony_on_the_rec`, rubric at revision `1bfa904c`

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. Correct as written. Sol initially suggested that this rubric was arguably defective, because it rewards models for making the point that the expert testimony in the record can support the Board's zoning approval. Sol's concern was that the expert testimony is not actually in the environment provided. While Sol correctly points out that the testimony is not in the environment, the Petition quotes portions of the testimony with citations to the transcripts, and other material in the environment does make clear that expert analysis was used to support the Board's decision. It's a relatively uncontroversial proposition, the relevance of which can be readily inferred from the environment, and my view is that any competent lawyer would consider adding that point (at least briefly) an improvement to the analysis.

- *Sources:* Petition.docx, quoting professional engineer Erik Howard of Johnson Engineering (Tr1 p. 36, ll. 1–4; p. 40, ll. 24–25; p. 41, ll. 11–21) and RVI Planner Crespo at the zoning hearing (Tr1 pp. 226–27), presenting "an exhibit generated by Johnson Engineering using that line work provided by the Port Authority"; Resolution.docx recitals (the Hearing Examiner "recommended APPROVAL of the request with conditions").

8. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 2: Brief Requesting Second-Tier Review

criterion `states_in_the_second_assignment_that_the_staff_report_recommendi`

Quoted from micro1 LegalReasonBench

> **Reward criterion, +5 points.** States in the second assignment that the staff report recommending approval is competent substantial evidence that supports the Board's approval of the rezoning.
>
> — micro1 LegalReasonBench, `states_in_the_second_assignment_that_the_staff_report_recommendi`, rubric at revision `1bfa904c`

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. Correct as written. This is a similar issue to item 7. The rubric wants models to make the point that the staff report is evidence supporting the Board's approval. Sol says the underlying staff report is not part of the environment, but it is referenced in other documents, and any competent lawyer in this situation would want the fact that such a report exists to be cited, as it's clearly going to be important evidence in support of the government's position.

- *Source:* Resolution.docx recitals (the Board "gave full and complete consideration to the recommendations of the staff, the Hearing Examiner, the documents on record and the testimony of all interested persons").

9. Viability of a DCCA Challenge, assignment 1: Record Mapping: Final Notice Against Award

criterion `criterion_8`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +28 points.** Identifies that the comparator memorandum's stated coverage years conflict with the dates of the cases it summarizes.
>
> — micro1 LegalReasonBench, `criterion_8`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Defective: Award ¶ 7 says "Two of the four comparators involved on-duty conduct," but the comparator memorandum shows three on duty (Guzman, Sadler and Voss; only Pietrzak was off duty), so the arbitrator's comparability finding misstates the record it rests on. The memorandum also says "The Commonwealth declined prosecution" of an off-duty shooting in the District.
AI auditors' reasoning	Correct
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. It's unclear how this rubric got created, since this seems like a straightforward mistake that even a nonlawyer would catch. The rubric rewards models for identifying that the "comparator memorandum's stated coverage years conflict with the dates of the cases it summarizes."

But both documents cited as the source for this have the correct dates. The memo says it reviewed "records covering calendar years 2018 through 2021." It then lists cases from August 2018, November 2019, May 2020, and March 2021, which were resolved in January 2019, June 2020, February 2021, and October 2021, respectively. There is no mismatch, so this rubric would reward models for hallucinating a mismatch that is not present in the environment.

- *Sources:* comparator_discipline_cases.pdf ¶¶ 1–5; final_notice_adverse_action_braddock.pdf ¶ 19 ("four comparator cases from 2018 through 2021").

10. Viability of a DCCA Challenge, assignment 2: Scope of Review

criterion `criterion_4`

QUOTED FROM MICRO1 LEGALREASONBENCH

> **Reward criterion, +26 points.** Identifies that no uploaded authority requires MPD to prove the comparability of its comparator cases.
>
> — micro1 LegalReasonBench, `criterion_4`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Defective: Award ¶ 7 says "Two of the four comparators involved on-duty conduct," but the comparator memorandum shows three on duty (Guzman, Sadler and Voss; only Pietrzak was off duty), so the arbitrator's comparability finding misstates the record it rests on. The memorandum also says "The Commonwealth declined prosecution" of an off-duty shooting in the District.
AI auditors' reasoning	Correct

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: met; Claude Sonnet 5: not met; Gemini 3.8 Flash: met. Result: met (2/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. The rubric wants an analysis of an employee disciplinary dispute to make the point that "no uploaded authority requires MPD to prove the comparability of its comparator cases." But *Douglas v. Veterans Administration*, 5 M.S.P.B. 313, 5 M.S.P.R. 280 (1981), one of the uploaded authorities, states that "it should be clear that the ultimate burden is upon the agency to persuade the Board [the Merit Systems Protection Board] of the appropriateness of the penalty imposed." *Id.* at 333. It adds that "[t]he selection of an appropriate penalty is a distinct element of the agency's decision, and therefore properly within its burden of persuasion," and that when the employee "challenges the severity of the penalty," as Braddock did (the stipulated arbitral issue was whether "the penalty of termination [is] appropriate"), "the agency will be called upon to present such further evidence." *Id.* at 334. That language comes just after a list of 12 factors that prior rulings have identified as relevant for determining the appropriateness of a penalty, which includes (as factor six) "consistency of the penalty with those imposed upon other employees for the same or similar offenses." *Id.* at 332. *Douglas* does note that "Not all of these factors will be pertinent in every case," so there is no bright-line requirement that the agency prove factor six in every case. But *Douglas* makes clear that the burden is on the agency to establish the appropriateness of the penalty and

comparability is expressly identified as an important factor in assessing appropriateness. Critically, in this specific case, MPD's Final Notice of Adverse Action expressly relied on comparable cases as an aggravating factor and a key factor for terminating Officer Braddock. Thus, in the context of this case, the proposition the rubric wants the AI to assert is wrong. If MPD cannot prove to the arbitrator that the comparator officers were similarly situated (as the Final Notice asserts, and as the arbitrator found MPD did not), that could be a basis for at least vacating the termination decision and remanding it for further consideration. Sol and Opus labeled this as an arguable defect, but under the specific circumstances of this case, I would view this assertion as objectively inaccurate. The stipulations do not change this. Stipulation (a), that reasonable minds can disagree about how comparable the four cases are, goes to whether MPD carried its burden, not the issue of whether MPD bore the burden in the first place. Stipulation (b), that Braddock never alleged disparate treatment before the hearing, matters only under the real-world rule that the employee must first make an initial showing, which this rubric itself treats as outside the uploaded materials (criterion 5). Failure to prove up the comparable cases does not automatically invalidate MPD's action, but it leaves MPD highly vulnerable to being overturned – so in this context, MPD does need to prove the comparable cases that it relied on in making its decision to terminate Braddock. This rubric appears to be lifted from a real case that matches the fact pattern in the task, *D.C. Metropolitan Police Department v. D.C. Public Employee Relations Board*, 282 A.3d 598 (D.C. 2022), and No. 24-CV-0573 (D.C. May 21, 2026), where MPD explicitly argued that it was not required to put forward comparable cases. But in the real case, the issue was that MPD had not identified any comparable cases where termination was the discipline imposed, and the terminated officer's union came forward with cases involving comparable conduct where the punishment was less severe. The task environment here was modified from the real case, so that in this fictional universe, MPD itself puts comparable cases at issue in its initial decision terminating the officer. In that context, the argument advanced in the real case no longer holds and the rubric is incorrect because it fails to take into account how this task environment was modified from the real case. The rubric is essentially importing a proposition of law from the real case without recognizing that the task environment was changed in a way that makes that proposition irrelevant on the facts being presented to the model.

- *Sources:* douglas_v_veterans_admin.pdf, printed pp. 332–34 (the twelve factors, "Not all of these factors will be pertinent in every case," the agency's "ultimate burden," and penalty selection as "within its burden of persuasion"); braddock_arbitration_award.pdf ¶ 2 (the stipulated issue); stipulations (a) and (b) in the assignment 2 prompt; final_notice_adverse_action_braddock.pdf ¶ 19 (MPD's factor-six finding, relying on the four comparator cases: "This factor is aggravating.") and ¶ 26 (termination after weighing all twelve factors, "each of which aggravates the penalty"); braddock_arbitration_award.pdf ¶¶ 7, 10 (the arbitrator finds MPD "did not prove that the comparator officers were similarly situated" and reduces the termination to a 45-day suspension); the real case: *MPD v. PERB*, No. 19-CV-1115 (D.C. Sept. 15, 2022), 282 A.3d 598, slip op. at 5–7 (MPD's panel "had not cited comparable cases"; MPD argued that its burden arose only on the officer's "initial showing" of disparate treatment), and *MPD v. PERB*, No. 24-CV-0573 (D.C. May 21, 2026), slip op. at 8–10 (FOP supplied three less-severe cases, including Officer Ford's 45-day suspension; the court said "MPD is correct about its burden under Douglas" but affirmed).

11. Viability of a DCCA Challenge, assignment 3: Statutory Exceptions

criterion `criterion_16`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +6 points.** Identifies that the PERB Decision misstates the length of the suspension imposed by the Award.
>
> — micro1 LegalReasonBench, `criterion_16`, rubric at revision `1bfa904c`

My assessment	
Criterion	Not defective
Task environment	Defective: Award ¶ 7 says "Two of the four comparators involved on-duty conduct," but the comparator memorandum shows three on duty (Guzman, Sadler and Voss; only Pietrzak was off duty), so the arbitrator's comparability finding misstates the record it rests on. The memorandum also says "The Commonwealth declined prosecution" of an off-duty shooting in the District.
AI auditors' reasoning	Wrong
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Rubric is correct. The rubric wants an internal memorandum to point out that the D.C. Public Employee Relations Board's decision incorrectly states that the officer was suspended for 40 days, when in fact the suspension was 45 days. Sol's point is that the prompt did not specifically request that analysis. But the prompt directs the associate to "[a]nchor each argument in specific findings of the Award and the PERB Decision," and a careful lawyer doing that would notice that the decision misstates the penalty it affirmed "in all respects." While I agree that it is not necessarily essential for a memo to call out what appears to be a typo in the Board's decision, doing so does improve an internal memo – and this is a situation where it would be desirable to train models to flag that discrepancy even if the user did not specifically prompt for it. The misstatement of fact in the Board's decision affirming the arbitration award is important information for the lawyer to be aware of, and a lawyer using AI to prepare this memo might not think to ask specifically about those kinds of discrepancies. So the rubric is a desirable training signal, and to the extent it asks the model to go somewhat beyond the prompt, this specific example is a good instance where you clearly would want the model to flag an issue that you might not have thought to ask about in your prompt.

- *Sources:* assignment 3 prompt ("Anchor each argument in specific findings of the Award and the PERB Decision, with page and paragraph citations"); perb_decision_braddock.pdf (Opinion No. 1904, PERB Case No. 25-A-09) ¶ 1 ("reduced the penalty of termination to a 40-day suspension without pay") and ¶ 7 (award "affirmed in all respects"); braddock_arbitration_award.pdf ¶ 10 ("reduced to a 45-day suspension without pay").

12. Plaintiff's Response - Forum Non Conveniens, assignment 1

criterion `criterion_2`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +5 points.** Identifies in Assignment #1 an affidavit by John Smith, M.D. does not indicate that Cook County would be an inconvenient forum for him to litigate this case as required under Langenhorst v. Norfolk S. Ry. Co., 219 Ill. 2d 430 (Ill. 2006).
>
> — micro1 LegalReasonBench, `criterion_2`, rubric at revision `1bfa904c`

My assessment	
Criterion	Not defective
Task environment	Defective: The partner's email says Baby Davis was transferred to "Lurie's Children's Hospital" and "almost passed away on September 16, 2024," but the Complaint says he was transferred to and died at Cook County Hospital (¶¶ 27, 29). The defendants' filed affidavits also carry unfilled placeholders: Dr. Smith's affidavit gives his residence and office as "[street address], [city]" (Ex. A ¶¶ 2–3), and the hospital's affiant is named only as "[Records Custodian]" (Ex. B).
AI auditors' reasoning	Wrong

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (1/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Rubric is correct as written. The case involves a request by defendants to move a case from Cook County (where Chicago is located) to DuPage County (a bordering county immediately west of Cook County, with Chicago O'Hare airport extending partially into DuPage County). One of the defendants, Dr. John Smith, submitted an affidavit arguing that DuPage County is more convenient than Cook County, but it does not expressly state that Cook County is "inconvenient." The rubric wants a memo analyzing the issue to point out that Dr. Smith does not state in his affidavit that it would be inconvenient for him to litigate in Cook County. Sol initially flagged the criterion as arguable, reasoning that because Dr. Smith swears DuPage would be "substantially more convenient" than Cook, "[a] careful reviewer could reasonably treat that as comparative inconvenience." But the rubric is correct because it's pointing out that the standard isn't which location is "more" convenient; a defendant seeking to move on forum non conveniens grounds needs to establish that the plaintiff's chosen forum is actually inconvenient to the defendant. The task environment contains a decision, *Langenhorst v. Norfolk Southern Railway Co.*, 219 Ill. 2d 430, 848 N.E.2d 927 (2006), that states the rule drawing this distinction between an argument that another location would be "more" convenient for one party and proof that the plaintiff's chosen location is actually inconvenient: "The defendant must show that the plaintiff's chosen forum is inconvenient to the defendant and that another forum is more convenient to all parties." 219 Ill. 2d at 444. So Sol's point is not well taken in this instance. The later reconciliation overturned Sol's flag on this ground.

- *Sources:* Defendants Motion to Transfer.pdf, Exhibit A ¶ 7 (Dr. Smith: DuPage "would be substantially more convenient for me, both as a party and as a fact witness, than trial in the

Circuit Court of Cook County"); Langenhorst v. Norfolk Southern.pdf, PDF p. 7, 219 Ill. 2d at 444 ("The defendant must show that the plaintiff's chosen forum is inconvenient to the defendant and that another forum is more convenient to all parties.").

13. Plaintiff's Response - Forum Non Conveniens, assignment 1

criterion `criterion_22`

Quoted from micro1 LegalReasonBench

Penalty criterion, −30 points. Mentions in Assignment #1 that Cook County lacks interest in adjudicating a controversy over care rendered by its DuPage based providers.

— micro1 LegalReasonBench, `criterion_22`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Defective: The partner's email says Baby Davis was transferred to "Lurie's Children's Hospital" and "almost passed away on September 16, 2024," but the Complaint says he was transferred to and died at Cook County Hospital (¶¶ 27, 29). The defendants' filed affidavits also carry unfilled placeholders: Dr. Smith's affidavit gives his residence and office as "[street address], [city]" (Ex. A ¶¶ 2–3), and the hospital's affiant is named only as "[Records Custodian]" (Ex. B).
AI auditors' reasoning	Correct
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: clear; Claude Sonnet 5: clear; Gemini 3.8 Flash: clear. Result: penalty not applied (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: clear; Claude Sonnet 5: clear; Gemini 3.8 Flash: clear. Result: penalty not applied (0/3)

My analysis. The rubric penalizes a memo that "Mentions . . . that Cook County lacks interest in adjudicating a controversy over care rendered by its DuPage based providers." The rubric is subtly misspecified in a way that can affect the grading. The reason the rubric is penalizing a memo that raises this point is that the environment contains a case, *Bradbury v. St. Mary's Hospital of Kankakee*, 273 Ill. App. 3d 555, 652 N.E.2d 1228 (1st Dist. 1995), that indicates that Cook County has an interest in situations like these if the patient died there – and the case at issue in the environment involves a patient death, so the theory underlying the rubric is sound. However, the defendants do argue in their motion that Cook County has limited interest in the death – and as written the rubric penalizes any memo that even "mentions" that Cook County lacks interest. So a grader could apply this penalty even to a memo that simply "mentions" the defendants' argument and then correctly explains why it's wrong under *Bradbury*. This penalty did not fire in our test runs, but a penalty in the same "Mentions" form in this assignment, criterion 25 ("Mentions in Assignment #1 that Cook County's overall clearance rate is 60.5%," -30) fired 3/3 against both Opus, which called the figure "the wrong statistic," and Luna, which listed it among the defendants' strengths as the prompt asked. Several criteria in this sample are penalties worded "Mentions" or "States," and the judges applied them inconsistently: sometimes literally, penalizing a response that raised the point only to rebut it, and sometimes as the rubric intends, to catch misstatements the model should not endorse (see items 16 and 22).

- *Sources:* Bradbury v. St. Mary's Hospital.pdf, PDF p. 4 ("Cook County has an interest because the decedent died here"); Defendants Motion to Transfer.pdf, PDF p. 5, Part V.A ("Cook County's only connection to this case is that Baby Davis was later transferred to, received treatment at, and died in, a Cook County hospital — a circumstance that, unlike in Bradbury, did not result from any Cook County conduct by these Defendants").
- *Native runs:* the penalty did not fire. Claude Opus 5.5 recounted the defendants' local-interest argument and rebutted it under *Bradbury*, and all three judges cleared it (0/3); GPT-6 Luna also cleared (0/3). The same wording problem did fire in this assignment's criterion 25 ("Mentions in Assignment #1 that Cook County's overall clearance rate is 60.5%," -30, not in the sample): 3/3 against Opus, which called it "the wrong statistic," and 3/3 against Luna, which listed it under "What is strong for Defendants" as the prompt asked.
- *Related items:* items 1, 13, 16 and 22 are penalties worded "States" or "Mentions" that can reach a response rebutting or setting aside the point; the native runs misfired on 16 and 22.

14. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_1`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +3 points.** States in Assignment #2 that the response is supported by the affidavit of Liam Vance, a physician who provided care to Baby Davis at Cook County Hospital.
>
> — micro1 LegalReasonBench, `criterion_1`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; arises from the task environment
Task environment	Defective: No affidavit from any of the five physicians is provided in the environment, and the prompt does not explicitly instruct the model to create one. The prompt also says, "If a document is not uploaded, it does not exist for purposes of this analysis."
AI auditors' reasoning	Correct
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. The rubric wants the plaintiff's opposition brief (in response to defendants' forum non conveniens motion) to state that the plaintiff's "response is supported by the affidavit of Liam Vance, a physician who provided care to Baby Davis at Cook County Hospital." But no such affidavit is provided in the environment, and the prompt does not explicitly instruct the model to create it. The prompt further says "Do not rely on, reference, or infer any legal authority, factual information, or data beyond what is contained in these uploaded files. If a document is not uploaded, it does not exist for purposes of this analysis." While the prompt elsewhere says, "Please include all affidavits and information you referenced from Assignment #1," that suggests the model should "include" documents that were referenced in the environment, not generate new ones or assume that they will be created separately. Although the assignment 1 rubric rewards recommending that the plaintiff

obtain affidavits from these physicians, and the assignment 2 prompt asks the model to include the affidavits referenced in assignment 1, I would still consider drafting a brief that assumes what third-party witnesses will sign to be hazardous and unreliable.

Drafting an opposition brief that cites to affidavits (without placeholders in the draft) is extremely hazardous behavior to train models on in a legal setting. In a real run, Opus was rewarded for drafting affidavits and including in the brief citations to them without placeholders, while Luna did not create the affidavits or cite to them. If forced to choose, I would consider Luna's approach more correct in this instance, as a matter of professional judgment. Although Luna's approach suffered from its own flaws (including that it incorrectly said the record does not reflect the physicians' willingness to testify), I would prefer the approach of not citing affidavits that don't yet exist. The contemplated affidavits are from nonparty physicians that the case team does not represent; while it is helpful that a lawyer on the case team spoke to these witnesses and obtained an understanding of their views that can be turned into a draft affidavit, ultimately a nonparty is entitled to decide exactly what it is willing to affirm in writing and a lawyer cannot simply assume that the witness will be willing to sign an affidavit based on the lawyer's understanding from an initial call. The normal approach in a situation like this would be to draft placeholder citations to an affidavit that make clear the language must be adjusted to conform to the witness's final testimony. Opus's draft affidavit includes one form for five physicians, and in practice, I would expect the physicians to object to some of what is included in the affidavit (and in fact getting all five to actually sign these affidavits in time would be a fairly significant undertaking, so I would not assume all five affidavits ultimately enter into the record as this assignment contemplates). The Opus work product creates a real risk that a human lawyer, through inadvertence or mistake, might file a brief that claims to be supported by affidavits that have not in fact been sworn, which is a serious ethical lapse.

Overall, the practical effect of this rubric seems to be to nudge the model in the direction of thinking it is OK to invent witness testimony that has not yet been approved by a witness. The issue could be avoided either by including an actual affidavit approved by the witness in the environment or asking the model to merely include bracketed text/placeholders that are clearly labeled as draft so that the draft is clear and explicit that the affidavits are contemplated or expected, rather than already secured, as Opus's draft of the brief implies.

This issue is distinct from items 7 and 8, where my view was that it was appropriate for a model to refer in a brief to expert analysis that was referred to in the environment but not supplied. In that case, we had clear indications that the expert analysis described really existed; the full details were simply not provided in the model prompt, so it is important to flag the existence of that evidence apparently already present in the record. Here, the environment makes clear that, for purposes of this assignment, the affidavits do not exist, yet the rubric rewards a model for creating a draft that could easily be filed mistakenly with the false assertion that these doctors have provided affidavits to support their views (which may or may not ultimately come to pass).

- *Sources:* assignment 2 prompt ("Please include all affidavits and information you referenced from Assignment #1"; "The names of the physicians from Cook County we have spoken to are"; "If a document is not uploaded, it does not exist for purposes of this analysis"); assignment 1 prompt (the five unnamed doctors "are willing to testify"); no uploaded document mentions any physician affidavit. Native runs: Opus 5.5 (low), 3/3, whose brief states "Plaintiff submits sworn affidavits from five named Cook County treating physicians" and "Each of the five treating physicians has executed an affidavit in the same form ... Each affidavit is separately signed," over a single unsigned form with "[Name]" and "[specialty]" blanks; GPT-6 Luna (xhigh), 0/3, which identifies the five as physicians "with whom counsel has spoken" and cites

no affidavit. Criteria 2–10 (the other four physicians and the five inconvenience statements) rest on the same missing affidavits.

- *Related items:* The missing-affidavit items (14, 15 and 17–20) share this defect.

15. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_11`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +5 points.** States in Assignment #2 that the response is supported by the affidavit of the plaintiff, Sara Davis.
>
> — micro1 LegalReasonBench, `criterion_11`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; arises from the task environment
Task environment	Defective: No affidavit from Sara Davis is provided in the environment, and the prompt does not explicitly instruct the model to create one. The only fact supplied about her circumstances is that she "resides in Kane County, Illinois."
AI auditors' reasoning	Correct

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: not met; Gemini 3.8 Flash: met. Result: met (2/3)

My analysis. This has the same problem as item 14 of asking the model to draft (apparently without placeholders) an assertion that a given point is supported by an affidavit that does not exist. This one is a closer call because the affiant is the client, so the lawyers preparing this opposition brief presumably can be more confident in what the client's testimony will be. But even a client needs to be comfortable that the facts they are swearing to are accurate and may tell the lawyer to change them, so preparing a complete sworn affidavit containing facts found nowhere in the record, and citing it in the brief as sworn fact, as Opus has done (and was fully rewarded for, on a 2–1 vote), is still hazardous.

There is a separate issue here where Sonnet misgraded the rubric because it refers to Assignment #2 in the actual rubric text that is supplied to the judge model, but that is a reference to the prompt the generative models receive and should not be in the output. Sonnet did the same thing on three other criteria in this run (criteria 4, 5 and 22), failing the brief only because it had no "Assignment #2" label; in each case the other two judges passed it, so the majority result did not change. For purposes of this exercise, I am not treating that as a rubric defect because all the rubrics in the land-use and forum non conveniens matters have the issue that they refer to a specific assignment number that is not legible to the judge grader. That is an overarching issue in the benchmark setup that can be addressed globally.

- *Sources:* assignment 1 prompt ("our client, Sara Davis, resides in Kane County, Illinois"); Defendants Motion to Transfer.pdf ¶ 9 (the complaint does not allege her residence); assignment 2 prompt ("Please include all affidavits and information you referenced from Assignment #1"). Opus 5.5 (low), 2/3: its Exhibit 1 is a complete sworn affidavit for Sara with facts found

nowhere in the record (¶ 7, "I had not consulted a lawyer, I was not considering a lawsuit"; ¶ 9, "Most of the medical and hospital expenses ... were incurred at Cook County Hospital"), which the brief cites as sworn fact. Sonnet 5's dissent: "The document contains no reference to an 'Assignment #2' or any assignment numbering scheme." GPT-6 Luna (xhigh), 0/3 cites no affidavit from Sara. Claude Sonnet 5 was the sole dissent, citing only the missing label, on criteria 4, 5, 11 and 22 of this Opus run. 90 of the 179 litigation criteria name an assignment (every criterion in the land-use and forum non conveniens matters; none in the DCCA or VICAR matters). Criterion 12, the companion requirement that her affidavit say DuPage County would be inconvenient, rests on the same missing affidavit; it was not drawn into this sample.

- *Related items:* The missing-affidavit items (14, 15 and 17–20) share this defect.

16. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_18`

Quoted from micro1 LegalReasonBench

Penalty criterion, −30 points. States in Assignment #2 that the 2023 Illinois Annual Report Statistics report a 60.5% clearance rate in Cook County, Illinois.

— micro1 LegalReasonBench, `criterion_18`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Defective: No affidavit from any of the five physicians or from Sara Davis is provided in the environment, and the prompt does not explicitly instruct the model to create one. The prompt also says, "If a document is not uploaded, it does not exist for purposes of this analysis."
AI auditors' reasoning	Correct
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: fires; Claude Sonnet 5: fires; Gemini 3.8 Flash: fires. Result: penalty applied (3/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: fires; Claude Sonnet 5: fires; Gemini 3.8 Flash: fires. Result: penalty applied (3/3)

My analysis. The rubric penalizes models for stating that the 2023 Illinois Annual Report Statistics report a 60.5% clearance rate in Cook County, Illinois. But that document is provided in the environment and does indeed report the rate given as the clearance rate for all cases. In our test runs, Opus and Luna were both penalized incorrectly for noting this figure while explaining why it deserved little weight: Opus contrasted it with the civil-case rates (97.9% and 96.3%), and Luna said it "expressly covers all case types." While some lawyers prefer not to raise bad facts that the other side has not raised, here, the defendants' opening brief squarely presents the 60.5% statistic to the Court – so any competent opposition would respond to that argument. There is no reason to penalize a model for doing so.

- *Sources:* 2023 Annual Report Statistical Summary.pdf, p. 21 (all-case clearance table: Cook County 60.5%) and p. 22 (civil category: Cook 97.9%, 18th Circuit 96.3%); Defendants Motion to Transfer.pdf, PDF p. 6 ("Defendants accordingly rely principally on the clearance-rate

comparison above"). Native runs, both penalized 3/3 (-30 of 73 positive points): Opus 5.5 (low), under the heading "Defendants rely on the wrong statistic," names the 60.5% all-case figure and then gives the civil rates (97.9% and 96.3%), earning +2 on criteria 19 and 20; GPT-6 Luna (xhigh) says the all-case figure gives Defendants "some support" but has "limited" weight because it "expressly covers all case types," and notes the report has no comparable Cook time-to-verdict data for civil jury cases over $50,000. The companion assignment 1 penalty (criterion 25, -30) has the same wording problem; its justification says the 60.5% figure "is not applicable to Civil Cases."

- *Related items:* items 1, 13, 16 and 22 are penalties worded "States" or "Mentions" that can reach a response rebutting or setting aside the point; the native runs misfired on 16 and 22.

17. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_3`

Quoted from micro1 LegalReasonBench

Reward criterion, +3 points. States in Assignment #2 that the response is supported by the affidavit of Caleb Montgomery, a physician who provided care to Baby Davis at Cook County Hospital.

— micro1 LegalReasonBench, `criterion_3`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; arises from the task environment
Task environment	Defective: No affidavit from any of the five physicians is provided in the environment, and the prompt does not explicitly instruct the model to create one. The prompt also says, "If a document is not uploaded, it does not exist for purposes of this analysis."
AI auditors' reasoning	Correct

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Same issue as item 14. The criterion rewards a brief that states that the response is supported by his affidavit, but no affidavit from Dr. Caleb Montgomery exists in the environment, and the prompt does not explicitly instruct the model to create one. This is one of the missing-affidavit items (14, 15 and 17–20), which all reward citing affidavits the record does not contain.

- *Sources:* see item 14. Native runs: Opus 5.5 (low) met this criterion 3/3 by citing its single unsigned physician form as an executed affidavit; GPT-6 Luna (xhigh) cited no physician affidavit and scored 0/3.

18. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_6`

Quoted from micro1 LegalReasonBench

Reward criterion, +3 points. Identifies in Assignment #2 that the affidavit by Liam

Vance indicates that DuPage County would be an inconvenient forum for him to litigate this case as required under Langenhorst v. Norfolk S. Ry. Co. 219 Ill. 2d 430 (Ill. 2006).

— micro1 LegalReasonBench, `criterion_6`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; arises from the task environment
Task environment	Defective: No affidavit from any of the five physicians is provided in the environment, and the prompt does not explicitly instruct the model to create one. The prompt also says, "If a document is not uploaded, it does not exist for purposes of this analysis."
AI auditors' reasoning	Correct
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Same issue as item 14. The criterion rewards a brief that states that his affidavit says DuPage County would be an inconvenient forum, but no affidavit from Dr. Liam Vance exists in the environment, and the prompt does not explicitly instruct the model to create one. This is one of the missing-affidavit items (14, 15 and 17–20), which all reward citing affidavits the record does not contain.

- *Sources:* see item 14. Native runs: Opus 5.5 (low) met this criterion 3/3 by citing its single unsigned physician form as an executed affidavit; GPT-6 Luna (xhigh) cited no physician affidavit and scored 0/3.

19. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_8`

Quoted from micro1 LegalReasonBench

Reward criterion, +3 points. Identifies in Assignment #2 that the affidavit by Caleb Montgomery indicates that DuPage County would be an inconvenient forum for him to litigate this case as required under Langenhorst v. Norfolk S. Ry. Co. 219 Ill. 2d 430 (Ill. 2006).

— micro1 LegalReasonBench, `criterion_8`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; arises from the task environment
Task environment	Defective: No affidavit from any of the five physicians is provided in the environment, and the prompt does not explicitly instruct the model to create one. The prompt also says, "If a document is not uploaded, it does not exist for purposes of this analysis."
AI auditors' reasoning	Correct
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Same issue as item 14. The criterion rewards a brief that states that his affidavit says DuPage County would be an inconvenient forum, but no affidavit from Dr. Caleb Montgomery exists in the environment, and the prompt does not explicitly instruct the model to create one. This is one of the missing-affidavit items (14, 15 and 17–20), which all reward citing affidavits the record does not contain.

- *Sources:* see item 14. Native runs: Opus 5.5 (low) met this criterion 3/3 by citing its single unsigned physician form as an executed affidavit; GPT-6 Luna (xhigh) cited no physician affidavit and scored 0/3.

20. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_9`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +3 points.** Identifies in Assignment #2 that the affidavit by Maya Linwood indicates that DuPage County would be an inconvenient forum for her to litigate this case as required under Langenhorst v. Norfolk S. Ry. Co. 219 Ill. 2d 430 (Ill. 2006).
>
> — micro1 LegalReasonBench, `criterion_9`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; arises from the task environment
Task environment	Defective: No affidavit from any of the five physicians is provided in the environment, and the prompt does not explicitly instruct the model to create one. The prompt also says, "If a document is not uploaded, it does not exist for purposes of this analysis."
AI auditors' reasoning	Correct
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Same issue as item 14. The criterion rewards a brief that states that her affidavit says DuPage County would be an inconvenient forum, but no affidavit from Dr. Maya Linwood

exists in the environment, and the prompt does not explicitly instruct the model to create one. This is one of the missing-affidavit items (14, 15 and 17–20), which all reward citing affidavits the record does not contain.

- *Sources:* see item 14. Native runs: Opus 5.5 (low) met this criterion 3/3 by citing its single unsigned physician form as an executed affidavit; GPT-6 Luna (xhigh) cited no physician affidavit and scored 0/3.

21. OTB Gang VICAR Trial Jury Instructions, assignment 1: Appeal brief re: justification instruction

criterion `criterion_21`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +15 points.** States that Patton's commission of conspiracy to commit a forcible felony precludes a justification defense.
>
> — micro1 LegalReasonBench, `criterion_21`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. The rubric rewards models for a legal error. The case involves federal charges against a gang member who shot at a rival gang member after following him from a mall. The defendant pled guilty to state charges of conspiracy to commit a forcible felony in connection with the incident. The rubric wants the prosecution's brief to argue that the commission of conspiracy to commit a forcible felony precludes the justification (self-defense) defense, under Iowa Code 704.6. But that statute only applies when the defendant commits a different forcible felony than the one he is alleging is justified. The paradigmatic case would be a robber who waves a hammer at a store clerk to take the cash in the register, intending only to frighten the clerk, and who cannot claim self-defense if the clerk responds by shooting at him (robbery is a forcible felony under the supplied Iowa Code § 702.11(1)). Iowa courts apply the bar in that setting. *See State v. Roby*, No. 22-1017, slip op. at 8 (Iowa Ct. App. Sept. 27, 2023) (defendant who brought a gun to an apartment and shot an occupant who drew his own weapon: "A justification defense was unavailable to Roby because he was participating in the forcible felony of robbery. See Iowa Code § 704.6(1)."). Here, the defendant is arguing that the shooting itself was an act of self-defense; there is no contention that he was committing a different forcible felony prior to the shooting that would negate his ability to claim self-defense for the shooting itself. (Although we did not find an Iowa case expressly holding that the forcible felony must be independent of the act claimed to be justified, judges in at least one other state with a similar statute have expressly held this. *See Martinez v. State*, 981 So. 2d 449 (Fla. 2008) (holding that for the forcible-felony bar to apply, "there must be an independent forcible felony other than the one which the defendant claims he or she committed in self-defense").

Obviously, the law could not be to the contrary because, if it were, the exception would swallow the rule; self-defense would never be available, because absent the defense, the use of force would almost always constitute a forcible felony.)

This environment is based on a real case, and in the real case, neither the Eighth Circuit nor the government's brief opposing certiorari relied on the argument (that the rubric requires) that the conspiracy to commit a forcible felony precludes the justification defense. Instead, the government argued that the defendant as a felon in possession had a duty to retreat safely and could have done so instead of shooting as he did. The evidence indicates this was a scenario where two rival gangs were actively seeking opportunities to shoot members of the other, not that an innocent person was attacked and then defended themselves.

- *Sources:* Patton State Plea.pdf, PDF pp. 4–5 (scanned; read from the visual transcription): Count VIII, "Conspiracy to Commit a Forcible Felony" (Iowa Code §§ 706.1, 706.3), Polk County No. FECR338165, admitting "On or about May 10, 2020 ... I did conspire with other individuals to commit a forcible felony by planning with others to engage in a shooting" (no target felony or victim named). 704.pdf, § 704.6(1) ("One who is participating in a forcible felony"); 702.11.pdf (the enumerated forcible felonies do not include conspiracy); Ellison.pdf (§ 704.6 lays out "the limited circumstances in which a person will have no justification defense"). Facts.pdf, PDF pp. 4–8 (retaliatory shootings between OTB and C-Block; Patton charged with attempted murder in aid of racketeering). Real case: *United States v. Mallory*, 104 F.4th 15 (8th Cir. 2024) (duty to retreat under § 704.1(3); cites only § 704.6(3)); the government's brief in opposition, No. 24-5503, rests on § 704.1(3), *Ellison* and *Baltazar*. The government's Eighth Circuit brief and the district court's instruction filings are not public. *Martinez v. State*, 981 So. 2d 449 (Fla. 2008); *State v. Roby*, No. 22-1017 (Iowa Ct. App. Sept. 27, 2023), slip op. at 8. The rubric's own justification for criteria 28 and 29 rejects the parallel theory as circular ("akin to arguing Patton can't get the justification instruction because he was engaged in the felonious assault of attempted murder").

22. OTB Gang VICAR Trial Jury Instructions, assignment 1: Appeal brief re: justification instruction

criterion `criterion_23`

> Quoted from micro1 LegalReasonBench
>
> **Penalty criterion, −5 points.** Mentions that Patton pled guilty in state court to intimidation with a dangerous weapon.
>
> — micro1 LegalReasonBench, `criterion_23`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: clear; Claude Sonnet 5: clear; Gemini 3.8 Flash: clear. Result: penalty not applied (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: fires; Claude Sonnet 5: fires; Gemini 3.8 Flash: clear. Result: penalty applied (2/3)

My analysis. The rubric penalizes mentioning that Patton pled guilty in state court to intimidation with a dangerous weapon, because that incident relates to a different event than the one at issue on his appeal. The rubric seems intended to test the model's ability to distinguish the events and understand that the earlier incident is unrelated. However, it is misspecified because it penalizes any mention of the issue. Opus was penalized merely for including a drafting note that flagged for the attorney that it had decided *not* to rely on the other incident because it was unrelated. A brief or drafting note that mentions the April 13 plea only to set it aside is correct work, so I do not think penalizing a model for merely mentioning it is correct. (The Count XX plea petition is dated November 18, 2020, after the May 10, 2020 shooting, so it is not a prior conviction.)

- *Sources:* Patton State Plea.pdf (scanned; read from the visual transcription): PDF p. 2, Count XX, intimidation with a dangerous weapon ("On or about April 13, 2020 ... I did discharging a gun into the vehicle of another person"); PDF p. 5, Count VIII, conspiracy for the May 10, 2020 shooting at issue (Facts.pdf, "on May 10, 2020, two more shootings occurred"). Rubric justification: the April 13 plea "involved a different shooting" and is "not relevant." Opus 5.5 (low), Drafting Note 2: the Count XX petition "covers a *different* event, a shooting on April 13, 2020. I rely only on the Count VIII admission"; penalty triggered 2/3 (-5). GPT-6 Luna (xhigh) did not mention Count XX; penalty clear.
- *Related items:* items 1, 13, 16 and 22 are penalties worded "States" or "Mentions" that can reach a response rebutting or setting aside the point; the native runs misfired on 16 and 22.

23. OTB Gang VICAR Trial Jury Instructions, assignment 2: Supplementing brief with new case

criterion `criterion_4`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +25 points.** Mentions that there was a lack of evidence proving who fired first on May 10, 2020.
>
> — micro1 LegalReasonBench, `criterion_4`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. This is another instance where the task environment differs from the real-world case on which it is based in a way that is material to this specific rubric – and a rubric that would be correct for the real-world case is incorrectly applied to this task environment. In the real case, *United States v. Mallory*, 104 F.4th 15 (8th Cir. 2024), there was a dispute concerning who fired first. This task environment was "closed book," however, and the model was given a statement of facts that says that Raysean Nelson "opened fire as the Trailblazer approached" and that "A rolling gun battle ensued." So there is no basis in the supplied environment for the model to assert that

there was a lack of evidence concerning who fired first; the supplied documents state that Nelson did so. The supplied facts are internally consistent; the error is the rubric's failure to follow the task's own changes to the real case.

- *Sources:* Facts.pdf, PDF p. 7 ("Nelson had his car parked near 30th and Clark Streets, and had hidden in the bushes. ... He opened fire as the Trailblazer approached. ... A rolling gun battle ensued."); assignment 2 prompt ("Be sure to limit your review for the Assignment to the attached documents"); *United States v. Mallory*, slip op. at 2 (8th Cir. June 12, 2024) ("though it is unclear who shot first"). The assignment 1 penalty justifications carry the same premise (criterion 25: "at a minimum, it is unclear which side fired first"). Related: item 10 is the other sampled criterion that carries a real-case rule onto changed task facts.

24. OTB Gang VICAR Trial Jury Instructions, assignment 2: Supplementing brief with new case

criterion `criterion_7`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +10 points.** Mentions that Nelson's car was already stopped when Patton fired his gun.
>
> — micro1 LegalReasonBench, `criterion_7`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. This benchmark does two assignments, sequentially, in separate model sessions that are in the same sandbox. In this task, assignment 2 was to build on the assignment 1 output by adding discussion of an additional case. The intent is for models to output response1.md in response to assignment 1, and to output response2.md in the same sandbox for assignment 2. The model is told, "Build on [the assignment 1] analysis as the prompt directs," but not clearly told that it is expected to output a full revised version of the brief (as opposed to supplemental or revised sections that can be inserted into a revised brief). Since assignment 2 is framed as a mere supplementation task, the model might expect that it should write only the new material. However, five of the assignment 1 rubrics are repeated on the assignment 2 output – so a model is penalized heavily if it interprets the instructions as requesting simply an update to the specific section to which the new authority is relevant. In effect, the rubric is testing for one particular response to somewhat ambiguous instructions. This rewards work beyond the scope the prompt states, and a complete revised brief is not what the prompt clearly requests.

- *Sources:* Assignment 2 prompt: "Please supplement your brief with relevant legal analysis of State v. Kuhse." Shared preamble: "Write the final complete legal deliverable for this assignment to `response2.md`. The verifier reads only that file" and "your prior deliverable

`response1.md` is also on disk. Build on that analysis as the prompt directs." Each assignment starts a fresh agent conversation in the same sandbox (run record). Repeated word for word from assignment 1: assignment 2 criteria 5–9 (windows rolled down; Patton hanging out the window; Nelson's car already stopped; the group could have driven away; Patton intended to fire regardless), 65 of assignment 2's 135 positive points (70 points in assignment 1). Native responses: Opus 5.5 (low) rewrote the full argument section (about 6,400 words); GPT-6 Luna (xhigh) wrote a "Supplemental response" (about 1,500 words) that reproduces its assignment 1 brief (about 1,100 words) nearly verbatim and adds sections on *Kuhse*. Neither run took the insert-only approach. On this criterion both runs scored 0/3, and both had also missed it in assignment 1 (criterion 12, +15); the Facts say Nelson's car was "parked" while he "had hidden in the bushes," and the word "stopped" tracks the real Eighth Circuit opinion.

- *Related items:* item 25 shares this defect.

25. OTB Gang VICAR Trial Jury Instructions, assignment 2: Supplementing brief with new case

criterion `criterion_8`

Quoted from micro1 LegalReasonBench

Reward criterion, +15 points. States that the OTB group could have driven away.

— micro1 LegalReasonBench, `criterion_8`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: not met. Result: met (2/3)

My analysis. This benchmark does two assignments, sequentially, in separate model sessions that are in the same sandbox. In this task, assignment 2 was to build on the assignment 1 output by adding discussion of an additional case. The intent is for models to output response1.md in response to assignment 1, and to output response2.md in the same sandbox for assignment 2. The model is told, "Build on [the assignment 1] analysis as the prompt directs," but not clearly told that it is expected to output a full revised version of the brief (as opposed to supplemental or revised sections that can be inserted into a revised brief). Since assignment 2 is framed as a mere supplementation task, the model might expect that it should write only the new material. However, five of the assignment 1 rubrics are repeated on the assignment 2 output – so a model is penalized heavily if it interprets the instructions as requesting simply an update to the specific section to which the new authority is relevant. In effect, the rubric is testing for one particular response to somewhat ambiguous instructions. This rewards work beyond the scope the prompt states, and a complete revised brief is not what the prompt clearly requests.

- *Sources:* Assignment 2 prompt: "Please supplement your brief with relevant legal analysis of State v. Kuhse." Shared preamble: "Write the final complete legal deliverable for this

assignment to `response2.md`. The verifier reads only that file" and "your prior deliverable `response1.md` is also on disk. Build on that analysis as the prompt directs." Each assignment starts a fresh agent conversation in the same sandbox (run record). Repeated word for word from assignment 1: assignment 2 criteria 5–9 (windows rolled down; Patton hanging out the window; Nelson's car already stopped; the group could have driven away; Patton intended to fire regardless), 65 of assignment 2's 135 positive points (70 points in assignment 1). Native responses: Opus 5.5 (low) rewrote the full argument section (about 6,400 words); GPT-6 Luna (xhigh) wrote a "Supplemental response" (about 1,500 words) that reproduces its assignment 1 brief (about 1,100 words) nearly verbatim and adds sections on *Kuhse*. Neither run took the insert-only approach. On this criterion Opus met 2/3 (+15) and Luna 0/3; neither met it in assignment 1 (criterion 13).

- *Related items:* item 24 shares this defect.