When Good Lawyering Gets a Bad Score: An Audit of Two Legal AI Benchmarks

John J. Hughes, III
Independent researcher
https://www.johnjhughesiii.com/research/

Working draft — November 9, 2026
Circulation draft for comment; please do not cite or redistribute

Note on this draft. This working draft paper is subject to revision and improvement. Text marked in red is a preliminary draft for which meaningful rewriting is planned; text marked in teal consists of passages under review by Harvey, which may change meaningfully based on its feedback.

Abstract

As frontier AI models have continued to gain capabilities at a rapid pace, they have saturated traditional benchmarks that posed simple tests of legal knowledge, requiring the industry to develop new benchmarks in 2026 that test models' ability to complete long-horizon legal agentic tasks. Harvey's Legal Agent Benchmark (LAB) and micro1's LegalReasonBench evaluate AI models' ability to draft legal work against lawyer-written rubrics. Some rubrics are used as reinforcement-learning rewards to train models, while held-out sets from both benchmarks are used to evaluate model performance. To assess these new benchmarks that aim to push the frontier of legal model capabilities, as a practicing lawyer who has 18 years of experience in litigation, I audited the publicly released litigation tasks from both benchmarks. AI auditors screened all of the public litigation criteria (2,858 in LAB's 52 litigation tasks; 179 in LegalReasonBench's 11 litigation assignments), and I reviewed a random sample of 25 criteria from each benchmark that the AI models had flagged as defective or potentially defective, under sampling protocols committed before the draws. I reviewed the task instructions, the environment or record the agents had been given, the governing law, and how each benchmark's own AI judges graded two model runs involving that criterion. I found that, in LAB, 16 of 25 sampled criteria were defective, 9 of them objective errors in which the record, the law, or the arithmetic contradicts the criterion. In LegalReasonBench, 21 of 25 were defective, 7 involving objective errors. At 95% confidence, at least 383 LAB criteria (13.4% of all litigation criteria) and at least 45 LegalReasonBench criteria (25.1%) are defective, and at least 171 and 11 of them, about 6% of each benchmark's litigation criteria, are objective errors. A separate AI check of the 3,618 citations in LAB's litigation environments flagged 436 citation errors, 76 of them citations to authorities the check could not locate, and my review of a sample established that at least 303 of the 436 are real errors.

The defects illustrate the difficulty of reducing professional practice to a series of rubrics and straightforward criteria—but the audit also revealed certain patterns in these benchmarks that suggest ways that these new approaches to evaluating frontier legal AI models can be improved. The defects fall into four groups: errors of law or fact, requirements the task does not justify, rewards for hallucinating (or making assertions unsupported by the supplied record), and criteria that would fail competent work product by making one choice on a debatable question of judgment the only acceptable answer. Some of these defects have the effect of rewarding worse lawyering. For example, LegalReasonBench included penalties that seemingly were intended

to discourage models from endorsing or accepting arguments made by a counterparty that are incorrect or misleading. But the rubric as written (and as applied by the benchmark's AI judges) actually penalizes any "mention" of those arguments, so models that correctly pointed out the flaws in an opposing party's argument in order to dismantle it were penalized for doing so.

The audit suggests that there are significant opportunities to improve legal benchmarks and environments, but raises questions about whether existing benchmarks are sufficiently reliable. The defect rates identified here exceed the scoring gaps that separate recently released frontier models on the published leaderboards, and the audit pointed to at least some evidence (although not statistically significant) that the defects could affect some models more than others, potentially impacting the scoring and ranking. The defects can also overstate how much room there is for models to improve on legal tasks: criteria that AI auditors flagged as defective or potentially defective resulted in failures 24.5% of the time, compared with a failure rate of 8.9% on criteria the auditors did not flag, suggesting that defects could be lowering every model's score and understating the existing model performance (although flagged criteria may also simply be harder than the rest). Many existing efforts to post-train models specifically for legal rely on evidence from held-out tasks from the same benchmark family that the models were trained on. While that shows that the training generalizes beyond the specific tasks used for training, it does not address the concerns of construct validity or establish that better performance on a benchmark corresponds to genuine improvement in legal capabilities. I conclude this paper with specific suggestions on how rubric-based benchmarks and training environments for legal could be improved with better quality control, rubric design, and environment design.

1 Introduction

Extraordinary gains in the quality of frontier models' legal capabilities over the past 2–3 years have saturated many older legal benchmarks, which tested models on the knowledge of well-settled propositions of law. This extraordinary growth in legal capabilities has led recently to a wave of new benchmarks (which often also serve as training environments) that aim to push frontier models' legal capabilities further. The leading example today is Harvey's Legal Agent Benchmark (LAB), which evaluates AI agents on a range of legal work product tasks, using LLM judges to score performance against expert-written rubrics. Harvey has trained its own legal AI models using environments that share LAB's structure, and Applied Compute has used LAB as a reinforcement-learning reward signal [9, 11, 12]. More recently, micro1, a data vendor that sells training environments to AI labs, built LegalReasonBench the same way [19].

The objective of these benchmarks, and their related training environments, is to create reward signals for evaluating or training models on high-level legal tasks. The expert-written rubrics aim to reduce at least a major portion of a lawyer's professional judgment to a list of criteria that can help the judge models distinguish good from bad legal work product. Recognizing that this process is imperfect, both micro1 and Harvey open-sourced a portion of their benchmark tasks; Harvey expressly invited practicing lawyers to audit the criteria its benchmarks use and to offer feedback on how these benchmarks can be improved [9]. Taking up that invitation, I audited the open-source litigation tasks from both benchmarks in order to understand how effectively the AI training techniques used at the frontier teach legal models legal judgment. I am a litigator with 18 years working in litigation, including sixteen years at top 10 law firms and two years clerking on federal appellate courts, and a broad generalist litigation practice that has given me exposure to everything from criminal trials and mass torts to securities class actions and bankruptcy; as such I audited the litigation tasks. After reviewing an initial sample, I gave AI auditors instructions on how to assess every criterion in LAB's 52 publicly available litigation tasks and in the 11 litigation assignments of LegalReasonBench's public sample. The AI auditors flagged criteria they believed

were defective or arguably defective under the standards I had provided. I then reviewed a simple random sample of 25 AI-flagged criteria from each benchmark, reading the task instructions, the supplied record, the governing law, and how each benchmark's own judges had graded two frontier models on the same criterion.

The results showed that measuring further gains in legal capabilities, from the already high base that frontier models have achieved, has become extraordinarily difficult. Both benchmarks, even though they were developed by industry leaders for use with cutting-edge reinforcement-learning techniques, had a substantial defect rate, including a mix of legal or factual errors, and criteria that trained models to take steps that I considered to be unreasonable professional judgments. The audit established, with 95% confidence, that at least 13.4% of LAB's litigation criteria and at least 25.1% of LegalReasonBench's are defective, and that at least 171 LAB criteria and 11 LegalReasonBench criteria, about 6% of each, are objective errors. To be clear, this audit does not suggest that these benchmarks are unusually flawed: they are the leading public examples of their kind, I have no comparable audit of any other benchmark to measure them against, and Harvey's research team welcomed the feedback and told me they would use it to improve future versions of the benchmark. Rather, the audit illustrates the challenges of translating lawyer judgment into machine-readable rubrics at scale and raises the question of how and whether AI researchers can be confident that measured gains in legal benchmarks correspond to meaningful improvements in legal capabilities.

One of LegalReasonBench's litigation assignments illustrates the issue. In the task environments, AI models are told they represent a plaintiff who has filed a medical malpractice case in Cook County (Chicago) and need to draft a brief opposing the defendants' motion to transfer venue to a different county where the defendants reside. The defendants' brief argues, as one reason for the transfer, that the Cook County courts have only a 60.5% clearance rate, but the report from which this statistic is taken is included in the environment and makes clear that the defendants' argument is misleading. It conflates criminal and other cases with the clearance rate for civil matters, which is the relevant rate for a medical malpractice case and is much higher. The expert-written rubric indicates that the model drafting the plaintiff's brief should be penalized if its brief "states" that the report shows a 60.5% clearance rate. The expert presumably meant to make the point that the model should not *endorse* or *adopt* defendants' position that the 60.5% clearance rate is the relevant comparison. The task presumably wants to check that the model realizes the defendants' statistic is misleading and thus refrains from endorsing it in the opposition brief. But the AI judges interpret "states" literally: one of the AI models I tested on the benchmark, Claude Opus 5.5, dealt with the defendants' argument exactly the way a good lawyer would have—only to be punished for it. Under the heading "Defendants rely on the wrong statistic," it stated the defendants' figure, explained that the statistic lumps together criminal, traffic, family, and juvenile matters with civil cases, and it cited the civil case statistic that actually matters: 97.9% for Cook County and 96.3% for the proposed forum—meaning, if anything, that on defendants' preferred measure, plaintiff's chosen forum of Cook County was better. But all three AI judges imposed a 30-point penalty on Opus for that response (tied for the heaviest penalty weight in the assignments I reviewed), because the rubric as written literally requires the agents not to mention the statistic at all, not even to refute it—even though that's exactly what a good lawyer should do. Holding every other grade fixed, that one penalty took the assignment score from 86% to 45% (Appendix A).

I observed several other kinds of defects in the audit. In several instances, criteria simply got the law wrong—or demanded that models assert a fact that was contradicted by the task environment provided to the model. Some criteria required models to take a particular position on a question of professional judgment where, in my view, the best lawyers either would reach a different position

from the one the rubric required or at least would consider a position that the rubric would fail to reward to be no less reasonable than the rubric's preferred position.

Section 2 reports the results of the audits. Section 3 sorts the defects I observed into four recurring groups and gives examples of each. Section 4 explores what implications this might have for legal-domain evaluations and training: what the scores can establish at the frontier, where the rewards favor worse lawyering, why penalizing errors is easier than rewarding judgment, and how legal environments could be improved. The appendices have a detailed write-up on each criterion I reviewed, with notes on how the AI models I tested performed on those criteria and a narrative description of my reasoning for each determination I made.

2 Audit design and results

2.1 The benchmarks

Harvey AI Corporation, the leading vertical legal AI company today, released LAB in May 2026. In recent model releases, LAB has been the most widely cited legal benchmark used by frontier labs to illustrate the legal capability of their current frontier models: Anthropic, Google, SpaceX, and Meta, among other labs, included LAB scores in their release posts or system cards in 2026 (OpenAI's GPT-6 model releases did not include a legal benchmark score)[1] [23–27]. The initial release had roughly 1,250 tasks across 24 practice areas (of which litigation, the area I audited, is just one) and more than 75,000 expert-written criteria. Harvey has said LAB uses real client matters handled by practicing lawyers as inspiration for AI-generated task documents and environments [9]. Harvey uses AI generation to get around the difficulty of getting clients to share real legal documents, which risks waiver of privilege over internal materials. Each LAB criterion is a binary pass or fail. Somewhat unusually, Harvey took the position that in legal work missing even one issue is a critical mistake, so a task is graded successful only if every criterion passes. A portion of the benchmark is open source on GitHub and available for lawyers to audit, while a separate private held-out set is used to report LAB results on the Artificial Analysis leaderboard [21]. Harvey says it started from real client matters, split them into tasks, and had the rubrics written and reviewed by experts; its launch post names Mercor and Snorkel AI as contributing research groups, and a later talk by Harvey's co-founder describes those providers as running expert review and verification of every generated LAB task [8, 9]. In comments on an earlier draft of this paper, Harvey told me that the outside providers had a more limited role on the public LAB tasks than on the held-out set used for the Artificial Analysis evaluation. A public post by one of LAB's authors describes review by more than 100 lawyers, at Harvey and through its partner vendors, totaling thousands of hours [10]. I audited the 52 tasks and 2,858 criteria in the public litigation and dispute resolution folder at Harvey commit `1dd81403`. Throughout, I use "criterion" for a single pass/fail or point-weighted test and "rubric" for the full set of criteria attached to a task or assignment, although the appendix worksheets sometimes call the one criterion under discussion the rubric; the units I screened, sampled, and labeled defective are criteria.

micro1 is a data and research lab that produces reinforcement-learning environments used by frontier labs to improve their models, in legal and several other domains, using experts such as practicing attorneys to help with the work. LegalReasonBench has 50 matters, 122 assignments, and 2,236 criteria, of which 353 are penalty criteria. micro1 states that a practicing lawyer created each assignment, the associated case file, and its rubrics [19]. It also reports a three-judge classification

[1] OpenAI's September 2026 announcement of Astra for Law, a legal product rather than a model release, reported results on Vals AI's Legal Research Bench instead.

of its criteria into IRAC categories with high agreement, which is a classification, not a substantive validation of each criterion. The open-source public sample on Hugging Face has 15 assignments and 221 criteria. I audited a litigation subset that covers 11 assignments, with 179 criteria across four matters. At least three of the four appear to be adapted from real cases with names and some facts changed. Each criterion has an associated points value; three AI judges decide, by majority vote, whether a given criterion was satisfied or not. The assignment score is calculated as the points earned on positive criteria, less penalties imposed under negative criteria, divided by the total points available on positive criteria; a negative score counts as zero. The dataset I reviewed is pinned on Hugging Face at revision `1bfa904c` [5]. Neither release permits reconstructing the review history of a disputed criterion.

There are two noteworthy differences between the benchmarks: LAB's all-pass scoring rule means that one failed criterion results in a zero on an otherwise excellent deliverable, while LegalReasonBench imposes a points-based penalty but does not zero an assignment for one failed criterion. Additionally, the task environments in Harvey's LAB are primarily synthetic with no recognizable connection to real cases, whereas micro1's LegalReasonBench tasks plainly are adapted from real litigation matters. There were tradeoffs and different failure modes associated with each of these approaches (Section 3.1 and Appendix C).

2.2 AI screening and two sampling frames

To conduct the audits, I gave GPT-6 Sol and Claude Opus 5.5 examples of defective criteria and asked them to flag any defective or arguably defective criteria. For LAB, Sol flagged 616 criteria and Opus flagged 650; after Opus reviewed Sol's findings, 835 criteria (29.2%) were flagged by one model or both. For LegalReasonBench, Sol flagged 51 criteria and Opus 43; between them they flagged 63 of 179 criteria (35.2%), and a later reconciliation pass retained 36 of the 63 [5]. In both audits an Opus pass was permitted to review both models' findings. This was intended to operate as two standalone reviews of these benchmarks, not a formal statistically valid comparison of the defect rate between them.

2.3 The standard I applied

After the AI auditors flagged defective or arguably defective criteria, I drew a simple random sample of 25 criteria from each benchmark's flagged set without replacement, with a fixed seed, using a sampling protocol that was committed before the draw [2, 6]. Criteria where the answer key contains legal or factual errors, or rests on objective discrepancies or errors in the task environment, were considered defective and labeled objectively incorrect.

An additional set of criteria was labeled defective as a matter of professional judgment; this was used in situations where the rubric required a decision that, in my view, top lawyers would consider inaccurate. Professional judgment defects also include a few instances where the rubric would fail a response that, in my view, most lawyers would see as a judgment call—and the rubric effectively forces the AI model to take one particular position on an issue where there is actually room for debate, which can itself be a kind of legal error. For example, LegalReasonBench required advising a governmental client who lost an administrative review petition in district court to both appeal the loss and redo the decision at the agency; that's a reasonable judgment, and probably the default path most lawyers would take, but in my view it's erroneous to require that recommendation in the internal memo to the client. The circumstances of that case are such that the client's objectives could be fully achieved by simply redoing the decision at the agency level and forgoing the appeal

(which likely would become moot once the agency reinstated its decision with new reasoning). Thus, in the context of the assignment for an internal memo, the rubric ideally would credit a memo that presents both options to the client and allows the client to decide whether the added expense of the appeal is worthwhile instead of simply assuming that appealing the adverse ruling is the best approach.

After preparing my initial draft of the verdicts, I asked AI models (GPT-6 Sol, GPT-6 Astra, Claude Opus 5.5, and Claude Fable 5.1) to stress test them and argue both sides of each call, then did a consistency pass through the full set. A handful of verdicts changed based on the AI-assisted stress testing and additional review that I did.

2.4 Native runs and citation checks

I ran both benchmarks on GPT-6 Luna (xhigh reasoning) and Claude Opus 5.5 (low reasoning), using each benchmark's own code (including the benchmark's harness, judge models, and scoring code). This was intended merely as a diagnostic, not a model comparison, to understand where a given criterion might pass for one model but fail for the other and whether that was based on genuine differences in the quality of the legal work product or pointed to potential defects in the rubrics (Appendix A).

The rubric audit revealed that some of the task documents on Harvey's LAB contained hallucinated case law; some of the LAB answer keys similarly referred to hallucinated cases. I had AI agents check every citation in the LAB litigation environments against CourtListener, then reviewed a random sample of the flagged citations myself under a pre-committed stopping rule (specifically, I drew 20 and committed in advance to stop after the first 10 if all 10 were confirmed errors, and otherwise to review all 20). A separate AI citation census covered the LegalReasonBench litigation documents (Appendix B).

2.5 Results

Table 1 reports the results. As noted, this analysis was designed as two distinct audits of the benchmarks, not a statistical comparison of defect rates. The one-sided 95% lower bounds of 383 and 45 are 13.4% and 25.1% of the respective audited populations, the point estimates 18.7% and 29.6%. Note that the bounds cover sampling uncertainty, not uncertainty about the accuracy of my verdicts, and in both benchmarks, there are cases where several sampled criteria share one underlying problem, so the verdicts are not truly independent. Regardless, the defect rates are of the same order as the differences in criteria pass rates that separate frontier models on these benchmarks, a comparison taken up in Section 4.1. The exhibits in Section 3 establish what the defects do.

In LAB, 9 of the 16 defective criteria are objective errors, such as these: the environment's time entries add up to $215,680, not the answer key value the rubric assumes of $193,372.50; the rubric states the opposite of Georgia tortious interference law; the rubric finds a conflict in a preservation notice that contains none. The other seven fail work an elite lawyer could consider at least as good as what the rubric requires. In LegalReasonBench, 7 of 21 are objective errors, and the remaining 14 are errors of professional judgment; 6 arise from the task materials rather than the criterion alone.

In cases where both AI auditors considered a criterion defective, they were almost always right. In LAB, 9 of the 10 criteria flagged by both models were defective, compared with 7 of 15 flagged by one; in LegalReasonBench, all 12 flagged by both blind reports were defective, compared with 9 of

	Harvey LAB public litigation tasks	LegalReasonBench public litigation assignments
Scope of audit	52 tasks	11 assignments in 4 matters
Criteria	2,858	179
AI-flagged as problematic	835	63
Lawyer-reviewed	25 (19 tasks)	25 (9 assignments)
Defective	16	21
of which objective error	9	7
of which professional judgment	7	14
Arguable	3	0
Not defective	6	4
In a task environment with a defect	14	12
Defective within audit scope, point estimate	534	53
Defective within audit scope, 95% lower bound	383	45
as a share of all audited criteria	13.4%	25.1%
Objective errors within audit scope, lower bound	171	11

Table 1: Lawyer review of a random sample of AI-flagged criteria. The point estimate shares of all audited criteria that have defects are 18.7% for LAB and 29.6% for LegalReasonBench.

13.

The two model runs I did with the benchmarks illustrate how much all-pass scoring (used in Harvey's LAB) impacts the results. Across the 52 LAB litigation tasks, averaging each judge's per-task pass fraction, GPT-6 Luna passed 83.8% of criteria and Claude Opus 5.5 passed 90.8%. The all-pass rate, by contrast, was 1.9% for Luna and 2.9% for Opus, and each model passed exactly one task outright with both judges. Under LAB's all-pass approach, one defective criterion is enough to fail a task. Thus, LAB's headline number is unusually sensitive to the correctness of each criterion.

I also did a sweep for citation errors, which found a substantial number in Harvey's LAB (likely owing to the synthetic data used to create it). Within LAB's litigation environments, an AI check of 3,618 citations found 436 seemingly unintentional errors in citations to purported real authority, including 76 citations to authorities the check could not locate (which I did not verify individually). I reviewed a random sample under a pre-committed stopping rule and confirmed that the first ten I reviewed were all true citation errors. Statistically, the one-sided 95% lower bound is that at least 303 of the 436 are real errors. While a similar AI audit yielded some instances where cases were mischaracterized in LegalReasonBench, it did not flag any hallucinated cases.

3 How the criteria misgrade legal work

The defects I identified, 16 in LAB and 21 in LegalReasonBench (LRB), are grouped here into four main categories: errors of law or fact in the rubrics; misspecified tasks or incorrect requirements; rewarding hallucinations or unsafe drafting; or errors of professional judgment. Table 2 shows the count by group, and Tables 3 through 6 list each defect in its group. The complete review of all 50 sampled criteria is in Appendices D and E.

These groups are somewhat overlapping. For example, the motion-to-dismiss criterion (LAB item 18) rewards models for raising a factual dispute on a motion to dismiss under Rule 12(b)(6), but courts are not permitted to resolve factual disputes under the rule. I consider that criterion an error of law, because it wants the model to cite facts from the defendant's internal documents that

are not incorporated into the complaint, which Rule 12(d) generally prohibits. *See Tellabs, Inc. v. Makor Issues & Rights, Ltd.*, 551 U.S. 308, 322 (2007). But it is also a strategic problem: raising that argument makes it easy for the opponent to point out to the judge that the movant is really trying to use the motion to dismiss to argue the facts, which makes it more likely the judge will simply deny the motion, even if there were other, better arguments. In that case, and others, a given criterion could fit under more than one of these groupings and I simply assigned it to the one that fit best. With one exception, the first group coincides with the 9 LAB and 7 LegalReasonBench criteria coded as unambiguous objective errors, and the other three groups are the professional-judgment defects. The exception is the parol-evidence criterion (LAB item 17), an objective error that I have placed in the third group because of what it rewards.

Group	LAB	LRB	Total	Listed in
Errors of law or fact	8	7	15	Table 3
Misspecified tasks or incorrect requirements	1	5	6	Table 4
Rewarding hallucinations or unsafe drafting	1	6	7	Table 5
Errors of professional judgment	6	3	9	Table 6
Total	16	21	37	

Table 2: The 16 LAB and 21 LegalReasonBench (LRB) criteria I found defective, by group. In the group tables, item numbers are the appendix entries (Appendices D and E); LAB criteria carry Harvey's identifiers, and LegalReasonBench criteria the dataset's numeric identifiers where they have one.

3.1 Errors of law or fact

The largest group contains criteria where the answer key is simply wrong. This could involve the answer key requiring models to misstate the applicable legal rule, to misread or misunderstand a document within the environment, or to reach a conclusion that is contradicted by facts supplied in the environment; this also applies if the rubric relies on arithmetic or numbers that do not add up. I subgrouped these into: errors of law, errors of fact and arithmetic, and errors in reading a provision together with its exceptions. Table 3 lists the criteria in this group.

Errors of law. There were several instances of straightforward errors of law, where the answer key is simply wrong.

- The jury-instruction task requires the memo to explain that under Georgia law tortious interference with a contract does not require an independently wrongful act while interference with business relations does. The AI auditors flagged that Georgia actually applies the same elements to both, and I did research to confirm they are correct. The fictional court order from which the rubric's law seems to be taken cites a case that cannot be found (LAB item 23).
- A LegalReasonBench land-use criterion rewards models that invoke Fla. Stat. § 163.3167(10). But that statute applies only if a development is approved under regulations, no one challenges the approval within the deadline, and the regulations at issue are later invalidated (for example, in a subsequent case involving a different development that results in a court challenge). In this case, the approval at issue was timely challenged and quashed by the court (item 5) so the statute cited is irrelevant. Other criteria in this same matter similarly try to apply a county rule that by its terms governs only remands from the Board to the hearing examiner, when

Item	Task	The problem
Errors of law		
LAB 23 (C-020)	Jury-instruction review	Georgia tortious interference law is the opposite of what the rubric states; the source order cites an apparently hallucinated case
LAB 18 (C-045)	Motion to dismiss brief	Requires raising fact disputes, which are impermissible on a Rule 12(b)(6) motion
LAB 13 (C-008)	Interrogatories	Fails interrogatories that omit a definition of "document," although the court's local rule already supplies one for all discovery requests and interrogatories need not use the term
LRB 1	Land use, assignment 1	Misapplies a county rule that is not applicable here and thus incorrectly penalizes models that cite a correct withdrawal date
LRB 6	Land use, assignment 1	Rewards stating an incorrect withdrawal date (which neither party asserted in the real litigation), under the same inapplicable rule
LRB 5	Land use, assignment 1	Misapplies a statute that protects unchallenged approvals to an approval that was timely challenged
LRB 10 (crit. 4)	Police discipline, assignment 2	Rewards the claim that no authority requires proving comparable cases, ignoring that *Douglas* and MPD's own notice do place that burden on MPD
LRB 21 (crit. 21)	Jury instructions, assignment 1	Misapplies an exception to the self-defense defense that does not apply to this fact pattern
Errors of fact and arithmetic		
LAB 1 (C-049)	Invoice staffing review	Required 17.2% overage is based on a subtotal contradicted by the invoice's time entries (30.7%)
LAB 3 (C-029)	Discovery gap memo	Miscounts the documents and treats five as "thin" evidence of damages
LAB 19 (C-028)	Interrogatory responses	Provides only calendar-year data and expects the model to conflate it with the contract-year data requested
LRB 9 (crit. 8)	Police discipline, assignment 1	Rewards asserting that there is a mismatch in dates, which the memorandum does not contain
LRB 23 (crit. 4)	Jury instructions, assignment 2	Rewards asserting that there's no evidence who fired first, ignoring the evidence in the environment that says who did
Errors in reading interacting provisions		
LAB 20 (C-002)	Preservation notice memo	Incorrectly treats an express broader date range for one request as somehow conflicting with other requests using a shorter date range
LAB 25 (C-042)	Invoice compliance review	Incorrectly treats business-class airfare expressly permitted by the policy in the environment as a borderline issue to be flagged

Table 3: The 15 criteria in the first group, errors of law or fact.

this case involves a remand from the examiner to staff (items 1 and 6).

- A jury-instruction task rewards arguing that the defendant's plea to conspiracy bars a self-defense defense under Iowa Code § 704.6, a provision that precludes a defendant from claiming justification if the defendant was committing a forcible felony (LRB item 21). But that means a different forcible felony than the one for which the defendant is claiming justification; the paradigmatic case would be that, say, an armed robber who tries to rob someone with a baseball bat cannot claim "self-defense" if the victim turns out to have a gun and puts the robber in fear for his life. It does not mean, as the rubric suggests, that a defendant who is attacked and uses force in response thereby negates his own claim of self-defense; that circular logic would eliminate self-defense completely, since self-defense is asserted precisely when the defendant claims to have used force to defend himself or others.
- LAB includes a motion to dismiss drafting assignment (LAB item 18), where the model is required to prepare a memorandum of law supporting a Rule 12(b)(6) motion to dismiss a

complaint. Criterion C-045 requires the brief to blame plaintiff's own failures (including late specifications and a vacancy in the project manager role) "as causes of the Go-Live delay." But this in effect demands that the model raise disputes over causation in a motion to dismiss, which are paradigmatic fact issues; moreover, the factual assertion that plaintiff caused the delays is grounded in the defendant's internal documents, not allegations made on the face of the complaint. On a motion to dismiss, the court must take the complaint's allegations as true so there is no legal basis for introducing the defendant's claimed version of events in the procedural posture of a motion to dismiss. The AI models that completed this task correctly left those issues out, and were failed on C-045 as a result, effectively penalizing them for a procedural judgment they got right.

In the course of discussing these results with people involved in creating these benchmarks, one argument raised was that environments that are internally inconsistent about the law are genuinely problematic, but that environments that are internally consistent are not problematic simply because they assert legal propositions that contradict the real-world legal rules. I have not accepted that position. Environments that reward models for applying some sort of fictitious, stylized version of the law teach the model that its knowledge of the law from pre-training is not useful for earning reward. In an adversarial legal setting, models should start with reasonably strong priors about the right answer and not blindly defer to evidence in the environment that contradicts their memory [40]. A good lawyer uses their extrinsic knowledge and training to identify claims or arguments that appear unsound or potentially debatable, and to investigate those before accepting them. Here, the reward seems to punish this desirable behavior. On LAB item 17, for example (Section 3.3), the model that checked the two cases, found they did not say what the memo claimed, and cited the right authority was failed. When given a task environment, the model has no way to know whether law supplied in the environment is a deliberate fiction that the environment's creators want the model to simply treat as true for purposes of a stylized exercise, or an error or misstatement of the law that the model is expected to identify and correct. While an environment can properly modify the specific facts of a given case (since these vary from case to case), the law itself is public and changes only slowly. While an environment can vary the law (e.g., introduce newly passed legislation to train the model to check for recent updates rather than relying blindly on its memory from pre-training in every case), environments should be designed so that models are rewarded for being skeptical of legal propositions that are completely at odds with their pre-training data, unless they are given a clear explanation (like a new authority that changed the law since the model's training cutoff).

Errors of fact and arithmetic. The invoice task rubric says the budget overage is 17.2%, which is what is shown on the invoice's summary page. But the 117 individual time entries sum up to $215,680, which puts the overage at 30.7%, far outside the rubric's tolerance (LAB item 1). The discovery-gap criterion requires the memo to say that only five documents relate to the counterclaim damages, when a sixth plainly does (item 3). The interrogatory criterion gives the model only calendar-year purchase data for a request that seeks contract-year purchases, which are not supplied (item 19). In LegalReasonBench's jury-instruction matter, 25 points are awarded for asserting that there's no evidence about who fired first. While that was true of the real case on which this environment is based, this task was constructed as a "closed universe" and includes a statement of facts that modifies the facts somewhat from the real-world case, including that the statement says explicitly who fired first (item 23).

Errors in reading interacting provisions. Reading a provision together with its qualifications and exceptions is a transferable skill, and two criteria do this incorrectly (one of them, LAB item 20, is discussed below). A travel policy permits business class on domestic flights of four hours or more, and the rubric treats a flight it describes as four hours and twenty minutes as a borderline issue that must be flagged for discussion, when the policy expressly says that's allowed (LAB item 25).

"No conflict" is an answer. Two criteria, one in each benchmark, require the model to find a discrepancy that is not there. In LAB's preservation notice task, a Department of Justice notice defines the relevant period as January 1, 2019 forward, "unless otherwise specified herein," and then provides that "[n]otwithstanding the Relevant Period defined in Paragraph 12," certain financial records must be preserved from January 1, 2017. Pursuant to Criterion C-002, the model is rewarded if the draft memo recommends asking the DOJ to provide clarification "regarding the conflicting date ranges." But there is no genuine conflict. The notice sets a time period applied to all requests by default, and then, for one, makes an exception and uses a longer period. Luna understood this correctly, but was failed on this criterion by both AI judges. Opus's memo was failed by one AI judge and passed by the other, on the strength of an unrelated clarification request about category numbering.

In LegalReasonBench's police-discipline matter, criterion 8 of the first assignment awards 28 of 71 available points for identifying "that the comparator memorandum's stated coverage years conflict with the dates of the cases it summarizes."

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +28 points.** Identifies that the comparator memorandum's stated coverage years conflict with the dates of the cases it summarizes.
>
> — micro1 LegalReasonBench, `criterion_8`, police discipline, assignment 1, rubric at revision `1bfa904c`

There does not seem to be any conflict. The memo says it covers calendar years 2018 through 2021. And the four cases it describes arose in August 2018, November 2019, May 2020, and March 2021 and were resolved between January 2019 and October 2021. So every date is inside the stated window. Luna and Opus correctly did not report a conflict and thus received zero of 28 points, but a model that hallucinated some sort of conflict would have received a reward for doing so.

3.2 Misspecified tasks or incorrect requirements

Several criteria appear to either be misspecified, in the sense that the criterion as written does not seem to capture what the environment's creator was intending to capture, or create a mismatch between the prompt/instructions the model received and the work product it is expected to create. Table 4 lists the criteria in this group.

Mention is not endorsement. Here are some examples:

- The 60.5% penalty described in the introduction of this paper is the clearest case.

  > Quoted from micro1 LegalReasonBench
  >
  > **Penalty criterion, −30 points.** States in Assignment #2 that the 2023 Illinois Annual Report Statistics report a 60.5% clearance rate in Cook County, Illinois.
  >
  > — micro1 LegalReasonBench, `criterion_18`, forum non conveniens, assignment 2, rubric at revision `1bfa904c`

Item	Task	The problem
Mentioning an argument mistaken for endorsing it		
LRB 16 (crit. 18)	Forum non conveniens, assignment 2	−30 for mentioning the opponent's 60.5% figure, even to rebut it
LRB 22 (crit. 23)	Jury instructions, assignment 1	−5 for mentioning an unrelated guilty plea, even to explain why it is not relevant
LRB 13 (crit. 22)	Forum non conveniens, assignment 1	−30 for mentioning the defendant's argument that Cook County lacks interest, even to refute it
Requirements the assignment does not justify		
LAB 6 (C-049)	Conflict check memo	Demands scope creep, with a conflicts memo for one engagement adding firm-wide process recommendations
LRB 24 (crit. 7)	Jury instructions, assignment 2	Grades a supplement as though a fully rewritten brief were clearly requested
LRB 25 (crit. 8)	Jury instructions, assignment 2	Same

Table 4: The 6 criteria in the second group, misspecified tasks or incorrect requirements.

The defendants' motion to transfer relied on the 60.5% clearance rate in Cook County as a reason to move the case elsewhere. The models drafting the response correctly realized that statistic was misleading because it includes a variety of other case types, not just civil cases, and made that point to rebut the defendants' argument. The rubric then punished them for that, because the rubric penalizes any mention of the issue. Presumably the environment's creator only meant to penalize endorsing or adopting the 60.5% statistic as the correct figure, rather than disputing its relevance.

- The same mechanism appears in the criminal appeal (LRB item 22), where a −5 penalty fires if the brief "mentions" that the defendant pled guilty in state court to intimidation with a dangerous weapon. Opus was penalized for including a drafting note to the user correctly explaining that it was *not* citing that other incident because it was from an unrelated case (the point the rubric was meant to test).
- The first forum non conveniens assignment had the same verbal formula ("Mentions ... that Cook County lacks interest"). Although it did not fire on my two test runs, I labeled that one defective as well because the formulation clearly can (depending on how the AI judge interprets it in a given run) punish models that simply recite defendants' argument that Cook County lacks an interest to rebut it.

Requirements the assignment does not justify. The second subgroup consists of cases where the prompt to the model does not reasonably convey what work the rubric expects the model to have done. One case involved a prompt that asked a model to write a memo analyzing a specific conflicts issue impacting a particular case; the rubric rewarded the model for making recommendations to improve the firm-wide conflicts database, an entirely different project that (even if the model could infer it might be desirable) belongs in a different task and memo.

LegalReasonBench uses sequential assignments, where the model completes two tasks in the same environment and sandbox. While in most cases the two tasks reviewed were sufficiently independent that they could be graded separately, in one instance, the second task was an assignment to supplement the analysis from the first task with a new case. The models' outputs appeared substantively competent and useful but lost points because the second task repeats the criteria from the first task, and the models did not appreciate that their output for the second task needed

to first repeat everything they had just output for the first task. The criteria reward a particular mechanical approach to revising documents that would normally be handled at the harness level and where the expectation is not clearly communicated to the model in the task environment. While not every requirement needs to be spelled out in detail, and it's appropriate to test the model's ability to infer things from the environment, the specific criteria here tested requirements that do not follow reasonably from the assignment and context provided (and where the rubric is not really rewarding better legal judgment, just a particular mechanical approach to editing documents within the environment).

3.3 Rewarding hallucinations or unsafe drafting

The third group rewards a model either for hallucinations, or for unsafe drafting practices.

Hallucinations. LAB's motion-to-dismiss task (discussed under item 18 above, Criterion C-037) requires the brief to cite "relevant Delaware authority on the parol evidence rule and integration clauses, such as SIGA Technologies v. PharmAthene or Eagle Industries v. DeVilbiss." The task environment contains a research memo with hallucinated quotations from both cases for language that does not actually appear in them. The rubric also conflates two distinct legal doctrines: under Delaware law the parol evidence rule and an integration clause are not the doctrines that bar a fraud claim on this fact pattern; rather, the issue is whether the plaintiff can reasonably allege reliance on pre-signing representations after agreeing contractually that it was not relying on any such representations. In its test run on this task, Opus correctly did not cite these cases and instead cited the correct doctrine from Delaware (the anti-reliance cases). That resulted in a failing grade.

Item	Task	The problem
LAB 17 (C-037)	Motion to dismiss brief	Demands authority on the wrong doctrine and names cases misquoted in the research memo; the brief that checked the cases and supplied appropriate anti-reliance authority was failed
LRB 14 (crit. 1)	Forum non conveniens, assignment 2	Rewards stating that the brief is supported by a physician's affidavit that does not exist
LRB 17 (crit. 3)	Forum non conveniens, assignment 2	Same, second physician
LRB 18 (crit. 6)	Forum non conveniens, assignment 2	Rewards saying a nonexistent affidavit shows DuPage County inconvenient for the affiant
LRB 19 (crit. 8)	Forum non conveniens, assignment 2	Same, second physician
LRB 20 (crit. 9)	Forum non conveniens, assignment 2	Same, third physician
LRB 15 (crit. 11)	Forum non conveniens, assignment 2	Rewards citing the plaintiff's own affidavit, which is not in the record

Table 5: The 7 criteria in the third group, rewarding hallucinations or unsafe drafting.

Unsafe drafting (presenting hoped-for evidence as if it had been obtained). LegalReasonBench had two sequential forum non conveniens tasks. The first states in the task environment that another lawyer on the case team has spoken with five Cook County physicians and learned they were willing to testify. But no affidavit from any of them (or from the plaintiff) is in the record, and the prompt states that a document not uploaded "does not exist for purposes of this analysis." Ten criteria (five of them in my sample) then award points for stating that the response is supported by

each physician's affidavit and that each affidavit shows DuPage County would be inconvenient for the affiant.

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +3 points.** States in Assignment #2 that the response is supported by the affidavit of Liam Vance, a physician who provided care to Baby Davis at Cook County Hospital.
>
> — micro1 LegalReasonBench, `criterion_1`, forum non conveniens, assignment 2, rubric at revision `1bfa904c`

This is an undesirable training signal, as currently formulated. While preliminary phone conversations indicating that third-party witnesses are willing to help are useful, actually obtaining executed affidavits from third-party witnesses that counsel does not control (in time to submit them in opposition to an already-filed motion to transfer) is a substantial undertaking. Several things can go wrong that result in the contemplated affidavits not being received or executed in time to support this opposition, as set forth more fully in the appendix entry for this criterion (Appendix E, item 14). There are several reasons to question whether it is strategically beneficial to submit any of these affidavits, let alone all five, and even if it is, there is no guarantee that the witnesses ultimately will follow through in time and sign affidavits that say exactly what counsel wants them to say. As a general matter, moreover, lawyers are extremely concerned about models hallucinating things that could result in sanctions, as a public database of court decisions had identified more than 2,200 cases worldwide, as of November 2026, in which lawyers or litigants were found to have relied on fictitious citations or other AI-generated material, nearly 900 of them involving lawyers [28]. Training models to just output a draft citing to affidavits that don't exist or that are still in draft form creates a genuine practical risk of human error that could lead to sanctions; the draft circulated within the firm might make a senior partner think that someone else has already obtained the necessary affidavits when in fact that's an open issue. If a careless junior associate or paralegal files the brief without checking, a brief that asserts the case team has affidavits that do not exist could hit the docket, creating a substantive issue in the case.

For these reasons, I think the approach of models that do not cite to these as of yet nonexistent documents is much safer and preferable in an age of AI drafting. Moreover, if an environment designer wanted to test a criterion like this, a much better approach would be to require that the model clearly indicate in draft brackets that the citation is to expected or contemplated testimony that the case team hopes to achieve but has not yet. The actual drafts that Opus prepared (which were rewarded) had no draft brackets around the citations so, by industry convention, they look like final citations to documents that actually exist. Having a citation of the form "[Proposed Vance Aff. ¶ ___ (not yet executed)]" allows testing the same substantive output but is much safer because it avoids the risk of miscommunication within the case team or an erroneous filing being submitted to the court.

The LAB interrogatory-response criterion (item 19, which I counted in the first group) involves an interrogatory that is seeking contract-year data but the environment contains only calendar-year figures, which are not the same. The rubric passes an answer that conceals this gap, by just providing the calendar-year data from the environment. (It appears that the mismatch was unintentional, but the consequence is that this rubric might train models to think that details like that don't matter, when they do.)

3.4 Errors of professional judgment

The last group is for criteria that take a position on a debatable question of professional judgment, and would fail a response that is nuanced, thoughtful, and at least as good as the rubric's preferred response. In some cases, these are criteria that require models to advance arguments or take litigation steps that are probably not necessary and which involve costs or drawbacks that may or may not be appropriate given the client's objectives. I ultimately labeled these criteria defective not because they are wrong in isolation but because they often incentivize the model to reach one particular result on a judgment call, where the better answer would be to acknowledge that there are two (or more) options, both of which entail costs and tradeoffs, so that the client can make an informed decision. Table 6 lists the criteria in this group.

Item	Task	The problem
Work whose cost exceeds its value		
LAB 16 (C-009)	Motion in limine	Fails a motion to exclude an expert unless it raises a minor rounding difference/typo in her damages total, a point a good advocate might omit; a companion criterion not in the sample (C-010) requires presenting it as a *Daubert* ground
LAB 4 (C-025)	Case assessment memo	Requires recommending a motion to dismiss that may not be in the client's interests, as it would not reduce exposure or narrow discovery
LAB 15 (C-028)	Discovery responses	Requires a questionable and toothless legal conclusion objection, then requires highlighting this minor point in the client strategy memo
LRB 3	Land use, assignment 1	Requires an appeal when curing the defects on remand alone could achieve the client's objective at lower cost
LRB 4	Land use, assignment 1	Requires remand "while an appeal is pending," instead of permitting a recommendation to remand in lieu of appeal
One reasonable choice made mandatory		
LAB 2 (C-012)	Privilege log	Fails "fully privileged" for counsel's meeting notes, the better answer
LAB 5 (C-030)	Conflict check memo	Requires screening an associate from a case, when informed client consent with a confidentiality protocol is permissible
LAB 9 (C-003)	Product-liability assessment	Requires a fatality to appear in the executive summary of a memo where it's more appropriately discussed elsewhere
LRB 2	Land use, assignment 1	−30 for recommending a pause in construction that the cited Florida Supreme Court precedent supports

Table 6: The 9 criteria in the fourth group, errors of professional judgment.

Three examples illustrate this point, all instances where the criteria require models to take a certain step in litigation or advance a certain argument. The rubrics seem to proceed on the assumption that it is desirable to have more arguments and points raised, but quite often, the opposite is true. As a Supreme Court justice and a leading legal writing professor have argued: "The most important—the very most important—step you will take in any presentation, whether before a trial court or an appellate court, is selecting the arguments that you'll advance" [41, § 12]. Less is often more in litigation. "Scattershot argument is ineffective," [41, § 12] and spends down the advocate's credibility with the court, while distracting from more important points. The Supreme Court has said the same in holding that appointed appellate counsel need not raise every nonfrivolous issue a client wants raised: "Experienced advocates since time beyond memory have emphasized the importance of winnowing out weaker arguments on appeal and focusing on one central issue if possible, or at most on a few key issues," *Jones v. Barnes*, 463 U.S. 745, 751–52 (1983), and, quoting Justice Jackson, "[l]egal contentions, like the currency, depreciate through over-issue," *id.* at 752.

Work whose cost exceeds its value. For this reason, I considered defective LAB's criterion that requires models to argue that an expert's testimony should be excluded under *Daubert* because the expert's report lists five categories of damage that sum to $14,211,000 and then claims total damages of $14,200,000. The task environment characterizes this as an "arithmetic error" and tries to suggest it is evidence that the expert's work cannot be trusted. But it's just rounding, and while rounding in an expert report is slightly unorthodox, it does not come close to the kind of fundamental reliability defect that a movant would need to show to get a ruling that the expert's opinion is completely inadmissible. The rounding is in movant's favor (the expert has rounded down the amount of damages the defendant might have to pay), and if the precise dollar amount were actually material for some reason, the non-rounded total could simply be clarified in a non-argumentative one-line footnote. It could be fine to train a model to flag this issue briefly in a user message or internal memo, but training the model to raise points like this as a basis for excluding testimony under *Daubert* could weaken the model's judgment.

A similar issue applies to a criterion that requires that the model recommend a motion to dismiss an unjust-enrichment count (LAB item 4) pleaded in the alternative to a contract claim that the rubric considers strong. In a situation where the defendant faces at least one clearly strong claim that is expected to survive a motion to dismiss no matter what, a motion becomes less attractive unless it will somehow reduce the financial exposure/settlement value (by taking higher-value claims off the table) or narrow the scope of discovery (by removing certain factual issues from the litigation). Neither appears to be the case from this environment, so rather than reflexively recommending another motion simply because the rules allow it, a better memo to the client would explain that the benefit of a motion is limited and may not be worth the costs and fees involved (and the tactical downside of spending credibility with the judge on weak arguments that do not advance the client's strategic objectives).

LegalReasonBench presented a similar tradeoff in the land-use matter, requiring that the county's memo recommend both an appeal and a remand to the hearing examiner. A protective appeal is reasonable and probably the default most lawyers would recommend, but in the context of an internal memo, the better answer would be to explain to the client that, while there are some advantages to at least filing a protective appeal, curing the defects on remand may be faster and more effective—and the models are rewarded here for the simpler answer. In the real case, the county did both, but its appellate filing explicitly acknowledged that the remand might end up mooting the appeal (items 3 and 4). While the advice those lawyers gave to their client is not public, the filings suggest they likely at least did discuss the limited value of the appeal here.

4 Implications for evaluation and training

The key result of this audit is that a meaningful minority of the audited criteria reward worse lawyering—for example, by penalizing models that properly rebut an adversary's arguments, or by rewarding models that miscite or misapply governing case law. This has three main implications. First, these two benchmarks, as released, may not reliably separate the models at the top of their leaderboards. The two benchmarks audited here were chosen precisely because they are industry-leading and likely among the best available metrics for long-horizon agentic legal tasks, but criteria that AI auditors flagged as defective accounted for a disproportionate share of the criterion failures in the model runs I did, and the two models failed the defective criteria in my sample at different rates (though not by a statistically significant margin), which at least raises questions about the stability of the ranking. Second, a large share of the specific defects identified are easily fixable. This

audit demonstrated that AI-led auditing techniques can be effective at identifying erroneous criteria, and Harvey has already made changes to LAB intended to address findings from this audit. Third, beyond fixing specific identified defects, this audit illustrates the enormous difficulty of reducing legal professional judgments to binary criteria. Section 4.1 considers what the scores can establish, Section 4.2 the rewards that favor worse lawyering, and Section 4.3 what the audit suggests might be effective methods for improving reward construction in the legal domain.

4.1 Legal benchmark rankings may be unreliable

The audit raises questions about how reliably current legal benchmarks can distinguish between frontier models, and how accurately they can depict capabilities. Table 7 shows that the defect rate found in this audit for litigation tasks is larger than the scoring gaps between the frontier models on these benchmarks (although the two panels measure different populations and a defect rate does not translate cleanly into a score error, as explained in the table notes). In the case of LegalReasonBench, even the conservative lower-bound estimate of the defect rate in litigation tasks (25.1%) is greater than the percentage point difference in scores between the highest- and lowest-ranked models (18.89 percentage points). Harvey's LAB has been used on a much larger set of models, including some older models released in 2025, so the gap between the highest and lowest score on LAB is 81.59 percentage points, much larger than the lower-bound estimated defect rate of 13.4%. But as an illustration, GPT-5.6 Luna (ranked 17th, and a different model from the GPT-6 Luna used in my runs) had a criteria pass rate that was 7.58 percentage points less than the highest-ranked model on Harvey's LAB, and the four highest-ranked models are within 1.93 percentage points of one another, so the difference between the top-scoring models and recently released small models is smaller than the defect rate observed in litigation tasks.

That means that, if even a portion of the defects end up penalizing one model more than another, correcting them could conceivably change the rankings. The audit does not show that any published ranking is wrong, and it does not estimate how the rankings would move if the defects were corrected.

There is at least some indication that the defects could differentially impact different models. In the audit, GPT-6 Luna failed 59.4% of the judge grades on the defective LAB criteria, while Opus failed 34.4%. That difference was not statistically significant ($p = 0.117$, exact paired permutation test over the 16 criteria, averaging the two judges per criterion). LegalReasonBench's forum non conveniens matter illustrates how inaccurate criteria can differentially affect models. Opus scored 53.2% while Luna earned only 1.0% (Appendix A). The models' scores on these two runs were separated by 105 of the 201 available rubric points, 35 of which came from twelve criteria that require citations to the affidavits discussed in Section 3.3: Opus drafted the affidavits and cited them, and Luna did not. (The remaining 70 points came from criteria that were not in my sample or that I did not judge defective, mostly on the first assignment, where the judges graded Opus's memorandum as covering far more of the rubric than Luna's.)

These issues also mean that benchmarks could understate model capabilities somewhat. The AI auditors' flags illustrate the impact. Although I reviewed only a sample of the criteria that AI auditors flagged as defective, across the full population of litigation criteria, criteria that AI auditors flagged as defective or potentially defective were graded as failures under LAB's scoring 24.5% of the time (819 of 3,340 judge grades across the two model runs), compared with 8.9% for criteria that no AI auditor flagged as potentially defective (723 of 8,092). The pattern holds for each model separately (Luna 28.5% versus 12.4%; Opus 20.5% versus 5.5%). The criteria that AI auditors flagged were only 29.2% of the population but accounted for 53.1% of all failed grades. While I

have not validated the assertion of AI auditors that those other criteria are defective, the results suggest that a meaningful share of failures in real-world grades could result from defective criteria, rather than deficiencies in the models being evaluated, although the comparison cannot rule out that the flagged criteria are simply harder ones.

Table 7: Criterion defect estimates and published leaderboard spreads. Defect shares count criteria; score gaps measure differences in model performance against the published rubrics.

	Harvey LAB	LegalReasonBench
A. Defect estimates: audited public litigation criteria		
Criteria in audited scope	2,858	179
Estimated defective share (defects within the AI-flagged frame)	18.7%	29.6%
Two-sided 95% confidence interval (flagged frame)	12.5–23.9%	24.0–32.4%
One-sided 95% lower bound	13.4%	25.1%
B. Published scores: full leaderboards		
Score metric	Criteria pass rate	Weighted score
Models listed	43	9
Highest score	95.48%	50.83%
Lowest score	13.89%	31.95%
Highest-to-lowest gap	81.59 pp	18.89 pp
First-to-fourth gap	1.93 pp	10.39 pp
First-to-tenth gap	4.24 pp	—
GPT-5.6 Luna rank	17 of 43	Not listed
GPT-5.6 Luna score	87.90%	—
Leader's gap to GPT-5.6 Luna	7.58 pp	—

Notes. pp = percentage points. Panel A and Panel B measure different populations, so this comparison is merely illustrative. Panel A reflects the estimated defect rate in litigation tasks, based on my audit of the litigation rubrics: the intervals are exact hypergeometric bounds within the AI-flagged frame, expressed as a share of all audited criteria and conditional on my verdicts, and they do not bound defects among unflagged criteria. The two-sided intervals correspond to 358–683 defective LAB criteria and 43–58 defective LRB criteria. Panel B shows scoring differences between models from their official scores on all tasks (not just litigation tasks), including held-out tasks; the defect shares are not estimates of score error. The LAB scores are taken from the October 7, 2026, snapshot of LAB-AA v1.0 and use each model's highest-scoring configuration (if the same model was tested at different reasoning settings, I use only that model's highest score to determine the ranking); the ranks are my calculation from a third-party recovery of the leaderboard [22]. LRB uses the October 26, 2026, leaderboard, captured November 8, with unrounded scores from its accompanying report; Gemini 4 Argon answered revised versions of 17 matters, so the gaps are descriptive, not matched-version comparisons [19, 20].

That said, there are limits to what this kind of audit can show: the audit was limited to the publicly available tasks, and both Harvey and micro1 use a private held-out set, which was not audited and which could differ in material respects from the audited population. Harvey has told me that outside expert review was more extensive on the held-out set (Section 2.1), so the defect rate there may be lower. Harvey continuously updates its benchmark and has made several updates since this audit (including changes intended to address the audit findings in particular), so this audit does not demonstrate the defect rate, if any, in the revised, operative v1.1 benchmark. And the audit was limited to litigation, which is one of LAB's 24 practice areas; the defect rate in non-litigation tasks could be higher or lower than the rate estimated for litigation tasks, and the effect of litigation defects on an all-practice-area leaderboard is diluted accordingly.

Nevertheless, the audit results suggest that legal benchmarks should be considered one signal and piece of evidence about the quality of models' legal AI work, not a definitive ranking as they are sometimes treated. The audit points to the need to consider a variety of data sources, not just one set of rubrics, given the challenges inherent to reducing legal judgment to rubrics.

4.2 Rubrics sometimes end up rewarding lower-quality lawyering

The audit found several places where the benchmarks unintentionally reward demonstrably worse lawyering, including:

- penalizing models that explained an opponent's argument in order to show why it was wrong (Section 3.2; LRB items 13, 16, and 22);
- rewarding models for misstating propositions of law or fact (Section 3.1; for example, LAB item 23 and LRB items 5, 10, and 21);
- rewarding models for asserting information that cannot be extracted from the environment (Section 3.3; LRB items 14, 15, and 17 through 20, and the nonexistent conflicts in LAB item 20 and LRB item 9).

These signals show that at least some criteria can be counterproductive. Since versions of Harvey's LAB have been used for training, and micro1 sells training environments of the same construction, this raises the question whether these environments could reinforce undesirable behaviors in some cases [11, 12, 19]. These defects are not necessarily noise that simply averages out across a large number of tasks; in some cases, the defects appear to constitute a consistent signal that creates undesirable incentives around a specific behavior (LegalReasonBench's penalties for mentioning the opponent's argument are one clear example of an inaccurate signal that several different task environments apply).

To be clear, the point should not be taken too far: the vast majority of criteria are likely to be sound, and thus are likely providing valuable training signal. (AI auditors did not raise any issue with the vast majority of criteria.) Models could still learn valuable skills from legal environments, even if some individual criteria have defects, if the positive training signal from the correct criteria outweighs the negative signal from defective ones. Litigation is also an especially challenging subdomain, because the setting is inherently adversarial and models need to approach evidence with due skepticism, distinguishing between reliable and unreliable information from the surrounding context.

Nevertheless, the audit's results suggest that it could be valuable to try to isolate the impact that legal-domain-specific training environments have on model quality, using benchmarks or techniques that have a meaningfully different construction than the training environment. Models are clearly getting better at legal tasks over time, but on the current publicly available evidence, it is impossible to determine how much of that gain comes from transfer from other task environments: teaching the models to do better work in easier-to-verify domains like math and coding could yield skills that improve the models' ability to handle legal tasks as well.

Exacerbating the difficulty of isolating the effect of legal-domain training is that many providers use one set of benchmark tasks as training signal, and then evaluate the model's improvement based largely on its performance on a held-out set from the same benchmark family. Examples are collected in Table 8. Evaluating on held-out tasks from the training benchmark is standard practice in machine learning, and Harvey appears in most rows only because it has published more legal post-training results than anyone else. This kind of evidence shows that legal-domain

Table 8: Notable reports of post-training gains on legal tasks in 2026, and whether any evidence outside the benchmark used for training was provided.

Report	Training signal	Gain reported on held-out tasks from the same benchmark	Other evidence of improvement presented	Source
Harvey and Baseten Research Qwen3.5-9B, May 2026	Reward "exclusively from the LAB rubric"	Criteria pass rate 42.5% to 63.0% on held-out LAB tasks	None	[13]
Harvey and Baseten Research Qwen3.5-27B, May 2026	Fine-tuned on rubric-passing LAB rollouts	Brought "into the closed-source frontier band" on held-out LAB tasks	None	[13]
Harvey and Trajectory Nemotron 3 Super, May 2026	Post-trained "on our Legal Agent Benchmark (LAB)"	All-pass rate matched GPT-5.5 on held-out LAB tasks	None	[15]
Harvey and Trajectory Nemotron 3 Ultra, June 2026	Same recipe as Nemotron 3 Super	All-pass rate 0% to 5.8% on held-out LAB tasks	None	[16]
Applied Compute and Harvey GLM-5.1, June 2026	LAB rubric grading as the RL reward	Criteria pass rate 0.853 to 0.913 on held-out LAB tasks	None	[12]
Harvey Tenet Kimi K3, August 2026	RL on environments that "share the structure of" LAB tasks	All-pass rate up 9 percentage points on the LAB held-out set	Gains on APEX Agents and Redline Bench, both rubric-graded; no regression on LegalBench, CUAD, MAUD, APEX v1	[11]
Mercor GLM-4.7, August 2026	RL on 1,800 APEX Agents tasks	Corporate-law score 4.7% to 26.6% on APEX Agents	Generalization to GDPval and APEX v1 asserted, without figures	[18]
Harvey and Baseten Research Qwen3.5-122B-A10B, September 2026	"Judged rubric criteria pass rate as the reward" on LAB Diligence	Criteria pass rate 29.9% to 63.0% on held-out LAB Diligence data rooms	None	[14]
Adaption Labs Qwen3.6-27B, September 2026	Trained on "successful end-to-end trajectories" from LAB tasks	Criteria pass rate 85.92% to 88.03% on held-out LAB tasks	None	[17]

Notes. RL = reinforcement learning. "None" means the published report evaluated the post-training gains exclusively using a held-out set of tasks from the same benchmark family used for training.

training changes the models' outputs in ways that generalize beyond the specific tasks that they were trained on. But this does not address the underlying problem of construct validity. The rubrics ultimately are only a proxy for gains in legal capabilities. Evaluating models only on a held-out set from the same family of tasks the models were trained on does not establish that there were real gains in underlying legal ability; just as a teacher might improve students' scores a few points by teaching them "tricks" to do well on a standardized test, training on a particular task family could make models better at scoring well on tasks from that specific family, without meaningfully improving the underlying legal capabilities. Several authors have previously noted the danger of relying on essentially the same rubrics as proof that a given training signal is useful [33–37]. And the phenomenon has been observed in legal RL specifically; LEGIT, which scores legal reasoning against issue trees built from court judgments, found that reinforcement learning raised a model's overall score on the benchmark's test set while reducing the share of legal issues the model addressed. The rubric penalized a wrong answer more heavily than an omission, so the model learned to address straightforward issues and skip more difficult ones to avoid losing points [32]. The benchmark's own score went up even though the output arguably was worse in that the model dodged difficult questions.

We ideally would want to see extrinsic evidence, from lawyer preference scores or perhaps a range of legal benchmarks that are constructed in a completely different way and that involve meaningfully different challenges, to be confident that the training improved underlying legal capabilities and not merely a proxy. In the legal domain, such evidence is currently scarce because there are so few rigorous benchmarks that challenge the models at the relevant frontier.

There is some reason to be optimistic that legal training environments, despite the challenges, are in fact improving model capabilities. Harvey has reported that its legal post-training work has resulted in capability gains that are observable in separately developed benchmarks (although the benchmarks were created in a similar way to LAB, with expert-annotated rubrics) [11]. Ultimately, it may be wise to look at a variety of signals. For example, Thomson Reuters reported practitioner preference scores in addition to results from rubric-based benchmarks as part of its model evaluation [38]. Evidence that post-training actually improved outputs in a way that lawyers consider improvements could be much stronger evidence that validates that a given set of rubrics is a good proxy for underlying capability gains.

4.3 Possible paths for improving legal RL environments

Although this audit found that some of the leading benchmarks have defects, rubric-based tasks and environments remain useful for several reasons. Rubrics are easy to audit and understand. This audit was possible precisely because Harvey and micro1 made a selection of their rubrics available for public inspection, and invited practitioners to audit them and point out ways they can be improved. The same approach would not work with preference comparison techniques because it may be difficult to understand why a given reviewer preferred one response over another (and such approaches have well-known biases, such as a tendency to sycophancy [44]). Rubrics can also be applied efficiently at scale: once created, rubrics can be run autonomously, with agents completing tasks and evaluating the results against rubrics, while approaches that rely on human intervention introduce bottlenecks and delays. And rubrics are the natural tool for open-ended work product, where no answer key exists; tasks with verifiable answers, such as the classification and extraction tasks in LegalBench, do not need them [43]. Accordingly, while rubrics have drawbacks and can encode errors, they clearly are an important component of benchmarking and training in the legal domain.

This audit suggests that there may be approaches that could improve rubric design and reward signals in evaluation and training data for legal work.

4.3.1 Quality control

Audit environments before training on them. Using frontier models to audit tasks and environments before using them as training signal could significantly improve quality. In this audit, instances where two distinct models (GPT-6 Sol and Claude Opus 5.5) both found a defect in a criterion were overwhelmingly real defects: 9 of 10 such criteria in the LAB sample and 12 of 12 in LegalReasonBench's. An automated citation check of the environment documents is cheaper still; the LAB check established at least 303 real citation errors and flagged 76 citations to authorities it could not locate, which I did not verify individually (Appendix B). Environment buyers may wish to have models adversarially inspect the environment and provided rubrics for defects, and to ask sellers to improve environments if the defect rate is unusually high.

Review rubrics against real model runs. Rubric authors should review how the rubrics apply in practice and confirm how well they are being applied by judges. For example, the recurring issue in LegalReasonBench where models were punished for rebutting an opponent's argument (Section 3.2) likely could have been caught by reviewing the scores from a real frontier model run, as I did in this audit. While it appears to be the case that, in many instances, one company (such as a data vendor) will create rubrics and a different one (a lab) actually uses them for training, the quality of these training signals could be improved significantly if one of the organizations involved in the process takes on the responsibility of validating rubrics end to end to ensure that their practical application provides the intended training signal. Otherwise, misspecifications can go undetected.

4.3.2 Rubric design

Rubrics can be highly effective at testing bright-line rules, instead of open-ended standards. Law professors often speak of the difference between bright-line rules and open-ended standards. The canonical example could be the difference between a posted speed limit of 55 miles per hour, versus a law that prohibits driving at "excessive speed" [39]. Rubrics in principle should be an excellent way of encoding bright-line rules, and there may be even more opportunities to develop rubrics that specifically exhaust the universe of bright-line rules that apply to a problem. For example, the motion to dismiss task environments I reviewed did not contain any overarching criteria that penalize models for citing evidence and facts outside the four corners of the complaint (or documents incorporated by reference therein or subject to judicial notice). But unless a model explicitly argues that the motion should be framed as a motion for summary judgment under Rule 12(d), it is almost always wrong to cite evidence that is not referenced in the complaint or subject to judicial notice. Rubrics that try to enforce all of the bright-line rules could improve the quality of legal environments. Similarly, rubrics can be used to penalize miscitations or misrepresentations of a cited source, as that is always unacceptable in effective advocacy, no matter how fluent or persuasive the document otherwise seems.

Rewarding good judgment is harder, but rubrics may be an effective way to enforce the permissible range of professional judgment. Many questions fall within a range of acceptable professional judgment, and reasonable lawyers can disagree about which is best; selecting among them is the heart of advocacy [41, § 12] and the hard part of evaluating open-ended legal answers

[30]. The misspecified requirements and the errors of professional judgment (Sections 3.2 and 3.4), and most of the unsafe-drafting rewards (Section 3.3), are standard-like: the rubric chose one answer where the law and good practice permit a range. Ideally, rubrics on questions of judgment or standards would be designed to track something like the abuse of discretion standard the appellate courts apply when reviewing a decision committed to the trial court's discretion: enforce hard constraints first, but allow any permissible approach within the accepted boundaries of discretion, failing those that exceed those boundaries (an "abuse of discretion," in judicial parlance). Many of the criteria I reviewed suggested that rubric authors were trying to capture a few seemingly obvious aspects of what a good legal output might entail in response to a given task, but often these involved something like a "local maximum" that occurred to the rubric author, not necessarily the optimal output overall. It may be more effective to conceive of strict rubrics like the ones Harvey and micro1 wrote as enforcing the outer limits on legal output—and using techniques like comparing briefs to decide which outputs are better amongst a category that complies with the basic requirements. Indeed, JudgmentBench reported that on 30 legal tasks, pairwise comparisons recovered a constructed quality ordering substantially better than rubric scores [31], and preference optimization can train on such comparisons [42].

General rubrics for all litigation documents. It may be desirable to develop a taxonomy of general rubrics that are not specific to a particular task environment, in addition to the rubrics that are more specific to a particular environment. One noteworthy aspect of LAB and LegalReasonBench was that the rubrics were almost all customized to a specific task and environment, so the rubrics tended to test case-specific knowledge. For example, in a random sample of 120 LAB litigation criteria that I had an AI agent classify, 104 were task-specific, and 178 of LegalReasonBench's 179 litigation criteria were task-specific [7]. No criterion in either benchmark mentions hallucination, and the one LAB criterion that addresses citation authenticity asks for "at least three case law authorities (whether real or plausibly cited)" (draft-motion-for-preliminary-injunction, C-051). (LAB v1.1, which I did not audit, adds hallucination checks.) Although the approach of focusing on task-specific criteria certainly makes sense, it was noteworthy to me that few of the rubrics enforce general principles of writing and argumentation that apply to a variety of different kinds of legal tasks. In principle, there could be value in carefully designing some overarching criteria that apply to, say, all litigation documents (e.g., a rubric that rewards citations for each contested factual or legal assertion made, and punishes any hallucinations or miscitations) and rubrics that apply to all briefs (e.g., a reward for sequencing the arguments roughly in order of their persuasive power and importance), in addition to task- or environment-specific rubrics (requirements grounded in the facts of a specific assignment).

4.3.3 Environment design

Adapt real cases with care. Environments grounded in real cases have advantages over purely AI-generated synthetic environments, but they must be adapted with great care, and AI can plainly contribute at multiple phases of the environment construction process. My impression from the audit was that the LRB environments, which were often grounded in real-world cases, were more realistic and more likely to train models on what a real litigation would look like. The audit was not designed to compare the two benchmarks' defect rates, however, and the real-case grounding introduced a failure mode of its own: the rubrics sometimes supplied the answer that would be accurate in the real case, but the task environment had been modified in a way that changed the answer (e.g., in one real-world case, there was a fact dispute over who opened fire first, and a criterion wanted the model to say that, but the task environment had been modified from the real-world case to specify who

fired first). When modifying task environments that have a real-world basis, environment authors need to take great care to consider how the modifications could change the task or its answers. Even simply "updating" the dates in a task by a few months could completely change the outcome because a new statute may have gone into effect or an old one been repealed with an effective date between the old and the new dates. It may be desirable to modify the task environment and the answers from the real-world case, so that a model cannot simply score well by repeating the answer to a case it learned in pre-training. But careful auditing is then required to ensure that the resulting task environment is accurate.

Other task formats. The environments I reviewed focused largely on drafting tasks, but there are a range of other tasks that could be useful. For example, it might be useful to give a model a brief where an important argument was ablated and have it identify and add the missing piece—or give it a longer version of the brief and test its ability to cut without losing important details from the original. In the context of training as opposed to evaluation, there may also be ways to combine reward signals from different sources. For example, LegalForecastBench asks models to predict a judicial disposition from the record the judge received [29]. With a fixed record and opening brief, models might be tested on the ability to generate candidate response briefs. After disqualifying briefs with fabrications or ethical violations (or concessions that compromise the client's interests in other matters), one could assign reward to the briefs that seem likely to maximize the chance of a favorable disposition, as scored by a well-calibrated prediction model. Such an approach would need to be tested to ascertain whether the forecasting model responds to real improvements in advocacy and to ensure that the training does not merely exploit the forecaster [33].

These suggestions are all indicative. The audit does not establish that they will work better than existing approaches, but their common purpose is to reward models for making legal work better, rather than for reproducing the arguments a rubric author expected to see.

5 Limitations and conclusion

There are several limits of this audit:

- I am the sole reviewer, as the author. Future audits like this could be made more rigorous by having multiple lawyers participate in the audit.
- This audit was limited to the publicly released litigation tasks for LAB and LegalReasonBench. The audit does not cover the private held-out datasets, or practice areas other than litigation.
- The samples I reviewed were drawn only from criteria that AI auditors flagged, so the confidence bounds do not establish the defect rate, if any, in the universe of criteria that the AI auditors did not identify as potentially problematic.
- The sampling method used for LAB and LegalReasonBench was slightly different. The two audits are meant to be independent, not a controlled comparison of the two.
- The confidence intervals cover only sampling error, not errors in my verdicts.
- Several sampled criteria share one underlying mistake or issue, so the numerical results could shift materially if the handling of a given issue were changed.
- My review was conducted with the AI auditors' findings, and the results of test model runs, in front of me. Although I endeavored to analyze the rubrics independently and often reached

my own conclusions, I could have been influenced by the AI models' arguments.

- The test benchmark runs used were a single run per model, with different reasoning settings, so these were intended as a diagnostic, not a comparison.
- Harvey commented on earlier drafts of this work. I believe its comments improved the paper, and I did not change any verdicts after sharing a draft with Harvey for feedback.

Despite these limits, the audit helps to illustrate some of the challenges of reducing legal judgment to rubrics as model capabilities have increased to the point where models are increasingly effective at long-horizon agentic tasks. These findings may help to improve not just the specific benchmarks but the process of legal environment task construction.

Acknowledgments

I extend my thanks to Julio Pereyra, who oversees the applied legal research team at Harvey, and who provided several substantive insights and reactions to earlier drafts that were incorporated herein. Mr. Pereyra and Niko Grupen, head of applied research at Harvey, also worked to address audit findings so that this work could contribute to improving future versions of Harvey's Legal Agent Benchmark. I also thank Prof. Dr. Daniel Katz, professor of law at Illinois Tech's Chicago-Kent College of Law, and Andreas Kirsch, an AI researcher, for their thoughtful comments on an earlier draft.

References

[1] GPT-6 Sol audit workers, Claude Opus 5.5 audit workers, Coordinating agents; individual provenance not fully recorded. Harvey LAB litigation and dispute resolution audit: GPT-6 Sol and Claude Opus 5.5. *Local research artifact, pinned September 29, 2026*, 2026. https://github.com/johnhughes3/LegalForecastBench/tree/58371a0e13494e591710e2eae3da8e428eebc359/docs/harvey-lab-audit/litigation-dispute-resolution.

[2] John Hughes. Human review of AI-flagged Harvey LAB litigation criteria: protocol, sample, and worksheet. *Research artifact*, 2026. https://github.com/johnhughes3/LegalForecastBench/tree/main/docs/harvey-lab-audit/litigation-dispute-resolution/human-review.

[3] Claude Sonnet 5.5 and Claude Opus 5.5 citation-check agents, coordinated by the author. Citation check of the Harvey LAB litigation and dispute resolution task environments. *AI-generated research artifact; a random sample attorney-reviewed*, 2026. https://github.com/johnhughes3/LegalForecastBench/tree/main/docs/harvey-lab-audit/litigation-dispute-resolution/citation-check.

[4] Harvey AI task authors; upstream authorship not individually listed. Pinned LAB task.json for analyze-counterparty-motion-to-dismiss. *Primary online source*, 2026. https://github.com/harveyai/harvey-labs/blob/1dd81403b2fbb60596f7aea3fcecafad7bf73143/tasks/litigation-dispute-resolution/analyze-counterparty-motion-to-dismiss/task.json.

[5] GPT-6 Sol and Claude Opus 5.5 audit workers, coordinated by the author. LegalReasonBench litigation rubric audit: AI screening, native model runs, and citation census, at dataset revision 1bfa904c. *Research artifact*, 2026. https://github.com/johnhughes3/LegalForecastBench/tree/main/docs/legalreasonbench-audit/litigation.

[6] John Hughes. Human review of AI-flagged LegalReasonBench litigation criteria: protocol, sample, worksheet, and stress tests. *Research artifact*, 2026. https://github.com/johnhughes3/LegalForecastBench/tree/main/docs/legalreasonbench-audit/litigation/human-review.

[7] AI classification agent, coordinated by the author. Are the LAB and LegalReasonBench litigation rubrics task-specific? Classification of a seeded random sample of 120 LAB criteria and all 179 LegalReasonBench criteria. *AI-generated research artifact*, November 8, 2026. https://github.com/johnhughes3/LegalForecastBench/blob/main/docs/papers/lab-audit/analysis/rubric-generality-check.md.

[8] Gabe Pereyra. How Harvey Built a Research Lab on a Budget. *Sequoia Capital, Own Your Intelligence*, August 11, 2026. Slide at 4:11: "Scale with expert data providers." https://www.youtube.com/watch?v=MGouk8W51v0&t=251s.

[9] Niko Grupen, Gabe Pereyra, Julio Pereyra. Introducing Harvey's Legal Agent Benchmark. *Harvey*, 2026. https://www.harvey.ai/blog/introducing-harveys-legal-agent-benchmark.

[10] Julio Pereyra. "Great write up by Rebecca Fordon on the data underlying LAB…" *LinkedIn post*, May 8, 2026, responding to Rebecca Fordon, "Reverse-Engineering How Harvey Built Legal Agent Benchmark," *LegalQuants Insights*, May 8, 2026. https://www.linkedin.com/posts/julio-pereyra-411738147_reverse-engineering-how-harvey-built-legal-activity-7458656036388880384-taQ2 (accessed November 9, 2026).

[11] Calvin Qi, Vasudha Rengarajan, Julio Pereyra, Niko Grupen, Gabe Pereyra. Harvey Tenet Research Preview. *Harvey*, August 20, 2026. https://www.harvey.ai/blog/post-training-update-harvey-tenet.

[12] Applied Compute. Training a State-of-the-Art Legal Agent with Harvey. *Case study*, June 22, 2026. https://www.appliedcompute.com/case-studies/harvey.

[13] Niko Grupen, Gabe Pereyra, Spencer Poff, Julio Pereyra, Charlie O'Neill, Mudith Jayasekara, Matthew Blau, Aaron Ellis-Bloor. Post-Training Open Legal Agents With Baseten Research. *Harvey*, May 27, 2026. https://www.harvey.ai/blog/post-training-open-legal-agents-with-baseten-research.

[14] Niko Grupen, Julio Pereyra, Gabe Pereyra, Charlie O'Neill, Mudith Jayasekara, Aaron Ellis-Bloor, Jonathon Liu, Matthew Blau. Post-Training RLM Agents for End-to-End M&A Diligence. *Harvey*, September 8, 2026. https://www.harvey.ai/blog/post-training-rlm-agents-for-m-and-a-diligence.

[15] Jerry Chan. The Sovereign Path for Continual Learning: Early Results on Harvey LAB with NVIDIA Nemotron. *Trajectory field notes*, May 29, 2026. https://www.trajectory.ai/field-notes/the-sovereign-path-for-continual-agent-learning-early-results-on-harvey-lab-with-nvidia-nemotron.

[16] Jerry Chan. Post-Training NVIDIA Nemotron 3 Ultra on Harvey LAB in Under 24 Hours. *Trajectory field notes*, June 9, 2026. https://www.trajectory.ai/field-notes/harvey-nemotron-3-ultra.

[17] Dhruv Naik, Sudip Roy. A Better Harness Can Unlock Smaller Models. *Adaption Labs*, September 9, 2026. https://adaptionlabs.ai/blog/a-better-harness-can-unlock-smaller-models.

[18] Brendan Foody. How RL Environments Are Built, and Why They're Your AI Moat. *Sequoia Capital, Own Your Intelligence*, August 2026. https://sequoiacap.com/podcast/how-rl-environments-are-built-and-why-they-re-your-ai-moat.

[19] micro1. LegalReasonBench: methodology and results. *micro1*, 2026. https://www.micro1.ai/benchmark/legal-reason-bench.

[20] micro1. LegalReasonBench: published leaderboard and paired comparisons. *Leaderboard labeled October 26, 2026; retrieved November 8, 2026.* Preserved numerical extract of the nine-model leaderboard and accompanying eight-model report, with original source links: https://github.com/johnhughes3/LegalForecastBench/tree/8c916227353e46248b4d03d528abec3dbf82bdb5/site/public/data/lab-audit/legalreasonbench-2026-11-08.

[21] Artificial Analysis. Harvey LAB-AA evaluation results (v1.1). *Leaderboard, accessed October 26, 2026.* https://artificialanalysis.ai/evaluations/harvey-lab-aa.

[22] Artificial Analysis. Harvey LAB-AA v1.0 criteria pass rates, October 7, 2026 snapshot. *Historical results recovered from Village LLM Atlas on November 8, 2026.* The original archived AA page was unavailable during retrieval. Preserved numerical extract, source links, and ranking calculation: https://github.com/johnhughes3/LegalForecastBench/tree/00a6c1aa76504d73b67ab0df27146751f2aace8e/site/public/data/lab-audit/lab-aa-v1-2026-10-07.

[23] xAI (SpaceXSI). Introducing Grok 4.7. *Release announcement*, September 21, 2026. https://x.ai/news/grok-4-7.

[24] Google DeepMind. Gemini 4 Argon: our next era of frontier intelligence. *Blog post and model evaluation methodology*, September 30, 2026. https://blog.google/innovation-and-ai/models-and-research/gemini-models/gemini-4-argon/.

[25] Anthropic. Claude Opus 5.5 System Card, § 8.14.2 (Legal Agent Benchmark). *System card*, September 22, 2026. https://anthropic.com/claude-opus-5-5-system-card.

[26] Meta. Muse Wonder: A Frontier Model for Personal and Professional AI Agents. *Release post and model card*, October 12, 2026. https://research.meta.ai/blog/muse-wonder.

[27] Mistral AI. Introducing Mistral Large 4. *Release announcement*, October 6, 2026. https://mistral.ai/news/mistral-large-4/.

[28] Damien Charlotin. AI Hallucination Cases Database. *Database of court and tribunal decisions, last updated November 5, 2026*. https://www.damiencharlotin.com/hallucinations/ (accessed November 9, 2026).

[29] John J. Hughes, III. Law as a Verifiable Domain: Using Real-World Judicial Outcomes to Evaluate AI Legal Reasoning. *Working paper*, 2026. https://www.legalforecastbench.org/paper/.

[30] Sebastian Nagl. SoK: Grading Open-Text Legal Answers, from Law School Exams to Trustworthy Rubric Judges. *arXiv:2610.19992*, 2026. https://arxiv.org/abs/2610.19992.

[31] Russell Yang, Ruishi Chen, Pierce Kelaita, Riya Ranjan, Sibo Ma, Charles Dickens, Matthew Guillod, Megan Ma, Julian Nyarko. JudgmentBench: Comparing Rubric and Preference Evaluation for Quality Assessment. *arXiv:2605.25240*, 2026. https://arxiv.org/abs/2605.25240.

[32] Jinu Lee, Kyoung-Woon On, Simeng Han, Arman Cohan, Julia Hockenmaier. Evaluating Legal Reasoning Traces with Legal Issue Tree Rubrics. *arXiv:2512.01020*, 2025. https://arxiv.org/abs/2512.01020.

[33] Leo Gao, John Schulman, Jacob Hilton. Scaling Laws for Reward Model Overoptimization. *Proceedings of the 40th International Conference on Machine Learning*, PMLR 202:10835–10866, 2023. https://proceedings.mlr.press/v202/gao23h.html.

[34] Robert Geirhos, Jörn-Henrik Jacobsen, Claudio Michaelis, Richard Zemel, Wieland Brendel, Matthias Bethge, Felix A. Wichmann. Shortcut Learning in Deep Neural Networks. *Nature Machine Intelligence* 2:665–673, 2020. https://arxiv.org/abs/2004.07780.

[35] Joar Skalse, Nikolaus H. R. Howe, Dmitrii Krasheninnikov, David Krueger. Defining and Characterizing Reward Hacking. In *Advances in Neural Information Processing Systems 35 (NeurIPS 2022)*. https://arxiv.org/abs/2209.13085.

[36] Anas Mahmoud, MohammadHossein Rezaei, Zihao Wang, Anisha Gunjal, Bing Liu, Yunzhong He. Reward Hacking in Rubric-Based Reinforcement Learning. *arXiv:2605.12474*, May 12, 2026. https://arxiv.org/abs/2605.12474.

[37] Minglai Yang, Xinyu Guo, Zhengliang Shi, Utkarsh Tyagi, Mian Zhang, Razvan Dumitru, Sunjie Hou, Yunzhong He, Daniel Yue Zhang, Dan Roth, Ying Liu. Rubric Dropout: A Simple Way to Mitigate Reward Hacking in Rubric-Based RL. *arXiv:2608.11669*, August 12, 2026 (revised October 21, 2026). https://arxiv.org/abs/2608.11669.

[38] Shengzhuang Chen, Jerrod Parker, Yejin Bang, et al. Thomson: Continual Learning of Frontier Models for SovereignAI. *arXiv:2608.27147*, 2026. https://arxiv.org/abs/2608.27147.

[39] Louis Kaplow. Rules versus Standards: An Economic Analysis. *Duke Law Journal* 42:557–629, 1992.

[40] Jian Xie, Kai Zhang, Jiangjie Chen, Renze Lou, Yu Su. Adaptive Chameleon or Stubborn Sloth: Revealing the Behavior of Large Language Models in Knowledge Conflicts. *International Conference on Learning Representations*, 2024. https://arxiv.org/abs/2305.13300.

[41] Antonin Scalia, Bryan A. Garner. Making Your Case: The Art of Persuading Judges. *Thomson/West*, 2008.

[42] Rafael Rafailov, Archit Sharma, Eric Mitchell, Stefano Ermon, Christopher D. Manning, Chelsea Finn. Direct Preference Optimization: Your Language Model is Secretly a Reward Model. *arXiv:2305.18290*, 2023. https://arxiv.org/abs/2305.18290.

[43] Neel Guha, Julian Nyarko, Daniel E. Ho, Christopher Ré, Adam Chilton, et al. LegalBench: A Collaboratively Built Benchmark for Measuring Legal Reasoning in Large Language Models. In *Advances in Neural Information Processing Systems 36 (NeurIPS 2023), Datasets and Benchmarks Track*, 2023. https://arxiv.org/abs/2308.11462.

[44] Mrinank Sharma, Meg Tong, Tomasz Korbak, David Duvenaud, Amanda Askell, et al. Towards Understanding Sycophancy in Language Models. In *International Conference on Learning Representations (ICLR 2024)*, 2024. https://arxiv.org/abs/2310.13548.

A Native execution and grading on both benchmarks

To understand how the criteria were applied in real benchmark runs, I ran both benchmarks with GPT-6 Luna (xhigh reasoning) and Claude Opus 5.5 (low reasoning), using each benchmark's own harness and judges, and considered the results in assessing criteria defects. The examples in Appendices D and E include these outputs, so I could consider how actual model outputs were being graded in evaluating the rubrics. The two models were run with different reasoning settings, so the difference in their scores is not a controlled comparison of the quality of those models; instead, it was merely intended as a diagnostic to help explore where differences in grading might arise. All deliverables, raw judge files, and every criterion's verdicts with the judges' reasoning are published unedited in the audit directories [1, 5].

Harvey LAB

LAB's litigation tasks were run at Harvey commit `1dd81403` with the original instructions, documents, and rubrics. LAB's designated AI judges, Claude Sonnet 4.6 and GPT-5.5, graded every output (Table 9).

Solver	Effort	Tasks	Criteria pass rate	All-pass cells	All-pass by both judges
GPT-6 Luna	xhigh	52	83.80%	2 of 104 (1.92%)	1 of 52 tasks
Claude Opus 5.5	low	52	90.81%	3 of 104 (2.88%)	1 of 52 tasks

Table 9: Native LAB AI-judge grades on the 52 public litigation tasks. For each task, each judge's fraction of criteria passed is averaged across the two judges and then across tasks with equal weights. An all-pass cell is one judge passing every criterion on one task; there are 104 such cells per solver.

Under LAB's all-pass rule, a task scores 1 only if every criterion is judged a pass; a single failed criterion on a task results in a score of zero on the entire task.

LegalReasonBench

The 11 litigation assignments were run using dataset revision `1bfa904c`. Assignments within a given matter were run in a persistent sandbox (consistent with the intended environment design) so that later assignments could read the deliverables produced at earlier steps. The benchmark's three-judge panel (GPT-5.6 Terra, Claude Sonnet 5, and Gemini 3.8 Flash) graded by majority vote (Table 10) [5].

	GPT-6 Luna (xhigh)		Claude Opus 5.5 (low)	
Matter	Signed points	Score	Signed points	Score
Land-use certiorari	50/320	15.6%	20/320	6.3%
Police discipline	160/332	48.2%	224/332	67.5%
Jury instructions	25/310	8.1%	60/310	19.4%
Forum non conveniens	2/201	1.0%	107/201	53.2%

Table 10: LegalReasonBench scores are calculated as signed points earned across assignments divided by total positive points. As an illustration, on the second forum non conveniens assignment, with 73 positive points available, the −30 penalty for explaining why the 60.5% statistic was misleading (Section 3.2) reduced Opus's score from 63 points (86.3%) to 33 (45.2%) and Luna's from 16 (21.9%) to −14.

B Citation checks of the task environments

Because the rubric audit turned up instances where the rubrics contained hallucinated case law (suggesting that whoever was responsible for verification may have relied on inaccurate AI-generated responses to prepare the answer key), I asked AI agents to conduct a separate citation check of the litigation and dispute resolution tasks to assess the frequency of AI hallucinations in the task environments or rubrics. A citation check is a standard pre-filing procedure where lawyers check every authority cited in a brief, regulatory submission, or other important document to ensure that the authority is cited accurately, contains any quotation attributed to it, and supports the legal proposition for which it is cited.

An initial AI-assisted check of every case citation in the 19 sampled LAB environments against CourtListener, with spot checks, found problems in 45 of 102 citations to purported real authority: 32 real cases cited for holdings, quotations, or facts they do not contain, 7 cases that could not be found, 5 garbled citations to real cases, and 1 citation pointing to a different case [2]. A broader occurrence-based check then covered all 52 tasks. After excluding errors in case law that were intentionally planted and then tested in rubrics (e.g., an error of law in an opponent's submission that the model was supposed to catch), the AI identified 436 errors that were apparently not intended to be part of the task environment [3] (Table 11).

Finding	All citations	Case citations
Verified	2,754	306
Verified, minor defect (pin cite, loose paraphrase, wrong year)	346	54
Intentional error tested by the rubric (excluded)	31	6
Unintentional errors		
Cited for a proposition it does not support	246	105
Not found by the check (unverified)	76	60
Garbled citation to a real authority	58	10
Quotation not found in the authority	40	31
Reporter citation belongs to a different case	16	12
Total unintentional errors	436	218
Source text not accessible	51	6
Total checked	3,618	590

Table 11: AI citation check of the 52 LAB litigation task environments. This table shows the AI findings; the attorney review below estimates how many are real.

I reviewed a simple random sample of the 436 errors personally. Before I reviewed any sampled finding, I drew 20 citation findings and pre-committed to a stopping rule that if the first ten were all confirmed errors, I would stop; otherwise I'd review ten more. I reviewed each citation finding and agreed that all ten involved errors a careful lawyer would correct before filing. The AI had correctly classified seven as cited for propositions they do not support and three as real authorities cited by the wrong number or subsection. Three of the ten involved miscitations to the Eleventh Circuit's pattern civil jury instructions. With ten of ten confirmed, a conservative one-sided 95% lower confidence bound is 303 of the 436 flagged errors (69.5%); the calculation accounts for both possible stopping points under the pre-committed rule by splitting the 5% error budget between them (a bound that ignored the stopping rule would be 325). I did not attempt to review or identify errors the AI may have missed among citations it marked as verified.

A separate AI audit reviewed for cite check issues. Of 2,312 authorities cited, it verified 1,957, and found 118 with minor defects, 73 problem occurrences (47 mischaracterized, 5 misquoted, and 21 wrong citations to real cases), and 164 that could not be verified (mostly historical local compilations and agency decisions that could not be retrieved through web search). There were 70 references to fictional authorities provided in the environment. In cases where an AI agent identified a problem with a citation, a second AI agent was given the same citation and asked to validate the issue [5]. The observed problem rate is 73 of 2,312 (3.2%). I did not sample these for attorney review, and the two rates should not be interpreted as a comparison between the benchmarks.

C A synthetic task environment: the motion-to-dismiss case study

In addition to auditing the accuracy of all 2,858 criteria, I did a deep-dive review of one substantive LAB litigation task. This review highlighted a different but equally important issue: the use of synthetic or AI-generated data may result in task environments that are unrealistic in a variety of subtle ways that may not be obvious to nonexperts but that are striking and clear to practicing attorneys. The task `analyze-counterparty-motion-to-dismiss` asks AI agents to prepare a list of issues for the opposition brief to a motion to dismiss, on a fictitious software service dispute.

I observed several ways in which the synthetically generated task environment seemed unrealistic and could potentially train models on a distorted view of what genuine practice looks like [4].

Planted issues. A number of the legal disputes in the task environment are obviously "planted" in a way that feels artificial and ultimately presents issues in a way that is meaningfully different from how real-world legal disputes manifest. For example, the contract the parties are suing over has a forum selection clause that requires any lawsuit to be filed in Travis County, Texas. The plaintiff literally quotes this entire provision in the complaint but nevertheless files suit in the Northern District of Georgia and offers no explanation for why they believe the forum selection clause they quoted does not apply or otherwise permits the suit in Georgia. It of course happens regularly that a lazy or incompetent litigant files suit in the wrong forum; a competent lawyer potentially could decide to ignore the clause and hope that their adversary overlooks the issue. But in both of those scenarios, the forum selection clause would not be explicitly flagged in the complaint as it is here; a real version of this legal task would require the lawyer to dig through a number of documents and realize that buried in one of them is a forum selection clause that allows the lawyer to have the case moved to Texas.

This is one example, but a significant number of the issues in this environment were similarly planted in a way that's artificial and unlike how disputes around these issues arise in real cases. Environments like this can distort an evaluation of a model's true capabilities, since simply looking for the obviously planted issues will satisfy the benchmark without testing whether the model would be able to identify this issue from more realistic source material akin to what it will see in real cases.

Environment documents violate court rules. Many of the documents in this environment do not resemble real lawyer work product. The environment provides counterparty litigation documents in docx format, and the native files from the git repository often cannot be opened, although agents and various CLI tools can recover the document text. In real litigation, lawyers rarely share their docx files with their opponent; the opponent's memorandum of law is ordinarily filed with the court as a PDF. The motion papers also lack a proper caption, include a stale table of contents, contain markdown formatting artifacts that render in Word as junk characters even in files that can be opened normally, lack citations for propositions that generally require citations, and are single spaced rather than double spaced in violation of the applicable court rules.

The motion that the agent is being asked to respond to is so far below professional norms that you would want any analysis of the motion to include consideration of whether the judge would strike the motion or require it to be refiled in compliance with the rules. That is not included in the rubric for this task, presumably because the quality control issues appear to be unintentional, not part of the test for the agent. But it is a genuine gap: if a document like this were filed in real litigation, the quality of the work product could itself become a substantive issue in the litigation and thus would need to be flagged given the goal motivating the all-pass rule, which was to ensure that the

AI does not overlook any important issues [9, section "Evaluating a task"].

There are valid reasons to prefer rubric-based reinforcement learning approaches. For example, while a benchmark like LegalForecastBench (which asks models to predict how federal judges ruled on motions to dismiss) can test a model's legal reasoning and analytical ability over a narrow task, it does not test the model's ability to generate new work product, which is an important consideration for a company like Harvey that is developing legal agents to prepare and revise different kinds of legal work product. But these task environments contain departures from practice that could weaken their value for training or evaluation. The model is in effect being tested on its ability to recognize a distinct environment, very unlike what it will encounter in production, and adhere to a stylized set of norms and rules that were artificially constructed in that environment but do not comport with the real norms or rules that govern real-world tasks. While creating LegalForecastBench illustrated that there are real challenges in working with real-world litigation data as well, real records contain contextual signals that this environment did not consistently preserve.

D Human review of AI-flagged LAB criteria

I asked two frontier models, GPT-6 Sol and Claude Opus 5.5, to audit the 52 LAB litigation tasks (specifically Harvey commit `1dd81403`). Table 12 reports how many criteria each pass flagged as defective or arguably defective. The combined row is Opus's later review of rubrics already flagged by either pass. A criterion counts as flagged if Sol flagged it or if Opus's final position after reconciliation flagged it; the 89 blind-pass flags that Opus withdrew on reconciliation, and 7 Sol entries that were not final flags, are not counted. That is 835 of the 2,858 criteria (29.2%) [1].

Pass	What the model saw	Criteria flagged
GPT-6 Sol	The tasks, without Opus's results	616
Claude Opus 5.5, first pass	The tasks, without Sol's results	650
Claude Opus 5.5, combined	The rubrics flagged by either pass	564
Flagged by either model	Share of all 2,858 criteria	835 (29.2%)

Table 12: AI audit of the 52 litigation tasks at Harvey commit `1dd81403`.

In order to validate the AI findings and determine how many of the AI-flagged rubrics were genuinely defective, I reviewed a simple random sample of 25 of the flagged rubrics, drawn with a fixed seed from 19 tasks. I committed the protocol for that review before it began [2]. For each rubric, I reviewed the grading rubric itself, the findings from both AI auditors, how LAB's AI judges graded two real model runs (GPT-6 Luna and Opus 5.5) on that rubric, and the relevant documents from the task environment. In a few cases I also did research into the relevant governing law. I marked a rubric as defective when it could produce an inaccurate grade, meaning: either the rubric could fail a response that a thoughtful litigator would consider appropriate, or the rubric could pass a response that a thoughtful litigator would consider wrong on the issue the rubric is testing. A rubric that forces one of several permissible answers is therefore defective, because it would fail the others.

During the review, I decided to distinguish between defects in the rubrics themselves and defects in the task environment. The line between the two was based on whether the defect could change the grade a model is given on the task. For example, one rubric wanted the work product to contain a specific metric sourced from the task environment, but the AI auditors pointed out that the task environment actually contained two contradictory numbers for that particular metric. Accordingly, an AI agent that relied on one of the numbers (which was not the one the rubric was expecting) would be failed incorrectly because of the underlying defect in the task environment. I treated that as a defect in both the rubric and the task environment. By contrast, a different task environment contained a notice of an OSHA regulatory violation from a state-level regulatory agency in Ohio that does not actually have jurisdiction over the violations charged. The rubric wanted work product to simply flag the existence of the charge in the work product. I coded that as a defect in the task environment but not as a defect in the rubric (because an AI agent handed that task could, and likely still would, pass the rubric by simply mentioning the charged violation).

Of the 25 criteria, I graded 16 rubrics defective in ways that affect grading, including 9 that I considered to be unambiguous objective errors in which the record, the law, the cited authority, or the arithmetic contradicts the rubric. 3 were arguable defects, and 6 were not defective. 14 of the 25 come from task environments that contain defects (10 of the 19 tasks). Only 4 had neither a defective nor an arguable criterion and came from an environment without a defect (Table 13).

Since the sample was a draw without replacement from the 835 criteria that AI agents had flagged as

defective, the one-sided 95% lower confidence bounds are: at least 383 flagged criteria are defective in ways that affect grading (point estimate 534; the bound is 13.4% of all criteria in the practice area); at least 171 are unambiguous objective errors (point estimate 301); and at least 319 come from task environments that contain defects (point estimate 468). At least 524 involve a defective criterion, a defective environment, or both (point estimate 668). In the sample, 9 of the 10 criteria flagged by both auditors were defective, against 7 of the 15 flagged by only one.

I did not attempt to estimate defects in criteria that were not flagged by the AI agents; it is possible there could be some additional defects that the agents did not flag. These estimates count criteria/rubrics, not the individual mistakes that give rise to the defects. In some cases, the defect was specific to the particular rubric at issue, but in other cases, a single error in a task environment could make several criteria defective. After I prepared an initial set of verdicts, I asked AI models (GPT-6 Sol and Opus 5.5) to stress test my answers and to draft additional arguments both for and against each call and a summary of those they would change. I changed a handful of verdicts after reflecting on these points, and did a second pass through the full set of my verdicts to ensure that they applied the principles I adopted during the review consistently to each call.

The audit of rubrics revealed instances where the task environment, and sometimes grading rubrics, appeared to rely on AI-generated hallucinated case law citations. The citation checks that followed are reported in Appendix B; several entries below cite their per-task results.

Criterion verdict	Environment defective	Environment correct	Total
Defective	10	6	16
Arguable	2	1	3
Not defective	2	4	6
Total	14	11	25

Table 13: My verdicts on the 25 sampled criteria, by whether the task environment itself contains a defect.

The sampled criteria

Each entry has four parts: (1) the criterion quoted verbatim from Harvey's task file (that text is from Harvey's data vendors, not me, and in many cases the text appears to be AI-generated), (2) my verdict on the criterion and task environment, (3) results from a sample run of LAB using GPT-6 Luna and Claude Opus 5.5, to assess how LAB's AI graders (Claude Sonnet 4.6 and GPT-5.5) graded outputs from Luna and Opus on the specific criteria at issue, and (4) my explanation of the verdict. Harvey task files are cited at commit `1dd81403`.

1. Assess Reasonableness of Staffing Levels on Litigation Invoice (C-049)

Quoted from Harvey LAB

> **Calculates budget overage percentage as approximately 17%.** PASS if the memo calculates the percentage by which the attorney fee subtotal ($193,372.50) exceeds the approved budget ceiling ($165,000) as approximately 17% (exact figure is 17.2%). FAIL if no percentage overage calculation is provided or if the percentage stated is materially incorrect (off by more than 2 percentage points).
>
> — Harvey LAB, criterion C-049, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error, deficiency in task environment
Task environment	Defective: the invoice's 117 time entries add up to $215,680.00, while its summary and subtotal show $193,372.50.
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric is testing the AI on a specific percentage calculation from an invoice, but the document in the environment is internally inconsistent. The invoice's summary page and its subtotal row show fees of $193,372.50 (17.2% over the $165,000 budget), but the 117 detailed time entries add up to $215,680.00, which puts the overage at $50,680, or 30.7% (13.5 points from the rubric's required 17.2% answer and far outside the 2-point tolerance).

2. Categorize Document Production Set by Relevance and Privilege — Privilege Log and Relevance Classification Report (C-012)

> Quoted from Harvey LAB
>
> **ISSUE_003: DOC_012 identified as Partially Privileged.** PASS if DOC_012 (chandrasekaran-meeting-notes-aug28.docx, Chandrasekaran's handwritten meeting notes from August 28, 2023) is classified as Partially Privileged. FAIL if classified as fully Privileged or Not Privileged.
>
> — Harvey LAB, criterion C-012, task file at commit `1dd81403`

My assessment	
Criterion	Defective — low-quality professional judgment, rubric is incorrect
Task environment	Correct
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric is wrong because it fails fully Privileged, which is the better answer here (even if Partially Privileged is defensible).

- The documents are handwritten lawyer notes from a fictitious meeting discussing problems with the R&D Department. The headings on individual topics seem intended to create the impression that there is a discussion of business topics (which would not be privileged) followed by a separate discussion of legal topics (privileged), so from a superficial glance, the rubric does seem to match the topic headings. However, the bullets under each heading make clear that the entire meeting is just dedicated to the topic of the company's desire to find a way to fire the head of R&D without getting sued. The business topics (which look to be nonprivileged) are discussions of how the R&D department is a bit of a disaster, such that management wants to fire the head. The legal sections then involve discussion of the fact that the R&D head has some sort of environmental whistleblower status, so termination is likely to trigger a lawsuit/whistleblower retaliation claim.

- The meeting was specifically convened because outside counsel was concerned about the retaliation claim, and even some of the 'ops' sections that are supposed to be nonprivileged reflect counsel's advice, so they cannot be cleanly treated as nonlegal business discussions.
- While partially privileged might be defensible, my preferred answer would be fully privileged, which the rubric incorrectly fails.

3. Compare Document Production Against Discovery Requests — Discovery Gap Analysis Memorandum (C-029)

Quoted from Harvey LAB

> **ISSUE_011: Notes only 5 documents actually relate to counterclaim damages.** PASS if the memorandum notes that only 5 of the 14 documents coded to RFP 27 (MLC-004056 through MLC-004102) actually relate to Pacific Corridor's volume shortfall and the counterclaim damages calculation, meaning the counterclaim damages support is thin. FAIL if this is not noted.
>
> — Harvey LAB, criterion C-029, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Defective: the production index codes 15 rows to RFP 27, but its summary tab gives a count of 14 for a range that contains 13.
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. This task requires models to make the point that "only" five documents in the bundle "actually relate to Pacific Corridor's volume shortfall and the counterclaim damages calculation, meaning the counterclaim damages support is thin." This rubric is flawed on multiple levels.

- First, it overlooks MLC-003781, a sixth document that plainly relates to the volume shortfall, so the count is objectively wrong (and the agents pointed to several other examples, at least some of which also seem to fit within the stated criteria). There is a document in the bundle that claims that only five documents concern the volume shortfall, but that assertion is undercut by other evidence supplied in the environment.
- Second, this rubric is someone's argumentative characterization of the documents and probably would not hold up under adversarial testing. For example, there is case law that the term "relate to" is fairly broad and flexible, so quite plausibly a number of other documents would apply.
- Third, the premise underlying this rubric—that "only" five documents is such "thin" evidence that it must be called out—is wrong. Even one document could constitute conclusive proof of damages (the general rule is that damages only need to be reasonably estimated; courts do not require great precision); to my mind, this is actually quite a lot of evidence at the fact stage compared to what one would normally see.

4. Draft Case Assessment Memorandum — Litigation Risk Analysis for Distribution Agreement Dispute (C-025)

Quoted from Harvey LAB

> **ISSUE_007: Recommends motion to dismiss on unjust enrichment.** PASS if the memo recommends filing a motion to dismiss (or motion for judgment on the pleadings, or similar dispositive motion) targeting the unjust enrichment claim (Count IV) on the ground that an express contract governs the parties' relationship. FAIL if no such motion recommendation is made regarding the unjust enrichment claim.
>
> — Harvey LAB, criterion C-025, task file at commit `1dd81403`

My assessment	
Criterion	Defective — low-quality professional judgment, rubric is incorrect
Task environment	Correct
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. A different rubric on this task concludes that plaintiff has a strong breach of contract claim. (See id: C-006.) The prompt instructs models to review case documents and "prepare a comprehensive case assessment memo," and models are dinged on this rubric if they do not "recommend[] filing a motion to dismiss" on the unjust enrichment claim, which is expressly pled in the alternative to the (strong) breach of contract claim.

- The rubric is wrong because a careful litigator could reasonably conclude that a motion to dismiss that claim is not a good idea: Count IV is pleaded in the alternative (Complaint ¶127) and seeks restitution of "no less than $22,521,312.50" (¶135), which is exactly the Count I contract damages ($17,721,312.50, ¶94) plus the Count III trade-secret damages ($4,800,000, ¶120). Counts I and III would survive the motion and double recovery is barred, so dismissing Count IV would not reduce the defendant's net exposure; and the factual underpinnings of both claims are the same so a successful motion would not reduce the scope of discovery. Filing a dispositive motion is costly and, on these facts, there seems to be no benefit to doing so. So the required recommendation would fail many thoughtful analyses that could discuss the pros and cons and recommend against filing such a motion.

5. Draft Conflict Check Memorandum for Litigation Engagement Clearance (C-030)

Quoted from Harvey LAB

> **ISSUE_005: Recommends screening or removal of Caleb Strand.** PASS if the memo recommends that Caleb Strand be screened from the engagement or removed from the proposed staffing. FAIL if no such recommendation is made.
>
> — Harvey LAB, criterion C-030, task file at commit `1dd81403`

My assessment	
Criterion	Defective — low-quality professional judgment, rubric is incorrect
Task environment	Correct
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric requires the model to recommend that an associate who shares an apartment with his sister (an in-house IP paralegal at TriPoint, the adverse party) be screened from the matter. Screening is generally used in certain specific situations where required by the professional conduct rules, but no rule requires screening in this scenario.

- I agree that a prudent memo would flag the fact that a junior team member is sharing an apartment with an employee of the adverse party, but that is covered by other rubrics (C-026 through C-029). This rubric specifically fails the work product (and results in a 0 for all work on this task) even if the model recommended, say, obtaining the client's informed consent, which would be perfectly permissible and appropriate here.
- Other rubrics in this task resolve similar kinds of conflicts with informed consent (C-024 accepts "clearance with conditions (such as obtaining informed consent)" for Lisa Chow's spousal conflict, and C-016 recommends informed consent for the Ridgeline conflict). Accordingly, informed consent plus a home confidentiality protocol could suffice as an alternative to screening.

6. Draft Conflict Check Memorandum for Litigation Engagement Clearance (C-049)

Quoted from Harvey LAB

> **ISSUE_011: Recommends integrating annual disclosure questionnaires with ConflictTracker.** PASS if the memo recommends that annual conflict-disclosure questionnaires be cross-referenced with or integrated into the ConflictTracker system. FAIL if no recommendation about integrating disclosure questionnaires with the conflicts database is made.
>
> — Harvey LAB, criterion C-049, task file at commit `1dd81403`

My assessment	
Criterion	Defective — requires specific outputs not expressly or impliedly requested in the prompt
Task environment	Correct
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The assignment in the prompt is to draft a "comprehensive conflict check memorandum for the proposed engagement." The rubric requires that the memo contain suggested firm-wide process improvements (memo must "recommend[] that annual conflict-disclosure questionnaires be cross-referenced with or integrated into the ConflictTracker system"). That's simply not the

assignment, and there are many reasons why it's undesirable to jam broad process suggestions into a conflicts memo for one specific case. This rubric actually failed work product from Luna that did everything a clearance memo needs on this point. The two judges split on that same Luna memo (Sonnet 4.6 failed it; GPT-5.5 passed it), which shows the criterion does not grade reliably.

7. Draft Counterclaim Against Plaintiff for Breach of Joint Development Agreement (C-045)

Quoted from Harvey LAB

> **Count: Patent Infringement (ISSUE_005).** PASS if the counterclaim includes a separate count for patent infringement under 35 U.S.C. § 271, alleging that the LumiSense 400 infringes U.S. Patent No. 11,234,567 covering the SensorCore™ architecture. FAIL if no patent infringement counterclaim is asserted.
>
> — Harvey LAB, criterion C-045, task file at commit `1dd81403`

My assessment	
Criterion	Arguable
Task environment	Defective: the record lacks the patent the counterclaim must plead, and the two expert reports give different titles for it.
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The task is to draft a counterclaim based on a handful of documents (including an expert report, some emails, and a counterclaim strategy memo that another lawyer prepared). The rubric dings the AI if the counterclaims do not include "a separate count for patent infringement under 35 U.S.C. § 271, alleging that the LumiSense 400 infringes U.S. Patent No. 11,234,567."

- The environment does not contain the allegedly infringed patent, which one would need to review in order to have a sufficient basis for drafting that counterclaim, so this environment appears to be incomplete. That said, a defensible approach might be for the agent to draft the counterclaim but flag in brackets that it requires verification.

8. Draft Counterclaim Against Plaintiff for Breach of Joint Development Agreement (C-046)

Quoted from Harvey LAB

> **Patent Infringement Count: Patent Number Cited.** PASS if the patent infringement count specifically identifies U.S. Patent No. 11,234,567. FAIL if the patent number is omitted or incorrect in the patent infringement count.
>
> — Harvey LAB, criterion C-046, task file at commit `1dd81403`

My assessment	
Criterion	Arguable
Task environment	Defective: same task environment as item 7.
AI auditors' reasoning	Partly correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. Same issue as prior rubric. This one is essentially derivative because it depends on pleading the patent count that was required in the prior rubric.

9. Draft Case Assessment Memorandum for Defective Industrial Equipment Product Liability Claim (C-003)

Quoted from Harvey LAB

> **Executive Summary mentions death of Marco Reyes.** PASS if the Executive Summary mentions the death of Marco Reyes. FAIL if the fatality is not mentioned in the Executive Summary.
>
> — Harvey LAB, criterion C-003, task file at commit `1dd81403`

My assessment	
Criterion	Defective — requires specific outputs not expressly or impliedly requested in the prompt
Task environment	Defective: same task environment as item 10 (Ohio state agency issuing OSHA citations).
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. The case involved a (fictitious) safety incident at a manufacturing plant that killed one employee and seriously injured two others. The rubric requires the fact that one employee died be discussed specifically in the "executive summary." When I tested Opus 5.5 and GPT-6 Luna on this benchmark, both discussed the death in several places in their memo, especially the factual background section, but the executive summary consisted of an overview of the legal analysis that did not restate the fact of the death. They were failed by one or more judges.

- This fails perfectly competent and appropriate work product. The company CEO who retained counsel to advise on this incident is the same person who delivered news of the (fictitious) tragic death to the decedent's widow and surviving family members, so the most senior management of the company is obviously very familiar with the incident and the agents' approach of focusing the executive summary on the legal exposure is appropriate.
- The two judges reached opposite results on Luna's memo: GPT-5.5 failed it, while Sonnet 4.6 passed it only by quoting a sentence from the memo's facts section as if it appeared in the executive summary.

10. Draft Case Assessment Memorandum for Defective Industrial Equipment Product Liability Claim (C-026)

Quoted from Harvey LAB

> **OSHA LOTO citation identified with specific violation and penalty.** PASS if the memo identifies the OSHA serious citation for LOTO deficiency under 29 CFR 1910.147(c)(4)(i) with a penalty of $18,900. FAIL if this specific citation, regulation, or penalty amount is not referenced.
>
> — Harvey LAB, criterion C-026, task file at commit `1dd81403`

My assessment	
Criterion	Not defective — defective environment, but rubric does not fail correct work product
Task environment	Defective: Ohio's Division of Safety & Hygiene cannot issue OSHA citations to this private employer. The criterion itself does not misgrade work (environment-only defect).
AI auditors' reasoning	Wrong
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. Although the specific rubric being tested does not fail work product incorrectly, the environment is defective in an unintended way. The environment contains an OSHA citation from Ohio's Division of Safety & Hygiene. But Ohio does not have an OSHA-approved State Plan, so only federal regulators have jurisdiction to issue OSHA citations against this company. The environment does not test that so it is not being presented as a "mistake" by one's adversary that the agents should pick up on.

- The rubric itself is not wrong because it simply wants a memo to address the purported citation, so an agent could write the memo while noticing that the task is obviously a fictitious evaluation with an incoherent environment.

11. Draft Discovery Plan Memorandum for Breach of Contract and Fraud Defense (C-036)

Quoted from Harvey LAB

> **Provides custodian priority rankings.** PASS if the memo provides priority rankings or tiers for custodians (e.g., high/medium/low priority, or Tier 1/Tier 2, or similar prioritization). FAIL if custodians are listed without any indication of relative priority.
>
> — Harvey LAB, criterion C-036, task file at commit `1dd81403`

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The AI agent is right that the rubric is really enforcing a stylistic choice, but the choice here is well-grounded in convention and an appropriate expectation for quality work product. The case file contains a court order where the judge encouraged the parties to discuss whether phased discovery would be appropriate at the Rule 26(f) conference, and there is some existing prioritization work in the environment. The rubric requires that the agent's discovery plan include some kind of prioritization or phasing of custodians, which is a common expectation in these cases. I think it's fair to judge anything that lacks any form of prioritization as deficient given the context.

12. Draft Federal Complaint for Breach of Contract and Fiduciary Duty — Placement Agent Dispute (C-025)

Quoted from Harvey LAB

> **ISSUE_012: Complaint requests injunctive relief.** PASS if the prayer for relief includes a request for preliminary and/or permanent injunctive relief — specifically to enforce the non-solicitation provision and prevent Graydon from continuing to divert Fund III investor prospects to competing funds. FAIL if no injunctive relief is requested.
>
> — Harvey LAB, criterion C-025, task file at commit `1dd81403`

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong

LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The client intake memo specifically describes an interest in seeking injunctive relief ("We may also wish to seek preliminary and permanent injunctive relief to enforce the twelve-month non-solicitation tail"), and the client's demand letter threatens it. While it's true that the intake memo flags the topic for discussion, the correct approach is to include the proposed prayer for relief in the draft, so the case team can see how it reads and discuss any downsides. As a practical matter, there's unlikely to be much downside to including this prayer for relief in the complaint. (Whether it's worth the cost of actually pursuing a preliminary injunction on this is debatable, so the rubric's approach makes sense.)

13. Draft First Set of Interrogatories to Defendant Veridian Health Systems in Trade Secret Misappropriation Case (C-008)

Quoted from Harvey LAB

> **Definitions section defines 'Document' or 'Documents'.** PASS if the Definitions section includes a definition of 'Document' or 'Documents' that is broad enough to encompass electronically stored information (ESI), emails, and physical documents. FAIL if there is no definition of 'Document' or it is so narrow as to exclude electronic records.
>
> — Harvey LAB, criterion C-008, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Defective: the supplied documents mischaracterize or miscite case law (citation sweep; peripheral to this criterion).
AI auditors' reasoning	Partly correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric fails any draft that does not supply a definition of "document" encompassing electronically stored information (ESI). There are two problems:

- First, these are interrogatories, not document requests. Well-drafted interrogatories do not necessarily need to use the word "document." The rubric does not really make sense as a requirement for this kind of work product.
- Second, even if we were dealing with document requests here (which we are not), the local rules of many courts, including the Western District of Texas (where this fictitious case is being litigated), define the term "document" for all discovery requests served by litigants in those cases. There is thus no requirement to define the term "documents," and in fact, supplying a definition could just trigger objections from the other side that the local rule controls. See Local Rule CV-26, available at https://www.txwd.uscourts.gov/court-information/lcr-civil-rules/ ("The term 'document' means any document or electronically stored information as described in Federal Rule of Civil Procedure 34(a). A draft of a nonidentical copy is a separate document within the meaning of this term.").
- The rubric does demonstrably fail rules-compliant work product. For example, Luna was failed for not including a definition of the term "Document" but that's actually what the local rules encourage in this Court.

14. Draft First Set of Interrogatories to Defendant Veridian Health Systems in Trade Secret Misappropriation Case (C-026)

Quoted from Harvey LAB

> **Interrogatory re: USB drive or CMT files on Veridian systems.** PASS if at least one interrogatory asks whether Oshiro's USB drive (referencing serial number CX-8827491 or otherwise describing the external storage device) was ever connected to, accessed from, or present on any Veridian system, device, or network, or asks whether any CMT files or materials were ever possessed, accessed, viewed, or copied by Veridian or its employees. FAIL if no interrogatory addresses the USB drive or CMT files on Veridian systems.
>
> — Harvey LAB, criterion C-026, task file at commit `1dd81403`

My assessment	
Criterion	Not defective
Task environment	Defective: same task environment as item 13 (environment-only defect).
AI auditors' reasoning	Wrong
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. This one does not really belong on the list, as GPT-6 Sol was complaining about a different rubric (C-048) which it argues is inconsistent with this one. It does not identify any issue in this rubric, standing alone, so this counts as an instance where Sol added something to the "flagged" list that does not belong.

15. Draft Discovery Responses and Objections to RFAs and RFPs in Breach of Supply Agreement Litigation (C-028)

Quoted from Harvey LAB

> **ISSUE_011: Identifies RFA No. 20 as calling for a legal conclusion.** PASS if the plan identifies that RFA No. 20 (asking Prismavale to admit the products were 'not merchantable as defined by the Uniform Commercial Code') calls for a legal conclusion regarding UCC § 2-314 merchantability and recommends objecting on that basis. FAIL if the plan does not identify the legal conclusion problem with RFA No. 20.
>
> — Harvey LAB, criterion C-028, task file at commit `1dd81403`

My assessment	
Criterion	Defective — low-quality professional judgment, rubric is incorrect
Task environment	Defective: the supplied Answer cites complaint paragraph ranges that do not match the complaint (e.g. Count III as ¶¶ 88–94; the complaint's Count III is ¶¶ 86–96).
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. Plaintiff served a request for admission on the defendant that says, "Admit that the SB-102 and PS-302 products shipped by Prismavale to Greenleaf were not merchantable as defined by the Uniform Commercial Code." The prompt asks the agent to "prepare a discovery strategy memo and formal RFA/RFP responses." The rubric requires "the plan" (i.e., the strategy memo) to recommend objecting to the request on the grounds that it calls for a legal conclusion. The rubric is wrong for two main reasons:

- First, the modern view is that requests for admission can call for application of law to facts. While it is within acceptable professional norms for a lawyer to object on the ground that such a request calls for a legal conclusion, and then proceed to answer, the better reading of the law is that this request is not objectionable such that the party should simply answer. On this legal conclusion in particular, there's really no reason why one would prefer to object rather than simply deny the requested admission. So this rubric fails competent work product

(and probably a better answer than the one the rubric calls for).

- Second, even assuming one should include the objection at all, the task fails the model if the issue is not expressly addressed in the strategy memo. Luna made the "legal conclusion" objection in the draft responses and then proceeded to a substantive denial (an acceptable approach), but did not discuss the "legal conclusion" objection in the memo. That was judged a failure. That's really requiring duplicative responses on a very minor point that probably shouldn't be in the drafts at all.

16. Draft Motion in Limine to Exclude Expert Testimony and Prejudicial Evidence in Commercial Breach of Contract Case (C-009)

Quoted from Harvey LAB

> **ISSUE_004: Identifies the $11,000 arithmetic discrepancy.** PASS if the motion identifies that the sum of Marchetti's five itemized damages categories totals $14,211,000, but her report states total damages as $14,200,000, creating an $11,000 discrepancy. FAIL if this arithmetic error is not identified.
>
> — Harvey LAB, criterion C-009, task file at commit `1dd81403`

My assessment	
Criterion	Defective — low-quality professional judgment, rubric is incorrect (a judgment call, not an objective error); defective environment (directs advocacy outside professional norms); companion rubric C-010 also incorrect (not sampled)
Task environment	Defective: the task directs advocacy outside professional norms (an $11,000 rounding point offered as a *Daubert* ground) and labels a rounding difference an "arithmetic error."
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The damages expert includes a summary table in her report listing various items of damages, which collectively add up to $14,211,000; the Total Damages line lists the total as $14,200,000. Opposing counsel questioned the expert about it at her deposition, and the expert calls it "a rounding discrepancy." An internal strategy memo highlights the issue as a ground for seeking to exclude the expert's testimony under *Daubert*, arguing "the arithmetic error supports a cumulative argument that the report was prepared without the rigor required of expert testimony in federal court."

- This task is poorly specified, and the recommended strategy of raising this with the judge is a poor one:
- First, the task environment repeatedly calls this an "arithmetic error." At her deposition the expert said it appeared to be "a rounding discrepancy" and that she would need her workpapers to confirm, and the table reads that way: the total is rounded to the nearest $100,000. That is somewhat unorthodox in an expert report (particularly since it is written out as $14,200,000 rather than $14.2 million), but numbers are often rounded in court filings, and a judge is unlikely to treat a rounded total as an error in the analysis.

- Second, even if the gap were an error rather than rounding, an $11,000 discrepancy in a $14.2 million calculation is generally not a basis for excluding the expert's testimony under *Daubert*. At most one could move to exclude the erroneous version and ask for more accurate testimony, but that does nothing to help the hypothetical client here; it just raises their damages exposure by the $11,000 that the expert was willing to round off.
- Third, the proposed framing comes close to a breach of decorum and carries a real risk of drawing a rebuke from the judge. The strategy memo wants the motion to argue the expert's 30 prior engagements were "in state courts and arbitrations that do not impose the same methodological gatekeeping requirements" as federal courts. While some state courts use the *Frye* standard instead of *Daubert*, many federal judges would take offense at an argument that state courts do not observe adequate methodological rigor. And the reality is that many busy judges would be annoyed at spending time on an argument that literally amounts to a rounding error. It's so trivial that raising it as a substantive point risks making the judge think you had no better argument to put forward, so your motion must be weak.
- This rubric is scored as Defective, as a professional-judgment call rather than an objective error. My view is that the better approach is to drop this issue entirely, and the sampled rubric (C-009) fails a motion that does so, even though a motion that merely mentions the discrepancy in a footnote would pass. The companion rubric, C-010, is in my view objectively wrong because it requires that this issue be presented as a basis for excluding the expert's testimony under *Daubert*. Luna was failed incorrectly on C-010 (by the Sonnet 4.6 judge) for downplaying this issue; in fact, Luna was right to do so, as it should not be raised at all. But since that's not the sampled criterion, I do not mark it here.

17. Draft Rule 12(b)(6) Motion to Dismiss Brief — Commercial Software Licensing Dispute (C-037)

Quoted from Harvey LAB

> **ISSUE_008 — Cites Delaware parol evidence authority.** PASS if the motion brief cites relevant Delaware authority on the parol evidence rule and integration clauses, such as SIGA Technologies v. PharmAthene or Eagle Industries v. DeVilbiss Health Care. FAIL if no Delaware parol evidence authority is cited.
>
> — Harvey LAB, criterion C-037, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error (including hallucinations), rubric is incorrect
Task environment	Defective: the supplied research memo cites fabricated or mischaracterized authority (citation sweep: 28 of 53 citations in this task have problems).
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. The rubric here involves a legal error, and the task environment contains hallucinations about case law. This one is defective for multiple reasons:

- First, the rubric is demanding case law on the wrong legal doctrine. Under Delaware law, the parol evidence rule and an integration clause do not preclude fraud claims. The issue here is whether the plaintiff can reasonably allege that they relied on their counterparty's pre-signing representations despite their agreement in the contract that they were not relying on any representations other than those written in the contract itself.
- Second, the cases in the environment appear to contain multiple hallucinations. The rubric refers to SIGA v. PharmAthene, but Opus noted that case does not address integration or parol evidence. The task environment contains a legal research memo that cites *SIGA Technologies, Inc. v. PharmAthene, Inc.*, 67 A.3d 330 (Del. 2013), and claims that it states (at page 344), "an integration clause is the clearest declaration that the parties intend the writing to be the complete and final expression of their agreement." The quoted language does not appear in the cited case. Page 344 of that case addresses the duty to negotiate in good faith, not integration clauses. Similarly, the memo cites *Eagle Industries, Inc. v. DeVilbiss Health Care, Inc.*, 702 A.2d 1228 (Del. 1997) and says: "The Delaware Supreme Court held that 'a party to a contract cannot promise, in a clear integration clause of a negotiated agreement, that it is not relying on promises or representations made by the other party outside of the agreement, and then assert a claim for fraud based on those very representations.' *Id.* at 1232." The quoted language does not appear in the cited case.
- Third, the rubric failed responses that stated the correct law on this point. In my tests, Opus reviewed SIGA and Eagle, correctly realized they were inapposite and declined to cite them, cited more on-point and accurate authority, and was then failed for it.

18. Draft Rule 12(b)(6) Motion to Dismiss Brief — Commercial Software Licensing Dispute (C-045)

Quoted from Harvey LAB

> **ISSUE_010 — References Arcadia's own delays contributing to Go-Live delay.** PASS if the motion brief identifies Arcadia's own contributing failures, including the 45-day Project Manager gap (Kevin Liu's departure) and/or the 75-day late API specifications, as causes of the Go-Live delay. FAIL if Arcadia's own contributory delays are not mentioned.
>
> — Harvey LAB, criterion C-045, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Defective: same task environment as item 17.
AI auditors' reasoning	Correct

LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. The rubric demands that the agent argue a factual dispute on a motion to dismiss, where the court is required to assume the plaintiff's allegations are true and cannot consider factual disputes. The 45-day and 75-day delays complained of come from the defendant's internal report, which the plaintiff's complaint disputes. Moreover, courts generally cannot resolve messy fact disputes over "causation" on a motion to dismiss.

- The criterion failed agents that correctly recognized that this argument was outside the scope of what can be argued on a motion to dismiss and thus decided to exclude it.

19. Draft Verified Responses and Objections to Plaintiff's First Set of Interrogatories in Commercial Breach of Contract and Fraud Action (C-028)

Quoted from Harvey LAB

> **Correct financial figures for Tri-Basin purchase history.** PASS if any purchase history figures cited in the responses are consistent with the canonical data: 2015: $9.2M; 2016: $10.1M; 2017: $10.8M; 2018: $12.4M; 2019: $13.6M; 2020: $9.7M; 2021: $15.3M; 2022: $16.8M; 2023 (Jan–Aug): $11.9M. Minor rounding differences are acceptable. Also PASS if the response invokes Rule 33(d) rather than citing specific figures. FAIL if specific purchase figures are cited but are materially incorrect.
>
> — Harvey LAB, criterion C-028, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error, deficiency in task environment; rubric is incorrect
Task environment	Defective: the purchase-history spreadsheet is internally inconsistent (three 2020 series; line items that do not sum to subtotals), and the record has no contract-year data.
AI auditors' reasoning	Partly correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. There are two issues:

- First, Interrogatory No. 4(b) asks for Tri-Basin's actual purchases for the 2020 contract year (the contract year runs from January 15 through January 14 of the following year; Interrogatory No. 14 separately asks for calendar-year totals), but the canonical data referred to in the rubric is calendar-year data, so the rubric's numbers are not correct for that request. The correct value cannot be discerned from the environment: the monthly data cannot be split at January 15, and nothing in the record supports assuming purchases are spread evenly within January (seasonal variation would make that unlikely). The rubric therefore grades wrong in the opposite direction from most of the other items: rather than failing good work product, it passes responses that supply the calendar-year figure for a contract-year request without acknowledging that the record does not contain the requested data. That is approving fundamentally deficient work product.
- Second, the data that is provided has some internal inconsistencies. For example, an "annual summary" tab lists 2020 purchases at $9.7 million, but the monthly detail lists 2020 three times, with totals of $8,700,000, $8,549,400 and $9,700,000. The units sold do not match up between those two sheets for 2021, 2022, 2023, and series 3 of 2020. One of the monthly rows does not add up (gross $1,210,000 + credits −$17,400 = $1,192,600, but the net shown is $1,192,500). The annual summary minimums do not add up. In the Product Line Breakdown, line-item net revenue never sums to the subtotal. So a response that relied on the line items would get graded as a failure.

- The environment contains law firm discovery guidelines that recommend that interrogatory responses match the client's Answer, but the Answer here contains only approximate figures, which is consistent with the Answer using calendar-year data as a rough approximation instead of providing the contract-year data requested.

20. Extract Key Obligations from Litigation Hold and Document Preservation Notice — Obligation Summary Memorandum (C-002)

Quoted from Harvey LAB

> **ISSUE_001: Flags financial records date discrepancy as requiring DOJ clarification.** PASS if the memo recommends seeking clarification from the DOJ regarding the conflicting date ranges (Jan 1, 2017 vs. Jan 1, 2019) for financial records. FAIL if no recommendation to seek clarification is made.
>
> — Harvey LAB, criterion C-002, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Correct
AI auditors' reasoning	Correct

LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: fail

My analysis. The task involves a document preservation notice from federal criminal authorities that defines the relevant period (for which the recipient must preserve documents) as January 1, 2019, through March 3, 2025 "unless otherwise specified herein." For one specific category of documents, the notice states that, "Notwithstanding the Relevant Period defined in Paragraph 12 of this Notice, Ridgeline shall preserve [certain kinds of] records . . . from January 1, 2017 to the present date."

- The rubric suggests that the agent should seek clarification of the "conflicting" date ranges but there is no conflict. The notice quite clearly imposed a longer period for one specific category of documents, while all the other categories are subject to a shorter period. This is a fairly common request.
- Both models were failed on this criterion (Luna by both judges, Opus by one), but the criterion makes no sense.

21. Identify Excessive or Duplicative Research Charges in Litigation Invoice (C-013)

Quoted from Harvey LAB

> **Issue 3: Identifies duplicative MTCA research across Takahashi and Wendt.** PASS if the memo identifies that multiple attorneys (Kenji Takahashi and Tyler Wendt) billed for overlapping MTCA liability/contractor liability/summary judgment research during July 2024, citing Section 6.4 (Duplicative Research Prohibition). FAIL if the MTCA duplication across timekeepers is not identified.
>
> — Harvey LAB, criterion C-013, task file at commit `1dd81403`

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. While Sol is right that the facts here are insufficient to establish that the client's billing guidelines were violated, for purposes of the task (a client-side memo identifying charges to dispute), it seems reasonable to flag that multiple attorneys were doing research into closely related issues. The task requested "a memo identifying all non-compliant or excessive charges with recommended adjustments," so flagging potential violations (even if they were not clearly proven from the face of the time entries) is a standard approach a client would take here.

22. Identify Issues in Litigation Matter Budget Proposal (C-029)

Quoted from Harvey LAB

> **Budget inconsistency rated as Critical or High severity.** PASS if the fee total reconciliation discrepancy ($116,050 difference between phase-based and staffing-model totals) is assigned a severity of Critical or High (or the highest or second-highest tier in whatever severity scale is used). FAIL if it is assigned a lower severity level or no severity level.
>
> — Harvey LAB, criterion C-029, task file at commit `1dd81403`

My assessment	
Criterion	Arguable — requires specific outputs not expressly or impliedly requested in the prompt
Task environment	Correct
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The prompt requested "a categorized issue memorandum" but didn't specify what the categories should be. The rubric fails the model if it does not break issues into severity levels or it does not assign a critical or high severity level to a specific issue. While it appears that models often do volunteer this kind of severity assessment, it's not requested by the prompt and competent work product could decline to include it.

23. Review Counterparty's Proposed Jury Instructions — Issue Memorandum for Trade Secrets Trial (C-020)

Quoted from Harvey LAB

> **Explains distinction between tortious interference with contract vs. business relations.** PASS if the memo explains that under Georgia law, tortious interference with an existing contract does not require a separately/independently wrongful act — the

knowing inducement of the breach is sufficient — and that the 'independently wrongful act' element applies to the distinct tort of tortious interference with business relations. FAIL if this legal distinction is not articulated.

— Harvey LAB, criterion C-020, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error (including hallucinations), rubric is incorrect
Task environment	Defective: the fictional summary-judgment order cites a case for the opposite of its holding and a case that cannot be found (citation sweep: 10 of 21 citations in this task have problems).
AI auditors' reasoning	Correct
LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. Multiple issues:

- First, this is a clear legal error. Georgia applies the same elements to interference with contractual and business relations. The rubric says the opposite of what Georgia law is. *Disaster Services, Inc. v. ERC P'ship*, 228 Ga. App. 739, 740–41 (1997) ("Tortious interference claims, whether asserting interference with contractual relations, business relations, or potential business relations, share certain common essential elements: (1) improper action or wrongful conduct by the defendant without privilege; (2) the defendant acted purposely and with malice with the intent to injure; (3) the defendant induced a breach of contractual obligations or caused a party or third parties to discontinue or fail to enter into an anticipated business relationship with the plaintiff; and (4) the defendant's tortious conduct proximately caused damage to the plaintiff.").
- Second, the fictional summary judgment order (from which the rubric's legal position seems to be taken) appears to contain several hallucinations. For example, the order paraphrases the third and fourth elements described in *Disaster Services* approximately correctly, but the first two elements are erroneously listed as "(1) the existence of a valid contract between the plaintiff and a third party; (2) the defendant's knowledge of the contract." The fictional order also cites *Valdosta Livestock, Inc. v. Furst*, 342 Ga. App. 25, 28 (2017), which is not found in any legal databases.

24. Review Litigation Invoice Against Outside Counsel Billing Guidelines — Compliance Deviation Report (C-024)

Quoted from Harvey LAB

Recommends full disallowance of black car service: $387.00. PASS if the report recommends disallowing the full $387.00 for black car service (or reducing it to a reasonable taxi/ride-share equivalent). FAIL if no disallowance or reduction is recommended.

— Harvey LAB, criterion C-024, task file at commit `1dd81403`

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong

LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The billing guidelines state that "Reimbursable ground transportation is limited to ride-share ... or taxi" and "Limousine services and black car services are not reimbursable under any circumstances." Sol's critique of the rubric is that the policy requires disallowing black car reimbursement, not reducing the reimbursement amount to what a taxi or ride share would have cost. The rubric's interpretation (that either disallowing the expense entirely or capping reimbursement at the amount a taxi would have cost is acceptable) seems reasonable and closer to the likely real-world application of a policy like this.

25. Review Litigation Invoice Against Outside Counsel Billing Guidelines — Compliance Deviation Report (C-042)

Quoted from Harvey LAB

> **Classifies business class airfare as borderline/requires discussion.** PASS if the report classifies the business class airfare issue (Messina's $2,847.00 flight) as a borderline or 'requires discussion' issue rather than a clear-cut mandatory reduction, given the 4-hour-20-minute flight duration and the 4-hour threshold. FAIL if it is classified as a clear mandatory reduction without acknowledging the ambiguity, OR if it is entirely dismissed as compliant.
>
> — Harvey LAB, criterion C-042, task file at commit `1dd81403`

My assessment	
Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Correct
AI auditors' reasoning	Correct

LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric's analysis is wrong. The travel policy included in the environment states, "For domestic flights with a scheduled duration of four (4) hours or more, business class may be booked without prior approval." The record does not state the flight's duration (the invoice lists only the route and fare class). The only basis for raising this seems to be in the rubric itself, which says the flight could take 4h20m, but even under that view, there's nothing to discuss, as the policy allows business class for that duration.

E Human review of AI-flagged LegalReasonBench criteria

Sol and Opus audited the 11 litigation assignments from micro1's public LegalReasonBench sample (dataset revision `1bfa904c`). They flagged a total of 63 criteria (out of the 179 they audited) as problematic or arguably problematic. A subsequent Opus pass reviewed both reports and retained 36 of the 63, but all 63 remained eligible for sampling. I reviewed a simple random sample of 25 of the 63 (seed 20261024, from 9 assignments) under a protocol committed prior to the draw. I applied the same standard used for the LAB review.

I did not finalize my verdicts until after an adversarial quality-control process in which AI models (Claude Opus 5.5, Claude Fable 5.1, GPT-6 Sol, and GPT-6 Astra) were invited to argue in favor of and against each verdict that I had reached. I updated the documentation after the audit QC process to more faithfully reflect the standards that I had applied during the audit. Most notably, I determined that rubrics that fail an answer a competent lawyer would consider the correct/best answer or a clearly reasonable one should be treated as defective, since they train models to reject reasonable approaches.

The results were that 21 of 25 reviewed rubrics were defective, with 7 of these constituting unambiguous objective errors and the remainder constituting instances where the rubric enforces an unreasonable or inaccurate position on a question of professional judgment. I found that 12 items came from defective environments.

Statistically this analysis implies that, at the 95% confidence level, at least 45 of the criteria that AI models flagged as defective are in fact defective (the point estimate would be 53 defective criteria) and at least 11 are objective errors (point estimate is 18).

Rubric verdict	Environment defective	Environment correct	Total
Defective	10	11	21
Arguable	0	0	0
Not defective	2	2	4
Total	12	13	25

Table 14: My verdicts on the 25 sampled LegalReasonBench criteria, by whether the task environment itself contains a defect.

1. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 1: Legal Memorandum to Board of County Commissioners

criterion `states_in_the_first_assignment_that_dci2018_10023_remained_activ`

Quoted from micro1 LegalReasonBench

> **Penalty criterion, −15 points.** States in the first assignment that DCI2018-10023 remained active until September 13, 2023.
>
> — micro1 LegalReasonBench, `states_in_the_first_assignment_that_dci2018_10023_remained_activ`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 1
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: clear; Claude Sonnet 5: clear; Gemini 3.8 Flash: clear. Result: penalty not applied (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: fires; Claude Sonnet 5: clear; Gemini 3.8 Flash: clear. Result: penalty not applied (1/3)

My analysis. Confirmed. This rubric penalizes the model if the memo it drafts says that the 2019 zoning case was still pending until September 13, 2023, which is what the task environment indicates (Petition.docx p. 19: "the 2019 zoning case was left open until September 13, 2023").

The rubric's theory is that a county rule, Lee County Admin. Code AC-2-6 § 4.1.B(1), which is not cited in the environment but can be found using web search (which models are allowed to use here), says that remanded hearings must occur within six months and that applications will be considered withdrawn after that. But the rubric ignores that the case is withdrawn only "[i]f the six month deadline is not met because of action or inaction of the Applicant/Appellant." In addition, AC-2-6 § 4.1 is prefaced with a sentence that states that it governs only applications "remanded to the Hearing Examiner by the Board." The case actually concerns a remand from the Hearing Examiner to the staff (Petition.docx: "we jointly request that DCI2018-10023 be remanded back to the staff"), so the regulation does not apply by its terms.

This case appears to be drawn from a real litigation, *Spahhn v. Lee County*, No. 24-CA-1479 (Fla. 20th Cir. Ct. Oct. 25, 2025) (order granting petition for writ of certiorari), and the plaintiffs' position is that the case remained live to September 13, 2023 (contradicting the rubric); the county says the case was withdrawn in March 2022. Elsewhere the rubric claims the case was withdrawn on February 13, 2020 (companion criterion `states_in_the_first_assignment_that_under_lee_county_fla_admin_c`), but neither party advanced that claim, so the rubric seems to be contradicted by the ground-truth evidence from the litigation.

2. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 1: Legal Memorandum to Board of County Commissioners

criterion `states_in_the_first_assignment_that_the_county_should_cease_cons`

> Quoted from micro1 LegalReasonBench
>
> **Penalty criterion, −30 points.** States in the first assignment that the County should cease construction activities at the subject property.
>
> — micro1 LegalReasonBench, `states_in_the_first_assignment_that_the_county_should_cease_cons`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 2
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: clear; Claude Sonnet 5: clear; Gemini 3.8 Flash: clear. Result: penalty not applied (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: fires; Claude Sonnet 5: fires; Gemini 3.8 Flash: clear. Result: penalty applied (2/3)

My analysis. Confirmed. The rubric deducts 30 points if the memo being prepared for a county in a zoning dispute recommends stopping construction.

The basis of this rubric is *Parkway Towers Condominium Ass'n v. Metropolitan Dade County*, 295 So. 2d 295 (Fla. 1974) (per curiam), which merely exercises discretion not to hear an appeal on certiorari. While it is true that the outcome in that particular case was that the county was allowed to proceed with construction, the court's reasoning, if anything, undercuts the rubric's reading: "Although we now consider our issuance of the writ of certiorari to have been improvident in this case, which results in our discharge thereof, it is nevertheless our view prospectively that zoning variations to accommodate county or municipal facility purposes should either have been anticipated in zoning ordinances before construction or operation of such facilities is commenced or, if this has not been done, construction should not be undertaken thereof until after due modification or change therefor is made in existing zoning ordinances." *Id.* at 295–96. In other words, the very case the rubric relies on essentially says the opposite of what it claims and expresses the Florida Supreme Court's view that zoning issues should be fixed before construction occurs.

The fact that some construction has begun is a relevant factor, but nothing in the environment establishes that it would be costly, impossible, or imprudent for the county to stop construction until the zoning issues are fixed properly to comply with the Court's opinion. Many prudent lawyers, relying on the Florida Supreme Court's guidance, might recommend addressing the zoning issue before continuing with construction, so this rubric fails a response many careful lawyers would prefer.

3. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 1: Legal Memorandum to Board of County Commissioners

criterion `states_in_the_first_assignment_that_the_county_should_file_a_tim`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +20 points.** States in the first assignment that the County should file a timely appeal with the 6th District Court of Appeals.
>
> — micro1 LegalReasonBench, `states_in_the_first_assignment_that_the_county_should_file_a_tim`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 3
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: not met. Result: met (2/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Rubrics 23 and 24 (criteria `states_in_the_first_assignment_that_the_county_should_file_a_tim` and `states_in_the_first_assignment_that_the_county_should_remand_the`, +20 each) require the government to simultaneously pursue an appeal and remand to a hearing examiner to fix the deficiencies the Court identified. I agree with Sol that the recommendation to take an appeal is not mandatory here and there are very good reasons not to appeal. While a party who loses in the trial court often wishes to appeal, the trial court's ruling here allows the government to reach the same conclusion it originally did after correcting some deficiencies in the process followed the first time. Often simply correcting the deficiencies identified by a trial court would be faster and more reliable for a government to proceed with a project than appealing.

Although the government in the real litigation did elect to simultaneously appeal and remand the issue to the hearing examiner (as the rubric recommends), the government's filing in the appeal acknowledged that the remand might moot the appeal. The government argued the appeal should still be considered because it wanted to undo the trial court's finding that the government had been biased in approving the initial rezoning, but that is a very debatable position (and arguably not a strong basis for an appeal).

The second level of certiorari review is also limited to whether the circuit court afforded procedural due process and applied the correct law, so the appeal appears to be somewhat weak on the merits (or at least not such an obvious winner that a lawyer in this spot would always advise bringing the appeal, rather than just trying to clean up the issues the lower court flagged). *See Haines City Cmty. Dev. v. Heggs*, 658 So. 2d 523, 530 (Fla. 1995). Filing a protective appeal is reasonable and probably the default most lawyers would implement, but in the case of a government entity, simply curing the identified defects on remand may work equally well. So this case does not seem to require an appeal—that's really a judgment call—and the rubric would incorrectly penalize a failure to pursue a strategy many lawyers would advise against.

4. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 1: Legal Memorandum to Board of County Commissioners

criterion `states_in_the_first_assignment_that_the_county_should_remand_the`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +20 points.** States in the first assignment that the County should remand the case to the Hearing Examiner to address the deficiencies identified in the Order Granting Writ of Certiorari while an appeal is pending.
>
> — micro1 LegalReasonBench, `states_in_the_first_assignment_that_the_county_should_remand_`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 4

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. With respect to the rubric recommending remand to the hearing examiner, I agree with the aspect of it that recommends remand: remand to the Hearing Examiner to fix issues identified by a court is the standard approach a government would follow in this spot. But as written, the criterion requires remand "while an appeal is pending," so it also requires an appeal, and it would therefore reject a remand-only recommendation that a careful litigator could adopt, for the reasons given for criterion 23 (item 3). A memo that recommends remand without an appeal would fail it. Sol flagged it (as arguable) on this same ground.

5. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 1: Legal Memorandum to Board of County Commissioners

criterion `states_in_the_first_assignment_that_under_fla_stat_163_3167_10_2`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +10 points.** States in the first assignment that under Fla. Stat § 163.3167(10) (2025) the development order issued for the project may not be invalidated by a subsequent judicial determination because the timeframe for filing an appeal of the development order has expired.
>
> — micro1 LegalReasonBench, `states_in_the_first_assignment_that_under_fla_stat_163_3167_10_2`, rubric at revision 1bfa904c

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 5

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. This rubric misreads the statutory provision it relies on. Fla. Stat. § 163.3167(10) (2025) protects the developer and allows a project to continue in a situation where no one challenges

that specific development, but a later case (typically involving a different development) invalidates the regulation under which the first development was approved. In that scenario, the developer whose project was approved (with no one challenging it) is entitled to proceed notwithstanding the fact that the regulations the first developer relied on were later invalidated in a *different* case. That principle has no relevance here, where the project at issue was timely challenged by a petition for a writ of certiorari, and a court invalidated the approval for the specific project that was timely challenged.

The statute also applies only when a later court decision rules the land development regulations themselves "invalid because of a deficiency in the approval standards." That is not what occurred here. The circuit court quashed the rezoning after finding defects in the Board's specific decision here, but did not invalidate Lee County's regulations. So even if the criterion were read to refer to some later development order issued under the rezoning, § 163.3167(10) would not address it.

6. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 1: Legal Memorandum to Board of County Commissioners

criterion `states_in_the_first_assignment_that_under_lee_county_fla_admin_c`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +5 points.** States in the first assignment that under Lee County, Fla. Admin. Code § 2-6, 4.1 B. (2025), DCI2018-10023 was considered withdrawn on February 13, 2020.
>
> — micro1 LegalReasonBench, `states_in_the_first_assignment_that_under_lee_county_fla_admin_c`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 6

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. Wrong for the same reasons as item 1. Neither party in the actual litigation contends that the six-month deadline was missed "because of action or inaction of the Applicant/Appellant." There's nothing in the environment to support that claim, and without a factual predicate for that, there's no factual basis for suggesting that Lee County Admin. Code AC-2-6 § 4.1.B(1) (2025) applies here. In addition, AC-2-6 § 4.1 is prefaced with a sentence that states that it governs only applications "remanded to the Hearing Examiner by the Board." The case actually concerns a remand from the Hearing Examiner to the staff (Petition.docx: "we jointly request that DCI2018-10023 be remanded back to the staff"), so the regulation does not apply by its terms.

7. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 2: Brief Requesting Second-Tier Review

criterion `states_in_the_second_assignment_that_expert_testimony_on_the_rec`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +5 points.** States in the second assignment that expert testimony on the record is competent substantial evidence that supports the Board's approval of the rezoning.
>
> — micro1 LegalReasonBench, `states_in_the_second_assignment_that_expert_testimony_on_the_rec`, rubric at revision `1bfa904c`

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong
Sources and run details	Review worksheet, item 7

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. Correct as written. Sol initially suggested that this rubric was arguably defective, because it rewards models for making the point that the expert testimony in the record can support the Board's zoning approval. Sol's concern was that the expert testimony is not actually in the environment provided. While Sol correctly points out that the testimony is not in the environment, the Petition quotes portions of the testimony with citations to the transcripts, and other material in the environment does make clear that expert analysis was used to support the Board's decision. It's a relatively uncontroversial proposition, the relevance of which can be readily inferred from the environment, and my view is that any competent lawyer would consider adding that point (at least briefly) an improvement to the analysis.

8. Petition for Writ of Certiorari for Quasi-judicial Proceeding, assignment 2: Brief Requesting Second-Tier Review

criterion `states_in_the_second_assignment_that_the_staff_report_recommendi`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +5 points.** States in the second assignment that the staff report recommending approval is competent substantial evidence that supports the Board's approval of the rezoning.
>
> — micro1 LegalReasonBench, `states_in_the_second_assignment_that_the_staff_report_recommendi`, rubric at revision `1bfa904c`

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong
Sources and run details	Review worksheet, item 8

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. Correct as written. This is a similar issue to item 7. The rubric wants models to make the point that the staff report is evidence supporting the Board's approval. Sol says the underlying staff report is not part of the environment, but it is referenced in other documents, and any competent lawyer in this situation would want the fact that such a report exists to be cited, as it's clearly going to be important evidence in support of the government's position.

9. Viability of a DCCA Challenge, assignment 1: Record Mapping: Final Notice Against Award

criterion `criterion_8`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +28 points.** Identifies that the comparator memorandum's stated coverage years conflict with the dates of the cases it summarizes.
>
> — micro1 LegalReasonBench, `criterion_8`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Defective: Award ¶ 7 says "Two of the four comparators involved on-duty conduct," but the comparator memorandum shows three on duty (Guzman, Sadler and Voss; only Pietrzak was off duty), so the arbitrator's comparability finding misstates the record it rests on. The memorandum also says "The Commonwealth declined prosecution" of an off-duty shooting in the District.
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 9

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. The rubric rewards models for identifying that the "comparator memorandum's stated coverage years conflict with the dates of the cases it summarizes." But both documents cited as the source for this have the correct dates. The memo says it reviewed "records covering calendar years 2018 through 2021." It then lists cases from August 2018, November 2019, May 2020, and

March 2021, which were resolved in January 2019, June 2020, February 2021, and October 2021, respectively. There is no mismatch, so this rubric would reward models for hallucinating a mismatch that is not present in the environment.

10. Viability of a DCCA Challenge, assignment 2: Scope of Review

criterion `criterion_4`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +26 points.** Identifies that no uploaded authority requires MPD to prove the comparability of its comparator cases.
>
> — micro1 LegalReasonBench, `criterion_4`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Defective: Award ¶ 7 says "Two of the four comparators involved on-duty conduct," but the comparator memorandum shows three on duty (Guzman, Sadler and Voss; only Pietrzak was off duty), so the arbitrator's comparability finding misstates the record it rests on. The memorandum also says "The Commonwealth declined prosecution" of an off-duty shooting in the District.
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 10
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: met; Claude Sonnet 5: not met; Gemini 3.8 Flash: met. Result: met (2/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. The rubric wants an analysis of an employee disciplinary dispute to make the point that "no uploaded authority requires MPD to prove the comparability of its comparator cases." But *Douglas v. Veterans Administration*, 5 M.S.P.B. 313, 5 M.S.P.R. 280 (1981), one of the uploaded authorities, states that "it should be clear that the ultimate burden is upon the agency to persuade the Board [the Merit Systems Protection Board] of the appropriateness of the penalty imposed." *Id.* at 333. It adds that "[t]he selection of an appropriate penalty is a distinct element of the agency's decision, and therefore properly within its burden of persuasion," and that when the employee "challenges the severity of the penalty," as Braddock did (the stipulated arbitral issue was whether "the penalty of termination [is] appropriate"), "the agency will be called upon to present such further evidence." *Id.* at 334.

That language comes just after a list of 12 factors that prior rulings have identified as relevant for determining the appropriateness of a penalty, which includes (as factor six) "consistency of the penalty with those imposed upon other employees for the same or similar offenses." *Id.* at 332. *Douglas* does note that "Not all of these factors will be pertinent in every case," so there is no bright-line requirement that the agency prove factor six in every case. But *Douglas* makes clear that the burden is on the agency to establish the appropriateness of the penalty and comparability is expressly identified as an important factor in assessing appropriateness.

Critically, in this specific case, MPD's Final Notice of Adverse Action expressly relied on comparable

cases as an aggravating factor and a key factor for terminating Officer Braddock. Thus, in the context of this case, the proposition the rubric wants the AI to assert is wrong. If MPD cannot prove to the arbitrator that the comparator officers were similarly situated (as the Final Notice asserts, and as the arbitrator found MPD did not), that could be a basis for at least vacating the termination decision and remanding it for further consideration. Sol and Opus labeled this as an arguable defect, but under the specific circumstances of this case, I would view this assertion as objectively inaccurate. Failure to prove up the comparable cases does not automatically invalidate MPD's action, but it leaves MPD highly vulnerable to being overturned—so in this context, MPD does need to prove the comparable cases that it relied on in making its decision to terminate Braddock.

This rubric appears to be lifted from a real case that matches the fact pattern in the task, *D.C. Metropolitan Police Department v. D.C. Public Employee Relations Board*, 282 A.3d 598 (D.C. 2022), and No. 24-CV-0573 (D.C. May 21, 2026), where MPD explicitly argued that it was not required to put forward comparable cases. But in the real case, the issue was that MPD had not identified any comparable cases where termination was the discipline imposed, and the terminated officer's union came forward with cases involving comparable conduct where the punishment was less severe. The task environment here was modified from the real case, so that in this fictional universe, MPD itself puts comparable cases at issue in its initial decision terminating the officer. In that context, the argument advanced in the real case no longer holds and the rubric is incorrect because it fails to take into account how this task environment was modified from the real case. The rubric is essentially importing a proposition of law from the real case without recognizing that the task environment was changed in a way that makes that proposition irrelevant on the facts being presented to the model.

11. Viability of a DCCA Challenge, assignment 3: Statutory Exceptions

criterion `criterion_16`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +6 points.** Identifies that the PERB Decision misstates the length of the suspension imposed by the Award.
>
> — micro1 LegalReasonBench, `criterion_16`, rubric at revision `1bfa904c`

My assessment	
Criterion	Not defective
Task environment	Defective: Award ¶ 7 says "Two of the four comparators involved on-duty conduct," but the comparator memorandum shows three on duty (Guzman, Sadler and Voss; only Pietrzak was off duty), so the arbitrator's comparability finding misstates the record it rests on. The memorandum also says "The Commonwealth declined prosecution" of an off-duty shooting in the District.
AI auditors' reasoning	Wrong
Sources and run details	Review worksheet, item 11

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Rubric is correct. The rubric wants an internal memorandum to point out that the D.C. Public Employee Relations Board's decision incorrectly states that the officer was suspended for 40 days, when in fact the suspension was 45 days.

Sol's point is that the prompt did not specifically request that analysis. But the prompt directs the associate to "[a]nchor each argument in specific findings of the Award and the PERB Decision," and a careful lawyer doing that would notice that the decision misstates the penalty it affirmed "in all respects." While I agree that it is not necessarily essential for a memo to call out what appears to be a typo in the Board's decision, doing so does improve an internal memo—and this is a situation where it would be desirable to train models to flag that discrepancy even if the user did not specifically prompt for it. The misstatement of fact in the Board's decision affirming the arbitration award is important information for the lawyer to be aware of, and a lawyer using AI to prepare this memo might not think to ask specifically about those kinds of discrepancies. So the rubric is a desirable training signal, and to the extent it asks the model to go somewhat beyond the prompt, this specific example is a good instance where you clearly would want the model to flag an issue that you might not have thought to ask about in your prompt.

12. Plaintiff's Response - Forum Non Conveniens, assignment 1

criterion `criterion_2`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +5 points.** Identifies in Assignment #1 an affidavit by John Smith, M.D. does not indicate that Cook County would be an inconvenient forum for him to litigate this case as required under Langenhorst v. Norfolk S. Ry. Co., 219 Ill. 2d 430 (Ill. 2006).
>
> — micro1 LegalReasonBench, `criterion_2`, rubric at revision `1bfa904c`

My assessment	
Criterion	Not defective
Task environment	Defective: The partner's email says Baby Davis was transferred to "Lurie's Children's Hospital" and "almost passed away on September 16, 2024," but the Complaint says he was transferred to and died at Cook County Hospital (¶¶ 27, 29). The defendants' filed affidavits also carry unfilled placeholders: Dr. Smith's affidavit gives his residence and office as "[street address], [city]" (Ex. A ¶¶ 2–3), and the hospital's affiant is named only as "[Records Custodian]" (Ex. B).
AI auditors' reasoning	Wrong
Sources and run details	Review worksheet, item 12

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (1/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Rubric is correct as written. The case involves a request by defendants to move a case from Cook County (where Chicago is located) to DuPage County (a bordering county immediately west of Cook County, with Chicago O'Hare airport extending partially into DuPage County). One of the defendants, Dr. John Smith, submitted an affidavit arguing that DuPage County is more

convenient than Cook County, but it does not expressly state that Cook County is "inconvenient." The rubric wants a memo analyzing the issue to point out that Dr. Smith does not state in his affidavit that it would be inconvenient for him to litigate in Cook County.

Sol initially flagged the criterion as arguable, reasoning that because Dr. Smith swears DuPage would be "substantially more convenient" than Cook, "[a] careful reviewer could reasonably treat that as comparative inconvenience." But the rubric is correct because it's pointing out that the standard isn't which location is "more" convenient; a defendant seeking to move on forum non conveniens grounds needs to establish that the plaintiff's chosen forum is actually inconvenient to the defendant.

The task environment contains a decision, *Langenhorst v. Norfolk Southern Railway Co.*, 219 Ill. 2d 430, 848 N.E.2d 927 (2006), that states the rule drawing this distinction between an argument that another location would be "more" convenient for one party and proof that the plaintiff's chosen location is actually inconvenient: "The defendant must show that the plaintiff's chosen forum is inconvenient to the defendant and that another forum is more convenient to all parties." 219 Ill. 2d at 444. So Sol's point is not well taken in this instance. The later reconciliation overturned Sol's flag on this ground.

13. Plaintiff's Response - Forum Non Conveniens, assignment 1

criterion `criterion_22`

> Quoted from micro1 LegalReasonBench
>
> **Penalty criterion, −30 points.** Mentions in Assignment #1 that Cook County lacks interest in adjudicating a controversy over care rendered by its DuPage based providers.
>
> — micro1 LegalReasonBench, `criterion_22`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Defective: The partner's email says Baby Davis was transferred to "Lurie's Children's Hospital" and "almost passed away on September 16, 2024," but the Complaint says he was transferred to and died at Cook County Hospital (¶¶ 27, 29). The defendants' filed affidavits also carry unfilled placeholders: Dr. Smith's affidavit gives his residence and office as "[street address], [city]" (Ex. A ¶¶ 2–3), and the hospital's affiant is named only as "[Records Custodian]" (Ex. B).
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 13
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: clear; Claude Sonnet 5: clear; Gemini 3.8 Flash: clear. Result: penalty not applied (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: clear; Claude Sonnet 5: clear; Gemini 3.8 Flash: clear. Result: penalty not applied (0/3)

My analysis. The rubric penalizes a memo that "Mentions . . . that Cook County lacks interest in adjudicating a controversy over care rendered by its DuPage based providers." The rubric is subtly misspecified in a way that can affect the grading. The reason the rubric is penalizing a memo that raises this point is that the environment contains a case, *Bradbury v. St. Mary's Hospital*

of Kankakee, 273 Ill. App. 3d 555, 652 N.E.2d 1228 (1st Dist. 1995), that indicates that Cook County has an interest in situations like these if the patient died there—and the case at issue in the environment involves a patient death, so the theory underlying the rubric is sound.

However, the defendants do argue in their motion that Cook County has limited interest in the death—and as written the rubric penalizes any memo that even "mentions" that Cook County lacks interest. So a grader could apply this penalty even to a memo that simply "mentions" the defendants' argument and then correctly explains why it's wrong under *Bradbury*.

In my test runs, the judges did not impose the penalty for this specific rubric, but penalties were triggered improperly on other rubrics that used the same "Mentions" format, such as criterion 25 ("Mentions in Assignment #1 that Cook County's overall clearance rate is 60.5%," −30), where all judges penalized the models for correctly raising an opponent's argument and rebutting it.

14. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_1`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +3 points.** States in Assignment #2 that the response is supported by the affidavit of Liam Vance, a physician who provided care to Baby Davis at Cook County Hospital.
>
> — micro1 LegalReasonBench, `criterion_1`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; arises from the task environment
Task environment	Defective: No affidavit from any of the five physicians is provided in the environment, and the prompt does not explicitly instruct the model to create one. The prompt also says, "If a document is not uploaded, it does not exist for purposes of this analysis."
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 14

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. The rubric wants the plaintiff's opposition brief (in response to defendants' forum non conveniens motion) to state that the plaintiff's "response is supported by the affidavit of Liam Vance, a physician who provided care to Baby Davis at Cook County Hospital." But no such affidavit is provided in the environment, and the prompt does not explicitly instruct the model to create it. The prompt further says "Do not rely on, reference, or infer any legal authority, factual information, or data beyond what is contained in these uploaded files. If a document is not uploaded, it does not exist for purposes of this analysis." While the prompt elsewhere says, "Please include all affidavits and information you referenced from Assignment #1," that suggests the model should "include" documents that were referenced in the environment, not generate new ones or assume that they will be created separately. Although the assignment 1 rubric rewards recommending that the plaintiff obtain affidavits from these physicians, and the assignment 2 prompt asks the model to include

the affidavits referenced in assignment 1, I would still consider drafting a brief that assumes what third-party witnesses will sign to be hazardous and unreliable.

Drafting an opposition brief that cites to affidavits (without placeholders in the draft) is hazardous behavior to train models on in a legal setting. In a real run, Opus was rewarded for drafting affidavits and including in the brief citations to them without placeholders, while Luna did not create the affidavits or cite to them. If forced to choose, I would consider Luna's approach more correct in this instance, as a matter of professional judgment. Although Luna's approach suffered from its own flaws (including that it incorrectly said the record does not reflect the physicians' willingness to testify), I would prefer the approach of not citing affidavits that don't yet exist.

The contemplated affidavits are from nonparty physicians that the case team does not represent; while it is helpful that a lawyer on the case team spoke to these witnesses and obtained an understanding of their views that can be turned into a draft affidavit, ultimately a nonparty is entitled to decide exactly what it is willing to affirm in writing and a lawyer cannot simply assume that the witness will be willing to sign an affidavit based on the lawyer's understanding from an initial call. The normal approach in a situation like this would be to draft placeholder citations to an affidavit that make clear the language must be adjusted to conform to the witness's final testimony. Opus's draft affidavit includes one form for five physicians, and in practice, I would expect the physicians to object to some of what is included in the affidavit (and in fact getting all five to actually sign these affidavits in time would be a fairly significant undertaking, so I would not assume all five affidavits ultimately enter into the record as this assignment contemplates). The Opus work product creates a real risk that a human lawyer, through inadvertence or mistake, might file a brief that claims to be supported by affidavits that have not in fact been sworn, which is a serious ethical lapse.

Overall, the practical effect of this rubric seems to be to nudge the model in the direction of thinking it is OK to invent witness testimony that has not yet been approved by a witness. The issue could be avoided either by including an actual affidavit approved by the witness in the environment or by asking the model to merely include bracketed text/placeholders that are clearly labeled as draft so that the draft is clear and explicit that the affidavits are contemplated or expected, rather than already secured, as Opus's draft of the brief implies.

The biggest issue here is that the affidavit does not actually exist, so this is training models to hallucinate. But there are also reasons why this would flunk professional judgment.

1. The environment indicates that the physicians work at hospitals, which almost certainly have their own in-house counsel. There's no indication in the environment that a lawyer has been consulted about the plan to inject the doctors' testimony into the forum issue. Once consulted, he may advise the physicians not to take a position on the forum issue even if their inclination was to be helpful.

2. A witness's agreement on a telephone call that they're willing to testify or be helpful is one thing, but actually getting them to sign an affidavit with the specific language you want is quite another. The affidavit needs to be supported by the witness's own personal knowledge, and the factual basis for each witness to assert inconvenience is unclear from this environment; only a conclusory assertion of inconvenience is made. On such a record, counsel may need to gather additional information to assess whether the assertion is well grounded and to explain, for each doctor, the claimed inconvenience.

3. There are also strategic considerations that may make it unwise to submit five identical affidavits, as the task environment proposes, since getting involved in this dispute could impact

the credibility of witnesses who will be needed for more important points later. Submitting five identical affidavits, as the rubric contemplates, will make it obvious that plaintiff's counsel recruited the doctors, drafted affidavits for them to sign, and coordinated with them on a disputed issue in the case. While not impermissible, that signals to the court that the doctors are effectively partisans supporting the plaintiff, which sacrifices some of their credibility when it comes to the merits of the case. In this spot, most lawyers would prefer to preserve the doctors' credibility by leaving them out of minor procedural disputes and presenting them as neutral experts when it comes to the merits. One would not want to reveal at this early stage that witnesses the court might otherwise view as neutral experts are in fact willing to sign whatever you put in front of them. A better strategy would be to leave them out of the minor procedural fight and make their only testimony in the case the merits testimony they will offer.

4. It is particularly strange to do so many affidavits here because it's somewhat incredible to suggest that it's truly "inconvenient" under the applicable standards for doctors to attend a case in adjoining jurisdictions. The defendant's own claim of inconvenience is not well supported and is the sort of conclusory assertion that the court can easily reject. So the better strategy for plaintiff in this spot is simply to point out that it's nonsense for defendant's doctor to argue that it's some grave inconvenience to drive an hour from DuPage County to Chicago for trial testimony, which is likely to be the extent of the doctor's in-person appearances. It's much easier to attack defendants' claim as bogus, rather than respond by putting in an equally bogus claim that the doctors on your side could not possibly drive an hour in the other direction. The more principled position is that neither jurisdiction can truly be inconvenient to the other side since they're only an hour's drive apart, which does not come close to proving *forum non conveniens*. It's defendants' burden to prove inconvenience. Although not absolute given that the plaintiff does not live in the county, the presumption is still that the case will be heard in plaintiff's forum.

5. Even if the plaintiff really wishes to make this "inconvenience" argument for some reason, five affidavits is overkill. The point could be made adequately by picking one or two doctors who have the strongest claim to genuine inconvenience (or perhaps the doctor whose testimony is less important at the merits stage) and submitting their affidavits. In the context of a case this size, there's no need for five affidavits.

This issue is distinct from items 7 and 8, where my view was that it was appropriate for a model to refer in a brief to expert analysis that was referred to in the environment but not supplied. In that case, I had clear indications that the expert analysis described really existed; the full details were simply not provided in the model prompt, so it is important to flag the existence of that evidence apparently already present in the record. Here, the environment makes clear that, for purposes of this assignment, the affidavits do not exist, yet the rubric rewards a model for creating a draft that could easily be filed mistakenly with the false assertion that these doctors have provided affidavits to support their views (which may or may not ultimately come to pass).

15. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_11`

Quoted from micro1 LegalReasonBench

Reward criterion, +5 points. States in Assignment #2 that the response is supported

by the affidavit of the plaintiff, Sara Davis.

— micro1 LegalReasonBench, `criterion_11`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; arises from the task environment
Task environment	Defective: No affidavit from Sara Davis is provided in the environment, and the prompt does not explicitly instruct the model to create one. The only fact supplied about her circumstances is that she "resides in Kane County, Illinois."
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 15
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: not met; Gemini 3.8 Flash: met. Result: met (2/3)

My analysis. This has the same problem as item 14 of asking the model to draft (apparently without placeholders) an assertion that a given point is supported by an affidavit that does not exist. This one is a closer call because the affiant is the client, so the lawyers preparing this opposition brief presumably can be more confident in what the client's testimony will be. But even a client needs to be comfortable that the facts they are swearing to are accurate and may tell the lawyer to change them, so preparing a complete sworn affidavit containing facts found nowhere in the record, and citing it in the brief as sworn fact, as Opus has done (and was fully rewarded for, on a 2–1 vote), is still hazardous.

There is a separate issue here where Sonnet misgraded the rubric because it refers to Assignment #2 in the actual rubric text that is supplied to the judge model, but that is a reference to the prompt the generative models receive and should not be in the output. Sonnet did the same thing on three other criteria in this run (criteria 4, 5 and 22), failing the brief only because it had no "Assignment #2" label; in each case the other two judges passed it, so the majority result did not change. For purposes of this exercise, I am not treating that as a rubric defect because all the rubrics in the land-use and forum non conveniens matters have the issue that they refer to a specific assignment number that is not legible to the judge grader. That is an overarching issue in the benchmark setup that can be addressed globally.

16. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_18`

Quoted from micro1 LegalReasonBench

Penalty criterion, −30 points. States in Assignment #2 that the 2023 Illinois Annual Report Statistics report a 60.5% clearance rate in Cook County, Illinois.

— micro1 LegalReasonBench, `criterion_18`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Defective: No affidavit from any of the five physicians or from Sara Davis is provided in the environment, and the prompt does not explicitly instruct the model to create one. The prompt also says, "If a document is not uploaded, it does not exist for purposes of this analysis."
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 16
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: fires; Claude Sonnet 5: fires; Gemini 3.8 Flash: fires. Result: penalty applied (3/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: fires; Claude Sonnet 5: fires; Gemini 3.8 Flash: fires. Result: penalty applied (3/3)

My analysis. The rubric penalizes models for stating that the 2023 Illinois Annual Report Statistics report a 60.5% clearance rate in Cook County, Illinois. But that document is provided in the environment and does indeed report the rate given as the clearance rate for all cases. In my test runs, Opus and Luna were both penalized incorrectly for noting this figure while explaining why it deserved little weight: Opus contrasted it with the civil-case rates (97.9% and 96.3%), and Luna said it "expressly covers all case types." While some lawyers prefer not to raise bad facts that the other side has not raised, here, the defendants' opening brief squarely presents the 60.5% statistic to the Court—so any competent opposition would respond to that argument. There is no reason to penalize a model for doing so.

17. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_3`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +3 points.** States in Assignment #2 that the response is supported by the affidavit of Caleb Montgomery, a physician who provided care to Baby Davis at Cook County Hospital.
>
> — micro1 LegalReasonBench, `criterion_3`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; arises from the task environment
Task environment	Defective: No affidavit from any of the five physicians is provided in the environment, and the prompt does not explicitly instruct the model to create one. The prompt also says, "If a document is not uploaded, it does not exist for purposes of this analysis."
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 17
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Same issue as item 14. The criterion rewards a brief for saying the response is supported by an affidavit that does not exist in the environment (and which the model is not instructed to create). This is one of several rubrics I deemed defective for encouraging models to cite to affidavits that do not exist or have not been settled/signed by the third-party witnesses whose views they purport to represent (14, 15 and 17–20).

18. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_6`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +3 points.** Identifies in Assignment #2 that the affidavit by Liam Vance indicates that DuPage County would be an inconvenient forum for him to litigate this case as required under Langenhorst v. Norfolk S. Ry. Co. 219 Ill. 2d 430 (Ill. 2006).
>
> — micro1 LegalReasonBench, `criterion_6`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; arises from the task environment
Task environment	Defective: No affidavit from any of the five physicians is provided in the environment, and the prompt does not explicitly instruct the model to create one. The prompt also says, "If a document is not uploaded, it does not exist for purposes of this analysis."
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 18
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Same issue as item 14. The criterion rewards a brief for saying an affidavit that does not exist in the environment (and which the model is not instructed to create) shows that DuPage County would be an inconvenient forum. This is one of several rubrics I deemed defective for encouraging models to cite to affidavits that do not exist or have not been settled/signed by the

third-party witnesses whose views they purport to represent (14, 15 and 17–20).

19. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_8`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +3 points.** Identifies in Assignment #2 that the affidavit by Caleb Montgomery indicates that DuPage County would be an inconvenient forum for him to litigate this case as required under Langenhorst v. Norfolk S. Ry. Co. 219 Ill. 2d 430 (Ill. 2006).
>
> — micro1 LegalReasonBench, `criterion_8`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; arises from the task environment
Task environment	Defective: No affidavit from any of the five physicians is provided in the environment, and the prompt does not explicitly instruct the model to create one. The prompt also says, "If a document is not uploaded, it does not exist for purposes of this analysis."
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 19

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Same issue as item 14. The criterion rewards a brief for saying an affidavit that does not exist in the environment (and which the model is not instructed to create) shows that DuPage County would be an inconvenient forum. This is one of several rubrics I deemed defective for encouraging models to cite to affidavits that do not exist or have not been settled/signed by the third-party witnesses whose views they purport to represent (14, 15 and 17–20).

20. Plaintiff's Response - Forum Non Conveniens, assignment 2

criterion `criterion_9`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +3 points.** Identifies in Assignment #2 that the affidavit by Maya Linwood indicates that DuPage County would be an inconvenient forum for her to litigate this case as required under Langenhorst v. Norfolk S. Ry. Co. 219 Ill. 2d 430 (Ill. 2006).
>
> — micro1 LegalReasonBench, `criterion_9`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; arises from the task environment
Task environment	Defective: No affidavit from any of the five physicians is provided in the environment, and the prompt does not explicitly instruct the model to create one. The prompt also says, "If a document is not uploaded, it does not exist for purposes of this analysis."
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 20
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. Same issue as item 14. The criterion rewards a brief for saying an affidavit that does not exist in the environment (and which the model is not instructed to create) shows that DuPage County would be an inconvenient forum. This is one of several rubrics I deemed defective for encouraging models to cite to affidavits that do not exist or have not been settled/signed by the third-party witnesses whose views they purport to represent (14, 15 and 17–20).

21. OTB Gang VICAR Trial Jury Instructions, assignment 1: Appeal brief re: justification instruction

criterion `criterion_21`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +15 points.** States that Patton's commission of conspiracy to commit a forcible felony precludes a justification defense.
>
> — micro1 LegalReasonBench, `criterion_21`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 21
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: met. Result: met (3/3)

My analysis. The rubric rewards models for a legal error. The case involves federal charges against a gang member who shot at a rival gang member after following him from a mall. The defendant pled guilty to state charges of conspiracy to commit a forcible felony in connection with the incident. The rubric wants the prosecution's brief to argue that the commission of conspiracy to commit a forcible felony precludes the justification (self-defense) defense, under Iowa Code § 704.6.

But that statute only applies when the defendant commits a different forcible felony than the one

he is alleging is justified. The paradigmatic case would be a robber who waves a hammer at a store clerk to take the cash in the register, intending only to frighten the clerk, and who cannot claim self-defense if the clerk responds by shooting at him (robbery is a forcible felony under the supplied Iowa Code § 702.11(1)). Iowa courts apply the bar in that setting. *See State v. Roby*, No. 22-1017, slip op. at 8 (Iowa Ct. App. Sept. 27, 2023) (defendant who brought a gun to an apartment and shot an occupant who drew his own weapon: "A justification defense was unavailable to Roby because he was participating in the forcible felony of robbery. See Iowa Code § 704.6(1).").

Here, the defendant is arguing that the shooting itself was an act of self-defense; there is no contention that he was committing a different forcible felony prior to the shooting that would negate his ability to claim self-defense for the shooting itself. (Although I did not find an Iowa case expressly holding that the forcible felony must be independent of the act claimed to be justified, judges in at least one other state with a similar statute have expressly held this. *See Martinez v. State*, 981 So. 2d 449 (Fla. 2008) (holding that for the forcible-felony bar to apply, "there must be an independent forcible felony other than the one which the defendant claims he or she committed in self-defense"). Obviously, the law could not be to the contrary because, if it were, the exception would swallow the rule; self-defense would never be available, because absent the defense, the use of force would almost always constitute a forcible felony.)

This environment is based on a real case, and in the real case, neither the Eighth Circuit nor the government's brief opposing certiorari relied on the argument (that the rubric requires) that the conspiracy to commit a forcible felony precludes the justification defense. Instead, the government argued that the defendant as a felon in possession had a duty to retreat safely and could have done so instead of shooting as he did. The evidence indicates this was a scenario where two rival gangs were actively seeking opportunities to shoot members of the other, not that an innocent person was attacked and then defended themselves.

22. OTB Gang VICAR Trial Jury Instructions, assignment 1: Appeal brief re: justification instruction

criterion `criterion_23`

> Quoted from micro1 LegalReasonBench
>
> **Penalty criterion, −5 points.** Mentions that Patton pled guilty in state court to intimidation with a dangerous weapon.
>
> — micro1 LegalReasonBench, `criterion_23`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 22

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: clear; Claude Sonnet 5: clear; Gemini 3.8 Flash: clear. Result: penalty not applied (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: fires; Claude Sonnet 5: fires; Gemini 3.8 Flash: clear. Result: penalty applied (2/3)

My analysis. The rubric penalizes mentioning that Patton pled guilty in state court to intimidation

with a dangerous weapon, because that incident relates to a different event than the one at issue on his appeal. The rubric seems intended to test the model's ability to distinguish the events and understand that the earlier incident is unrelated. However, it is misspecified because it penalizes any mention of the issue. Opus was penalized merely for including a drafting note that flagged for the attorney that it had decided *not* to rely on the other incident because it was unrelated. A brief or drafting note that mentions the April 13 plea only to set it aside is correct work, so I do not think penalizing a model for merely mentioning it is correct. (The Count XX plea petition is dated November 18, 2020, after the May 10, 2020 shooting, so it is not a prior conviction.)

23. OTB Gang VICAR Trial Jury Instructions, assignment 2: Supplementing brief with new case

criterion `criterion_4`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +25 points.** Mentions that there was a lack of evidence proving who fired first on May 10, 2020.
>
> — micro1 LegalReasonBench, `criterion_4`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — unambiguous objective error; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 23

LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. This is another instance where the task environment differs from the real-world case on which it is based in a way that is material to this specific rubric—and a rubric that would be correct for the real-world case is incorrectly applied to this task environment. In the real case, *United States v. Mallory*, 104 F.4th 15 (8th Cir. 2024), there was a dispute concerning who fired first. This task environment was "closed book," however, and the model was given a statement of facts that says that Raysean Nelson "opened fire as the Trailblazer approached" and that "A rolling gun battle ensued." So there is no basis in the supplied environment for the model to assert that there was a lack of evidence concerning who fired first; the supplied documents state that Nelson did so. The supplied facts are internally consistent; the error is the rubric's failure to follow the task's own changes to the real case.

24. OTB Gang VICAR Trial Jury Instructions, assignment 2: Supplementing brief with new case

criterion `criterion_7`

> Quoted from micro1 LegalReasonBench
>
> **Reward criterion, +10 points.** Mentions that Nelson's car was already stopped

when Patton fired his gun.

— micro1 LegalReasonBench, `criterion_7`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 24
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)

My analysis. This benchmark does two assignments, sequentially, in separate model sessions that are in the same sandbox. In this task, assignment 2 was to build on the assignment 1 output by adding discussion of an additional case. The intent is for models to output response1.md in response to assignment 1, and to output response2.md in the same sandbox for assignment 2.

The model is told, "Build on [the assignment 1] analysis as the prompt directs," but not clearly told that it is expected to output a full revised version of the brief (as opposed to supplemental or revised sections that can be inserted into a revised brief). Since assignment 2 is framed as a mere supplementation task, the model might expect that it should write only the new material.

However, five of the assignment 1 rubrics are repeated on the assignment 2 output—so a model is penalized heavily if it interprets the instructions as requesting simply an update to the specific section to which the new authority is relevant. In effect, the rubric is testing for one particular response to somewhat ambiguous instructions. This rewards work beyond the scope the prompt states, and a complete revised brief is not what the prompt clearly requests.

25. OTB Gang VICAR Trial Jury Instructions, assignment 2: Supplementing brief with new case

criterion `criterion_8`

Quoted from micro1 LegalReasonBench

Reward criterion, +15 points. States that the OTB group could have driven away.

— micro1 LegalReasonBench, `criterion_8`, rubric at revision `1bfa904c`

My assessment	
Criterion	Defective — professional judgment; does not arise from the task environment
Task environment	Correct
AI auditors' reasoning	Correct
Sources and run details	Review worksheet, item 25
LegalReasonBench's own grading of two model runs (its judges, not my assessment)	
GPT-6 Luna (xhigh)	GPT-5.6 Terra: not met; Claude Sonnet 5: not met; Gemini 3.8 Flash: not met. Result: not met (0/3)
Claude Opus 5.5 (low)	GPT-5.6 Terra: met; Claude Sonnet 5: met; Gemini 3.8 Flash: not met. Result: met (2/3)

My analysis. This benchmark does two assignments, sequentially, in separate model sessions that are in the same sandbox. In this task, assignment 2 was to build on the assignment 1 output by adding discussion of an additional case. The intent is for models to output response1.md in response to assignment 1, and to output response2.md in the same sandbox for assignment 2.

The model is told, "Build on [the assignment 1] analysis as the prompt directs," but not clearly told that it is expected to output a full revised version of the brief (as opposed to supplemental or revised sections that can be inserted into a revised brief). Since assignment 2 is framed as a mere supplementation task, the model might expect that it should write only the new material.

However, five of the assignment 1 rubrics are repeated on the assignment 2 output—so a model is penalized heavily if it interprets the instructions as requesting simply an update to the specific section to which the new authority is relevant. In effect, the rubric is testing for one particular response to somewhat ambiguous instructions. This rewards work beyond the scope the prompt states, and a complete revised brief is not what the prompt clearly requests.