

LegalForecastBench: Forecasting Judicial Decisions as a Test of Legal Reasoning Ability

John J. Hughes, III
Independent researcher
<https://www.johnjhughesiii.com/research/>

Working draft — October 7, 2026

Abstract

Reinforcement learning (RL) techniques, in domains that can offer objectively verifiable rewards, have led to incredible and noteworthy recent gains in the capabilities of frontier AI models. But not all domains are equally verifiable. Many high-level legal judgment tasks are difficult to verify, and model capability gains on those tasks have not been nearly as noteworthy as model gains in domains that are susceptible to automated verification like math and coding.

Existing benchmarks used to evaluate high-level legal tasks rely largely on human-written rubrics that test model output in response to fictitious, AI-generated environments. But an AI-assisted audit of one of the most widely used such benchmarks, checked by a practicing litigator’s review of a random sample of the flagged criteria, found evidence of hundreds of erroneous rubrics in just one subpart of the benchmark, many of them in task environments that are themselves misspecified. The errors in these environments illustrate the challenges facing even the top data providers and research labs trying to improve frontier model capabilities, at a time when many previously challenging benchmarks and tasks have become saturated.

This paper argues that the heavy reliance on artificially constructed environments fails to take advantage of key properties of the common law judicial tradition, which gives precedential effect and the force of law to reasoned judicial decisions issued by judges in specific cases or controversies. For AI researchers, those judicial dispositions supply an external signal that allows for evaluation of some high-level legal judgment tasks against an objectively verifiable real-world ground truth. In the legal system, the opinions of purported “legal experts” or data providers as to what the law is have no authoritative or precedential value. The pronouncements of a tribunal, by contrast, carry the force of law (and often precedential effect).

LegalForecastBench illustrates one mechanism by which the signal from judicial decisions can be used to objectively evaluate the legal reasoning capabilities of frontier models. Models receive the same briefs that a federal judge received when ruling on a motion to dismiss, and are asked to estimate the probability that the judge dismissed each claim challenged by a defendant.

The beta includes 91 cases and 387 scored claim–defendant units. In the initial run, GPT-6 Sol has the lowest micro-Brier loss (0.118) and the highest accuracy (84.5%) of seventeen models tested. We describe how more training signal could potentially be extracted from real-world litigation data, and suggest this approach may be more effective than relying on fictitious environments created by some data providers.

1 Introduction

Reinforcement learning (RL) techniques have led to significant recent gains in frontier model capabilities, but the most noteworthy gains have been in domains like mathematics and coding that

are susceptible to automated verification. Reinforcement learning with verifiable rewards (RLVR) has allowed frontier models to gain capabilities by receiving reward signals from automated correctness checks (like mathematical answer matching, executable tests, or formal proof checking) [1, 2]. Earlier self-play systems demonstrated the powerful capabilities that automated outcome feedback could provide to models playing games with explicit rules and automatable scoring mechanisms [3].

Frontier models have also made noteworthy gains on legal tasks that are objectively verifiable; for example, tests that ask models bar exam-style questions about established legal propositions are approaching saturation. Neel Guha’s legal benchmark dashboard reports 98.3% accuracy (115/117) for Gemini 3 Flash on a publicly released subset of Multistate Bar Examination questions, with several other frontier models exceeding 90% [8]. This increase in legal capabilities has forced researchers to develop new tasks that challenge today’s frontier models. The literature illustrates at least two approaches being used recently: preference evaluation (comparing which of two outputs practicing attorneys prefer) and rubric-based reinforcement learning (grading model outputs against specified criteria to produce a reward) [4]. While these approaches are important and useful, our audit of one of the most widely used rubric-based legal benchmarks (which utilized rubrics supplied by leading data companies) found that more than 13.4% of the criteria in the litigation practice area (95% CI) contained material errors that could impact the grading on those rubrics, and more than 11.1% come from task environments with errors or inconsistencies. This illustrates the extraordinary difficulty of obtaining reliable, high-quality training data that can improve frontier models’ legal abilities at current capability levels.

This paper aims to demonstrate that judicial dispositions in litigation matters provide a valuable external signal that can make high-level legal judgment tasks verifiable against an objective ground truth. This approach rests on a simple reality of the legal system: “It is emphatically the province and duty of the Judicial Department to say what the law is.” *Marbury v. Madison*, 5 U.S. 137, 177 (1803). Every litigated judgment has the force of law as between the parties, and definitively resolves any disputes between the parties on the interpretation of the law or its application to a specific fact pattern.

LegalForecastBench illustrates one specific use of this signal, using judges’ rulings on motions to dismiss to evaluate the predictive abilities of frontier models. The benchmark focuses on district court level motions because lower courts are heavily constrained by precedent and largely decide garden-variety disputes in which the ability to correctly analyze and apply the law to the alleged facts will be the primary way to predict the outcome. Prediction at this level is likely to serve as a reasonable proxy for legal analytical ability (whereas the U.S. Supreme Court is not nearly as constrained by precedent, since it is free to overrule any of its own prior precedential decisions, and prediction at that level would focus to a much greater extent on modeling the preferences or likely vote of individual justices). Obviously, not every judicial disposition is accurate, and there are some politically salient cases even in the lower courts where different judges might have different leanings, so the prediction task is not a pure test of legal reasoning ability.

But even so, frontier models’ prediction ability is valuable and important, because prediction is a high-level legal judgment task often performed by the most senior and highly compensated lawyer on the case team because of its enormous consequences for the client’s legal strategy. In many kinds of cases (such as securities fraud class actions), whether the motion to dismiss is granted or denied can have an enormous impact on the practical economic value of a case; the vast majority of cases in which pre-trial dismissal motions are denied end in settlement, whereas a case that is dismissed at the outset usually results in no recovery to the plaintiff [11, p. 16]. Prior research shows that even parties who have no intention of litigating negotiate legal resolutions “in the shadow of the

law” [12]. When deciding whether to settle a dispute or litigate, each side’s expected value from litigation (net of the costs and risks involved) can function as a kind of reservation value; if the settlement offer is worth less than the expected value from litigation, then a party will rationally reject settlement and litigate instead. [32, 33]. This can make settlement difficult or impossible if both sides are overly optimistic about their chances of success in litigation; models that help lawyers and their clients better calibrate their expectations could help unlock the enormous surplus that litigants achieve when they settle consensually (avoiding the costs, risks, and distraction brought on by protracted, adversarial litigation), by helping both sides to see their case more realistically, which should make them more willing to offer and accept the fair value of the case.

2 Related work

2.1 Legal outcome prediction

Using AI models to predict litigation outcomes has a long lineage, although the incredible recent growth in frontier model capabilities has created opportunities for more powerful approaches than those available to prior researchers. One early experiment involved the 2002 Term of the U.S. Supreme Court; Ruger and colleagues predicted the outcome of cases that Term using hand-coded metadata and statistical techniques [20]. Later researchers used more sophisticated machine-learning and natural language processing techniques. Aletras and colleagues in 2016 used a text-based approach to predict the outcomes of European Court of Human Rights decisions [22]. Katz, Bommarito, and Blackman in 2017 used a random-forest classifier over structured Supreme Court data to predict Justice votes and case outcomes. [21]. CAIL2018 used a dataset of more than 2.6 million Chinese criminal cases to predict charges and sentences, with text-classification models providing initial baselines. [23]. As the transformer architecture became dominant, models like BERT and T5 were increasingly used. CLC-UKET, released in 2024, compared fine-tuned BERT and T5 models to zero- and few-shot GPT-3.5 and GPT-4 to predict U.K. Employment Tribunal cases. [24].

2.2 Legal AI benchmarks and evaluation

The growth in frontier model capabilities has similarly coincided with increasingly ambitious benchmarks. LegalBench, released in 2023, is a major collaborative benchmark to which 40 authors contributed 162 tasks intended to assess legal domain expertise. The benchmark deliberately tests relatively clean legal propositions that can be reduced to a binary answer. [7].

Most recent benchmarks have sought to test higher-level judgment tasks, which usually lack an objective ground truth. JudgmentBench, released in May 2026, and Harvey’s LAB targeted open-ended professional legal work, like drafting, analysis, and advice, for which there are potentially multiple defensible outputs and no independently verifiable ground truth [4, 26]. These benchmarks take different approaches to reward specification. JudgmentBench compared rubric-based scoring to comparative judgment (pairwise preference judgments collected from practicing attorneys), and found that comparative judgments recover the intended quality ordering substantially better than rubrics. LAB relies on a rubric-based approach, to allow rapid and scalable scoring of new model releases.

2.3 Verifiable rewards, forecasting, and outcome feedback

Frontier models are trained on several different forms of feedback, as described in prior work, including Reinforcement Learning with Verifiable Rewards (RLVR). Recent work has used real-world (nonlegal) events to supply an objective training signal for RLVR. For example, Autocast, released in 2022, gathered forecasting questions from tournaments and a date-indexed news corpus, then had language models make predictions based on the information a human would have had at the time. [27]. Halawi and colleagues in 2024 used a retrieval-augmented forecasting system that could search for information, and generate and aggregate multiple forecasts [28]. ForecastBench, introduced in 2024, used this approach to create a dynamic benchmark that continually collects questions about unresolved future events and scores models once the event outcome is known [29].

Outcome-Based Reinforcement Learning to Predict the Future showed in 2025 that resolved binary forecasting questions could serve as reinforcement-learning rewards, training a 14-billion-parameter model on more than 100,000 events using adaptations of GRPO and ReMax [13]. OpenForecaster similarly explored training language models directly on resolved forecasting questions [14]. Some research in China has also used frontier models to predict criminal sentences [15].

3 LegalForecastBench: Dataset and evaluation protocol

3.1 Prediction target and information set

Let $c \in \{1, \dots, C\}$ index cases and $u \in \mathcal{U}_c$ index scored units within case c . Each unit identifies a challenged claim or a separately resolvable pleaded subclaim that a plaintiff has asserted against a moving defendant (or a group of similarly situated moving defendants). The “units” (i.e., the claims that each set of defendant(s) is moving to dismiss) are identified by labeling models that were given access to the motion papers.

When the benchmark is run, a model receives the same docket and motion papers that the judge received X_c and the specification of the units that are being challenged U_{cu} and returns

$$p_{cu} \in [0, 1], \quad p_{cu} \approx \Pr(Y_{cu} = 1 \mid X_c, U_{cu}), \quad (1)$$

where $Y_{cu} = 1$ means that the judge dismissed that specific claim-defendant unit in its entirety, in the target motion’s first written disposition. That event is not necessarily a final victory in the lawsuit: a plaintiff could be given leave to amend the complaint (in which case the defendant typically will file a renewed motion to dismiss), or other claims might survive even if some are dismissed. The plaintiff also has a right to appeal final dismissal orders.

First written disposition of the frozen unit	Primary label
Entire claim dismissed, with or without leave to amend	1: fully dismissed
Only one nonseparable theory or remedy dismissed; claim remains	0: not fully dismissed
Claim dismissed against one defendant but survives against another	Separate labels for separately frozen units
Disposition does not reliably resolve the frozen target	Excluded from primary scoring

Table 1: Illustrative label semantics.

3.2 Cohort construction and labeling

The analyzed set consists of 91 U.S. federal cases with written decisions dated June 30–August 7, 2026. The cases, as labeled, contained 409 claim-defendant units that could be forecasted. Of those, 387 are scored. The unscored units involved situations where the plaintiff voluntarily withdrew the frozen claim, the order did not cleanly grant or deny the frozen claim, or (for four units) we determined after freezing the unit that the frozen defendant was the wrong target or the motion did not address the frozen claim. Of the scored claim-defendant units, 285 were fully dismissed and 102 were not. The sample was not random; our pipeline favored cases where some of the necessary documents could be obtained for free from online databases, although we typically still had to purchase a number of documents to obtain a complete set of the substantive motion papers that were before the judge.

The case construction and labeling pipeline used models as well as human legal oversight from me. The pipeline maintained separation between *unitization* and *outcome coding*. To avoid the possibility that the prediction units could leak information about the judge’s ultimate outcome to the models that were performing prediction, AI models first were asked to identify the challenged claim-defendant units based only on the motion and the docket (without the decision). A later stage was given access to the judge’s ruling, to interpret the disposition of each frozen claim-defendant unit. I resolved disputed units (where two models did not agree on the disposition) and corrected labels as necessary. The process of creating the pipeline was iterative, and we did not track a blinded human-audit error estimate.

Models predicting outcomes were not given access to the judge’s ruling. The pipeline also attempted to screen dockets and filing text to avoid outcome leakage. Models had access to the motion to dismiss papers (through document text tools) in a sandbox with no Internet access. Model API calls were made with provider-side web search disabled. Provider-specific reasoning settings are retained in the run registries. Appendix A lists the configuration settings. The study population is the seventeen configurations in Table 2.

The decisions used in the benchmark were decided by the federal judges and became public, only after the reported knowledge or training cutoffs for all of the models that disclose such cutoffs. Kimi and Muse do not publish training cutoffs, but based on the model release dates, it appears unlikely that those models were exposed to the decisions in training.

The acquisition and labeling pipeline is designed to create new cohorts in the future, as new judicial decisions become available. This enables fresh comparisons even after models have been trained on the specific decisions at issue here. The downside of this approach is that refreshing the benchmark with a new cohort of cases changes the case mix and difficulty, so scores from different versions cannot be interpreted as a time series of model improvement.

3.3 Scoring and statistical inference

For $N = \sum_c |\mathcal{U}_c| = 387$ and $C = 91$, the primary metric is the micro-Brier loss

$$B_{\text{micro}} = \frac{1}{N} \sum_{c=1}^C \sum_{u \in \mathcal{U}_c} (p_{cu} - Y_{cu})^2. \quad (2)$$

The equal-case score changes the estimand by giving each case the same weight:

$$B_{\text{case}} = \frac{1}{C} \sum_{c=1}^C \frac{1}{|\mathcal{U}_c|} \sum_{u \in \mathcal{U}_c} (p_{cu} - Y_{cu})^2. \quad (3)$$

Lower is better for both. Micro-Brier treats each scored claim–defendant unit equally, so cases that involve a motion to dismiss many separate claims receive more weight than cases where only one claim and defendant was subject to a motion to dismiss. Equal-case Brier weights each prediction unit in inverse proportion to the number of units from that case.

The Brier rule is strictly proper for the binary target. If the true conditional event probability is q , then

$$\mathbb{E}[(p - Y)^2 \mid X, U] = (p - q)^2 + q(1 - q), \quad (4)$$

which is minimized by reporting $p = q$. The second term is irreducible uncertainty for that information set. A set of observable judicial events therefore creates a verifiable score across the entire probability distribution, even though a single realized outcome standing alone does not reveal whether the probability a model assigned to that outcome was correct *ex ante* [9, 10]. Accuracy at threshold $p \geq 0.5$ is reported as a secondary descriptive metric.

The observed dismissal frequency is $\hat{\pi} = 285/387 = 0.73643$. Always predicting dismissal has accuracy 73.64%, so a constant forecast $p = \hat{\pi}$ has micro-Brier $\hat{\pi}(1 - \hat{\pi}) = 0.19410$. The constant is fitted retrospectively on this cohort, so it is a descriptive reference, not a baseline derived from earlier cases. The uninformative forecast $p = 0.5$ has Brier 0.25.

Log loss is a complementary diagnostic:

$$L = -\frac{1}{N} \sum_{c,u} [Y_{cu} \log p_{cu} + (1 - Y_{cu}) \log(1 - p_{cu})]. \quad (5)$$

For a wrong forecast assigning 0.99 to the predicted class, Brier loss is 0.9801 and log loss is approximately 4.605; at confidence 0.90 the corresponding losses are 0.81 and 2.303. Thus log loss imposes a large penalty on instances where the model assigned a very high probability to the outcome that did not occur.

Claim–defendant units from the same case share facts and legal issues; distinct claims in a single case may suffer from closely related legal defects, so the outcomes are not independent. We therefore measure within-case dependence using an observed-scale intraclass correlation coefficient (ICC) [35].

Across the 387 scored units in 91 cases, the observed-scale outcome ICC was 0.76 (95% CI: 0.61–0.87). Prediction losses also clustered within case because if a model failed to account, for example, for the possibility that a given motion might be denied on procedural grounds, all of the claims in a case that it expected to be dismissed could survive: GPT-6 Sol’s Brier-loss ICC was 0.42 (95% CI: 0.24–0.58). These intervals use 20,000 whole-case bootstrap resamples, and assume independence between cases.

4 Results

4.1 Model performance and unconditional references

Table 2 reports the seventeen configurations on the ranked leaderboard. GPT-6 Sol achieves the best observed micro-Brier (0.11763), equal-case Brier (0.12949), and accuracy (327/387, 84.50%). All seventeen beat the retrospectively fitted constant micro-Brier of 0.19410; GPT-6 Sol’s micro-Brier is lower than that constant by about 39.4%.

Sixteen configurations also beat majority-class accuracy (the 73.64% accuracy of predicting dismissal for every unit). Muse Spark 1.3 does not: 279/387 correct forecasts is 72.09%. Gemini 3.1 Pro

Configuration	Micro Brier	Case Brier	Correct	Acc. (%)	Cost (\$)
GPT-6 Sol	0.11763	0.12949	327/387	84.50	11.01
Claude Fable 5.1	0.12877	0.13653	318/387	82.17	83.18
Grok 4.6	0.13368	0.13480	312/387	80.62	15.72
Claude Opus 5	0.13987	0.13548	310/387	80.10	45.89
GPT-5.6 Sol	0.14009	0.13319	313/387	80.88	24.91
Kimi K3	0.14130	0.15367	303/387	78.29	16.87
GPT-5.6 Luna	0.14322	0.15084	313/387	80.88	1.48
Claude Opus 5.5	0.14540	0.14762	299/387	77.26	20.59
GPT-6.1 Sol	0.14948	0.13885	304/387	78.55	10.68
GPT-6 Astra	0.15221	0.14275	305/387	78.81	59.74
Gemini 3.8 Flash	0.15331	0.13900	310/387	80.10	35.61
Grok 4.7	0.15833	0.16454	298/387	77.00	33.00
Claude Sonnet 5	0.15848	0.15964	293/387	75.71	24.68
GPT-6 Luna	0.16368	0.17379	296/387	76.49	0.62
Claude Sonnet 5.5	0.17298	0.17183	295/387	76.23	6.99
Muse Spark 1.3	0.18121	0.17845	279/387	72.09	8.78
Gemini 3.1 Pro Preview	0.18345	0.18080	302/387	78.04	59.04
Cohort-rate reference	0.19410	—	285/387	73.64	—
Constant $p = 0.5$	0.25000	0.25000	—	—	—

Table 2: Benchmark results, sorted by micro-Brier loss.

Preview has the highest micro-Brier in the panel (0.18345) and still beats that accuracy baseline (302/387, 78.04%).

4.2 Confidence and calibration

Define confidence $h_{cu} = \max(p_{cu}, 1 - p_{cu})$ and the high-confidence subset $H_m = \{(c, u) : h_{cu} \geq 0.9\}$ for configuration m . Comparing the mean confidence in H_m with classification accuracy in H_m gives a useful conditional calibration diagnostic. Of course, this is not a full reliability curve; the subset differs by model, and its units remain clustered by case.

OpenAI models appeared more likely than others to make high-confidence forecasts. GPT-5.6 Sol, GPT-5.6 Luna, and Astra make high-confidence forecasts on 58.9%, 47.0%, and 53.5% of units respectively, compared with 20.9% for Fable 5.1. Their mean confidence in those subsets is approximately 95–97%, while realized accuracy is approximately 89–90%, suggesting some overconfidence, particularly for Astra, which missed 21 of 207 such predictions. Gemini 3.8 Flash also had relatively high sharpness, making 276 high-confidence forecasts (71.3% of units) and missing 35. Gemini 3.1 Pro Preview is sharper still, making 320 high-confidence forecasts and missing 57. GPT-6 Sol provides a useful counterexample within OpenAI, missing 8 of 169 high-confidence forecasts, for 95.3% subset accuracy. GPT-6.1 Sol is sharper and less accurate in that subset, missing 18 of 199. GPT-6 Luna misses 19 of 183. Claude Opus 5.5 misses none of 65.

Other high-performing models were less sharp. For example, Fable has no errors among 81 high-confidence units, while Kimi and Grok 4.6 each miss one. Claude Sonnet 5.5 makes 63 high-confidence forecasts and misses 5. The practical finding is that confidence varies substantially and it materially affects proper scores. Astra’s underperformance relative to cheaper models like Sol, for example, appears largely due to several high-confidence misses, which resulted in large Brier losses.




Figure 1: Micro- and equal-case Brier scores for the same seventeen configurations.

4.3 Inference cost

The normalized costs range from \$0.62 for GPT-6 Luna to \$83.18 for Fable. These are not actual invoice charges. For example, Fable was initially run without caching, resulting in an actual cost of \$189.70. The normalized price was reconstructed after the fact using turn-by-turn token counts, with the reconstruction validated against real billing data from other Anthropic model runs. In several other cases, provider discounts that we received were normalized to standard API rates. On these standardized estimates, GPT-6 Luna, GPT-5.6 Luna, and GPT-6 Sol form the observed micro-Brier cost–score frontier.

4.4 Evaluating Jev, a non-autoregressive discriminative decision model

Jev—described by its developer as a probabilistic “System One” model—was evaluated as a separate experiment because its context window and lack of agentic tool calling differ from the document-enabled configurations used with frontier models. [31]. GPT-5.6 Luna summarized the motion documents, without commentary or analysis. Jev received the summaries, and the unit questions, in a single request per case, without document tools or multiple turns. In a second condition, Grok 4.6 prepared a separate set of shorter summaries of the same 482 documents, which Jev received under the same protocol.

Jev was not well-calibrated for this task. Jev’s micro-Brier of 0.36754 is worse than the 0.25 loss of always reporting $p = 0.5$. Its mean predicted dismissal probability was approximately 0.305, while the observed frequency in this data set was 0.736. This resulted in severe underprediction in the aggregate. Its 32.6% threshold accuracy is also below majority-class prediction.

We tested Luna using the same summaries that Jev received in order to determine whether Jev’s underperformance was attributable to receiving only summaries of the litigation documents, rather than the full documents used in our primary evaluation. Luna performed much better than Jev,




Figure 2: Conditional confidence diagnostic on predictions with $\max(p, 1 - p) \geq 0.9$. Each row uses the subset selected by that model, rather than a common set of units. Wrong/selected counts: GPT-6 Sol: 8/169; Fable 5.1: 0/81; Grok 4.6: 1/112; Opus 5: 4/121; GPT-5.6 Sol: 26/228; Kimi K3: 1/92; GPT-5.6 Luna: 20/182; Opus 5.5: 0/65; GPT-6.1 Sol: 18/199; GPT-6 Astra: 21/207; Gemini 3.8 Flash: 35/276; Grok 4.7: 4/98; Sonnet 5: 5/60; GPT-6 Luna: 19/183; Sonnet 5.5: 5/63; Muse Spark 1.3: 15/135; Gemini 3.1 Pro: 57/320. This is not a full reliability diagram or an estimate of statistical significance.

even when it only had access to the same summary Jev received. GPT-5.6 Luna with the identical summaries achieves micro-Brier 0.15489 and accuracy 78.6%, demonstrating that the summaries contained useful predictive signal that Jev was unable to exploit (although Luna performed better in the primary evaluation, where it had full document access, which could mean that the summaries lost some relevant signal or that the interactive environment allows for better performance).

We also tested whether a different summary pipeline would change Jev’s result. With the Grok summaries, Jev improved to micro-Brier 0.32725, equal-case Brier 0.30505, and 38.2% accuracy; the improvement over Jev with Luna summaries is significant on all three metrics (paired case-cluster bootstrap, Bonferroni-corrected across the six pairs of the four summary conditions and three metrics). Jev nevertheless remained significantly worse than both Luna summary controls and worse than the constant $p = 0.5$ forecast; its mean predicted dismissal probability rose only to about 0.356. Because Grok wrote different summaries and no Luna control received them, this compares

Condition	Micro Brier	Case Brier	Acc. (%)
GPT-5.6 Luna full documents	0.14322	0.15084	80.88
GPT-5.6 Luna summary + high reasoning	0.15489	0.16119	78.55
GPT-6 Luna summary + no reasoning	0.17152	0.18694	74.16
Jev Luna summary + one-shot	0.36754	0.34416	32.56
Jev Grok summary + one-shot	0.32725	0.30505	38.24

Table 3: Results from the Jev experiment, with Luna evaluating the same Luna summaries as a control. No control received the Grok summaries.



Figure 3: Probability quality with a shared summary cache. Both Luna conditions exploit predictive signal in the summaries, while Jev scores worse than the constant $p = 0.5$ forecast with either set of summaries. The full-document reference changes both information access and interaction. Connecting segments indicate weighting sensitivity, not uncertainty intervals.

pipelines; it does not isolate summary quality or model ability.

Summary preparation cost an estimated \$10.32 for the Luna summaries and \$29.99 for the Grok summaries, both reconstructed from recorded usage. An earlier Grok preparation under the standard summary prompt cost a further \$2.00, and \$4.11 in uncertain Grok preparation charges is held separately and excluded from these totals. Successful inference estimates are approximately \$0.044 for Jev with Luna summaries, \$0.020 for Jev with Grok summaries (provider-reported), \$1.62 for GPT-5.6 Luna (with high reasoning), and \$0.103 for GPT-6 Luna (which was run without reasoning as a control).

5 Potential advantages of using real-world litigation data as ground truth instead of synthetic data or rubric-based rewards

Using real-world litigation data as ground truth for evaluating frontier models has several advantages over other leading approaches. The leading vertical legal AI company, Harvey AI Corporation, released Legal Agent Benchmark (LAB), in May 2026. Now widely cited by frontier research labs as a key evaluation of agentic legal capabilities for frontier models, LAB uses AI-generated task environments to test models on pre-specified rubrics. Harvey has publicly disclosed that it works with Mercor and Snorkel, leading third-party data providers, to conduct expert review and verification of every generated LAB task [25]. Harvey open-sourced the benchmark (except for a held-out test set), and invited practicing lawyers to validate and improve it by reviewing existing tasks and auditing rubrics. [26] I took up that invitation and did a deep-dive review of one of the substantive litigation tasks (drafting an issues list for a motion to dismiss opposition brief) and an audit of all the litigation rubrics. This analysis underscored the challenges involved in creating suitable reinforcement learning environments that can advance the capabilities frontier in the legal domain—and illustrated why extracting signal from real-world litigation data could be more effective than existing approaches.

5.1 Audit of LAB’s litigation rubrics

Several LAB rubrics appeared to contain anomalous results, so I conducted an audit of all rubrics on litigation and dispute resolution tasks (my area of specialty for 18 years). I supplied GPT-6 Sol and Claude Opus 5.5 with examples of the kinds of rubrics that I consider defective and asked them to iterate through each task to audit the rubrics using the high-level standards I supplied (which allowed them to raise issues that were not specifically on my list).

Of the 52 tasks and 2,858 rubrics included in the litigation and dispute resolution category, GPT-6 Sol flagged 616 criteria as defective or arguably defective, Claude Opus 5.5 flagged 564 (with the ability to review GPT-6 Sol’s findings). Between them, the models flagged 835 criteria, leaving 2,023 unflagged by either report. Of 345 overlapping flags, both reviewers called 83 problematic and 135 arguable; 127 received different levels of concern. Opus was allowed to see Sol’s findings when reaching its own conclusions, so agreement in the final reports is not independent confirmation [16].

To validate these findings and ascertain a statistically valid estimate of the number of rubrics that are defective, I reviewed a simple random sample of 25 of the 835 flagged criteria, drawn with a fixed seed from 19 tasks (Appendix C). I treated a rubric as defective if, applied as written to the task as given, it would fail a response a careful litigator would consider correct or pass one that is wrong. 16 of the 25 were defective, and I assessed 9 of those to be unambiguous objective errors where the accuracy of the rubric did not depend on my exercise of professional judgment but there was indisputably an error or problem in the task of some kind. Another 3 were arguable and 6 were not defective.

These results demonstrate that at least 383 of the 835 flagged criteria are defective (one-sided 95% bound; point estimate 534), or more than 13.4% of all 2,858 criteria in the practice area; at least 171 are unambiguous objective errors. Agreement between the two AI auditors was informative: 9 of the 10 sampled criteria flagged by both were defective, against 7 of the 15 flagged by only one.

Examples of defective rubrics, drawn from the sampled criteria in Appendix C, included instances where the rubric misstates the law (e.g., requiring the answer that is the opposite of settled law in the jurisdiction at issue), mathematical or numerical errors (e.g., the task environment contained mathematically impossible documents, or contained a rubric that required models to assume that one of two conflicting numbers was accurate while ignoring the conflicting number), or instances where I assessed the rubric’s recommended approach to contravene widely respected professional norms.

I separately counted defects in the task environment relevant to the rubric (as opposed to defects that made the rubric’s grading criteria themselves incorrect). 14 of the 25 criteria related to defective task environments, for example, documents that did not sum to the stated figures; these span 10 of the 19 sampled tasks, and only 4 of the 25 criteria had neither a defective nor an arguable criterion and came from an environment without a defect. This means that at least 319 flagged criteria (more than 11.1% of all criteria) come from such environments.

These errors are highly significant relative to the margins that separate frontier models on LAB’s scoring criteria. On Artificial Analysis’s Harvey LAB-AA leaderboard, Muse Spark 1.3 leads with a criterion pass rate of 95.5%, followed by Kimi K3 at 94.6%, Claude Fable 5.1 at 93.3%, Claude Opus 5.5 at 91.2%, and GPT-5.6 Luna at 87.9% [34]. If criteria elsewhere in LAB are defective at anything like the rate we found (more than 13.4% of the criteria in this practice area), the rate of defective rubrics is as large as or larger than the reported gaps between leading models on criterion pass rate, and it is unclear how much of those gaps the defects could explain. All-pass scoring

compounds the problem, because a single defective criterion can fail an otherwise acceptable task.

The flaws in these environments raise questions about whether general-purpose data vendors have the qualifications or ability to prepare verified and accurate RL environments for legal tasks. Harvey retained leading data vendors to review and verify the tasks and rubrics [25]. But the verification services offered by these third-party vendors did not reliably detect defects like rubrics that misstated the law. These general-purpose data vendors provide a variety of different kinds of reinforcement learning environments, many of them for nonlegal tasks, and the companies utilized do not purport to be specialized in legal work. Some companies in this space hire freelance lawyers at contract attorney rates to verify tasks and rubrics. At an earlier phase of frontier model development, that approach may have provided useful signal. But as the capabilities of frontier models have improved, and benchmarks have become more ambitious, the results here suggest that this kind of freelancer data could have defects that end up undercutting the evaluation or training signal the dataset is intended to provide. Even sophisticated buyers doing frontier AI research in the legal domain may find it difficult to catch these errors without effectively reproducing much of the work that the data vendors are paid to do. This suggests that further improvement in the legal domain might require efforts to increase the quality of the data used to evaluate and train models on legal tasks.

The challenges involved in developing and maintaining acceptable benchmarks for legal tasks call to mind some of the challenges frontier labs have had developing new benchmarks for coding tasks. In February 2026, OpenAI publicly criticized SWE-bench Verified (which OpenAI had helped develop), noting that its audit revealed problems involving test requirements and task specification [30]. The rapid capability growth has meant that errors that once might have been unnoticeable now show up in the form of frontier models failing tasks incorrectly because the model is smarter than the benchmark. That dynamic could require greater investment in evaluations to keep pace with the strength of the underlying models.

5.2 Cite Check of LAB’s Task Environment and Rubrics

Because the rubric audit turned up instances where the rubrics contained hallucinated case law (suggesting that the external experts who were responsible for verification may have relied on inaccurate AI-generated responses to prepare the answer key), I asked AI agents to conduct a separate citation check of all the litigation and dispute resolution tasks to assess the frequency of AI hallucinations in the task environments or rubrics. A citation check is a standard pre-filing procedure where lawyers check every authority cited in a brief, regulatory submission, or other important document to ensure that the authority is cited accurately, contains any quotation attributed to it, and supports the legal proposition for which it is cited.

A separate AI-assisted check of every case citation in the 19 sampled environments against CourtListener, with spot checks, found problems in 45 of 102 citations to purported real authority: 32 real cases cited for holdings, quotations, or facts they do not contain, 7 cases that could not be found, 5 garbled citations to real cases, and 1 citation pointing to a different case [17].

After excluding errors in case law that were intentionally planted and then tested in rubrics (e.g., an error of law in an opponent’s submission that the model was supposed to catch), the AI identified 436 errors that were apparently not intended to be part of the task environment [18]:

[[Placeholder: To add results of attorney audit of the results, the hallucinated cases issue was identified only on 10/4/2026, the day before our original intended publication date. We delayed publication to understand how widespread the issue was, and initial spot checks suggest there are a meaningful number of hallucinated citations in the task environments so we need to present that

Finding	All citations	Case citations
Verified	2,754	306
Verified, minor defect (pin cite, loose paraphrase, wrong year)	346	54
Intentional error tested by the rubric (excluded)	31	6
<i>Unintentional errors</i>		
Cited for a proposition it does not support	246	105
Not found (likely fabricated)	76	60
Garbled citation to a real authority	58	10
Quotation not found in the authority	40	31
Reporter citation belongs to a different case	16	12
Total unintentional errors	436	218
Source text not accessible	51	6
Total checked	3,618	590

Table 4: AI citation check of the 52 LAB litigation task environments. This table shows the AI findings.

with some validation so the issue is properly quantified.]]

5.3 A motion-to-dismiss case study

In addition to auditing the accuracy of all 2,858 rubrics, I did a deep dive review of one substantive litigation task. This review highlighted a different but equally important issue: the use of synthetic or AI-generated data may result in task environments that are unrealistic in a variety of subtle ways that may not be obvious to nonexperts but that are striking and clear to practicing attorneys. The task `analyze-counterparty-motion-to-dismiss` asks AI agents to prepare a list of issues for the opposition brief to a motion to dismiss, on a fictitious software service dispute.

I observed several ways in which the synthetically generated task environment seemed unrealistic and could potentially train models on a distorted view of what genuine practice looks like. [19]

Planted issues. A number of the legal disputes in the task environment are obviously “planted” in a way that feels artificial and ultimately presents issues in a way that is meaningfully different from how real-world legal disputes manifest. For example, the contract the parties are suing over has a forum selection clause that requires any lawsuit to be filed in Travis County, Texas. The plaintiff literally quotes this entire provision in the complaint but nevertheless files suit in the Northern District of Georgia and offers no explanation for why they believe the forum selection clause they quoted does not apply or otherwise permits the suit in Georgia. It of course happens regularly that a lazy or incompetent litigant files suit in the wrong forum; a competent lawyer potentially could decide to ignore the clause and hope that their adversary overlooks the issue. But in both of those scenarios, the forum selection clause would not be explicitly flagged in the complaint as it is here; a real version of this legal task would require the lawyer to dig through a number of documents and realize that buried in one of them is a forum selection clause that allows the lawyer to have the case moved to Texas.

This is one example, but a significant number of the tasks in these environments were similarly planted in a way that’s artificial and unlike how disputes around these issues arise in real cases. Environments like this can distort an evaluation of a model’s true capabilities, since simply looking for the obviously planted issues will satisfy the benchmark without testing whether the model would

be able to identify this issue from more realistic source material akin to what it will see in real cases.

Environment Documents Violate Court Rules. Many of the documents in these environments do not resemble real lawyer work product. The environment provides counterparty litigation documents in docx format, and the native files from the git repository often cannot be opened, although agents and various CLI tools can recover the document text. In real litigation, lawyers rarely share their docx files with their opponent; the opponent’s memorandum of law is ordinarily filed with the court as a PDF. The motion papers also lack a proper caption, include a stale table of contents, contain markdown formatting artifacts that render in Word as junk characters even in files that can be opened normally, lack citations for propositions that generally require citations, and are single spaced rather than double spaced in violation of the applicable court rules.

The motion that the agent is being asked to respond to is so far below professional norms that you would want any analysis of the motion to include consideration of whether the judge would strike the motion completely, or potentially discipline or sanction the counsel that filed that deficient work product. That is not included in the rubric for this task because the quality control issues appear to be unintentional, not part of the test for the agent. But it is a genuine gap: if a document like this were filed in real litigation, the quality of the work product could itself become a substantive issue in the litigation and thus would need to be flagged given the goal motivating the all-pass criterion, which was to ensure that the AI does not overlook any important issues [26, section “Evaluating a task”].

There are valid reasons to prefer rubric-based reinforcement learning approaches. For example, while a benchmark like LegalForecastBench can test a model’s legal reasoning and analytical ability over a narrow task, it does not test the model’s ability to generate new work product, which is an important consideration for a company like Harvey that is developing legal agents to prepare and revise different kinds of legal work product. But reviewing these environments suggests that the current state of the art in synthetic data generation introduces drift on a variety of dimensions that could undermine the efficacy of the task environment as a training or evaluation signal. The model is in effect being tested on its ability to recognize a distinct environment, very unlike what it will encounter in production, and adhere to a stylized set of norms and rules that were artificially constructed in that environment but do not comport with the real norms or rules that govern real-world tasks. While creating LegalForecastBench illustrated that there are real challenges in working with real-world litigation data as well, for a variety of tasks that data has a lot of rich, implicit signal that data providers generating synthetic data are unable to capture.

6 Conclusion

LegalForecastBench is one illustration of how real-world litigation data can be used as a reward signal with an objective ground truth to evaluate frontier models on high-level legal analytical tasks. Next steps in my research include testing whether these signals are also useful as training; some early results suggest that merely training on the Brier score loss function may not be a sufficiently dense reward signal to improve model performance, but we are testing techniques that might exploit more of the judge’s written reasoning to create a richer task, where the model is asked to make several predictions, to create potentially denser rewards.

The experiment also illustrated that AI research has a long way to go to bridge the gap between current training and evaluation environments used to train models, and the messy real-world tasks that models are asked to perform. LAB is notable as the first large-scale, open-source set of legal

task environments that try to test high-level legal judgment tasks where there is often no objective ground truth. But LAB illustrates the challenges of creating effective environments, because even leading data vendors that took on the responsibility of expert verification of the LAB tasks were unable to create realistic task environments or accurate rubrics. Harvey has explicitly invited practicing attorneys to help collaborate on improving LAB, and building models that can automate more of the high-level legal work that lawyers do clearly will require careful collaboration to obtain training signal that can reliably improve the capabilities of, and distinguish between the performance of, increasingly capable frontier models.

The good news is that this experiment suggests that there are still lots of sources of training signal that have not been exploited and that could meaningfully improve the evaluation of, and performance of, frontier models.

Reproducibility, access, and AI use disclosure

Code is maintained in the [LegalForecastBench repository](#); documents used in the benchmark evaluation, including the documents I purchased in order to complete the evaluation, have been made available to the public through [CourtListener](#). I am also happy to provide copies of the litigation documents as they were shown to models to interested researchers.

As I am a practicing attorney and do not have any formal training in artificial intelligence or machine learning, the use of AI models was essential to allow me to make a research contribution in this space that otherwise would have been beyond my capabilities. My personal contribution included the benchmark design, auditing a sample of the LAB rubrics by hand to statistically validate the AI audit findings, labeling cases for the benchmark in instances where frontier models reached differing conclusions about the correct label, conducting a deep dive into one of LAB’s litigation tasks, and orchestrating AI coding agents that helped write and execute the code that ran the benchmark evaluation, starting on May 18, 2026, and continuing through at least October 6, 2026. I wrote most of this paper personally, but discussed the results at several points with AI agents and considered AI-generated analysis as an input into my thinking as I drafted. The major exception to this was sections 3 and 4 and other places where numerical results or mathematical notations are included; I relied on AI agents to write code that performs the necessary mathematical calculations and to prepare the LaTeX source summarizing those mathematical results. I used several different models from Anthropic, OpenAI, SpaceXAI, and Google to check the work. I also met with experienced researchers in the space to validate the experimental design and results, and obtain feedback on the draft. I am grateful to several AI researchers in industry positions for helpful discussions, and to Prof. Dr. Daniel Katz, professor of law at Illinois Tech’s Chicago-Kent College of Law, for reviewing an early draft and discussing relevant findings from his prior Supreme Court prediction research. This version of the paper is intended as a working draft, and feedback from researchers, lawyers, and others on how to improve this benchmark or make our findings more useful is welcome.

References

- [1] Daya Guo, Dejian Yang, Haowei Zhang et al. DeepSeek-R1 incentivizes reasoning in LLMs through reinforcement learning. *Nature* 645, 633–638 (2025); *arXiv:2501.12948v2*, 2025. <https://www.nature.com/articles/s41586-025-09422-z>.
- [2] Zhiyuan Zeng, Hamish Ivison, Yiping Wang et al. RLVE: Scaling Up Reinforcement Learning for Language Models with Adaptive Verifiable Environments. *ICML 2026, Proceedings of Machine Learning Research* 306:153406–153428, 2026. <https://proceedings.mlr.press/v306/zeng26j.html>.
- [3] David Silver, Thomas Hubert, Julian Schrittwieser et al. Mastering Chess and Shogi by Self-Play with a General Reinforcement Learning Algorithm. *arXiv:1712.01815* (2017), preprint; precursor to the distinct 2018 *Science* article, 2017. <https://arxiv.org/abs/1712.01815>.
- [4] Russell Yang, Ruishi Chen, Pierce Kelaita et al. JudgmentBench: Comparing Rubric and Preference Evaluation for Quality Assessment. *arXiv:2605.25240v3*, preprint, 2026. <https://arxiv.org/abs/2605.25240>.
- [5] Niko Grupen, Julio Pereyra, Gabe Pereyra et al. Training a Legal Agent with Applied Compute. *Primary online source*, 2026. <https://www.harvey.ai/blog/training-a-legal-agent-with-applied-compute>.
- [6] Calvin Qi, Vasudha Rengarajan, Julio Pereyra et al. Harvey Tenet Research Preview. *August 20, 2026 company research report*, 2026. <https://www.harvey.ai/blog/post-training-update-harvey-tenet>.
- [7] Neel Guha, Julian Nyarko, Daniel E. Ho et al. LegalBench: A Collaboratively Built Benchmark for Measuring Legal Reasoning in Large Language Models. *NeurIPS 2023, Datasets and Benchmarks*, 2023. <https://arxiv.org/abs/2308.11462>.
- [8] Neel Guha. Legal LLM Leaderboard: Tracking LLM Performance on Academic Legal Benchmarks. *Public evaluation dashboard; results generated April 8, 2026, accessed October 6, 2026*. BarExam: 117 publicly released Multistate Bar Examination questions. <https://neelguha.github.io/legal-benchmark-dashboard/>.
- [9] Glenn W. Brier. Verification of Forecasts Expressed in Terms of Probability. *Monthly Weather Review* 78(1) (1950), 1–3, 1950. https://journals.ametsoc.org/view/journals/mwre/78/1/1520-0493_1950_078_0001_vofeit_2_0_co_2.xml.
- [10] Tilmann Gneiting, Adrian E. Raftery. Strictly Proper Scoring Rules, Prediction, and Estimation. *Journal of the American Statistical Association* 102 (2007), 359–378, 2007. <https://doi.org/10.1198/016214506000001437>.
- [11] Cornerstone Research and Stanford Law School Securities Class Action Clearinghouse. Securities Class Action Filings—2024 Year in Review. *Research report*, 2025. p. 16, status of core federal securities class action filings as of the end of 2024. <https://securities.stanford.edu/research-reports/1996-2024/Securities-Class-Action-Filings-2024-Year-in-Review.pdf>.
- [12] Robert H. Mnookin and Lewis Kornhauser. Bargaining in the Shadow of the Law: The Case of Divorce. *Yale Law Journal*, 88(5):950–997, 1979. <https://doi.org/10.2307/795824>.

- [13] Benjamin Turtel, Danny Franklin, Kris Skotheim et al. Outcome-based Reinforcement Learning to Predict the Future. *Transactions on Machine Learning Research (November 2025)*; *arXiv:2505.17989v4*, 2025. <https://arxiv.org/abs/2505.17989>.
- [14] Nikhil Chandak, Shashwat Goel, Ameya Prabhu et al. Scaling Open-Ended Reasoning To Predict the Future. *arXiv:2512.25070v2 (January 2026)*; *preprint*, 2026. <https://arxiv.org/abs/2512.25070>.
- [15] Junjie Chen, Haitao Li, Qilei Zhang et al. Multi-Source Retrieval and Reasoning for Legal Sentencing Prediction. *arXiv:2602.04690v3*; *arXiv record comments 'EMNLP 2026 main' (publisher proceedings record not independently checked)*, 2026. <https://arxiv.org/abs/2602.04690>.
- [16] GPT-6 Sol audit workers, Claude Opus 5.5 audit workers, Coordinating agents; individual provenance not fully recorded. Harvey LAB litigation and dispute resolution audit: GPT-6 Sol and Claude Opus 5.5. *Local research artifact, pinned September 29, 2026*, 2026. <https://github.com/johnhughes3/LegalForecastBench/tree/58371a0e13494e591710e2ae3da8e428eebc359/docs/harvey-lab-audit/litigation-dispute-resolution>.
- [17] John Hughes. Human review of AI-flagged Harvey LAB litigation criteria: protocol, sample, and worksheet. *Research artifact*, 2026. <https://github.com/johnhughes3/LegalForecastBench/tree/main/docs/harvey-lab-audit/litigation-dispute-resolution/human-review>.
- [18] Claude Sonnet 5.5 and Claude Opus 5.5 citation-check agents, coordinated by the author. Citation check of the Harvey LAB litigation and dispute resolution task environments. *AI-generated research artifact, not attorney-reviewed*, 2026. <https://github.com/johnhughes3/LegalForecastBench/tree/main/docs/harvey-lab-audit/litigation-dispute-resolution/citation-check>.
- [19] Harvey AI task authors; upstream authorship not individually listed. Pinned LAB task.json for analyze-counterparty-motion-to-dismiss. *Primary online source*, 2026. <https://github.com/harveyai/harvey-labs/blob/1dd81403b2fbb60596f7aea3fcecfaad7bf73143/tasks/litigation-dispute-resolution/analyze-counterparty-motion-to-dismiss/task.json>.
- [20] Theodore W. Ruger, Pauline T. Kim, Andrew D. Martin, Kevin M. Quinn. The Supreme Court Forecasting Project: Legal and Political Science Approaches to Predicting Supreme Court Decision-Making. *Columbia Law Review* 104 (2004), 1150–1210, 2004. [Author-hosted article PDF](#).
- [21] Daniel Martin Katz, Michael J. Bommarito II, Josh Blackman. A General Approach for Predicting the Behavior of the Supreme Court of the United States. *PLOS ONE* 12(4): e0174698 (2017), 2017. <https://journals.plos.org/plosone/article?id=10.1371%2Fjournal.pone.0174698>.
- [22] Nikolaos Aletras, Dimitrios Tsarapatsanis, Daniel Preotiuc-Pietro, Vasileios Lampos. Predicting Judicial Decisions of the European Court of Human Rights: A Natural Language Processing Perspective. *PeerJ Computer Science* 2:e93 (2016), 2016. <https://peerj.com/articles/cs-93/>.
- [23] Chaojun Xiao, Haoxi Zhong, Zhipeng Guo et al. CAIL2018: A Large-Scale Legal Dataset for Judgment Prediction. *arXiv:1807.02478 (2018)*; *CAIL challenge dataset paper*, 2018. <https://arxiv.org/abs/1807.02478>.

- [24] Huiyuan Xie, Felix Steffek, Joana Ribeiro de Faria et al. The CLC-UKET Dataset: Benchmarking Case Outcome Prediction for the UK Employment Tribunal. *Proceedings of the Natural Legal Language Processing Workshop (ACL)*, 2024, pp. 81–96, 2024. <https://aclanthology.org/2024.nllp-1.7/>.
- [25] Gabe Pereyra. How Harvey Built a Research Lab on a Budget. *Sequoia Capital, Own Your Intelligence*, August 11, 2026. Slide at 4:11: “Scale with expert data providers.” <https://www.youtube.com/watch?v=MGouk8W51v0&t=251s>.
- [26] Niko Grupen, Gabe Pereyra, Julio Pereyra. Introducing Harvey’s Legal Agent Benchmark. *Primary online source*, 2026. <https://www.harvey.ai/blog/introducing-harveys-legal-agent-benchmark>.
- [27] Andy Zou, Tristan Xiao, Ryan Jia et al. Forecasting Future World Events with Neural Networks. *NeurIPS 2022, Datasets and Benchmarks*, 2022. <https://arxiv.org/abs/2206.15474>.
- [28] Danny Halawi, Fred Zhang, Chen Yueh-Han, Jacob Steinhardt. Approaching Human-Level Forecasting with Language Models. *ICML 2024, Proceedings of Machine Learning Research 235*, 2024. <https://arxiv.org/abs/2402.18563>.
- [29] Ezra Karger, Houtan Bastani, Chen Yueh-Han et al. ForecastBench: A Dynamic Benchmark of AI Forecasting Capabilities. *ICLR2025*, 2025. <https://arxiv.org/abs/2409.19839>.
- [30] OpenAI. Why SWE-bench Verified no longer measures frontier coding capabilities. *Primary online source*, 2026. <https://openai.com/index/why-we-no-longer-evaluate-swe-bench-verified/>.
- [31] TypeSafe AI. Models: Jev and the System One Interface. *Product documentation*, accessed September 29, 2026, 2026. <https://docs.typesafe.ai/models>.
- [32] George L. Priest, Benjamin Klein. The Selection of Disputes for Litigation. *Journal of Legal Studies* 13 (1984), 1–55, 1984. <https://www.journals.uchicago.edu/doi/10.1086/467732>.
- [33] Lucian A. Bebchuk. Litigation and Settlement Under Imperfect Information. *RAND Journal of Economics* 15(3) (1984), 404–415, 1984. <https://hls.harvard.edu/bibliography/litigation-and-settlement-under-imperfect-information/>.
- [34] Artificial Analysis. Harvey LAB-AA evaluation results. *Leaderboard*, accessed October 5, 2026, 2026. <https://artificialanalysis.ai/evaluations/harvey-lab-aa>.
- [35] Tak K. Mak. Analysing Intraclass Correlation for Dichotomous Variables. *Journal of the Royal Statistical Society, Series C (Applied Statistics)*, 37(3):344–352, 1988. <https://doi.org/10.2307/2347309>.

A Model configuration and temporal evidence

“K” denotes a reported knowledge cutoff, “T” a reported training cutoff, and “?” no established cutoff. Month-only dates are compared as the first day of the reported month. The reported cutoffs for fifteen configurations precede June 30, 2026. Kimi and Muse are treated as “qualified,” meaning we cannot be certain that those models were not trained on the decisions at issue (but based on the release dates that appears unlikely).

Configuration	Reasoning	Reported cutoff	Relation to window
GPT-6 Sol	High	2026-04-20 (K)	Precedes window
Claude Fable 5.1	Adaptive/default	2026-06 (T)	Precedes window
Grok 4.6	High	2026-02-01 (K)	Precedes window
Claude Opus 5	Adaptive/default	2026-05 (T)	Precedes window
GPT-5.6 Sol	High	2026-02-16 (K)	Precedes window
Kimi K3	High	Unknown (?)	Unknown
GPT-5.6 Luna	High	2026-02-16 (K)	Precedes window
Claude Opus 5.5	High	2026-06 (T)	Precedes window
GPT-6.1 Sol	High	2026-04-30 (K)	Precedes window
GPT-6 Astra	High	2026-04-30 (K)	Precedes window
Gemini 3.8 Flash	High thinking	2026-03 (K)	Precedes window
Grok 4.7	High	2026-05 (K)	Precedes window
Claude Sonnet 5	Adaptive/default	2026-01 (T)	Precedes window
GPT-6 Luna	High	2026-05-18 (K)	Precedes window
Claude Sonnet 5.5	High	2026-06 (T)	Precedes window
Muse Spark 1.3	High	Unknown (?)	Unknown
Gemini 3.1 Pro Preview	High thinking	2025-01 (K)	Precedes window

Table 5: Configuration and cutoff evidence retained by the benchmark. “Precedes window” is a date comparison under the maintained rule, not an independently verified training-history claim.

Claude Fable 5.1, Claude Opus 5, and Claude Sonnet 5 used adaptive thinking at provider defaults. Claude Opus 5.5 and Claude Sonnet 5.5 used high effort. The other configurations use the high reasoning or thinking setting.

B Result provenance

The primary cohort identifier is `cycle-1-91-2026-09-08-luna-r5`. Costs in the main table are standardized estimates from recorded usage and frozen rates; Fable’s estimate reconstructs cached execution from turn-by-turn usage rather than its original uncached charges, as described in Section 4.3.

The checked source revisions are [LegalForecastBench 58371a0e](#), [LegalForecastCorpus b392ddaf](#), and [LegalForecastGym 4cc0f528](#). The LAB source snapshot is [1dd81403](#). These identify inspected versions; they do not certify every execution against the current source tree.

C Human review of AI-flagged LAB criteria

After noticing instances in which LAB rubrics appeared to contain errors, I asked two frontier models, GPT-6 Sol and Claude Opus 5.5, to audit the 52 tasks in LAB’s litigation-and-dispute-resolution

practice area at Harvey commit 1dd81403. The goal was to identify rubric criteria that misstate the law, and environments or task specifications containing errors or deficiencies that would make the tasks unreliable for benchmarking or for training models. Harvey has used LAB for both [5, 6]; [confirm the source for public tasks used in training and a held-out set used for benchmarking](#).

The models flagged 835 of the 2,858 criteria (29.2%) as defective or arguably defective. GPT-6 Sol flagged 616. Claude Opus 5.5, working without access to Sol’s results, flagged 650; it then reviewed every flag from both passes against the record and the governing law, ending with 564. The 835 figure is the union of Sol’s flags and Opus’s final flags [16].

In order to then establish how many of the AI-generated flags reflected genuine defects, I reviewed a simple random sample of 25 of the 835 flagged criteria, drawn with a fixed seed under a protocol committed before the review began [17]. The sample spans 19 tasks. For each criterion I examined the task instructions, the rubric, both auditors’ findings, how LAB’s AI judges graded two model runs on it, and the relevant documents (and in a few cases the relevant governing law). I treated a criterion as defective when, applied as written to the task as given, it would grade a competent response wrong: failing a response a careful litigator would consider correct, or passing one that is wrong. During the review, I adopted two applications of that standard: a criterion that requires one of several permissible approaches is defective when it rejects an approach a careful litigator would adopt given that record. And an unplanted error in the supplied record, such as figures that do not add up or misdescribed case law, makes a criterion defective when it changes what a competent response can get right. (For example, in one case, a rubric wanted work product to contain a specific metric sourced from the task environment, but during the review, I discovered that the metric appeared in two different places with different numbers for each, seemingly a mistake in the task environment rather than an intentional inconsistency that was part of the test.) Separately, I recorded whether each task environment itself contains a defect, even if that defect did not affect the sampled criterion. (For example, one task contained a notice of regulatory violations, but the notice emanated from a state-level regulatory agency that, in the specific state at issue, does not have jurisdiction over the violations charged.)

Of the 25 criteria, I found 16 defective in ways that affect grading, 9 of them unambiguous objective errors in which the record, the law, the cited authority, or the arithmetic contradicts the rubric. 3 were arguable defects, and 6 were not defective. 14 of the 25 come from task environments that contain defects (10 of the 19 tasks). Only 4 had neither a defective nor an arguable criterion and came from an environment without a defect (Table 6).

Treating the sample as a draw without replacement from the 835 flagged criteria, the exact one-sided 95% lower confidence bounds are: at least 383 flagged criteria are defective in ways that affect grading (point estimate 534; the bound is 13.4% of all criteria in the practice area); at least 171 are unambiguous objective errors (point estimate 301); and at least 319 come from task environments that contain defects (point estimate 468). At least 524 involve a defective criterion, a defective environment, or both (point estimate 668). In the sample, 9 of the 10 criteria flagged by both auditors were defective, against 7 of the 15 flagged by only one.

These estimates cover only criteria the AI auditors flagged; defects among unflagged criteria are not estimated. The estimates count criteria/rubrics, not each individual mistake; one error in a task environment can make several criteria defective. After I prepared an initial set of verdicts, based on my professional judgment as a litigator, I asked AI models to stress test my answers and to provide me with additional arguments both for and against each call. I changed a handful of verdicts upon further reflection, and did a second pass through my verdicts to ensure that the principles I adopted

during the review were applied consistently to each call.

The audit of rubrics revealed instances where the task environment, and sometimes grading rubrics, appeared to rely on AI-generated hallucinated case law citations. As a result, I had AI agents (Opus 5.5) conduct a separate AI-assisted check of every case citation in the 19 sampled task environments against CourtListener. This utilized an ultracode workflow where a separate agent reviewed each task, with multiple passes by separate agents to verify findings. This process found problems in 45 of 102 citations to real authority, including 7 cases that could not be found and 32 real cases cited for propositions they do not contain [17].

Criterion verdict	Environment defective	Environment correct	Total
Defective	10	6	16
Arguable	2	1	3
Not defective	2	4	6
Total	14	11	25

Table 6: My verdicts on the 25 sampled criteria, by whether the task environment itself contains a defect.

The sampled criteria

Each entry has four parts: (1) the criterion quoted verbatim from Harvey’s task file (that text is from Harvey’s data vendors, not me, and in many cases the text appears to be AI-generated), (2) my verdict on the criterion and task environment, (3) results from a sample run of LAB using GPT-6 Luna and Claude Opus 5.5, to assess how LAB’s AI graders (Claude Sonnet 4.6 and GPT-5.5) graded outputs from Luna and Opus on the specific criteria at issue, (4) my explanation of the verdict. Harvey task files are cited at commit 1dd81403.

1. Assess Reasonableness of Staffing Levels on Litigation Invoice (C-049)

QUOTED FROM HARVEY LAB

Calculates budget overage percentage as approximately 17%. PASS if the memo calculates the percentage by which the attorney fee subtotal (\$193,372.50) exceeds the approved budget ceiling (\$165,000) as approximately 17% (exact figure is 17.2%). FAIL if no percentage overage calculation is provided or if the percentage stated is materially incorrect (off by more than 2 percentage points).

— Harvey LAB, criterion C-049, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — unambiguous objective error, deficiency in task environment
Task environment	Defective: the invoice’s 117 time entries add up to \$215,680.00, while its summary and subtotal show \$193,372.50.
AI auditors’ reasoning	Correct

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric is testing the AI on a specific percentage calculation from an invoice, but

the document in the environment is internally inconsistent. The invoice’s summary page and its subtotal row show fees of \$193,372.50 (17.2% over the \$165,000 budget), but the 117 detailed time entries add up to \$215,680.00, which puts the overage at \$50,680, or 30.7% (13.5 points from the rubric’s required 17.2% answer and far outside the 2-point tolerance).

2. Categorize Document Production Set by Relevance and Privilege — Privilege Log and Relevance Classification Report (C-012)

QUOTED FROM HARVEY LAB

ISSUE_003: DOC_012 identified as Partially Privileged. PASS if DOC_012 (chandrasekaran-meeting-notes-aug28.docx, Chandrasekaran’s handwritten meeting notes from August 28, 2023) is classified as Partially Privileged. FAIL if classified as fully Privileged or Not Privileged.

— Harvey LAB, criterion C-012, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — low-quality professional judgment, rubric is incorrect
Task environment	Correct
AI auditors’ reasoning	Correct

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric is wrong because it fails fully Privileged, which is the better answer here (even if Partially Privileged is defensible).

- The documents are handwritten lawyer notes from a fictitious meeting discussing problems with the R&D Department. The headings on individual topics seem intended to create the impression that there is a discussion of business topics (which would not be privileged) followed by a separate discussion of legal topics (privileged), so from a superficial glance, the rubric does seem to match the topic headings. However, the bullets under each heading make clear that the entire meeting is just dedicated to the topic of the company’s desire to find a way to fire the head of R&D without getting sued. The business topics (which look to be nonprivileged) are discussions of how the R&D department is a bit of a disaster, such that management wants to fire the head. The legal sections then involve discussion of the fact that the R&D head has some sort of environmental whistleblower status, so termination is likely to trigger a lawsuit/whistleblower retaliation claim.
- The meeting was specifically convened because outside counsel was concerned about the retaliation claim, and even some of the ‘ops’ sections that are supposed to be nonprivileged reflect counsel’s advice, so they cannot be cleanly treated as nonlegal business discussions. For example, the ‘Ops’ sections record HR “compiling the file” of performance documentation and lay out a “Performance improvement plan first” / “Documentation period” sequence, which carries out outside counsel’s recommendation, recorded in the legal section, to build a “contemporaneous performance record to insulate against retaliation” claims.
- While partially privileged might be defensible, my preferred answer would be fully privileged, which the rubric incorrectly fails.

3. Compare Document Production Against Discovery Requests — Discovery Gap Analysis Memorandum (C-029)

QUOTED FROM HARVEY LAB

ISSUE_011: Notes only 5 documents actually relate to counterclaim damages.

PASS if the memorandum notes that only 5 of the 14 documents coded to RFP 27 (MLC-004056 through MLC-004102) actually relate to Pacific Corridor’s volume shortfall and the counterclaim damages calculation, meaning the counterclaim damages support is thin. FAIL if this is not noted.

— Harvey LAB, criterion C-029, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Defective: the production index codes 15 rows to RFP 27, but its summary tab gives a count of 14 for a range that contains 13.
AI auditors’ reasoning	Correct

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. This task requires models to make the point that “only” five documents in the bundle “actually relate to Pacific Corridor’s volume shortfall and the counterclaim damages calculation, meaning the counterclaim damages support is thin.” This rubric is flawed on multiple levels.

- First, it overlooks MLC-003781, a sixth document that plainly relates to the volume shortfall, so the count is objectively wrong (and the agents pointed to several other examples, at least some of which also seem to fit within the stated criteria). There is a document in the bundle that claims that only five documents concern the volume shortfall, but that assertion is undercut by other evidence supplied in the environment.
- Second, this rubric is someone’s argumentative characterization of the documents and probably would not hold up under adversarial testing. For example, there is case law that the term “relate to” is fairly broad and flexible, so quite plausibly a number of other documents would apply.
- Third, the premise underlying this rubric—that “only” five documents is such “thin” evidence that it must be called out—is wrong. Even one document could constitute conclusive proof of damages (the general rule is that damages only need to be reasonably estimated; courts do not require great precision); to my mind, this is actually quite a lot of evidence at the fact stage compared to what one would normally see.

4. Draft Case Assessment Memorandum — Litigation Risk Analysis for Distribution Agreement Dispute (C-025)

QUOTED FROM HARVEY LAB

ISSUE_007: Recommends motion to dismiss on unjust enrichment. PASS

if the memo recommends filing a motion to dismiss (or motion for judgment on the pleadings, or similar dispositive motion) targeting the unjust enrichment claim (Count IV) on the ground that an express contract governs the parties’ relationship. FAIL if no

such motion recommendation is made regarding the unjust enrichment claim.

— Harvey LAB, criterion C-025, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — low-quality professional judgment, rubric is incorrect
Task environment	Correct
AI auditors’ reasoning	Correct

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. A different rubric on this task concludes that plaintiff has a strong breach of contract claim. (See id: C-006.) The prompt instructs models to review case documents and “prepare a comprehensive case assessment memo,” and models are dinged on this rubric if they do not “recommend[] filing a motion to dismiss” on the unjust enrichment claim, which is expressly pled in the alternative to the (strong) breach of contract claim.

- The rubric is wrong because a careful litigator could reasonably conclude that a motion to dismiss that claim is not a good idea: Count IV is pleaded in the alternative (Complaint ¶127) and seeks restitution of “no less than \$22,521,312.50” (¶135), which is exactly the Count I contract damages (\$17,721,312.50, ¶94) plus the Count III trade-secret damages (\$4,800,000, ¶120). Counts I and III would survive the motion and double recovery is barred, so dismissing Count IV would not reduce the defendant’s net exposure; and the factual underpinnings of both claims are the same so a successful motion would not reduce the scope of discovery. Filing a dispositive motion is costly and, on these facts, there seems to be no benefit to doing so. So the required recommendation would fail many thoughtful analyses that could discuss the pros and cons and recommend against filing such a motion.

5. Draft Conflict Check Memorandum for Litigation Engagement Clearance (C-030)

QUOTED FROM HARVEY LAB

ISSUE_005: Recommends screening or removal of Caleb Strand. PASS if the memo recommends that Caleb Strand be screened from the engagement or removed from the proposed staffing. FAIL if no such recommendation is made.

— Harvey LAB, criterion C-030, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — low-quality professional judgment, rubric is incorrect
Task environment	Correct
AI auditors’ reasoning	Correct

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric requires the model to recommend that an associate who shares an apartment with his sister (an in-house IP paralegal at TriPoint, the adverse party) be screened from the matter. Screening is generally used in certain specific situations where required by the professional conduct rules, but no rule requires screening in this scenario.

- I agree that a prudent memo would flag the fact that a junior team member is sharing an apartment with an employee of the adverse party, but that is covered by other rubrics (C-026 through C-029). This rubric specifically fails the work product (and results in a 0 for all work on this task) even if the model recommended, say, obtaining the client’s informed consent, which would be perfectly permissible and appropriate here.
- Other rubrics in this task resolve similar kinds of conflicts with informed consent (C-024 accepts “clearance with conditions (such as obtaining informed consent)” for Lisa Chow’s spousal conflict, and C-016 recommends informed consent for the Ridgeline conflict). Accordingly, informed consent plus a home confidentiality protocol could suffice as an alternative to screening.

6. Draft Conflict Check Memorandum for Litigation Engagement Clearance (C-049)

QUOTED FROM HARVEY LAB

ISSUE_011: Recommends integrating annual disclosure questionnaires with ConflictTracker. PASS if the memo recommends that annual conflict-disclosure questionnaires be cross-referenced with or integrated into the ConflictTracker system. FAIL if no recommendation about integrating disclosure questionnaires with the conflicts database is made.

— Harvey LAB, criterion C-049, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — requires specific outputs not expressly or impliedly requested in the prompt
Task environment	Correct
AI auditors’ reasoning	Correct

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The assignment in the prompt is to draft a “comprehensive conflict check memorandum for the proposed engagement.” The rubric requires that the memo contain suggested firm-wide process improvements (memo must “recommend[] that annual conflict-disclosure questionnaires be cross-referenced with or integrated into the ConflictTracker system”). That’s simply not the assignment, and there’s many reasons why it’s undesirable to jam broad process suggestions into a conflicts memo for one specific case. This rubric actually failed work product from Luna that did everything a clearance memo needs on this point. The two judges split on that same Luna memo (Sonnet 4.6 failed it; GPT-5.5 passed it), which shows the criterion does not grade reliably.

7. Draft Counterclaim Against Plaintiff for Breach of Joint Development Agreement (C-045)

QUOTED FROM HARVEY LAB

Count: Patent Infringement (ISSUE_005). PASS if the counterclaim includes a separate count for patent infringement under 35 U.S.C. § 271, alleging that the LumiSense 400 infringes U.S. Patent No. 11,234,567 covering the SensorCore™ architecture. FAIL

if no patent infringement counterclaim is asserted.

— Harvey LAB, criterion C-045, [task file at commit 1dd81403](#)

My assessment

Criterion	Arguable
Task environment	Defective: the record lacks the patent the counterclaim must plead, and the two expert reports give different titles for it.
AI auditors' reasoning	Correct

LAB's own grading of two model runs (Harvey's judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The task is to draft a counterclaim based on a handful of documents (including an expert report, some emails, and a counterclaim strategy memo that another lawyer prepared). The rubric dings the AI if the counterclaims do not include “a separate count for patent infringement under 35 U.S.C. § 271, alleging that the LumiSense 400 infringes U.S. Patent No. 11,234,567.”

- The environment does not contain the allegedly infringed patent, which one would need to review in order to have a sufficient basis for drafting that counterclaim, so this environment appears to be incomplete. That said, a defensible approach might be for the agent to draft the counterclaim but flag in brackets that it requires verification.

8. Draft Counterclaim Against Plaintiff for Breach of Joint Development Agreement (C-046)

QUOTED FROM HARVEY LAB

Patent Infringement Count: Patent Number Cited. PASS if the patent infringement count specifically identifies U.S. Patent No. 11,234,567. FAIL if the patent number is omitted or incorrect in the patent infringement count.

— Harvey LAB, criterion C-046, [task file at commit 1dd81403](#)

My assessment

Criterion	Arguable
Task environment	Defective: same task environment as item 7.
AI auditors' reasoning	Partly correct

LAB's own grading of two model runs (Harvey's judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. Same issue as prior rubric. This one is essentially derivative because it depends on pleading the patent count that was required in the prior rubric.

9. Draft Case Assessment Memorandum for Defective Industrial Equipment Product Liability Claim (C-003)

QUOTED FROM HARVEY LAB

Executive Summary mentions death of Marco Reyes. PASS if the Executive Summary mentions the death of Marco Reyes. FAIL if the fatality is not mentioned in

the Executive Summary.

— Harvey LAB, criterion C-003, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — requires specific outputs not expressly or impliedly requested in the prompt
Task environment	Defective: same task environment as item 10 (Ohio state agency issuing OSHA citations).
AI auditors’ reasoning	Correct

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. The case involved a (fictitious) safety incident at a manufacturing plant that killed one employee and seriously injured two others. The rubric requires the fact that one employee died be discussed specifically in the “executive summary.” When we tested Opus 5.5 and GPT-6 Luna on this benchmark, both discussed the death in several places in their memo, especially the factual background section, but the executive summary consisted of an overview of the legal analysis that did not restate the fact of the death. They were failed by one or more judges.

- This fails perfectly competent and appropriate work product. The company CEO who retained counsel to advise on this incident is the same person who delivered news of the (fictitious) tragic death to the decedent’s widow and surviving family members, so the most senior management of the company is obviously very familiar with the incident and the agents’ approach of focusing the executive summary on the legal exposure is appropriate.
- The two judges reached opposite results on Luna’s memo: GPT-5.5 failed it, while Sonnet 4.6 passed it only by quoting a sentence from the memo’s facts section as if it appeared in the executive summary.

10. Draft Case Assessment Memorandum for Defective Industrial Equipment Product Liability Claim (C-026)

QUOTED FROM HARVEY LAB

OSHA LOTO citation identified with specific violation and penalty. PASS if the memo identifies the OSHA serious citation for LOTO deficiency under 29 CFR 1910.147(c)(4)(i) with a penalty of \$18,900. FAIL if this specific citation, regulation, or penalty amount is not referenced.

— Harvey LAB, criterion C-026, [task file at commit 1dd81403](#)

My assessment	
Criterion	Not defective — defective environment, but rubric does not fail correct work product
Task environment	Defective: Ohio’s Division of Safety & Hygiene cannot issue OSHA citations to this private employer. The criterion itself does not misgrade work (environment-only defect).
AI auditors’ reasoning	Wrong

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. Although the specific rubric being tested does not fail work product incorrectly, the environment is defective in an unintended way. The environment contains an OSHA citation from Ohio’s Division of Safety & Hygiene. But Ohio does not have an OSHA-approved State Plan, so only federal regulators have jurisdiction to issue OSHA citations against this company. The environment does not test that so it is not being presented as a “mistake” by one’s adversary that the agents should pick up on.

- The rubric itself is not wrong because it simply wants a memo to address the purported citation, so an agent could write the memo while noticing that the task is obviously a fictitious evaluation with an incoherent environment.

11. Draft Discovery Plan Memorandum for Breach of Contract and Fraud Defense (C-036)

QUOTED FROM HARVEY LAB

Provides custodian priority rankings. PASS if the memo provides priority rankings or tiers for custodians (e.g., high/medium/low priority, or Tier 1/Tier 2, or similar prioritization). FAIL if custodians are listed without any indication of relative priority.

— Harvey LAB, criterion C-036, [task file at commit 1dd81403](#)

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors’ reasoning	Wrong

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The AI agent is right that the rubric is really enforcing a stylistic choice, but the choice here is well-grounded in convention and an appropriate expectation for quality work product. The case file contains a court order where the judge encouraged the parties to discuss whether phased discovery would be appropriate at the Rule 26(f) conference, and there is some existing prioritization work in the environment. The rubric requires that the agent’s discovery plan include some kind of prioritization or phasing of custodians, which is a common expectation in these cases. I think it’s fair to judge anything that lacks any form of prioritization as deficient given the context.

12. Draft Federal Complaint for Breach of Contract and Fiduciary Duty — Placement Agent Dispute (C-025)

QUOTED FROM HARVEY LAB

ISSUE_012: Complaint requests injunctive relief. PASS if the prayer for relief includes a request for preliminary and/or permanent injunctive relief — specifically to enforce the non-solicitation provision and prevent Graydon from continuing to divert Fund III investor prospects to competing funds. FAIL if no injunctive relief is requested.

— Harvey LAB, criterion C-025, [task file at commit 1dd81403](#)

My assessment

Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong

LAB's own grading of two model runs (Harvey's judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The client intake memo specifically describes an interest in seeking injunctive relief (“We may also wish to seek preliminary and permanent injunctive relief to enforce the twelve-month non-solicitation tail”), and the client’s demand letter threatens it. While it’s true that the intake memo flags the topic for discussion, the correct approach is to include the proposed prayer for relief in the draft, so the case team can see how it reads and discuss any downsides. As a practical matter, there’s unlikely to be much downside to including this prayer for relief in the complaint. (Whether it’s worth the cost of actually pursuing a preliminary injunction on this is debatable, so the rubric’s approach makes sense.)

13. Draft First Set of Interrogatories to Defendant Veridian Health Systems in Trade Secret Misappropriation Case (C-008)

QUOTED FROM HARVEY LAB

Definitions section defines 'Document' or 'Documents'. PASS if the Definitions section includes a definition of 'Document' or 'Documents' that is broad enough to encompass electronically stored information (ESI), emails, and physical documents. FAIL if there is no definition of 'Document' or it is so narrow as to exclude electronic records.

— Harvey LAB, criterion C-008, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Defective: the supplied documents mischaracterize or miscite case law (citation sweep; peripheral to this criterion).
AI auditors' reasoning	Partly correct

LAB's own grading of two model runs (Harvey's judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric fails any draft that does not supply a definition of “document” encom-

passing electronically stored information (ESI). There are two problems:

- First, these are interrogatories, not document requests. Well-drafted interrogatories do not necessarily need to use the word “document.” The rubric does not really make sense as a requirement for this kind of work product.
- Second, even if we were dealing with document requests here (which we are not), the local rules of many courts, including the Western District of Texas (where this fictitious case is being litigated), define the term “document” for all discovery requests served by litigants in those cases. There is thus no requirement to define the term “documents,” and in fact, supplying a definition could just trigger objections from the other side that the local rule controls. See Local Rule CV-26, available at <https://www.txwd.uscourts.gov/court-information/lcr-civil-rules/> (“The term ‘document’ means any document or electronically stored information as described in Federal Rule of Civil Procedure 34(a). A draft of a nonidentical copy is a separate document within the meaning of this term.”).
- The rubric does demonstrably fail rules-compliant work product. For example, Luna was failed for not including a definition of the term “Document” but that’s actually what the local rules encourage in this Court.

14. Draft First Set of Interrogatories to Defendant Veridian Health Systems in Trade Secret Misappropriation Case (C-026)

QUOTED FROM HARVEY LAB

Interrogatory re: USB drive or CMT files on Veridian systems. PASS if at least one interrogatory asks whether Oshiro’s USB drive (referencing serial number CX-8827491 or otherwise describing the external storage device) was ever connected to, accessed from, or present on any Veridian system, device, or network, or asks whether any CMT files or materials were ever possessed, accessed, viewed, or copied by Veridian or its employees. FAIL if no interrogatory addresses the USB drive or CMT files on Veridian systems.

— Harvey LAB, criterion C-026, [task file at commit 1dd81403](#)

My assessment

Criterion	Not defective
Task environment	Defective: same task environment as item 13 (environment-only defect).
AI auditors’ reasoning	Wrong

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. This one does not really belong on the list, as GPT-6 Sol was complaining about a different rubric (C-048) which it argues is inconsistent with this one. It does not identify any issue in this rubric, standing alone, so this counts as an instance where Sol added something to the “flagged” list that does not belong.

15. Draft Discovery Responses and Objections to RFAs and RFPs in Breach of Supply Agreement Litigation (C-028)

QUOTED FROM HARVEY LAB

ISSUE_011: Identifies RFA No. 20 as calling for a legal conclusion. PASS if the plan identifies that RFA No. 20 (asking Prismavale to admit the products were 'not merchantable as defined by the Uniform Commercial Code') calls for a legal conclusion regarding UCC § 2-314 merchantability and recommends objecting on that basis. FAIL if the plan does not identify the legal conclusion problem with RFA No. 20.

— Harvey LAB, criterion C-028, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — low-quality professional judgment, rubric is incorrect
Task environment	Defective: the supplied Answer cites complaint paragraph ranges that do not match the complaint (e.g. Count III as ¶¶ 88–94; the complaint's Count III is ¶¶ 86–96).
AI auditors' reasoning	Correct

LAB's own grading of two model runs (Harvey's judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. Plaintiff served a request for admission on the defendant that says, “Admit that the SB-102 and PS-302 products shipped by Prismavale to Greenleaf were not merchantable as defined by the Uniform Commercial Code.” The prompt asks the agent to “prepare a discovery strategy memo and formal RFA/RFP responses.” The rubric requires “the plan” (i.e., the strategy memo) to recommend objecting to the request on the grounds that it calls for a legal conclusion. The rubric is wrong for two main reasons:

- First, the modern view is requests for admissions can call for application of law to facts. While it is within acceptable professional norms for a lawyer to object on the ground that such a request calls for a legal conclusion, and then proceed to answer, the better reading of the law is that this request is not objectionable such that the party should simply answer. On this legal conclusion in particular, there's really no reason why one would prefer to object rather than simply deny the requested admission. So this rubric fails competent work product (and probably the better answer than the one the rubric calls for).
- Second, even assuming one should include the objection at all, the task fails the model if the issue is not expressly addressed in the strategy memo. Luna made the “legal conclusion” objection in the draft responses and then proceeded to a substantive denial (an acceptable approach), but did not discuss the “legal conclusion” objection in the memo. That was judged a failure. That's really requiring duplicative responses on a very minor point that probably shouldn't be in the drafts at all.

16. Draft Motion in Limine to Exclude Expert Testimony and Prejudicial Evidence in Commercial Breach of Contract Case (C-009)

QUOTED FROM HARVEY LAB

ISSUE_004: Identifies the \$11,000 arithmetic discrepancy. PASS if the motion identifies that the sum of Marchetti's five itemized damages categories totals \$14,211,000,

but her report states total damages as \$14,200,000, creating an \$11,000 discrepancy. FAIL if this arithmetic error is not identified.

— Harvey LAB, criterion C-009, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — low-quality professional judgment, rubric is incorrect (a judgment call, not an objective error); defective environment (directs advocacy outside professional norms); companion rubric C-010 also incorrect (not sampled)
Task environment	Defective: the task directs advocacy outside professional norms (an \$11,000 rounding point offered as a <i>Daubert</i> ground) and labels a rounding difference an “arithmetic error.”
AI auditors’ reasoning	Correct

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The damages expert includes a summary table in her report listing various items of damages, which collectively add up to \$14,211,000; the Total Damages line lists the total as \$14,200,000. Opposing counsel questioned the expert about it at her deposition, and the expert calls it “a rounding discrepancy.” An internal strategy memo highlights the issue as a ground for seeking to exclude the expert’s testimony under *Daubert*, arguing “the arithmetic error supports a cumulative argument that the report was prepared without the rigor required of expert testimony in federal court.”

- This task is poorly specified, and the recommended strategy of raising this with the judge is a poor one:
- First, the task environment repeatedly states that this is an “arithmetic error.” That is false. The expert has rounded the total in the table to the nearest \$100,000. While that is somewhat unorthodox in an expert report (particularly since it is written out as \$14,200,000 not \$14.2 million), numbers often are rounded in court filings so the judge is unlikely to agree that this is an error.
- Second, even if there were an error (which there is not), that is generally not a basis for excluding the expert’s testimony under *Daubert*. At most one could move to exclude the erroneous version and ask for more accurate testimony, but that does nothing to help the hypothetical client here; it just raises their damages exposure by the \$11,000 that the expert was willing to round off.
- Third, the proposed framing of the argument comes very close to a breach of decorum and carries a real risk of drawing a rebuke from the judge. The strategy memo wants the motion to argue the expert’s 30 prior engagements were “in state courts and arbitrations that do not impose the same methodological gatekeeping requirements” as federal courts. While some state courts use the *Frye* standard instead of *Daubert*, many federal judges would take offense at the suggestion that state courts do not observe adequate methodological rigor. Many busy judges could be annoyed at spending time on an argument that literally just amounts to a rounding error.
- This rubric is scored as Defective, as a professional-judgment call rather than an objective

error. Our view is that the better approach is to drop this issue entirely, and the sampled rubric (C-009) fails a motion that does so, even though a motion that merely mentions the discrepancy in a footnote would pass. The companion rubric, C-010, is in my view objectively wrong because it requires that this issue be presented as a basis for excluding the expert’s testimony under *Daubert*. Luna was failed incorrectly on C-010 (by the Sonnet 4.6 judge) for downplaying this issue; in fact, Luna was right to do so, as it should not be raised at all. But since that’s not the sampled criterion, we do not mark it here.

17. Draft Rule 12(b)(6) Motion to Dismiss Brief — Commercial Software Licensing Dispute (C-037)

QUOTED FROM HARVEY LAB

ISSUE_008 — Cites Delaware parol evidence authority. PASS if the motion brief cites relevant Delaware authority on the parol evidence rule and integration clauses, such as *SIGA Technologies v. PharmAthene* or *Eagle Industries v. DeVilbiss Health Care*. FAIL if no Delaware parol evidence authority is cited.

— Harvey LAB, criterion C-037, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — unambiguous objective error (including hallucinations), rubric is incorrect
Task environment	Defective: the supplied research memo cites fabricated or mischaracterized authority (citation sweep: 28 of 53 citations in this task have problems).
AI auditors’ reasoning	Correct

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. The rubric here involves a legal error, and the task environment contains hallucinations about case law. This one is extremely defective for multiple reasons:

- First, the rubric is demanding case law on the wrong legal doctrine. Under Delaware law, the parol evidence rule and an integration clause do not preclude fraud claims. The issue here is whether the plaintiff can reasonably allege that they relied on their counterparty’s pre-signing representations despite their agreement in the contract that they were not relying on any representations other than those written in the contract itself.
- Second, the cases in the environment appear to contain multiple hallucinations. The rubric refers to *SIGA v. PharmAthene*, but Opus noted that case does not address integration or parol evidence. The task environment contains a legal research memo that cites *SIGA Technologies, Inc. v. PharmAthene, Inc.*, 67 A.3d 330 (Del. 2013), and claims that it states (at page 344), “an integration clause is the clearest declaration that the parties intend the writing to be the complete and final expression of their agreement.” The quoted language does not appear in the cited case. Page 344 of that case addresses the duty to negotiate in good faith, not integration clauses. Similarly, the memo cites *Eagle Industries, Inc. v. DeVilbiss Health Care, Inc.*, 702 A.2d 1228 (Del. 1997) and says: “The Delaware Supreme Court held that ‘a party to a contract cannot promise, in a clear integration clause of a negotiated agreement, that it is not relying on promises or representations made by the other party outside of the

agreement, and then assert a claim for fraud based on those very representations.’ *Id.* at 1232.” The quoted language does not appear in the cited case.

- Third, the rubric failed responses that stated the correct law on this point. In our tests, Opus reviewed SIGA and Eagle, correctly realized they were inapposite and declined to cite them, cited more on-point and accurate authority, and was then failed for doing a better job than the rubric’s author did at this task.

18. Draft Rule 12(b)(6) Motion to Dismiss Brief — Commercial Software Licensing Dispute (C-045)

QUOTED FROM HARVEY LAB

ISSUE_010 — References Arcadia’s own delays contributing to Go-Live delay. PASS if the motion brief identifies Arcadia’s own contributing failures, including the 45-day Project Manager gap (Kevin Liu’s departure) and/or the 75-day late API specifications, as causes of the Go-Live delay. FAIL if Arcadia’s own contributory delays are not mentioned.

— Harvey LAB, criterion C-045, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Defective: same task environment as item 17.
AI auditors’ reasoning	Correct

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. The rubric makes a rookie mistake, demanding that the agent argue a factual dispute on a motion to dismiss, where the court is required to assume the plaintiff’s allegations are true and cannot consider factual dispute. The 45-day and 75-day delays complained of come from the defendant’s internal report, which the plaintiff’s complaint disputes. Moreover, courts generally cannot resolve messy fact disputes over “causation” on a motion to dismiss.

- The criterion failed agents that correctly recognized that this argument was outside the scope of what can be argued on a motion to dismiss and thus decided to exclude it.

19. Draft Verified Responses and Objections to Plaintiff’s First Set of Interrogatories in Commercial Breach of Contract and Fraud Action (C-028)

QUOTED FROM HARVEY LAB

Correct financial figures for Tri-Basin purchase history. PASS if any purchase history figures cited in the responses are consistent with the canonical data: 2015: \$9.2M; 2016: \$10.1M; 2017: \$10.8M; 2018: \$12.4M; 2019: \$13.6M; 2020: \$9.7M; 2021: \$15.3M; 2022: \$16.8M; 2023 (Jan–Aug): \$11.9M. Minor rounding differences are acceptable. Also PASS if the response invokes Rule 33(d) rather than citing specific figures. FAIL if specific purchase figures are cited but are materially incorrect.

— Harvey LAB, criterion C-028, [task file at commit 1dd81403](#)

My assessment	
Criterion	Defective — unambiguous objective error, deficiency in task environment; rubric is incorrect
Task environment	Defective: the purchase-history spreadsheet is internally inconsistent (three 2020 series; line items that do not sum to subtotals), and the record has no contract-year data.
AI auditors' reasoning	Partly correct

LAB's own grading of two model runs (Harvey's judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. There are two issues:

- First, Interrogatory No. 4(b) asks for Tri-Basin's actual purchases for the 2020 contract year (the contract year runs from January 15 through January 14 of the following year; Interrogatory No. 14 separately asks for calendar-year totals), but the canonical data referred to in the rubric is calendar year data, so the rubric's numbers are not correct for that request. The correct value cannot be discerned from the environment: the monthly data cannot be split at January 15, and nothing in the record supports assuming purchases are spread evenly within January (seasonal variation would make that unlikely). The rubric therefore grades wrong in the opposite direction from most of the other items: rather than failing good work product, it passes responses that supply the calendar-year figure for a contract-year request without acknowledging that the record does not contain the requested data. That is approving fundamentally deficient work product.
- Second, the data that is provided has some internal inconsistencies. For example, an "annual summary" tab lists 2020 purchases at \$9.7 million, but the monthly detail lists 2020 three times, with totals of \$8,700,000, \$8,549,400 and \$9,700,000. The units sold do not match up between those two sheets for 2021, 2022, 2023, and series 3 of 2020. One of the monthly rows does not add up (gross \$1,210,000 + credits −\$17,400 = \$1,192,600, but the net shown is \$1,192,500). The annual summary minimums do not add up. In the Product Line Breakdown, line-item net revenue never sums to the subtotal. So a response that relied on the line items would get graded as a failure.
- The environment contains law firm discovery guidelines that recommend that interrogatory responses match the client's Answer, but the Answer here contains only approximate figures, which is consistent with the Answer using calendar year data as a rough approximation instead of providing the contract year data requested.

20. Extract Key Obligations from Litigation Hold and Document Preservation Notice — Obligation Summary Memorandum (C-002)

QUOTED FROM HARVEY LAB

ISSUE_001: Flags financial records date discrepancy as requiring DOJ clarification. PASS if the memo recommends seeking clarification from the DOJ regarding the conflicting date ranges (Jan 1, 2017 vs. Jan 1, 2019) for financial records. FAIL if no recommendation to seek clarification is made.

— Harvey LAB, criterion C-002, [task file at commit 1dd81403](#)

My assessment	
Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Correct
AI auditors' reasoning	Correct

LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: fail

My analysis. The task involves a document preservation notice from federal criminal authorities that defines the relevant period (for which the recipient must preserve documents) as January 1, 2019, through March 3, 2025 “unless otherwise specified herein.” For one specific category of documents, the notice states that, “Notwithstanding the Relevant Period defined in Paragraph 12 of this Notice, Ridgeline shall preserve [certain kinds of] records . . . from January 1, 2017 to the present date.”

- The rubric suggests that the agent should seek clarification of the “conflicting” date ranges but there is no conflict. The notice quite clearly imposed a longer period for one specific category of documents, while all the other categories are subject to a shorter period. This is a fairly common request.
- All the models were failed on this criterion, but the criterion makes no sense.

21. Identify Excessive or Duplicative Research Charges in Litigation Invoice (C-013)

QUOTED FROM HARVEY LAB

Issue 3: Identifies duplicative MTCA research across Takahashi and Wendt.

PASS if the memo identifies that multiple attorneys (Kenji Takahashi and Tyler Wendt) billed for overlapping MTCA liability/contractor liability/summary judgment research during July 2024, citing Section 6.4 (Duplicative Research Prohibition). FAIL if the MTCA duplication across timekeepers is not identified.

— Harvey LAB, criterion C-013, [task file at commit 1dd81403](#)

My assessment	
Criterion	Not defective
Task environment	Correct
AI auditors' reasoning	Wrong

LAB's own grading of two model runs (Harvey's judges, not my assessment)	
GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. While Sol is right that the facts here are insufficient to establish that the client's billing guidelines were violated, for purposes of the task (a client-side memo identifying changes to dispute), it seems reasonable to flag that multiple attorneys were doing research into closely related issues. The task requested “a memo identifying all non-compliant or excessive charges with recommended adjustments,” so flagging potential violations (even if they were not clearly proven from the face of the time entries) is a standard approach a client would take here.

22. Identify Issues in Litigation Matter Budget Proposal (C-029)

QUOTED FROM HARVEY LAB

Budget inconsistency rated as Critical or High severity. PASS if the fee total reconciliation discrepancy (\$116,050 difference between phase-based and staffing-model totals) is assigned a severity of Critical or High (or the highest or second-highest tier in whatever severity scale is used). FAIL if it is assigned a lower severity level or no severity level.

— Harvey LAB, criterion C-029, [task file at commit 1dd81403](#)

My assessment

Criterion	Arguable — requires specific outputs not expressly or impliedly requested in the prompt
Task environment	Correct
AI auditors’ reasoning	Correct

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The prompt requested “a categorized issue memorandum” but didn’t specify what the categories should be. The rubric fails the model if it does not break issues into severity levels or it does not assign a critical or high severity level to a specific issue. While it appears that models often do volunteer this kind of severity assessment, it’s not requested by the prompt and competent work product could decline to include it.

23. Review Counterparty’s Proposed Jury Instructions — Issue Memorandum for Trade Secrets Trial (C-020)

QUOTED FROM HARVEY LAB

Explains distinction between tortious interference with contract vs. business relations. PASS if the memo explains that under Georgia law, tortious interference with an existing contract does not require a separately/independently wrongful act — the knowing inducement of the breach is sufficient — and that the ‘independently wrongful act’ element applies to the distinct tort of tortious interference with business relations. FAIL if this legal distinction is not articulated.

— Harvey LAB, criterion C-020, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — unambiguous objective error (including hallucinations), rubric is incorrect
Task environment	Defective: the fictional summary-judgment order cites a case for the opposite of its holding and a case that cannot be found (citation sweep: 10 of 21 citations in this task have problems).
AI auditors’ reasoning	Correct

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail
Claude Opus 5.5	Sonnet 4.6 judge: fail; GPT-5.5 judge: fail

My analysis. Multiple issues:

- First, this is a clear legal error. Georgia applies the same elements to interference with contractual and business relations. The rubric says the opposite of what Georgia law is. *Disaster Services, Inc. v. ERC P’ship*, 228 Ga. App. 739, 740–41 (1997) (“Tortious interference claims, whether asserting interference with contractual relations, business relations, or potential business relations, share certain common essential elements: (1) improper action or wrongful conduct by the defendant without privilege; (2) the defendant acted purposely and with malice with the intent to injure; (3) the defendant induced a breach of contractual obligations or caused a party or third parties to discontinue or fail to enter into an anticipated business relationship with the plaintiff; and (4) the defendant’s tortious conduct proximately caused damage to the plaintiff.”).
- Second, the fictional summary judgment order (from which the rubric’s legal position seems to be taken) appears to contain several hallucinations. For example, the order paraphrases the third and fourth element described in *Disaster Services* approximately correctly, but the first two elements are erroneously listed as “(1) the existence of a valid contract between the plaintiff and a third party; (2) the defendant’s knowledge of the contract.” The fictional order also cites *Valdosta Livestock, Inc. v. Furst*, 342 Ga. App. 25, 28 (2017), which is not found in any legal databases.

24. Review Litigation Invoice Against Outside Counsel Billing Guidelines — Compliance Deviation Report (C-024)

QUOTED FROM HARVEY LAB

Recommends full disallowance of black car service: \$387.00. PASS if the report recommends disallowing the full \$387.00 for black car service (or reducing it to a reasonable taxi/ride-share equivalent). FAIL if no disallowance or reduction is recommended.

— Harvey LAB, criterion C-024, [task file at commit 1dd81403](#)

My assessment

Criterion	Not defective
Task environment	Correct
AI auditors’ reasoning	Wrong

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The billing guidelines state that “Reimbursable ground transportation is limited to ride-share ... or taxi” and “Limousine services and black car services are not reimbursable under any circumstances.” Sol’s critique of the rubric is that the policy requires disallowing black car reimbursement, not reducing the reimbursement amount to what a taxi or ride share would have cost. The rubric’s interpretation (that either disallowing the expense entirely or capping reimbursement at the amount a taxi would have cost) seems reasonable and closer to the likely real-world application of a policy like this.

25. Review Litigation Invoice Against Outside Counsel Billing Guidelines — Compliance Deviation Report (C-042)

QUOTED FROM HARVEY LAB

Classifies business class airfare as borderline/requires discussion. PASS if the report classifies the business class airfare issue (Messina’s \$2,847.00 flight) as a borderline or ‘requires discussion’ issue rather than a clear-cut mandatory reduction, given the 4-hour-20-minute flight duration and the 4-hour threshold. FAIL if it is classified as a clear mandatory reduction without acknowledging the ambiguity, OR if it is entirely dismissed as compliant.

— Harvey LAB, criterion C-042, [task file at commit 1dd81403](#)

My assessment

Criterion	Defective — unambiguous objective error, rubric is incorrect
Task environment	Correct
AI auditors’ reasoning	Correct

LAB’s own grading of two model runs (Harvey’s judges, not my assessment)

GPT-6 Luna	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass
Claude Opus 5.5	Sonnet 4.6 judge: pass; GPT-5.5 judge: pass

My analysis. The rubric’s analysis is wrong. The travel policy included in the environment states, “For domestic flights with a scheduled duration of four (4) hours or more, business class may be booked without prior approval.” The record does not state the flight’s duration (the invoice lists only the route and fare class). The only basis for raising this seems to be in the rubric itself, which says the flight could take 4h20m, but even under that view, there’s nothing to discuss, as the policy allows business class for that duration.