

# LegalForecastBench: Forecasting Judicial Decisions as a Test of Legal Reasoning Ability

John J. Hughes, III  
Independent researcher  
<https://www.johnjhughes3.com/research/>

Working draft — October 1, 2026

## Abstract

Frontier agentic models have recently demonstrated enormous gains in capabilities in a number of domains, but particularly in domains like math and coding where model outputs on even high-level tasks can be verified automatically. Capabilities in the legal domain have advanced as well, but noticeably more slowly than capability levels in verifiable domains, with human lawyers continuing to outperform frontier models on many higher-level legal judgment tasks. The demonstrated importance of verifiable rewards to improving frontier model capabilities raises the question of whether legal judgment tasks can be decomposed into objectively verifiable tasks that could be used to evaluate or train models on higher-level legal reasoning capabilities.

Existing benchmarks for evaluating high-level legal output largely use human-written rubrics that test model outputs. When models were still acquiring basic legal knowledge, simple, bar exam-style questions provided a bounded way to evaluate that knowledge. As capability levels have increased, one approach has been to create new human-written rubrics; Harvey’s Legal Agent Benchmark (LAB), for example, contains 1,200 legal tasks and more than 75,000 rubric criteria that are intended to test what an ideal output would look like. But my preliminary AI-assisted audit of LAB’s litigation tasks (my area of specialty as a practicing lawyer) flagged hundreds of criteria for possible legal errors or misspecification; verification of those findings is ongoing. Moreover, my deep dive into a specific task environment identified concerns about the realism of the synthetic documents and the alignment between the task and its rubric. For example, the supplied documents include programmatic-generation metadata and a stale table-of-contents instruction, and certain criteria prescribe legal conclusions without adequately distinguishing stronger from weaker arguments. These observations raise questions about whether the rewarded conduct will transfer to real-world litigation. They do not establish a benchmark-wide error rate or show that rubric-based training is ineffective.

I argue that judicial decisions in real-world litigation matters supply an external signal that allows evaluation of some high-level legal judgment tasks against an objectively verifiable real-world ground truth. While lawyers and litigants may disagree on what the law is on a certain point, or how it applies to a specific fact pattern, a court’s ruling authoritatively resolves such disputes. Unlike the opinions of expert annotators creating synthetic environments, moreover, judicial decisions have the force of law (and analyzing and understanding such decisions is one of the primary mechanisms by which human lawyers are trained). LegalForecastBench is an example of one way judicial dispositions can be used for more objective evaluation of frontier models’ legal reasoning capabilities. Models receive the available pre-decision litigation records for a motion to dismiss, and are asked to estimate the probability that the presiding judge will dismiss each challenged claim against each moving defendant. The beta includes 91 cases and 387 scored claim–defendant units. In the initial run, GPT-6 Sol has the lowest micro-Brier loss (0.118) and the highest accuracy (84.5%) of seventeen model configurations tested. All seventeen configurations achieved lower micro-Brier loss than a retrospectively fitted

constant-probability reference, but confidence and calibration vary widely, and several strong models make consequential high-confidence errors. We describe how the approach could be extended to model training; this study reports evaluation results, not a completed training experiment.

## 1 Introduction

Reinforcement learning (RL) techniques have led to significant recent gains in frontier model capabilities, but the most noteworthy gains have been in domains like mathematics and coding that are susceptible to automated verification. Reinforcement learning with verifiable rewards (RLVR) has allowed frontier models to gain capabilities by receiving reward signals from automated correctness checks (like mathematical answer matching, executable tests, or formal proof checking) [1, 2]. Earlier self-play systems demonstrated that automated outcome feedback, in games with explicit rules and automatable scoring mechanisms, allowed models to achieve superhuman performance with sufficient training [3]. Recent announcements from frontier labs have demonstrated the incredible power of these techniques in solving hundreds of longstanding mathematical problems that eluded human solutions.

Frontier models have also made impressive gains in the legal domain, but progress there appears to have been limited to capabilities that are susceptible of some degree of automated verification—meaning that models today cannot reliably perform high-level legal judgment tasks that are difficult to verify. Frontier models have gained excellent legal knowledge. Existing benchmarks that test knowledge of well-settled legal propositions, or their application to simple fact patterns, are easily verifiable (and frontier models likely gain significant legal knowledge in pre-training), so frontier models have become excellent at answering legal questions. And various techniques to ensure that answers to legal research tasks are grounded in sources have largely eliminated earlier problems with hallucinations and made frontier models excellent legal research assistants. But models have increasingly saturated legal benchmarks that focus on bounded tasks with objectively scored answers, like the Multistate Bar Examination [13]. To train models on higher-level legal tasks, researchers have tried at least two approaches: preference evaluation (comparing which of two outputs practicing attorneys prefer) and rubric-based reinforcement learning (grading model outputs against specified criteria to produce a reward). [7]. There are limits to both approaches, as we will discuss in Section 6.4.

This paper aims to demonstrate that judicial decisions in litigation matters provide a useful external signal that can create high-level legal judgment tasks that are objectively verifiable against an external ground truth. This approach rests on a simple fact about the legal system: "It is emphatically the province and duty of the Judicial Department to say what the law is." *Marbury v. Madison*, 5 U.S. 137, 177 (1803). Litigated matters almost always involve disputes about how to interpret, or apply, abstract legal principles to a specific fact pattern. A court's judgment on such a dispute represents the resolution of the matter, and has the force of law. LegalForecastBench attempts to illustrate one specific example of how this external signal can be utilized to create evaluations that assess the legal reasoning and analytical abilities of frontier models against an objectively verifiable ground truth.

As discussed herein, the alternatives of preference evaluation or rubric-based reinforcement learning are useful techniques for certain kinds of training and evaluation, but both approaches rest on human labels applied by experts whose opinions have no legal authoritative or precedential value. Moreover, as discussed herein, after noticing legal errors in several rubrics, I used frontier models

to audit the litigation-related tasks in Harvey’s Legal Agent Benchmark, and models identified hundreds of potential errors in rubrics or task specifications. While my work to verify the models’ findings is ongoing, my review of a small sample identified genuine legal errors in the rubrics or task specifications. These issues illustrate some of the practical challenges of trying to train models on legal tasks using artificially constructed environments.

LegalForecastBench demonstrates that real-world litigation matters can be used to provide models with realistic, challenging high-level judgment tasks that have an objective resolution. Although this paper uses such tasks only for model evaluation, I have developed a separate training data set and am separately exploring whether training on this data can also help improve models’ legal reasoning abilities.

## 2 Related work

The approach here is consistent with previous research that (i) used older AI techniques to build smaller, custom models for predicting the outcome of litigation (especially U.S. Supreme Court decisions), and (ii) more recent efforts to use frontier models to predict the outcomes of events defined in prediction markets.

### 2.1 Legal outcome prediction

Using AI models and other techniques to predict litigation outcomes has a long lineage. Ruger and colleagues used hand-coded metadata and statistical techniques to predict the 2002 Term of the U.S. Supreme Court [35]. Later work incorporated more sophisticated machine-learning and natural language processing techniques. Aletras and colleagues in 2016 used a text-based approach to predict the outcomes of European Court of Human Rights decisions [37]. Katz, Bommarito, and Blackman in 2017 used a random-forest classifier over structured Supreme Court data to predict Justice votes and case outcomes. [36]. CAIL2018 used a dataset of more than 2.6 million Chinese criminal cases to predict charges and sentences, with text-classification models providing initial baselines. [38]. As the transformer architecture became dominant, later legal prediction systems increasingly used models like BERT and T5. CLC-UKET, released in 2024, compared fine-tuned BERT and T5 models to zero- and few-shot GPT-3.5 and GPT-5 to predict U.K. Employment Tribunal cases. [40]. The field has thus progressed from statistical models over hand-coded metadata through conventional MLP and supervised machine learning to transformer-based classifiers and, more recently, large language models. Much of this literature differs from the approach studied here, because the models operate on structured metadata or facts and claims extracted from judicial opinions. The approach here simply asks frontier models to reason about the actual pre-decision litigation record, essentially the documents filed in court before the judge who decided the motion, without any access to statistical data or models showing prior decisions.

### 2.2 Legal AI benchmarks and evaluation

One of the most impressive academic benchmarks is LegalBench, a collaborative benchmark to which 40 authors contributed 162 tasks that aimed to assess legal domain expertise. Although some of the legal issues tested could give rise to disputes in real litigation (for example, there are tests of whether a given statement is hearsay or not), the benchmark deliberately uses relatively clean examples that can be reduced to a binary answer. [12]. The hearsay tasks, for example, explicitly articulate what proposition a given statement is being used to prove or disprove, and then ask

whether that use is hearsay; in real-world litigation, what proposition the statement is being offered to prove may itself be disputed so providing a clean prompt on that point narrows the task to a relatively mechanical application of an explicit rule.

Most recent benchmarks test higher-level judgment tasks that lack an objective ground truth. JudgmentBench, released in May 2026, and Harvey’s LAB targeted open-ended professional legal work, like drafting, analysis, and advice, for which there are potentially multiple defensible outputs and no independently verifiable ground truth. [7, 41]. These benchmarks take different approaches to reward specification. JudgmentBench explicitly compared rubric-based scoring to comparative judgment (pairwise preference judgments collected from practicing attorneys), and found that comparative judgments recover the intended quality ordering substantially better than rubrics. LAB relies on a rubric-based approach.

### 2.3 Verifiable rewards, forecasting, and outcome feedback

Frontier models are trained on several different forms of feedback, as described in prior work. InstructGPT in 2022 provided an influential demonstration of reinforcement learning from human feedback [11]. In 2023, Lightman and colleagues showed that, on mathematical reasoning tasks, process supervision (assigning feedback to individual reasoning steps) outperformed outcome supervision, which gives rewards based on whether the final answer is correct [42].

The efficacy of these techniques in verifiable domains has been exceptional. During reinforcement learning, DeepSeek-R1-Zero’s average pass@1 on AIME 2024 rose from 15.6% to 77.9%; the final DeepSeek-R1 system, which used a broader training pipeline, reached the 96.3rd percentile on Codeforces [1]. AlphaProof coupled reinforcement learning with Lean’s formal proof checker and proved three problems from the 2024 International Mathematical Olympiad, including its hardest problem, solved by only five of 609 human contestants [4]. And OpenAI’s September 2026 announcement that it had solved the Navier–Stokes Millennium Prize Problem reportedly relied on an internal model developed through large-scale reinforcement learning, followed by Lean formalization and verification. [5, 6].

More recent work has used externally resolved real-world events to supply an analogous signal outside domains like mathematics. For example, Autocast, released in 2022, assembled forecasting questions from tournaments and a date-indexed news corpus to allow language models to make predictions based on the information a human would have had at the time. [43]. Halawi and colleagues in 2024 used a retrieval-augmented forecasting system that could search for information, and generate and aggregate multiple forecasts; the system approached the performance of competitive human forecasters [44]. ForecastBench, introduced in 2024, extended this approach into a dynamic benchmark that continually collects questions about unresolved future events and scores models once the event outcome is known [45].

Outcome-Based Reinforcement Learning to Predict the Future showed in 2025 that resolved binary forecasting questions could serve as reinforcement-learning rewards, training a 14-billion-parameter model on more than 100,000 events using adaptations of GRPO and ReMax [24]. OpenForecaster similarly explored training language models directly on resolved forecasting questions [25]. Some research in China has also used frontier models to predict criminal sentences [26]. LegalForecastBench is an effort to apply a similar approach to litigation outcomes, at least as an evaluation metric for now.

### 3 Judicial outcomes as a target for legal AI evaluation

#### 3.1 Judicial outcomes as externally verifiable events

Litigation converts a messy legal dispute as to which reasonable minds can differ into a concrete and fixed disposition. Courts prize finality, and final judgments are explicitly written with the goal of achieving a clear and objective disposition of the dispute between the parties that leaves no room for further debate or relitigation of the issues the court has resolved. A final judgment on the merits, with appeals exhausted, has preclusive effect and the force of law, even if a litigant could show that the judgment was based on a legal error [17].

While LegalForecastBench does not directly test models' ability to predict the court's final judgment (instead, testing only their ability to predict the first written disposition of a specified motion), that first-level prediction task is a good initial target because many more cases reach the motion to dismiss stage than the final judgment stage (a litigated final judgment is comparatively rare). And the outcome of the motion to dismiss itself is of critical importance in many disputes and potentially the most economically consequential single stage of litigation. In securities fraud class actions, for example, the motion to dismiss can strongly influence whether the defendant will ultimately have to make a settlement payment to resolve the case. The Stanford Law School Securities Class Action Clearinghouse and Cornerstone Research report that, among core federal filings from 2014 to 2021, 52% had been dismissed, 41% had settled, and 6% remained ongoing as of the end of 2024 [21, p. 16]. These cohort outcomes do not isolate the effect of the motion to dismiss, but illustrate the importance of dismissal and settlement as alternative resolutions.

#### 3.2 Forecasting as a component of legal judgment

Litigation prediction is a key subcomponent of legal work, even when litigation is not at all contemplated or likely. While there are many other tasks involved in giving strategic advice to a client (like understanding the client's objectives), outcome prediction is a key subtask, consistent with the MacCrate skills framework and the New York State Judicial Institute's account of counseling and negotiation [23]. Past research shows that even parties who do not intend to litigate with each other bargain "in the shadow of the law" [22]. When parties are negotiating, their best alternative to a negotiated agreement is often to litigate. As a result, the parties' predictions of how a court will resolve their dispute constrain the bargaining range; neither party is likely to accept an agreement that is less favorable than the result they believe they could obtain in court (after taking into account the costs and risks of litigation).

A simple risk-neutral model illustrates this [51]. Suppose a plaintiff would receive a fixed award  $J > 0$  if successful and zero otherwise. Let  $q_P$  and  $q_D$  be the plaintiff's and defendant's respective estimates of the probability of that success, and let  $k_P, k_D \geq 0$  be their remaining litigation costs. Assuming each side bears its own costs, ignoring settlement costs, and restricting attention to monetary payoffs, the plaintiff's expected litigation payoff is  $q_P J - k_P$ , while the defendant's expected trial cost is  $q_D J + k_D$ . A payment  $s$  from defendant to plaintiff is weakly acceptable to both when

$$q_P J - k_P \leq s \leq q_D J + k_D. \quad (1)$$

The interval is nonempty when

$$(q_P - q_D)J \leq k_P + k_D. \quad (2)$$

This is an illustrative reservation-value calculation [52, 53]. Where the parties agree about the probability of success, they can enjoy a surplus from settling the case, equal to the avoided costs. The

possibility of settlement disappears when the plaintiff’s estimate of success exceeds the defendant’s estimate of the plaintiff’s success by enough that  $(q_P - q_D)J > k_P + k_D$ . Even when some interval exists where both parties believe they are better off settling, private information or game theoretic considerations/strategic behavior might prevent agreement [53]. Nonmonetary goals, risk aversion, fee shifting, and uncertainty around the actual damages that would be awarded at trial are additional complications that would require a more complex model. Nevertheless, this simple example illustrates how forecasting litigation outcomes can impact even out-of-court negotiations in a broad swath of commercial disputes or transactions.

Notably, better AI forecasts could yield more surplus to litigants, systematically, if such forecasts help parties’ expectations to converge; simply put, if an AI system becomes sufficiently accurate at predicting litigation outcomes that parties are willing to adjust their expectations based on its output, then the parties are more likely to agree on the likely outcome of a trial. Where  $q_P > q_D$ , narrowing the gap between their estimates reduces  $(q_P - q_D)J$ , making the settlement condition easier to satisfy and potentially creating a mutually acceptable payment where none previously existed. If the parties adopt a common forecast,  $q_P = q_D$ , the settlement interval has width  $k_P + k_D$ : the full amount of the avoided litigation costs. While modern litigation is much better than many historical options used to resolve disputes (such as feud, self-help, trial by ordeal, etc.), reaching a mutually acceptable agreement is virtually always preferable to the costs, risks, and distraction brought on by protracted, adversarial litigation. Improved forecasting of outcomes can thus help parties more realistically understand the weaknesses and risks of their case, and increase the likelihood of settlement early in the matter.

One major limitation of LegalForecastBench, relative to alternatives, is that it focuses primarily on the models’ analytical ability and does not directly evaluate their ability to prepare legal work product. But litigation prediction is at least a subcomponent of many drafting tasks, and could improve models’ skills.

A basic principle observed in the legal profession is that lawyers must confine their argument to the strongest points, not raise every possible point, and generally should put the strongest arguments first in the brief, moving on to points that are less likely to be dispositive later. In *Jones v. Barnes*, the Supreme Court expressly recognized the value of selecting stronger appellate arguments and the risk that adding weaker ones can dilute them [54]. In order to prepare good work product, then, models need to distinguish the potentially dispositive issues from objections or points that are technically valid but less likely to drive the ultimate outcome.

### 3.3 What judicial forecasting measures

**Prediction is not legal correctness.** Judicial outcomes reflect doctrine and record facts, but also procedure, judicial behavior, and institutional context. Judge identity remains visible. A metadata-only comparison, judge masking, and a human forecasting baseline are important future tests. Neither fluent rationales nor a correct forecast demonstrate valid legal reasoning. Training to imitate outcomes can reproduce historical disparities or errors; professional advice also has obligations that prediction scores do not capture.

## 4 LegalForecastBench: Dataset and evaluation protocol

### 4.1 Prediction target and information set

Let  $c \in \{1, \dots, C\}$  index cases and  $u \in \mathcal{U}_c$  index scored units within case  $c$ . Each unit identifies a challenged claim or a separately resolvable pleaded subclaim against a moving defendant (or a group of similarly situated defendants, for whom the outcome is predicted together since the legal and factual arguments for dismissal are materially identical as to all of them). We used models to identify appropriate units based on the arguments articulated in each motion; an affirmative defense, an element of a legal cause of action, a requested remedy, or a distinct factual theory was not automatically treated as a separate unit. Rather, units generally tracked claims asserted in the complaint (in some cases with a single claim asserted against multiple defendants being treated as different units, if different defendants advanced different grounds for dismissal). The model receives the pre-decision record that the judge received  $X_c$  and the unit specification  $U_{cu}$  and returns

$$p_{cu} \in [0, 1], \quad p_{cu} \approx \Pr(Y_{cu} = 1 \mid X_c, U_{cu}), \quad (3)$$

where  $Y_{cu} = 1$  means that the judge dismissed that claim-defendant unit in its entirety, in the target motion’s first written disposition. That event is not a final victory in the lawsuit, since a plaintiff might be given an opportunity to file an amended complaint with additional allegations supporting a dismissed claim, other claims might survive even if some are dismissed, and the plaintiff normally has a right to appeal dismissal and seek a de novo adjudication of the issue from the circuit court of appeals.

A full dismissal with leave to amend still has label 1 (i.e., a dismissal is coded the same way, even if it is without prejudice and the plaintiff may seek leave to amend the complaint and try again). A partial dismissal of some claims or theories that leaves the frozen unit alive in a material respect has label 0, meaning the claim-defendant unit was not entirely dismissed. Later amendment, settlement, reconsideration, and appellate reversal do not change the target; the model is specifically asked to predict only the trial judge’s ruling on the specific motion to dismiss. Dispositions that cannot be reliably mapped to a unit were excluded from primary scoring. These choices sacrifice some procedural detail for a stable, objectively interpretable forecast event. Even with these limitations, label construction is a substantive legal task in its own right as the labeling models had to analyze the filed motions to determine what claims were being challenged and ascertain from the judge’s ruling how those were resolved. Two labeling models had to agree on the units, and had the option to escalate ambiguous cases to me for review; in practice, the labeling process was iterative in this preliminary run, as we discovered a number of edge cases that needed to be resolved. The process started with 100 cases but seven were eliminated during the labeling process because the order did not decide the frozen claims, one was eliminated because it was a stipulated dismissal, and one was a plaintiffs’ Rule 12(c) motion for judgment on the pleadings rather than a defendant’s motion to dismiss under Rule 12(b).

### 4.2 Cohort construction and labeling

The analyzed set consists of 91 U.S. federal cases with written decisions dated June 30–August 7, 2026. The cases, as labeled, contained 409 claim-defendant units that could be forecasted. Of those, 387 are scored. Of the scored claim-defendant units, 285 were fully dismissed and 102 were not. The shared-summary preparation inventory contains 482 selected documents. The sample was not random; our pipeline favored cases where some of the necessary documents could be obtained from online databases, although we typically had to purchase a number of documents from the federal

| First written disposition of the frozen unit                       | Primary label                               |
|--------------------------------------------------------------------|---------------------------------------------|
| Entire claim dismissed, with or without leave to amend             | 1: fully dismissed                          |
| Only one nonseparable theory or remedy dismissed; claim remains    | 0: not fully dismissed                      |
| Claim dismissed against one defendant but survives against another | Separate labels for separately frozen units |
| Disposition does not reliably resolve the frozen target            | Excluded from primary scoring               |

**Table 1:** Illustrative label semantics, not additional dataset observations. The documented protocol defines unit boundaries before consulting the outcome.

PACER system to ensure that the task included all the substantive filings on the motion that were before the judge.

Of the 409 claim-defendant units, 387 are scored; 22 units therefore do not enter the headline metrics. The unscored units involved situations where the plaintiff voluntarily withdrew the frozen claim, the order did not cleanly grant or deny the frozen claim, or (for four units) we determined after freezing the unit that the frozen defendant was the wrong target or the motion did not address the frozen claim.


The case construction and labeling pipeline used models as well as human legal oversight from me. We did take pains to maintain separation between *unitization* and *outcome coding*. In particular, we wished to avoid the possibility that the prediction units could leak information about the judge’s ultimate outcome to the models that were performing prediction. AI models first were asked to identify the claim-defendant units that a motion sought to dismiss, based only on the pre-decision materials (the motion and the docket). A later stage was given access to the claim-defendant units and the judge’s order disposing of the motion, to interpret the disposition for each frozen target. The labeling pipeline we used requires structured outcome agreement and supporting disposition excerpts for automatic acceptance; disagreements between models, ambiguity, and structural defects in the cases were routed to me for lawyer review and a decision. I resolved disputed units and corrected labels as necessary. For purposes of the initial release, this process was iterative and we did not track a blinded human-audit error estimate.

### 4.3 Information separation and contamination

The models predicting outcomes were not given access to the target decision or the outcome. The construction pipeline also screens docket and filing text to avoid outcome leakage. Models had access to document text tools in a sandbox with no Internet access. Model API calls were made with provider-side web search tools disabled where applicable.

The decisions used in the benchmark were handed down after the reported knowledge or training cutoffs for fifteen of the seventeen tested models. Kimi and Muse did not publish training cutoffs, but based on the model release dates, it appears unlikely that those models were exposed to the decisions being tested in training.

The acquisition and labeling pipeline is designed to create new cohorts in the future, when new decisions become available. This enables fresh comparisons even after models might have been trained on the decisions at issue here. However, the downside is that refreshing a benchmark changes its case mix and difficulty; scores from different versions cannot be interpreted as a time series of model improvement.



**Figure 1:** Information separation in the forecasting task. Models inspect pre-decision material and frozen targets. The subsequent judicial disposition is used in the labeling and scoring path, not in the model-visible packet. In the retrospective beta, this is an information boundary rather than a claim that forecasts were made before the decisions occurred.

#### 4.4 Model execution

The main comparison uses a common document-access interface. Models can read documents using tools and produce a probability for each unit, together with a short rationale. Provider-specific reasoning settings are retained in the run registries. Appendix A lists the configuration settings. The study population is the seventeen configurations in Table 2.

#### 4.5 Scoring and statistical inference

For  $N = \sum_c |\mathcal{U}_c| = 387$  and  $C = 91$ , the primary metric is the micro-Brier loss

$$B_{\text{micro}} = \frac{1}{N} \sum_{c=1}^C \sum_{u \in \mathcal{U}_c} (p_{cu} - Y_{cu})^2. \quad (4)$$

The equal-case score changes the estimand by giving each case the same weight:

$$B_{\text{case}} = \frac{1}{C} \sum_{c=1}^C \frac{1}{|\mathcal{U}_c|} \sum_{u \in \mathcal{U}_c} (p_{cu} - Y_{cu})^2. \quad (5)$$

Lower is better for both. Micro-Brier treats each scored claim-defendant unit equally, so cases that involve a motion to dismiss many separate claims receive more weight than cases where only one claim and defendant was subject to a motion to dismiss. Equal-case Brier weights each prediction unit in inverse proportion to the number of units from that case.

The Brier rule is strictly proper for the binary target. If the true conditional event probability is  $q$ , then

$$\mathbb{E}[(p - Y)^2 \mid X, U] = (p - q)^2 + q(1 - q), \quad (6)$$

which is minimized by reporting  $p = q$ . The second term is irreducible uncertainty for that information set. A set of observable judicial events therefore create a verifiable score across the entire probability distribution, even though a single realized outcome standing alone does not reveal whether the individual probability a model assigned to that outcome was correct ex ante [18, 19].

We report accuracy at threshold  $p \geq 0.5$  as a secondary descriptive metric; in other words, accuracy percentage is simply the percentage of units where the model assigned a 50

The observed dismissal frequency is  $\hat{\pi} = 285/387 = 0.73643$ . Always predicting dismissal has accuracy 73.64%, so a constant forecast  $p = \hat{\pi}$  has micro-Brier  $\hat{\pi}(1 - \hat{\pi}) = 0.19410$ . The constant is fitted retrospectively on this cohort, so it is a descriptive reference, not a baseline derived from earlier cases. The uninformative forecast  $p = 0.5$  has Brier 0.25.

Log loss is a complementary diagnostic:

$$L = -\frac{1}{N} \sum_{c,u} [Y_{cu} \log p_{cu} + (1 - Y_{cu}) \log(1 - p_{cu})]. \quad (7)$$

For a wrong forecast assigning 0.99 to the predicted class, Brier loss is 0.9801 and log loss is approximately 4.605; at confidence 0.90 the corresponding losses are 0.81 and 2.303. Thus log loss imposes a large penalty on extreme errors, specifically, instances where the model assigned a very high probability to the outcome that did not occur.

Claim-defendant units from the same case share facts and legal issues; distinct claims in a single case may suffer from closely related legal defects, so the outcomes are not independent. We therefore measure within-case dependence using an observed-scale intraclass correlation coefficient (ICC) [70].

Across the 387 scored units in 91 cases, the observed-scale outcome ICC was 0.76 (95% CI: 0.61–0.87). Prediction losses also clustered within case because if a model failed to account, for example, for the possibility that a given motion might be denied on procedural grounds, all of the claims in a case that it expected to be dismissed could survive: GPT-6 Sol’s Brier-loss ICC was 0.42 (95% CI: 0.24–0.58). These intervals use 20,000 whole-case bootstrap resamples, and assume independence between cases.

The observed dependence supports resampling whole cases, with model predictions kept paired, when estimating uncertainty in score differences [71]. The seventeen-configuration table is descriptive and does not claim a significance ordering. It does not measure variation across repeated model runs [46].

## 5 Results

### 5.1 Frontier models substantially outperform unconditional references

Table 2 reports the seventeen configurations on the ranked leaderboard. GPT-6 Sol achieves the best observed micro-Brier (0.11763), equal-case Brier (0.12949), and accuracy (327/387, 84.50%). All seventeen beat the retrospectively fitted constant micro-Brier of 0.19410; GPT-6 Sol’s micro-Brier is lower than that constant by about 39.4%.

Sixteen configurations also beat majority-class accuracy (the 73.64% accuracy of predicting dismissal for every unit). Muse Spark 1.3 does not: 279/387 correct forecasts is 72.09%. Gemini 3.1 Pro Preview has the highest micro-Brier in the panel (0.18345) and still beats that accuracy baseline (302/387, 78.04%).

More expensive or newer configurations do not uniformly score better:

- Astra’s micro-Brier is worse than those of GPT-5.6 Sol and Luna; Astra improves on Luna’s equal-case Brier but not its accuracy.
- Grok 4.7 has worse observed scores than Grok 4.6.

| Configuration          | Micro Brier | Case Brier | Correct | Acc. (%) | Cost (\$) |
|------------------------|-------------|------------|---------|----------|-----------|
| GPT-6 Sol              | 0.11763     | 0.12949    | 327/387 | 84.50    | 11.01     |
| Claude Fable 5.1       | 0.12877     | 0.13653    | 318/387 | 82.17    | 83.18     |
| Grok 4.6               | 0.13368     | 0.13480    | 312/387 | 80.62    | 15.72     |
| Claude Opus 5          | 0.13987     | 0.13548    | 310/387 | 80.10    | 45.89     |
| GPT-5.6 Sol            | 0.14009     | 0.13319    | 313/387 | 80.88    | 24.91     |
| Kimi K3                | 0.14130     | 0.15367    | 303/387 | 78.29    | 16.87     |
| GPT-5.6 Luna           | 0.14322     | 0.15084    | 313/387 | 80.88    | 1.48      |
| Claude Opus 5.5        | 0.14540     | 0.14762    | 299/387 | 77.26    | 20.59     |
| GPT-6.1 Sol            | 0.14948     | 0.13885    | 304/387 | 78.55    | 10.68     |
| GPT-6 Astra            | 0.15221     | 0.14275    | 305/387 | 78.81    | 59.74     |
| Gemini 3.8 Flash       | 0.15331     | 0.13900    | 310/387 | 80.10    | 35.61     |
| Grok 4.7               | 0.15833     | 0.16454    | 298/387 | 77.00    | 33.00     |
| Claude Sonnet 5        | 0.15848     | 0.15964    | 293/387 | 75.71    | 24.68     |
| GPT-6 Luna             | 0.16368     | 0.17379    | 296/387 | 76.49    | 0.62      |
| Claude Sonnet 5.5      | 0.17298     | 0.17183    | 295/387 | 76.23    | 6.99      |
| Muse Spark 1.3         | 0.18121     | 0.17845    | 279/387 | 72.09    | 8.78      |
| Gemini 3.1 Pro Preview | 0.18345     | 0.18080    | 302/387 | 78.04    | 59.04     |
| Cohort-rate reference  | 0.19410     | —          | 285/387 | 73.64    | —         |
| Constant $p = 0.5$     | 0.25000     | 0.25000    | —       | —        | —         |


**Table 2:** Full-document results, ordered by micro-Brier loss. Each model row covers 91 cases and 387 scored units. The last column reports the cost estimate under standard pricing, as detailed in the text. The cohort-rate reference is retrospectively fitted, and its accuracy column uses always-dismiss classification.

- GPT-6.1 Sol is worse than GPT-6 Sol on micro-Brier (0.14948 versus 0.11763).
- Claude Sonnet 5.5 is worse than Claude Sonnet 5 on micro-Brier (0.17298 versus 0.15848), though its unit accuracy is slightly higher (76.23% versus 75.71%).
- Claude Opus 5.5 is worse than Claude Opus 5 on micro-Brier (0.14540 versus 0.13987) and on accuracy (77.26% versus 80.10%).
- GPT-6 Luna’s full-document scores are worse than GPT-5.6 Luna’s: micro-Brier 0.16368 versus 0.14322, and accuracy 76.49% versus 80.88%.

## 5.2 Confidence and calibration

Define confidence  $h_{cu} = \max(p_{cu}, 1 - p_{cu})$  and the high-confidence subset  $H_m = \{(c, u) : h_{cu} \geq 0.9\}$  for configuration  $m$ . Comparing the mean confidence in  $H_m$  with classification accuracy in  $H_m$  gives a useful conditional calibration diagnostic. Of course, this is not a full reliability curve; the subset differs by model, and its units remain clustered by case.

OpenAI models appeared more likely than others to make high-confidence forecasts. GPT-5.6 Sol, GPT-5.6 Luna, and Astra make high-confidence forecasts on 58.9%, 47.0%, and 53.5% of units respectively, compared with 20.9% for Fable 5.1. Their mean confidence in those subsets is approximately 95–97%, while realized accuracy is approximately 89–90%, suggesting some overconfidence, particularly for Astra, which missed 21 of 207 such predictions. Gemini 3.8 Flash also had relatively high sharpness, making 276 high-confidence forecasts (71.3% of units) and missing 35. Gemini 3.1 Pro Preview is sharper still, making 320 high-confidence forecasts and missing 57. GPT-6 Sol provides a useful counterexample within OpenAI, missing 8 of 169 high-confidence



**Figure 2:** Micro- and equal-case Brier scores for the same seventeen configurations. The connecting segments show weighting sensitivity, not uncertainty intervals. The dashed constant-probability reference applies only to micro-Brier.

forecasts, for 95.3% subset accuracy. GPT-6.1 Sol is sharper and less accurate in that subset, missing 18 of 199. GPT-6 Luna misses 19 of 183. Claude Opus 5.5 misses none of 65.


Other high-performing models were less sharp. For example, Fable has no errors among 81 high-confidence units, while Kimi and Grok 4.6 each miss one. Claude Sonnet 5.5 makes 63 high-confidence forecasts and misses 5. The practical finding is that confidence varies substantially and it materially affects proper scores. Astra’s underperformance relative to cheaper models like Sol, for example, appears largely due to several high-confidence misses, which resulted in large Brier losses.

### 5.3 Inference cost

The normalized successful-workload costs range from \$0.62 for GPT-6 Luna to \$83.18 for Fable. These are not actual invoice charges. For example, Fable was initially run without caching, resulting in an actual cost of \$189.70. The normalized price was reconstructed after the fact using turn-by-turn token counts, with the reconstruction validated against real billing data from other Anthropic model runs. In several other cases, provider discounts that we received were normalized to standard API rates. On these standardized estimates, GPT-6 Luna, GPT-5.6 Luna, and GPT-6 Sol form the observed micro-Brier cost–score frontier.

### 5.4 Evaluating Jev, a non-autoregressive discriminative decision model

Jev is described by its developer as a probabilistic “System One” model, exposing native Boolean probabilities rather than generating text or using tools agentially as frontier models do [49]. Because Jev’s context window and lack of agentic tool calling differ from the document-enabled configurations used with frontier models, Jev was evaluated as a separate experiment.




**Figure 3:** Conditional confidence diagnostic on predictions with  $\max(p, 1 - p) \geq 0.9$ . Each row uses the subset selected by that model, rather than a common set of units. Wrong/selected counts: GPT-6 Sol: 8/169; Fable 5.1: 0/81; Grok 4.6: 1/112; Opus 5: 4/121; GPT-5.6 Sol: 26/228; Kimi K3: 1/92; GPT-5.6 Luna: 20/182; Opus 5.5: 0/65; GPT-6.1 Sol: 18/199; GPT-6 Astra: 21/207; Gemini 3.8 Flash: 35/276; Grok 4.7: 4/98; Sonnet 5: 5/60; GPT-6 Luna: 19/183; Sonnet 5.5: 5/63; Muse Spark 1.3: 15/135; Gemini 3.1 Pro: 57/320. This is not a full reliability diagram or an estimate of statistical significance.

The full set of documents relevant to motions to dismiss is frequently much longer than Jev’s context window, so GPT-5.6 Luna summarized each document with instructions to preserve material facts, arguments, authorities, and procedural posture without commentary or analysis. Luna received only the motion documents that forecasting models received, to prepare summaries. Jev (and two Luna comparator configurations) then received the summaries prepared by the prior Luna session, and the unit questions, in a single request per case, without document tools or multiple turns. The conditions compare Jev through the `typesafe-ai/jev` Gateway alias, GPT-5.6 Luna with high reasoning, and GPT-6 Luna without reasoning.

Jev was not well-calibrated for this task. Jev’s micro-Brier of 0.36754 is worse than the 0.25 loss of always reporting  $p = 0.5$ . Its mean predicted dismissal probability was approximately 0.305, while the observed frequency in this data set was 0.736. This resulted in severe underprediction in the aggregate. Its 32.6% threshold accuracy is also below majority-class prediction.

| Condition                             | Micro Brier | Case Brier | Acc. (%) |
|---------------------------------------|-------------|------------|----------|
| GPT-5.6 Luna full documents           | 0.14322     | 0.15084    | 80.88    |
| GPT-5.6 Luna summary + high reasoning | 0.15489     | 0.16119    | 78.55    |
| GPT-6 Luna summary + no reasoning     | 0.17152     | 0.18694    | 74.16    |
| Jev summary + one-shot                | 0.36754     | 0.34416    | 32.56    |

**Table 3:** Shared-summary results and the full-document Luna reference. All rows use the same 91 cases and 387 scored units. Only the last three rows share a one-request summary interface; the first row is a descriptive full-document reference.



**Figure 4:** Probability quality with a shared summary cache. Both Luna conditions exploit predictive signal in the summaries, while Jev scores worse than the constant  $p = 0.5$  forecast. The full-document reference changes both information access and interaction. Connecting segments indicate weighting sensitivity, not uncertainty intervals.

We tested Luna using the same summaries that Jev received in order to determine how much of Jev’s underperformance was attributable to receiving only summaries of the litigation documents, rather than the full documents that we used in our primary evaluation. However, Luna performed much better than Jev even on this comparison. GPT-5.6 Luna with the identical summaries achieves micro-Brier 0.15489 and accuracy 78.6%. The paired case-cluster bootstrap gives a Luna-minus-Jev difference of  $-0.21265$ , with a 99.44% Bonferroni-adjusted interval (95% familywise coverage) of  $[-0.32644, -0.06433]$  across the three model pairs and three metrics. The result remains exploratory because potential cross-case dependence was not assessed.

The shared-summary control shows that the summaries contain useful predictive signal that the Jev configuration does not exploit. Of course, that does not indicate that every important detail survived the summarization process. Luna performed better in the primary evaluation, where it had full document access, than it did in this comparison, with only summaries. That could mean either that the summaries lost some relevant signal, or that the interactive environment (where Luna could choose which documents to read) allowed better performance, or some mix of the two.

Although this experiment indicates that Jev is inferior to frontier models on this particular task, it does not establish that probabilistic or non-generative architectures are unsuitable for litigation prediction or that they will necessarily underperform. It is possible that Jev was not exposed to sufficient training data for this task, or that its architecture or interface was not well-suited to the task, or some combination of the three.

Summary preparation cost an estimated \$10.32. Successful inference estimates are approximately \$0.044 for Jev, \$1.62 for GPT-5.6 Luna (with high reasoning), and \$0.103 for GPT-6 Luna (which was run without reasoning as a control).

## 6 Alternative evaluation signals: preferences and rubric-based rewards

### 6.1 Evaluating open-ended legal work

JudgmentBench compares two ways of evaluating open-ended legal work: asking practicing attorneys which of two outputs they prefer, and asking them to score outputs against specified criteria. Across

30 legal tasks with outputs at three constructed quality levels, comparative judgments recovered the intended quality ordering substantially better than rubric scores and required less than half the annotation time [7]. The constructed quality levels and task selection limit what this establishes about professional work more broadly. Together with judicial outcomes, these approaches provide three distinct evaluation signals: what a court did, which work product an expert prefers, and whether an output satisfies specified requirements.

Preferences and rubrics can assess aspects of advice, analysis, and drafting that a forecast score does not directly measure. Rubrics make the requirements explicit; comparative judgments let reviewers weigh the work as a whole without an exhaustive checklist. Either can supply evaluation data, and either can be used to construct training rewards. Collecting preference judgments for evaluation is not itself reinforcement learning.

Harvey’s Legal Agent Benchmark (LAB) provides a large-scale case study of the rubric approach. Launched on May 6, 2026, it contains 1,200 tasks spanning 24 practice areas and more than 75,000 criteria, with LLM judges scoring model outputs [41, 57]. The LAB-AA partnership with Artificial Analysis extends its use in evaluation [58].<sup>1</sup> Harvey and Applied Compute also report full-parameter reinforcement learning of a legal agent using rubric-based rewards [9]. Harvey’s August 2026 Tenet research preview reports gains on LAB and other legal benchmarks after training a model based on Kimi K3 [10]. These are company-reported training results; they make the quality of the criteria consequential for learning as well as evaluation.

## 6.2 Preliminary audit of LAB

A rubric can assess richer work than a judicial forecast, but that broader coverage is useful only if the rubric specifies the task correctly and the grader applies it reliably. Errors can arise in the requested work, the supplied record, the legal answer prescribed by a criterion, or the grader’s treatment of a model response. These are distinct failure modes. The underlying question is one of construct validity: whether the score measures the capability the benchmark claims to assess [30].

Our retained AI-assisted review covers the pinned litigation-and-dispute-resolution subset: 52 tasks and 2,858 criteria. GPT-6 Sol flagged 616 criteria; Claude Opus 5.5’s final reconciliation flagged 564. Their union contains 835 flagged criteria, leaving 2,023 unflagged by either final report. Of 345 overlapping flags, both reviewers called 83 problematic and 135 arguable; 127 received different levels of concern. The union represents 29.2% of the 2,858 criteria; the other 70.8% were not flagged, which does not establish that they are correct. These are review outputs, not a verified error rate. Opus’s reconciliation saw Sol’s findings, so agreement in the final reports is not independent confirmation [27]. A human review of a random sample of 25 flagged criteria (Appendix C) found 16 defective in ways that affect grading, implying with 95% confidence that at least 383 of the 835 flagged criteria are defective.

If a defective criterion supplies a training reward, it also creates a problem of reward misspecification. Reward hacking is a possible learned response; reward-model overoptimization is a possible divergence between an improving proxy score and the underlying target [31, 32]. Our audit examines tasks and criteria, not a training trajectory. It therefore cannot establish that rubric-based RL is ineffective, that synthetic tasks generally harm performance, or that training on litigation tasks

---

<sup>1</sup>Harvey’s product announcements report LAB all-pass results for Opus 4.8 (10.4%, May 28), Fable 5 (13.3%, June 9), and Sonnet 5 (5.8%, June 30) [59–61]. These are Harvey’s evaluations of third-party models, not evidence that each developer independently adopted LAB as its definitive legal benchmark. Prominence does not establish evaluation validity.

caused degradation.

OpenAI’s 2026 analysis of SWE-bench Verified provides a useful methodological analogy: its selected audit of 138 difficult tasks reported problems involving test requirements and task specification [48]. That finding supports examining alignment between tasks and their scoring rules; it provides no defect estimate for LAB. A stronger legal study would compare matched environments built from synthetic and authentic records, independently adjudicate their criteria, and evaluate transfer to held-out professional tasks.

### 6.3 A motion-to-dismiss case study

The task `analyze-counterparty-motion-to-dismiss` requests an opposition issues memorandum in a fictional software-services dispute involving alleged pre-contractual representations, implementation defects, and jurisdictional, venue, and pleading defenses. The AI reviewers flagged 16 of its 34 criteria. My targeted review uses the task and supplied documents at the pinned upstream revision recorded by the local audit [29]. The analysis below distinguishes direct record observations from legal interpretations and unverified AI flags; it does not treat every flag or solver failure as a grading defect.

**Forum selection and litigation strategy.** Complaint paragraph 48 sets out a clause selecting state or federal courts in Travis County, Texas, while the action is filed in the Northern District of Georgia. Paragraphs 18–19 assert statutory venue in Georgia, and paragraph 123 argues that the statutory deceptive-practices claim falls outside the clause. Those arguments may be weak, but the complaint is not wholly silent about the forum issue. A rigorous task should ask whether transfer is likely, which claims the clause covers, and whether a procedural objection changes the client’s position. Under *Atlantic Marine*, a valid clause selecting another federal forum is generally enforced through section 1404(a), while a state or foreign forum implicates forum non conveniens [62]. Correcting an opponent’s procedural vehicle is not the same as establishing that the chosen Georgia forum will retain the dispute. A criterion that emphasizes the former without measuring the latter can reward issue identification without strategic assessment.

**Accrual and pleaded dates.** Criteria C-013–C-015 require an accrual dispute and alternative discovery dates. The complaint alleges discovery of integration problems in August 2022, latency problems by November 2022, an acknowledgment in January 2023, and an unsuccessful patch on April 15, 2023. Paragraph 83 alleges that the plaintiff could not know the full extent of the breach until the patch failed, claims reliance on a promised fix, and states a March 22, 2024 filing date. These allegations require analyzing knowledge sufficient for accrual, distinct breaches, tolling or estoppel, and enforceability of the one-year contract term. Texas law itself restricts contractual limitations periods, subject to statutory exceptions [63].

The appropriate critique is therefore not that the dates mechanically prove the entire claim time-barred. Nor is it enough to announce that knowledge is always a fact issue immune from dismissal. A plaintiff’s allegations can establish an affirmative defense on the face of the pleading, while a genuine factual or legal issue may prevent that result [64]. The rubric should require applying the governing rule to the actual chronology and contract. It should allow the model to distinguish stronger from weaker opposition arguments instead of requiring the conclusion that every accrual question is unresolved.

**Integration versus non-reliance.** Criterion C-012 requires the model to characterize the merger-clause defense as legally incorrect. MSA section 14.7 also contains express non-reliance language, which is analytically distinct from an ordinary integration clause. Texas authority distinguishes these contractual concepts and examines the language and surrounding circumstances [65]. The source record therefore does not justify a categorical rubric requiring the ordinary-merger-clause answer without that analysis. This observation identifies an overbroad criterion; it does not establish that the particular disclaimer is enforceable or that the fraud claim must fail.

**Hidden output requirements.** C-028 requires concrete opposition recommendations for at least eight issues, although the prompt supplies no minimum count. C-022–C-025 prescribe high-severity treatment of specified issues. These criteria permit equivalent severity language; it would be inaccurate to say they necessarily require literal labels. C-026 prescribes organization tied to the motion’s framework and expressly permits omitting Rule 12(b)(1) when no genuine issue exists. Whether those requirements improve usefulness to counsel is a measurement question. They should not be confused with independently established legal errors.

**Document construction and limits of the review.** The supplied documents have programmatic-generation metadata and include a stale table-of-contents instruction in the agreement. The motion nevertheless contains a caption, table of contents, and table of authorities. We have not established a court-rule violation warranting striking it, proven LLM authorship, or estimated how often such artifacts occur across LAB. Synthetic design can be useful for controlled testing; the concern is whether the rewarded conduct transfers to actual matters. The local audit also records genuine solver mistakes and unresolved legal disagreements. A benchmark-wide conclusion would require a representative expert review of source documents, criteria, model outputs, and grading decisions.

## 6.4 Comparing supervision signals

**Judicial outcomes.** Outcomes provide an external reference that the evaluator does not choose. The task designer must still define the event, separate pre-decision information from later evidence, and interpret the disposition correctly. Once that target is fixed, a proper scoring rule evaluates the forecast against what happened without requiring an expert to prescribe the best legal argument. The coverage is narrow: predicting a disposition does not directly evaluate the quality of a brief, the usefulness of advice, or the faithfulness of an explanation. A successful forecast may combine doctrinal analysis with institutional priors rather than identify the legally correct answer.

**Rubrics.** Rubrics can directly assess several dimensions of professional work, including factual support, legal analysis, completeness, and compliance with a client’s instructions. Their difficulty is specification: the criteria must distinguish indispensable analysis from optional presentation and allow defensible alternative answers. All-pass aggregation magnifies this problem because one defective criterion can make an otherwise useful work product fail. Some requirements are indispensable, so all-pass scoring is not inherently inappropriate; the question is whether the required items are valid. Nor does learning from rubric rewards necessarily produce only rubric compliance: Gunjal and colleagues report gains on the multiple-choice GPQA-Diamond benchmark relative to direct Likert-based judge rewards, although that finding does not establish transfer to legal work [8].

**Preferences.** Comparative judgments can capture professional assessments that are difficult to specify in advance, including whether a memorandum prioritizes the issues that matter. They

also depend on reviewer expertise, objectives, and consistency. Disagreement may reflect different defensible judgments or measurement noise, and a preferred output need not be the most accurate or useful one. Research outside law shows that human preferences can favor persuasive agreement with a user’s beliefs over factual correctness, and that optimizing a learned reward can amplify such biases [33, 34]. In legal evaluation, a corresponding concern is that a reviewer might prefer confident or congenial advice to a better-supported adverse assessment. This is a risk to investigate, not a finding that JudgmentBench’s attorney judgments exhibit sycophancy. Preference evaluation and preference-based training must be assessed separately.

These signals are complementary. Authentic case records can support an outcome forecast, a rubric checking necessary legal and factual requirements, and a blinded expert comparison of the resulting advice or drafting. Reporting these assessments separately would make disagreements visible: a model might forecast accurately while giving poor advice, or produce polished work that misses the issue most likely to determine the result. The research question is which signals, individually or together, support useful legal judgment beyond the tasks on which they are scored.

## 7 Future work: from outcome evaluation to outcome feedback

### 7.1 Outcome feedback as a training signal

LegalForecastGym implements a separate, single-step forecasting environment intended for post-training experiments. It renders a static litigation packet and receives a structured action containing a dismissal probability and cited evidence spans. Training loaders admit only explicitly designated training collections and reject locked test matters. Episode assembly keeps the label out of the observation. The current phase does not implement a multi-turn litigation simulator or train an agent to take procedural actions.

The checked Gym artifacts contain no completed parameter-update experiment or held-out before/after training comparison. They do contain a live sampling comparison on 50 synthetic development episodes, used to assess output format and select an open-weight sampling model. That run is infrastructure evidence, not evidence of legal forecasting improvement. A separately recorded Gym export of 100 matters and 425 scoreable units is also distinct from the 91-case, 387-unit benchmark analyzed here; those populations are not pooled.

The intended learning study asks whether earlier outcome feedback improves predictions on later, disjoint matters and whether any gains extend beyond forecasting. Parameter learning, in-context feedback, and changes to the surrounding agent system are different interventions. Each needs a matched control and a held-out endpoint. Forecasting literature makes such learning plausible, and OpenForecaster reports calibration transfer to other question-answering benchmarks; those findings do not establish improved litigation strategy [24, 25]. Appendix D specifies a concrete prospective design.

### 7.2 Testing transfer beyond forecasting

A reduction in forecasting loss on later cases would be evidence of improvement on the prediction task. Transfer would require improvement on additional legal tasks assessed independently of that score. The prospective study in Appendix D would therefore compare the trained model with a matched control on issue prioritization, strategic advice, and drafting using later, disjoint matters. The relevant questions are whether the model selects the strongest arguments, relates its advice to

the client’s objectives and litigation risks, and produces work supported by the record and applicable law. Explanations should be evaluated separately from the probabilities they accompany.

Practicing lawyers would assess paired outputs with model identity and training condition concealed and presentation order randomized. Both models would receive the same pre-decision record and client instructions; reviewers would initially assess the work without seeing the eventual disposition to limit hindsight effects. Expert comparisons would be supplemented by checks of factual and legal support, with reviewer disagreement reported. Improvement on these independent assessments would provide evidence of transfer; a better forecast score alone would not. LegalForecastGym provides infrastructure for testing this hypothesis, not a completed learning or transfer result.

## 8 Limitations

**Selection into adjudication.** The target population is a set of accessible cases with written motion decisions, not all disputes, all filed cases, or all contemplated motions. Parties settle, abandon claims, amend pleadings, or choose not to sue. Cases reaching a published disposition may differ in uncertainty, stakes, representation, and bargaining conditions. Calibration conditional on reaching this dataset does not establish calibration for counterfactual outcomes in the broader dispute population [52]. At the same time, selection does not prove that learned skills cannot transfer: that is an empirical question requiring an independently defined evaluation population.

**Judge, court, and party priors.** We have not established how much performance comes from close legal analysis rather than judge, court, party, or case-type priors.

**Labels and completeness.** Unitization can omit, merge, or split legally distinct targets; dispositions can be ambiguous or procedural. Human oversight helps but does not quantify residual error. The absence of a complete blinded label audit and acquisition/exclusion census limits claims about both label accuracy and representativeness. Excluding 22 unscored units preserves an interpretable target but may also change difficulty. These limitations deserve the same scrutiny as disputed rubrics in another benchmark.

**Statistical and execution limits.** One successful prediction per unit does not characterize run-to-run variance, operational reliability, or expected total cost. The cohort is small and imbalanced. Case clustering does not resolve unverified dependence among related cases. Reported scores compare model–harness–provider configurations, not isolated model weights.

**Temporal and transfer limits.** A provider-reported cutoff and disabled browsing reduce specific leakage risks without proving contamination absent. New benchmark versions change the case distribution. Generalization to other jurisdictions, procedural stages, settlement decisions, drafting, or strategic advice is not established by this study.

## 9 Conclusion

LegalForecastBench operationalizes an uncertain but consequential legal judgment as a probabilistic forecast of a later judicial event. In the beta, all seventeen document-enabled configurations outperform a cohort-fitted constant in micro-Brier, while confidence behavior and model ordering differ in informative ways. GPT-6 Sol has the strongest observed scores, and the shared-summary

comparison exposes a large gap between Jev and Luna that cannot be explained by the summaries containing no signal.

Real outcomes provide an external scoring reference, but useful training signals still require careful task definition, information boundaries, and legal labeling. The next research question is whether feedback on such forecasts improves performance on later cases and transfers to independently assessed legal work. The present evidence motivates that question; it does not answer it.

## **Reproducibility, access, and disclosure**

This manuscript reports the seventeen-configuration comparison and checks its numerical claims against the retained repository artifacts as of September 30, 2026. The accompanying aggregate data and figure script regenerate the plots without court-document bytes. Code and result artifacts are maintained in the [LegalForecastBench repository](#); access to restricted source documents and private adjudication material is a separate audit process. Availability of code is not equivalent to unrestricted source-data replication.

The author designed the benchmark and supplied domain judgments. AI assistants supported this manuscript’s drafting, literature retrieval, and numerical and source checks. The LAB audit discussed above is explicitly an AI-assisted preliminary analysis; the author’s targeted task review is distinguished from a completed independent audit. No model-weight reinforcement-learning experiment is claimed unless explicitly identified as a result of an external cited study.

## References

- [1] Daya Guo, Dejian Yang, Haowei Zhang et al. DeepSeek-R1 incentivizes reasoning in LLMs through reinforcement learning. *Nature* 645, 633–638 (2025); *arXiv:2501.12948v2*, 2025. <https://www.nature.com/articles/s41586-025-09422-z>.
- [2] Zhiyuan Zeng, Hamish Ivison, Yiping Wang et al. RLVE: Scaling Up Reinforcement Learning for Language Models with Adaptive Verifiable Environments. *ICML 2026, Proceedings of Machine Learning Research* 306:153406–153428, 2026. <https://proceedings.mlr.press/v306/zeng26j.html>.
- [3] David Silver, Thomas Hubert, Julian Schrittwieser et al. Mastering Chess and Shogi by Self-Play with a General Reinforcement Learning Algorithm. *arXiv:1712.01815 (2017)*, preprint; precursor to the distinct 2018 *Science* article, 2017. <https://arxiv.org/abs/1712.01815>.
- [4] Thomas Hubert, Rishi Mehta, Laurent Sartran et al. Olympiad-level formal mathematical reasoning with reinforcement learning. *Nature* 651, 607–613; version of record 13 March 2026, 2026. <https://www.nature.com/articles/s41586-025-09833-y>.
- [5] OpenAI. Finite Time Blowup for Navier–Stokes. *Technical manuscript*, September 2026. <https://cdn.openai.com/pdf/32d9f210-8b73-45e0-91bc-82a30aef8a9a/navier-stokes.pdf>.
- [6] OpenAI. On the Navier–Stokes Millennium Prize Problem. *Research announcement*, September 8, 2026; updated September 10, 2026. <https://openai.com/index/navier-stokes-solution/>.
- [7] Russell Yang, Ruishi Chen, Pierce Kelaita et al. JudgmentBench: Comparing Rubric and Preference Evaluation for Quality Assessment. *arXiv:2605.25240v3*, preprint, 2026. <https://arxiv.org/abs/2605.25240>.
- [8] Anisha Gunjal, Anthony Wang, Elaine Lau, Vaskar Nath, Yunzhong He, Bing Liu, and Sean Hendryx. Rubrics as Rewards: Reinforcement Learning Beyond Verifiable Domains. In *International Conference on Learning Representations*, 2026. First posted as *arXiv:2507.17746* in 2025. [https://proceedings.iclr.cc/paper\\_files/paper/2026/hash/cfd7bee7a651ee9af525098ef67a9e45-Abstract-Conference.html](https://proceedings.iclr.cc/paper_files/paper/2026/hash/cfd7bee7a651ee9af525098ef67a9e45-Abstract-Conference.html).
- [9] Niko Grupen, Julio Pereyra, Gabe Pereyra et al. Training a Legal Agent with Applied Compute. *Primary online source*, 2026. <https://www.harvey.ai/blog/training-a-legal-agent-with-applied-compute>.
- [10] Calvin Qi, Vasudha Rengarajan, Julio Pereyra et al. Harvey Tenet Research Preview. *August 20, 2026 company research report*, 2026. <https://www.harvey.ai/blog/post-training-update-harvey-tenet>.
- [11] Long Ouyang, Jeffrey Wu, Xu Jiang et al. Training language models to follow instructions with human feedback. *NeurIPS 2022*, 2022. <https://arxiv.org/abs/2203.02155>.
- [12] Neel Guha, Julian Nyarko, Daniel E. Ho et al. LegalBench: A Collaboratively Built Benchmark for Measuring Legal Reasoning in Large Language Models. *NeurIPS 2023, Datasets and Benchmarks*, 2023. <https://arxiv.org/abs/2308.11462>.
- [13] National Conference of Bar Examiners. About the Multistate Bar Examination. Official examination description, n.d.; accessed September 29, 2026. <https://www.ncbex.org/exams/mbe/about-mbe>.

- [14] H. L. A. Hart. Positivism and the Separation of Law and Morals. *Harvard Law Review*, 71(4):593–629, 1958. Especially pp. 607–608. <https://www.jstor.org/stable/1338225>.
- [15] Lon L. Fuller. Positivism and Fidelity to Law—A Reply to Professor Hart. *Harvard Law Review*, 71(4):630–672, 1958. Especially pp. 661–664. <https://www.jstor.org/stable/1338226>.
- [16] United States Congress. Declaratory Judgment Act, 28 U.S.C. § 2201(a). Text consulted September 29, 2026. <https://www.law.cornell.edu/uscode/text/28/2201>.
- [17] Supreme Court of the United States. Federated Department Stores, Inc. v. Moitie, 452 U.S. 394, 398–399 (1981). <https://www.govinfo.gov/content/pkg/USREPORTS-452/pdf/USREPORTS-452-394.pdf>.
- [18] Glenn W. Brier. Verification of Forecasts Expressed in Terms of Probability. *Monthly Weather Review* 78(1) (1950), 1–3, 1950. [https://journals.ametsoc.org/view/journals/mwre/78/1/1520-0493\\_1950\\_078\\_0001\\_vofeit\\_2\\_0\\_co\\_2.xml](https://journals.ametsoc.org/view/journals/mwre/78/1/1520-0493_1950_078_0001_vofeit_2_0_co_2.xml).
- [19] Tilmann Gneiting, Adrian E. Raftery. Strictly Proper Scoring Rules, Prediction, and Estimation. *Journal of the American Statistical Association* 102 (2007), 359–378, 2007. <https://doi.org/10.1198/016214506000001437>.
- [20] Oliver Wendell Holmes, Jr. The Path of the Law. *Harvard Law Review*, 10(8):457–478, 1897. Especially pp. 457, 460–461. <https://www.jstor.org/stable/1322028>.
- [21] Cornerstone Research and Stanford Law School Securities Class Action Clearinghouse. Securities Class Action Filings—2024 Year in Review. *Research report*, 2025. p. 16, status of core federal securities class action filings as of the end of 2024. <https://securities.stanford.edu/research-reports/1996-2024/Securities-Class-Action-Filings-2024-Year-in-Review.pdf>.
- [22] Robert H. Mnookin and Lewis Kornhauser. Bargaining in the Shadow of the Law: The Case of Divorce. *Yale Law Journal*, 88(5):950–997, 1979. <https://doi.org/10.2307/795824>.
- [23] Paul C. Saunders. Legal Skills and Professional Values: A Handbook. *Journal of the New York State Judicial Institute on Professionalism in the Law*, 7(1), Fall 2019. Reproducing and discussing the MacCrate skills framework. <https://www.nycourts.gov/LegacyPDFS/ip/JIPL/pdf/JIPL%20-%20FINAL%20-%20Skills%20and%20Professional%20Values.pdf>.
- [24] Benjamin Turtel, Danny Franklin, Kris Skotheim et al. Outcome-based Reinforcement Learning to Predict the Future. *Transactions on Machine Learning Research (November 2025)*; *arXiv:2505.17989v4*, 2025. <https://arxiv.org/abs/2505.17989>.
- [25] Nikhil Chandak, Shashwat Goel, Ameya Prabhu et al. Scaling Open-Ended Reasoning To Predict the Future. *arXiv:2512.25070v2 (January 2026)*; *preprint*, 2026. <https://arxiv.org/abs/2512.25070>.
- [26] Junjie Chen, Haitao Li, Qilei Zhang et al. Multi-Source Retrieval and Reasoning for Legal Sentencing Prediction. *arXiv:2602.04690v3*; *arXiv record comments 'EMNLP 2026 main' (publisher proceedings record not independently checked)*, 2026. <https://arxiv.org/abs/2602.04690>.
- [27] GPT-6 Sol audit workers, Claude Opus 5.5 audit workers, Coordinating agents; individual provenance not fully recorded. Harvey LAB litigation and dispute resolution audit: GPT-6 Sol

- and Claude Opus 5.5. *Local research artifact, pinned September 29, 2026*, 2026. <https://github.com/johnhughes3/LegalForecastBench/tree/58371a0e13494e591710e2ae3da8e428eebc359/docs/harvey-lab-audit/litigation-dispute-resolution>.
- [28] John Hughes. Human review of AI-flagged Harvey LAB litigation criteria: protocol, sample, and worksheet. *Research artifact*, 2026. <https://github.com/johnhughes3/LegalForecastBench/tree/main/docs/harvey-lab-audit/litigation-dispute-resolution/human-review>.
  - [29] Harvey AI task authors; upstream authorship not individually listed. Pinned LAB task.json for analyze-counterparty-motion-to-dismiss. *Primary online source*, 2026. <https://github.com/harveyai/harvey-labs/blob/1dd81403b2fbb60596f7aea3fcecfaad7bf73143/tasks/litigation-dispute-resolution/analyze-counterparty-motion-to-dismiss/task.json>.
  - [30] Abigail Z. Jacobs and Hanna Wallach. Measurement and Fairness. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, pp. 375–385, 2021. <https://doi.org/10.1145/3442188.3445901>.
  - [31] Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. Concrete Problems in AI Safety. *arXiv:1606.06565v2*, 2016. <https://arxiv.org/abs/1606.06565>.
  - [32] Leo Gao, John Schulman, and Jacob Hilton. Scaling Laws for Reward Model Overoptimization. In *Proceedings of the 40th International Conference on Machine Learning*, PMLR 202:10835–10866, 2023. <https://proceedings.mlr.press/v202/gao23h.html>.
  - [33] Mrinank Sharma, Meg Tong, Tomasz Korbak et al. Towards Understanding Sycophancy in Language Models. *ICLR 2024; first posted as arXiv:2310.13548 in 2023*, 2024. <https://arxiv.org/abs/2310.13548>.
  - [34] Itai Shapira, Gerdus Benade, and Ariel D. Procaccia. How RLHF Amplifies Sycophancy. *ICML 2026; arXiv:2602.01002*, 2026. <https://arxiv.org/abs/2602.01002>.
  - [35] Theodore W. Ruger, Pauline T. Kim, Andrew D. Martin, Kevin M. Quinn. The Supreme Court Forecasting Project: Legal and Political Science Approaches to Predicting Supreme Court Decision-Making. *Columbia Law Review* 104 (2004), 1150–1210, 2004. [Author-hosted article PDF](#).
  - [36] Daniel Martin Katz, Michael J. Bommarito II, Josh Blackman. A General Approach for Predicting the Behavior of the Supreme Court of the United States. *PLOS ONE* 12(4): e0174698 (2017), 2017. <https://journals.plos.org/plosone/article?id=10.1371%2Fjournal.pone.0174698>.
  - [37] Nikolaos Aletras, Dimitrios Tsarapatsanis, Daniel Preotiuc-Pietro, Vasileios Lamos. Predicting Judicial Decisions of the European Court of Human Rights: A Natural Language Processing Perspective. *PeerJ Computer Science* 2:e93 (2016), 2016. <https://peerj.com/articles/cs-93/>.
  - [38] Chaojun Xiao, Haoxi Zhong, Zhipeng Guo et al. CAIL2018: A Large-Scale Legal Dataset for Judgment Prediction. *arXiv:1807.02478 (2018); CAIL challenge dataset paper*, 2018. <https://arxiv.org/abs/1807.02478>.
  - [39] Masha Medvedeva, Pauline McBride. Legal Judgment Prediction: If You Are Going to Do It, Do It Right. *Proceedings of the Natural Legal Language Processing Workshop (ACL)*, 2023, pp. 73–84, 2023. <https://aclanthology.org/2023.nllp-1.9/>.

- [40] Huiyuan Xie, Felix Steffek, Joana Ribeiro de Faria et al. The CLC-UKET Dataset: Benchmarking Case Outcome Prediction for the UK Employment Tribunal. *Proceedings of the Natural Legal Language Processing Workshop (ACL)*, 2024, pp. 81–96, 2024. <https://aclanthology.org/2024.nllp-1.7/>.
- [41] Niko Grupen, Gabe Pereyra, Julio Pereyra. Introducing Harvey’s Legal Agent Benchmark. *Primary online source*, 2026. <https://www.harvey.ai/blog/introducing-harveys-legal-agent-benchmark>.
- [42] Hunter Lightman, Vineet Kosaraju, Yura Burda et al. Let’s Verify Step by Step. *ICLR 2024; arXiv:2305.20050*, 2024. <https://arxiv.org/abs/2305.20050>.
- [43] Andy Zou, Tristan Xiao, Ryan Jia et al. Forecasting Future World Events with Neural Networks. *NeurIPS 2022, Datasets and Benchmarks*, 2022. <https://arxiv.org/abs/2206.15474>.
- [44] Danny Halawi, Fred Zhang, Chen Yueh-Han, Jacob Steinhardt. Approaching Human-Level Forecasting with Language Models. *ICML 2024, Proceedings of Machine Learning Research 235*, 2024. <https://arxiv.org/abs/2402.18563>.
- [45] Ezra Karger, Houtan Bastani, Chen Yueh-Han et al. ForecastBench: A Dynamic Benchmark of AI Forecasting Capabilities. *ICLR2025*, 2025. <https://arxiv.org/abs/2409.19839>.
- [46] Evan Miller. Adding Error Bars to Evals: A Statistical Approach to Language Model Evaluations. *arXiv preprint*, 2024. <https://arxiv.org/abs/2411.00640>.
- [47] Carlos E. Jimenez, John Yang, Alexander Wettig et al. SWE-bench: Can Language Models Resolve Real-World GitHub Issues?. *ICLR 2024*, 2024. <https://arxiv.org/abs/2310.06770>.
- [48] OpenAI. Why SWE-bench Verified no longer measures frontier coding capabilities. *Primary online source*, 2026. <https://openai.com/index/why-we-no-longer-evaluate-swe-bench-verified/>.
- [49] TypeSafe AI. Models: Jev and the System One Interface. *Product documentation, accessed September 29, 2026*, 2026. <https://docs.typesafe.ai/models>.
- [50] Supreme Court of the United States. Dobbs v. Jackson Women’s Health Organization, 597 U.S. 215 (2022). Decided June 24, 2022. [https://www.supremecourt.gov/opinions/21pdf/19-1392\\_6j37.pdf](https://www.supremecourt.gov/opinions/21pdf/19-1392_6j37.pdf).
- [51] Steven Shavell. Economic Analysis of Litigation and the Legal Process. *NBER Working Paper 9697*, 2003. Chapter 17 of the draft *Foundations of Economic Analysis of Law*. <https://www.nber.org/papers/w9697>.
- [52] George L. Priest, Benjamin Klein. The Selection of Disputes for Litigation. *Journal of Legal Studies* 13 (1984), 1–55, 1984. <https://www.journals.uchicago.edu/doi/10.1086/467732>.
- [53] Lucian A. Bebchuk. Litigation and Settlement Under Imperfect Information. *RAND Journal of Economics* 15(3) (1984), 404–415, 1984. <https://hls.harvard.edu/bibliography/litigation-and-settlement-under-imperfect-information/>.
- [54] Supreme Court of the United States. Jones v. Barnes, 463 U.S. 745, 751–753 (1983). <https://tile.loc.gov/storage-services/service/ll/usrep/usrep463/usrep463745/usrep463745.pdf>.

- [55] Jane Goodman-Delahunty, Pär Anders Granhag, Maria Hartwig, and Elizabeth F. Loftus. Insightful or Wishful: Lawyers’ Ability to Predict Case Outcomes. *Psychology, Public Policy, and Law*, 16(2):133–157, 2010. <https://doi.org/10.1037/a0019060>.
- [56] Harvey Team. How AI Agents are Changing the Way Lawyers do Legal Work. *Primary online source*, 2026. <https://www.harvey.ai/blog/ai-agents-for-legal-work>.
- [57] Niko Grupen, Gabe Pereyra, Julio Pereyra. Legal Agent Benchmark: Initial Results. *Primary online source*, 2026. <https://www.harvey.ai/blog/legal-agent-benchmark-initial-results>.
- [58] Artificial Analysis. Harvey LAB-AA: New Benchmark Evaluating AI on Legal Tasks. *July 7, 2026, benchmark announcement*, 2026. <https://artificialanalysis.ai/articles/harvey-lab-aa>.
- [59] Harvey Team. Opus 4.8, Now Live in Harvey. *Primary online source*, 2026. <https://www.harvey.ai/blog/opus-4-8-now-live-in-harvey>.
- [60] Harvey Team. Fable 5, Now Available in Harvey. *Primary online source*, 2026. <https://www.harvey.ai/blog/fable-5-now-available-in-harvey>.
- [61] Harvey Team. Claude Sonnet 5, Now Live in Harvey. *Primary online source*, 2026. <https://www.harvey.ai/blog/sonnet-5-in-harvey>.
- [62] United States Supreme Court. Atlantic Marine Construction Co. v. United States District Court. *571 U.S.49,59–66 (2013)*, 2013. <https://www.law.cornell.edu/supremecourt/text/12-929>.
- [63] Texas Legislature. Texas Civil Practice and Remedies Code, Chapter 16. *Texas Civil Practice and Remedies Code § 16.070, accessed September 29, 2026*, 2026. <https://statutes.capitol.texas.gov/Docs/CP/htm/CP.16.htm#16.070>.
- [64] Supreme Court of the United States. Jones v. Bock. *549 U.S.199,215*, 2007. <https://tile.loc.gov/storage-services/service/ll/usrep/usrep549/usrep549199/usrep549199.pdf>.
- [65] Supreme Court of Texas. Italian Cowboy Partners, Ltd. v. Prudential Insurance Co. of America. *341 S.W.3d 323, 334–37 (Tex. 2011)*, 2011. <https://www.txcourts.gov/media/819896/OpinionsFY2011.pdf>.
- [66] Clay Mathematics Institute. Navier–Stokes Announcement. September 11, 2026. <https://www.claymath.org/news/navier-stokes-announcement/>.
- [67] Matthew Stubenberg, Chloe Berridge, Thomas Smith, and Joshua Casey. How AI Stacks Up Against the Multistate Bar Exam. *University of Hawai’i Richardson School of Law Research Paper No. 5291811; SSRN working paper*, May 12, 2025. <https://doi.org/10.2139/ssrn.5291811>.
- [68] Matthew Stubenberg. AI model achieves 100% on MBE bar exam questions. *Researcher’s online benchmark update, LinkedIn*, 2026; accessed September 30, 2026. [https://www.linkedin.com/posts/matthew-stubenberg\\_well-it-finally-happened-an-ai-model-correctly-activity-7482410775534448641-urAX](https://www.linkedin.com/posts/matthew-stubenberg_well-it-finally-happened-an-ai-model-correctly-activity-7482410775534448641-urAX).
- [69] Vals AI. LegalBench Leaderboard and Methodology. *Benchmark-provider results and commentary*, updated September 27, 2026; accessed September 30, 2026. [https://www.vals.ai/benchmarks/legal\\_bench](https://www.vals.ai/benchmarks/legal_bench).

- [70] Tak K. Mak. Analysing Intraclass Correlation for Dichotomous Variables. *Journal of the Royal Statistical Society, Series C (Applied Statistics)*, 37(3):344–352, 1988. <https://doi.org/10.2307/2347309>.
- [71] A. Colin Cameron and Douglas L. Miller. A Practitioner’s Guide to Cluster-Robust Inference. *Journal of Human Resources*, 50(2):317–372, 2015. <https://doi.org/10.3368/jhr.50.2.317>.

## A Model configuration and temporal evidence

The table records the maintained local evidence, rather than certifying provider training histories. “K” denotes a reported knowledge cutoff, “T” a reported training cutoff, and “?” no established cutoff. Month-only dates are compared as the first day of the reported month. The reported cutoffs for fifteen configurations precede June 30, 2026. Kimi and Muse remain qualified. A release date is not used to fill an unknown cutoff.

| Configuration          | Reasoning        | Reported cutoff | Relation to window |
|------------------------|------------------|-----------------|--------------------|
| GPT-6 Sol              | High             | 2026-04-20 (K)  | Precedes window    |
| Claude Fable 5.1       | Adaptive/default | 2026-06 (T)     | Precedes window    |
| Grok 4.6               | High             | 2026-02-01 (K)  | Precedes window    |
| Claude Opus 5          | Adaptive/default | 2026-05 (T)     | Precedes window    |
| GPT-5.6 Sol            | High             | 2026-02-16 (K)  | Precedes window    |
| Kimi K3                | High             | Unknown (?)     | Unknown            |
| GPT-5.6 Luna           | High             | 2026-02-16 (K)  | Precedes window    |
| Claude Opus 5.5        | High             | 2026-06 (T)     | Precedes window    |
| GPT-6.1 Sol            | High             | 2026-04-30 (K)  | Precedes window    |
| GPT-6 Astra            | High             | 2026-04-30 (K)  | Precedes window    |
| Gemini 3.8 Flash       | High thinking    | 2026-03 (K)     | Precedes window    |
| Grok 4.7               | High             | 2026-05 (K)     | Precedes window    |
| Claude Sonnet 5        | Adaptive/default | 2026-01 (T)     | Precedes window    |
| GPT-6 Luna             | High             | 2026-05-18 (K)  | Precedes window    |
| Claude Sonnet 5.5      | High             | 2026-06 (T)     | Precedes window    |
| Muse Spark 1.3         | High             | Unknown (?)     | Unknown            |
| Gemini 3.1 Pro Preview | High thinking    | 2025-01 (K)     | Precedes window    |

**Table 4:** Configuration and cutoff evidence retained by the benchmark. “Precedes window” is a date comparison under the maintained rule, not an independently verified training-history claim.

The common document interface standardizes access, not every inference setting. Claude Fable 5.1, Claude Opus 5, and Claude Sonnet 5 use adaptive thinking at provider defaults. Claude Opus 5.5 and Claude Sonnet 5.5 pin high effort. The other configurations use the high reasoning or thinking setting recorded in their registries. Provider-default sampling and serving updates may still introduce variation. Full-document predictions and the one-request shared-summary conditions are kept separate throughout the analysis.

## B Result provenance

The primary cohort identifier is `cycle-1-91-2026-09-08-luna-r5`. Costs in the main table are the totals recorded on the run receipts from provider-reported token usage and the frozen rates.

The checked source revisions are [LegalForecastBench 58371a0e](#), [LegalForecastCorpus b392ddaf](#), and [LegalForecastGym 4cc0f528](#). The LAB source snapshot is [1dd81403](#). These identify inspected versions; they do not certify every execution against the current source tree.

## C Human review of AI-flagged LAB criteria

After noticing instances in which LAB tasks appeared to contain errors, I asked two frontier models, GPT-6 Sol and Claude Opus 5.5, to audit the 52 tasks in LAB’s litigation-and-dispute-resolution practice area at Harvey commit 1dd81403. The goal was to identify rubric criteria that misstate the law, and environments or task specifications containing errors or deficiencies that would make the tasks unreliable for benchmarking or for training models. Harvey has used LAB for both [9, 10]; [confirm the source for public tasks used in training and a held-out set used for benchmarking](#).

The models flagged 835 of the 2,858 criteria (29.2%) as problematic or arguable. GPT-6 Sol flagged 616. Claude Opus 5.5, working without access to Sol’s results, flagged 650; it then reviewed every flag from both passes against the record and the governing law, ending with 564. The 835 figure is the union of Sol’s flags and Opus’s final flags [27].

To establish how many of these flags reflect genuine defects, I reviewed a simple random sample of 25 of the 835 flagged criteria, drawn with a fixed seed under a protocol committed before the review began [28]. The sample spans 19 tasks. For each criterion I examined the task instructions, the supplied documents, the governing law, both auditors’ findings, and how LAB’s own judges graded two model runs on it. A criterion is defective when, applied as written to the task as given, it would grade a competent response wrong: failing a response a careful litigator would consider correct, or passing one that is wrong. Two applications of that standard were made explicit during the review. A criterion that requires one of several permissible approaches is defective when it rejects an approach a careful litigator would adopt on the record. And an unplanted error in the supplied record, such as figures that do not add up or misdescribed case law, makes a criterion defective when it changes what a competent response can get right. Separately, I recorded whether each task environment itself contains a defect, whether or not that defect affects the sampled criterion.

Of the 25 criteria, I found 16 defective in ways that affect grading, 9 of them unambiguous objective errors in which the record, the law, the cited authority, or the arithmetic contradicts the rubric. 3 were arguable and 6 were not defective. 14 of the 25 come from task environments that contain defects (10 of the 19 tasks). Only 4 involved neither a defective criterion nor a defective environment (Table 5).

Treating the sample as a draw without replacement from the 835 flagged criteria, the exact one-sided 95% lower confidence bounds are: at least 383 flagged criteria are defective in ways that affect grading (point estimate 534; the bound is 13.4% of all criteria in the practice area); at least 171 are unambiguous objective errors (point estimate 301); and at least 319 come from task environments that contain defects (point estimate 468). At least 524 involve a defective criterion, a defective environment, or both (point estimate 668). In the sample, 9 of the 10 criteria flagged by both auditors were defective, against 7 of the 15 flagged by only one.

These estimates cover only criteria the auditors flagged; defects among unflagged criteria are not estimated. They count criteria, not independent mistakes: one error in a task environment can make several criteria defective. Each verdict is my professional judgment as a litigator, made after an AI stress test that argued both for and against each call, and no second human reviewer has re-decided the sample. A separate AI-assisted check of every case citation in the 19 sampled task environments against CourtListener, with spot checks, found problems in 45 of 102 citations to real authority, including 7 cases that could not be found and 32 real cases cited for propositions they do not contain [28].

| Criterion verdict | Environment defective | Environment correct | Total |
|-------------------|-----------------------|---------------------|-------|
| Defective         | 10                    | 6                   | 16    |
| Arguable          | 2                     | 1                   | 3     |
| Not defective     | 2                     | 4                   | 6     |
| Total             | 14                    | 11                  | 25    |

**Table 5:** Reviewer verdicts on the 25 sampled criteria, by whether the task environment itself contains a defect.

## The sampled criteria

Each entry gives the criterion as written, the reviewer’s verdicts, how LAB’s two native judges (Claude Sonnet 4.6 and GPT-5.5) graded the GPT-6 Luna and Claude Opus 5.5 runs on that criterion, and the reviewer’s analysis. Harvey task files are cited at commit 1dd81403.

### 1. Assess Reasonableness of Staffing Levels on Litigation Invoice (C-049)

*Criterion.* **Calculates budget overage percentage as approximately 17%.** PASS if the memo calculates the percentage by which the attorney fee subtotal (\$193,372.50) exceeds the approved budget ceiling (\$165,000) as approximately 17% (exact figure is 17.2%). FAIL if no percentage overage calculation is provided or if the percentage stated is materially incorrect (off by more than 2 percentage points).

*Verdict.* Defective (unambiguous objective error, deficiency in task environment). *Environment.* Defective: the invoice’s 117 time entries add up to \$215,680.00, while its summary and subtotal show \$193,372.50. *AI auditors’ reasoning (reviewer’s assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: F/F; Claude Opus 5.5: P/P.

*Analysis.* The rubric is testing the AI on a specific percentage calculation from an invoice, but the document in the environment is internally inconsistent. The invoice’s summary page and its subtotal row show fees of \$193,372.50 (17.2% over the \$165,000 budget), but the 117 detailed time entries add up to \$215,680.00, which puts the overage at \$50,680, or 30.7% (13.5 points from the rubric’s required 17.2% answer and far outside the 2-point tolerance).

### 2. Categorize Document Production Set by Relevance and Privilege — Privilege Log and Relevance Classification Report (C-012)

*Criterion.* **ISSUE\_003: DOC\_012 identified as Partially Privileged.** PASS if DOC\_012 (chandrasekaran-meeting-notes-aug28.docx, Chandrasekaran’s handwritten meeting notes from August 28, 2023) is classified as Partially Privileged. FAIL if classified as fully Privileged or Not Privileged.

*Verdict.* Defective (low-quality professional judgment, rubric is incorrect). *Environment.* Correct. *AI auditors’ reasoning (reviewer’s assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/P; Claude Opus 5.5: P/P.

*Analysis.* The rubric is wrong because it fails fully Privileged, which is the better answer here (even if Partially Privileged is defensible).

- The documents are handwritten lawyer notes from a fictitious meeting discussing problems with the R&D Department. The headings on individual topics seem intended to create the

impression that there is a discussion of business topics (which would not be privileged) followed by a separate discussion of legal topics (privileged), so from a superficial glance, the rubric does seem to match the topic headings. However, the bullets under each heading make clear that the entire meeting is just dedicated to the topic of the company’s desire to find a way to fire the head of R&D without getting sued. The business topics (which look to be nonprivileged) are discussions of how the R&D department is a bit of a disaster, such that management wants to fire the head. The legal sections then involve discussion of the fact that the R&D head has some sort of environmental whistleblower status, so termination is likely to trigger a lawsuit/whistleblower retaliation claim.

- The meeting was specifically convened because outside counsel was concerned about the retaliation claim, and even some of the ‘ops’ sections that are supposed to be nonprivileged reflect counsel’s advice, so they cannot be cleanly treated as nonlegal business discussions. For example, the ‘Ops’ sections record HR “compiling the file” of performance documentation and lay out a “Performance improvement plan first” / “Documentation period” sequence, which carries out outside counsel’s recommendation, recorded in the legal section, to build a “contemporaneous performance record to insulate against retaliation” claims.
- While partially privileged might be defensible, my preferred answer would be fully privileged, which the rubric incorrectly fails.

### 3. Compare Document Production Against Discovery Requests — Discovery Gap Analysis Memorandum (C-029)

*Criterion.* **ISSUE\_011: Notes only 5 documents actually relate to counterclaim damages.** PASS if the memorandum notes that only 5 of the 14 documents coded to RFP 27 (MLC-004056 through MLC-004102) actually relate to Pacific Corridor’s volume shortfall and the counterclaim damages calculation, meaning the counterclaim damages support is thin. FAIL if this is not noted.

*Verdict.* Defective (unambiguous objective error, rubric is incorrect). *Environment.* Defective: the production index codes 15 rows to RFP 27, but its summary tab gives a count of 14 for a range that contains 13. *AI auditors’ reasoning (reviewer’s assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: F/P; Claude Opus 5.5: F/F.

*Analysis.* This task requires models to make the point that “only” five documents in the bundle “actually relate to Pacific Corridor’s volume shortfall and the counterclaim damages calculation, meaning the counterclaim damages support is thin.” This rubric is flawed on multiple levels.

- First, it overlooks MLC-003781, a sixth document that plainly relates to the volume shortfall, so the count is objectively wrong (and the agents pointed to several other examples, at least some of which also seem to fit within the stated criteria). There is a document in the bundle that claims that only five documents concern the volume shortfall, but that assertion is undercut by other evidence supplied in the environment.
- Second, this rubric is someone’s argumentative characterization of the documents and probably would not hold up under adversarial testing. For example, there is case law that the term “relate to” is fairly broad and flexible, so quite plausibly a number of other documents would apply.
- Third, the premise underlying this rubric—that “only” five documents is such “thin” evidence that it must be called out—is wrong. Even one document could constitute conclusive proof of damages (the general rule is that damages only need to be reasonably estimated; courts do

not require great precision); to my mind, this is actually quite a lot of evidence at the fact stage compared to what one would normally see.

#### 4. Draft Case Assessment Memorandum — Litigation Risk Analysis for Distribution Agreement Dispute (C-025)

*Criterion.* **ISSUE\_007: Recommends motion to dismiss on unjust enrichment.** PASS if the memo recommends filing a motion to dismiss (or motion for judgment on the pleadings, or similar dispositive motion) targeting the unjust enrichment claim (Count IV) on the ground that an express contract governs the parties' relationship. FAIL if no such motion recommendation is made regarding the unjust enrichment claim.

*Verdict.* Defective (low-quality professional judgment, rubric is incorrect). *Environment.* Correct. *AI auditors' reasoning (reviewer's assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: F/F; Claude Opus 5.5: P/P.

*Analysis.* A different rubric on this task concludes that plaintiff has a strong breach of contract claim. (See id: C-006.) The prompt instructs models to review case documents and “prepare a comprehensive case assessment memo,” and models are dinged on this rubric if they do not “recommend[] filing a motion to dismiss” on the unjust enrichment claim, which is expressly pled in the alternative to the (strong) breach of contract claim.

- The rubric is wrong because a careful litigator could reasonably conclude that a motion to dismiss that claim is not a good idea: Count IV is pleaded in the alternative (Complaint ¶127) and seeks restitution of “no less than \$22,521,312.50” (¶135), which is exactly the Count I contract damages (\$17,721,312.50, ¶94) plus the Count III trade-secret damages (\$4,800,000, ¶120). Counts I and III would survive the motion and double recovery is barred, so dismissing Count IV would not reduce the defendant's net exposure; and the factual underpinnings of both claims are the same so a successful motion would not reduce the scope of discovery. Filing a dispositive motion is costly and, on these facts, there seems to be no benefit to doing so. So the required recommendation would fail many thoughtful analyses that could discuss the pros and cons and recommend against filing such a motion.

#### 5. Draft Conflict Check Memorandum for Litigation Engagement Clearance (C-030)

*Criterion.* **ISSUE\_005: Recommends screening or removal of Caleb Strand.** PASS if the memo recommends that Caleb Strand be screened from the engagement or removed from the proposed staffing. FAIL if no such recommendation is made.

*Verdict.* Defective (low-quality professional judgment, rubric is incorrect). *Environment.* Correct. *AI auditors' reasoning (reviewer's assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/P; Claude Opus 5.5: P/P.

*Analysis.* The rubric requires the model to recommend that an associate who shares an apartment with his sister (an in-house IP paralegal at TriPoint, the adverse party) be screened from the matter. Screening is generally used in certain specific situations where required by the professional conduct rules, but no rule requires screening in this scenario.

- I agree that a prudent memo would flag the fact that a junior team member is sharing an apartment with an employee of the adverse party, but that is covered by other rubrics (C-026 through C-029). This rubric specifically fails the work product (and results in a 0 for all work

on this task) even if the model recommended, say, obtaining the client’s informed consent, which would be perfectly permissible and appropriate here.

- Other rubrics in this task resolve similar kinds of conflicts with informed consent (C-024 accepts “clearance with conditions (such as obtaining informed consent)” for Lisa Chow’s spousal conflict, and C-016 recommends informed consent for the Ridgeline conflict). Accordingly, informed consent plus a home confidentiality protocol could suffice as an alternative to screening.

## 6. Draft Conflict Check Memorandum for Litigation Engagement Clearance (C-049)

*Criterion.* **ISSUE\_011: Recommends integrating annual disclosure questionnaires with ConflictTracker.** PASS if the memo recommends that annual conflict-disclosure questionnaires be cross-referenced with or integrated into the ConflictTracker system. FAIL if no recommendation about integrating disclosure questionnaires with the conflicts database is made.

*Verdict.* Defective (requires specific outputs not expressly or impliedly requested in the prompt).

*Environment.* Correct. *AI auditors’ reasoning (reviewer’s assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: F/P; Claude Opus 5.5: P/P.

*Analysis.* The assignment in the prompt is to draft a “comprehensive conflict check memorandum for the proposed engagement.” The rubric requires that the memo contain suggested firm-wide process improvements (memo must “recommend[] that annual conflict-disclosure questionnaires be cross-referenced with or integrated into the ConflictTracker system”). That’s simply not the assignment, and there’s many reasons why it’s undesirable to jam broad process suggestions into a conflicts memo for one specific case. This rubric actually failed work product from Luna that did everything a clearance memo needs on this point. The two judges split on that same Luna memo (Sonnet 4.6 failed it; GPT-5.5 passed it), which shows the criterion does not grade reliably.

## 7. Draft Counterclaim Against Plaintiff for Breach of Joint Development Agreement (C-045)

*Criterion.* **Count: Patent Infringement (ISSUE\_005).** PASS if the counterclaim includes a separate count for patent infringement under 35 U.S.C. § 271, alleging that the LumiSense 400 infringes U.S. Patent No. 11,234,567 covering the SensorCore™ architecture. FAIL if no patent infringement counterclaim is asserted.

*Verdict.* Arguable. *Environment.* Defective: the record lacks the patent the counterclaim must plead, and the two expert reports give different titles for it. *AI auditors’ reasoning (reviewer’s assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/P; Claude Opus 5.5: P/P.

*Analysis.* The task is to draft a counterclaim based on a handful of documents (including an expert report, some emails, and a counterclaim strategy memo that another lawyer prepared). The rubric dings the AI if the courtclaims do not include “a separate count for patent infringement under 35 U.S.C. § 271, alleging that the LumiSense 400 infringes U.S. Patent No. 11,234,567.”

- The environment does not contain the allegedly infringed patent, which one would need to review in order to have a sufficient basis for drafting that counterclaim, so this environment appears to be incomplete. That said, a defensible approach might be for the agent to draft the counterclaim but flag in brackets that it requires verification.

## 8. Draft Counterclaim Against Plaintiff for Breach of Joint Development Agreement (C-046)

*Criterion.* **Patent Infringement Count: Patent Number Cited.** PASS if the patent infringement count specifically identifies U.S. Patent No. 11,234,567. FAIL if the patent number is omitted or incorrect in the patent infringement count.

*Verdict.* Arguable. *Environment.* Defective: same task environment as item 7. *AI auditors' reasoning (reviewer's assessment).* Partly correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/P; Claude Opus 5.5: P/P.

*Analysis.* Same issue as prior rubric. This one is essentially derivative because it depends on pleading the patent count that was required in the prior rubric.

## 9. Draft Case Assessment Memorandum for Defective Industrial Equipment Product Liability Claim (C-003)

*Criterion.* **Executive Summary mentions death of Marco Reyes.** PASS if the Executive Summary mentions the death of Marco Reyes. FAIL if the fatality is not mentioned in the Executive Summary.

*Verdict.* Defective (requires specific outputs not expressly or impliedly requested in the prompt). *Environment.* Defective: same task environment as item 10 (Ohio state agency issuing OSHA citations). *AI auditors' reasoning (reviewer's assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/F; Claude Opus 5.5: F/F.

*Analysis.* The case involved a (fictitious) safety incident at a manufacturing plant that killed one employee and seriously injured two others. The rubric requires the fact that one employee died be discussed specifically in the “executive summary.” When we tested Opus 5.5 and GPT-6-Luna on this benchmark, both discussed the death in several places in their memo, especially the factual background section, but the executive summary consisted of an overview of the legal analysis that did not restate the fact of the death. They were failed by one or more judges.

- This fails perfectly competent and appropriate work product. The company CEO who retained counsel to advise on this incident is the same person who delivered news of the (fictitious) tragic death to the decedent’s widow and surviving family members, so the most senior management of the company is obviously very familiar with the incident and the agents’ approach of focusing the executive summary on the legal exposure is appropriate.
- The two judges reached opposite results on Luna’s memo: GPT-5.5 failed it, while Sonnet 4.6 passed it only by quoting a sentence from the memo’s facts section as if it appeared in the executive summary.

## 10. Draft Case Assessment Memorandum for Defective Industrial Equipment Product Liability Claim (C-026)

*Criterion.* **OSHA LOTO citation identified with specific violation and penalty.** PASS if the memo identifies the OSHA serious citation for LOTO deficiency under 29 CFR 1910.147(c)(4)(i) with a penalty of \$18,900. FAIL if this specific citation, regulation, or penalty amount is not referenced.

*Verdict.* Not defective (defective environment, but rubric does not fail correct work product).

*Environment.* Defective: Ohio’s Division of Safety & Hygiene cannot issue OSHA citations to this private employer. The criterion itself does not misgrade work (environment-only defect). *AI auditors’ reasoning (reviewer’s assessment).* Wrong. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: F/F; Claude Opus 5.5: P/P.

*Analysis.* Although the specific rubric being tested does not fail work product incorrectly, the environment is defective in an unintended way. The environment contains an OSHA citation from Ohio’s Division of Safety & Hygiene. But Ohio does not have an OSHA-approved State Plan, so only federal regulators have jurisdiction to issue OSHA citations against this company. The environment does not test that so it is not being presented as a “mistake” by one’s adversary that the agents should pick up on.

- The rubric itself is not wrong because it simply wants a memo to address the purported citation, so an agent could write the memo while noticing that the task is obviously a fictitious evaluation with an incoherent environment.

## **11. Draft Discovery Plan Memorandum for Breach of Contract and Fraud Defense (C-036)**

*Criterion.* **Provides custodian priority rankings.** PASS if the memo provides priority rankings or tiers for custodians (e.g., high/medium/low priority, or Tier 1/Tier 2, or similar prioritization). FAIL if custodians are listed without any indication of relative priority.

*Verdict.* Not defective. *Environment.* Correct. *AI auditors’ reasoning (reviewer’s assessment).* Wrong. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/P; Claude Opus 5.5: P/P.

*Analysis.* The AI agent is right that the rubric is really enforcing a stylistic choice, but the choice here is well-grounded in convention and an appropriate expectation for quality work product. The case file contains a court order where the judge encouraged the parties to discuss whether phased discovery would be appropriate at the Rule 26(f) conference, and there is some existing prioritization work in the environment. The rubric requires that the agent’s discovery plan include some kind of prioritization or phasing of custodians, which is a common expectation in these cases. I think it’s fair to judge anything that lacks any form of prioritization as deficient given the context.

## **12. Draft Federal Complaint for Breach of Contract and Fiduciary Duty — Placement Agent Dispute (C-025)**

*Criterion.* **ISSUE\_012: Complaint requests injunctive relief.** PASS if the prayer for relief includes a request for preliminary and/or permanent injunctive relief — specifically to enforce the non-solicitation provision and prevent Graydon from continuing to divert Fund III investor prospects to competing funds. FAIL if no injunctive relief is requested.

*Verdict.* Not defective. *Environment.* Correct. *AI auditors’ reasoning (reviewer’s assessment).* Wrong. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/P; Claude Opus 5.5: P/P.

*Analysis.* The client intake memo specifically describes an interest in seeking injunctive relief (“We may also wish to seek preliminary and permanent injunctive relief to enforce the twelve-month non-solicitation tail”), and the client’s demand letter threatens it. While it’s true that the intake memo flags the topic for discussion, the correct approach is to include the proposed prayer for relief

in the draft, so the case team can see how it reads and discuss any downsides. As a practical matter, there's unlikely to be much downside to including this prayer for relief in the complaint. (Whether it's worth the cost of actually pursuing a preliminary injunction on this is debatable, so the rubric's approach makes sense.)

### **13. Draft First Set of Interrogatories to Defendant Veridian Health Systems in Trade Secret Misappropriation Case (C-008)**

*Criterion.* **Definitions section defines 'Document' or 'Documents'.** PASS if the Definitions section includes a definition of 'Document' or 'Documents' that is broad enough to encompass electronically stored information (ESI), emails, and physical documents. FAIL if there is no definition of 'Document' or it is so narrow as to exclude electronic records.

*Verdict.* Defective (unambiguous objective error, rubric is incorrect). *Environment.* Defective: the supplied documents mischaracterize or miscite case law (citation sweep; peripheral to this criterion). *AI auditors' reasoning (reviewer's assessment).* Partly correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: F/F; Claude Opus 5.5: P/P.

*Analysis.* The rubric fails any draft that does not supply a definition of "document" encompassing electronically stored information (ESI). There are two problems:

- First, these are interrogatories, not document requests. Well-drafted interrogatories do not necessarily need to use the word "document." The rubric does not really make sense as a requirement for this kind of work product.
- Second, even if we were dealing with document requests here (which we are not), the local rules of many courts, including the Western District of Texas (where this fictitious case is being litigated), define the term "document" for all discovery requests served by litigants in those cases. There is thus no requirement to define the term "documents," and in fact, supplying a definition could just trigger objections from the other side that the local rule controls. See Local Rule CV-26, available at <https://www.txwd.uscourts.gov/court-information/lcr-civil-rules/> ("The term 'document' means any document or electronically stored information as described in Federal Rule of Civil Procedure 34(a). A draft of a nonidentical copy is a separate document within the meaning of this term.").
- The rubric does demonstrably fail rules-compliant work product. For example, Luna was failed for not including a definition of the term "Document" but that's actually what the local rules encourage in this Court.

### **14. Draft First Set of Interrogatories to Defendant Veridian Health Systems in Trade Secret Misappropriation Case (C-026)**

*Criterion.* **Interrogatory re: USB drive or CMT files on Veridian systems.** PASS if at least one interrogatory asks whether Oshiro's USB drive (referencing serial number CX-8827491 or otherwise describing the external storage device) was ever connected to, accessed from, or present on any Veridian system, device, or network, or asks whether any CMT files or materials were ever possessed, accessed, viewed, or copied by Veridian or its employees. FAIL if no interrogatory addresses the USB drive or CMT files on Veridian systems.

*Verdict.* Not defective. *Environment.* Defective: same task environment as item 13 (environment-only defect). *AI auditors' reasoning (reviewer's assessment).* Wrong. *Native grades* (Sonnet

4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/P; Claude Opus 5.5: P/P.

*Analysis.* This one does not really belong on the list, as GPT-6-Sol was complaining about a different rubric (C-048) which it argues is inconsistent with this one. It does not identify any issue in this rubric, standing alone, so this counts as an instance where Sol added something to the “flagged” list that does not belong.

## 15. Draft Discovery Responses and Objections to RFAs and RFPs in Breach of Supply Agreement Litigation (C-028)

*Criterion. ISSUE\_011: Identifies RFA No. 20 as calling for a legal conclusion.* PASS if the plan identifies that RFA No. 20 (asking Prismavale to admit the products were ‘not merchantable as defined by the Uniform Commercial Code’) calls for a legal conclusion regarding UCC § 2-314 merchantability and recommends objecting on that basis. FAIL if the plan does not identify the legal conclusion problem with RFA No. 20.

*Verdict.* Defective (low-quality professional judgment, rubric is incorrect). *Environment.* Defective: the supplied Answer cites complaint paragraph ranges that do not match the complaint (e.g. Count III as ¶¶ 88–94; the complaint’s Count III is ¶¶ 86–96). *AI auditors’ reasoning (reviewer’s assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: F/F; Claude Opus 5.5: P/P.

*Analysis.* Plaintiff served a request for admission on the defendant that says, “Admit that the SB-102 and PS-302 products shipped by Prismavale to Greenleaf were not merchantable as defined by the Uniform Commercial Code.” The prompt asks the agent to “prepare a discovery strategy memo and formal RFA/RFP responses.” The rubric requires “the plan” (i.e., the strategy memo) to recommend objecting to the request on the grounds that it calls for a legal conclusion. The rubric is wrong for two main reasons:

- First, the modern view is requests for admissions can call for application of law to facts. While it is within acceptable professional norms for a lawyer to object on the ground that such a request calls for a legal conclusion, and then proceed to answer, the better reading of the law is that this request is not objectionable such that the party should simply answer. On this legal conclusion in particular, there’s really no reason why one would prefer to object rather than simply deny the requested admission. So this rubric fails competent work product (and probably the better answer than the one the rubric calls for).
- Second, even assuming one should include the objection at all, the task fails the model if the issue is not expressly addressed in the strategy memo. Luna made the “legal conclusion” objection in the draft responses and then proceeded a substantive denial (an acceptable approach), but did not discuss the “legal conclusion” objection in the memo. That was judged a failure. That’s really requiring duplicative responses on a very minor point that probably shouldn’t be in the drafts at all.

## 16. Draft Motion in Limine to Exclude Expert Testimony and Prejudicial Evidence in Commercial Breach of Contract Case (C-009)

*Criterion. ISSUE\_004: Identifies the \$11,000 arithmetic discrepancy.* PASS if the motion identifies that the sum of Marchetti’s five itemized damages categories totals \$14,211,000, but her report states total damages as \$14,200,000, creating an \$11,000 discrepancy. FAIL if this arithmetic error is not identified.

*Verdict.* Defective (low-quality professional judgment, rubric is incorrect (a judgment call, not an objective error); defective environment (directs advocacy outside professional norms); companion rubric C-010 also incorrect (not sampled)). *Environment.* Defective: the task directs advocacy outside professional norms (an \$11,000 rounding point offered as a *Daubert* ground) and labels a rounding difference an “arithmetic error.” *AI auditors’ reasoning (reviewer’s assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/P; Claude Opus 5.5: P/P.

*Analysis.* The damages expert includes a summary table in her report listing various items of damages, which collectively add up to \$14,211,000; the Total Damages line lists the total as \$14,200,000. Opposing counsel questioned the expert about it at her deposition, and the expert calls it “a rounding discrepancy.” An internal strategy memo highlights the issue as a ground for seeking to exclude the expert’s testimony under *Daubert*, arguing “the arithmetic error supports a cumulative argument that the report was prepared without the rigor required of expert testimony in federal court.”

- This task is poorly specified, and the recommended strategy of raising this with the judge is a poor one:
- First, the task environment repeatedly states that this is an “arithmetic error.” That is false. The expert has rounded the total in the table to the nearest \$100,000. While that is somewhat unorthodox in an expert report (particularly since it is written out as \$14,200,000 not \$14.2 million), numbers often are rounded in court filings so the judge is unlikely to agree that this is an error.
- Second, even if there were an error (which there is not), that is generally not a basis for excluding the expert’s testimony under *Daubert*. At most one could move to exclude the erroneous version and ask for more accurate testimony, but that does nothing to help the hypothetical client here; it just raises their damages exposure by the \$11,000 that the expert was willing to round off.
- Third, the proposed framing of the argument comes very close to a breach of decorum and carries a real risk of drawing a rebuke from the judge. The strategy memo wants the motion to argue the expert’s 30 prior engagements were “in state courts and arbitrations that do not impose the same methodological gatekeeping requirements” as federal courts. While some state courts use the *Frye* standard instead of *Daubert*, many federal judges would take offense at the suggestion that state courts do not observe adequate methodological rigor. Many busy judges could be annoyed at spending time on an argument that literally just amounts to a rounding error.
- This rubric is scored as Defective, as a professional-judgment call rather than an objective error. Our view is that the better approach is to drop this issue entirely, and the sampled rubric (C-009) fails a motion that does so, even though a motion that merely mentions the discrepancy in a footnote would pass. The companion rubric, C-010, is in my view objectively wrong because it requires that this issue be presented as a basis for excluding the expert’s testimony under *Daubert*. Luna was failed incorrectly on C-010 (by the Sonnet 4.6 judge) for downplaying this issue; in fact, Luna was right to do so, as it should not be raised at all. But since that’s not the sampled criterion, we do not mark it here.

## 17. Draft Rule 12(b)(6) Motion to Dismiss Brief — Commercial Software Licensing Dispute (C-037)

*Criterion.* **ISSUE\_008 — Cites Delaware parol evidence authority.** PASS if the motion brief cites relevant Delaware authority on the parol evidence rule and integration clauses, such as *SIGA Technologies v. PharmAthene* or *Eagle Industries v. DeVilbiss Health Care*. FAIL if no Delaware parol evidence authority is cited.

*Verdict.* Defective (unambiguous objective error (including hallucinations), rubric is incorrect).

*Environment.* Defective: the supplied research memo cites fabricated or mischaracterized authority (citation sweep: 28 of 53 citations in this task have problems). *AI auditors' reasoning (reviewer's assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: F/F; Claude Opus 5.5: F/F.

*Analysis.* The rubric here involves a legal error, and the task environment contains hallucinations about case law. This one is extremely defective for multiple reasons:

- First, the rubric is demanding case law on the wrong legal doctrine. Under Delaware law, the parol evidence rule and an integration clause do not preclude fraud claims. The issue here is whether the plaintiff can reasonably allege that they relied on their counterparty's pre-signing representations despite their agreement in the contract that they were not relying on any representations other than those written in the contract itself.
- Second, the cases in the environment appear to contain multiple hallucinations. The rubric refers to *SIGA v. PharmAthene*, but Opus noted that case does not address integration or parol evidence. The task environment contains a legal research memo that cites *SIGA Technologies, Inc. v. PharmAthene, Inc.*, 67 A.3d 330 (Del. 2013), and claims that it states (at page 344), "an integration clause is the clearest declaration that the parties intend the writing to be the complete and final expression of their agreement." The quoted language does not appear in the cited case. Page 344 of that case addresses the duty to negotiation in good faith, not integration clauses. Similarly, the memo cites *Eagle Industries, Inc. v. DeVilbiss Health Care, Inc.*, 702 A.2d 1228 (Del. 1997) and says: "The Delaware Supreme Court held that 'a party to a contract cannot promise, in a clear integration clause of a negotiated agreement, that it is not relying on promises or representations made by the other party outside of the agreement, and then assert a claim for fraud based on those very representations.' *Id.* at 1232." The quoted language does not appear in the cited case.
- Third, the rubric failed responses that stated the correct law on this point. In our tests, Opus reviewed *SIGA* and *Eagle*, correctly realized they were inapposite and declined to cite them, cited more on-point and accurate authority, and was then failed for doing a better job than the rubric's author did at this task.

## 18. Draft Rule 12(b)(6) Motion to Dismiss Brief — Commercial Software Licensing Dispute (C-045)

*Criterion.* **ISSUE\_010 — References Arcadia's own delays contributing to Go-Live delay.** PASS if the motion brief identifies Arcadia's own contributing failures, including the 45-day Project Manager gap (Kevin Liu's departure) and/or the 75-day late API specifications, as causes of the Go-Live delay. FAIL if Arcadia's own contributory delays are not mentioned.

*Verdict.* Defective (unambiguous objective error, rubric is incorrect). *Environment.* Defective: same task environment as item 17. *AI auditors' reasoning (reviewer's assessment).* Correct. *Native*

*grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: F/F; Claude Opus 5.5: F/F.

*Analysis.* The rubric makes a rookie mistake, demanding that the agent argue a factual dispute on a motion to dismiss, where the court is required to assume the plaintiff's allegations are true and cannot consider factual dispute. The 45-day and 75-day delays complained of come from the defendant's internal report, which the plaintiff's complaint disputes. Moreover, courts generally cannot resolve messy fact disputes over "causation" on a motion to dismiss.

- The criterion failed agents that correctly recognized that this argument was outside the scope of what can be argued on a motion to dismiss and thus decided to excluded it.

## **19. Draft Verified Responses and Objections to Plaintiff's First Set of Interrogatories in Commercial Breach of Contract and Fraud Action (C-028)**

*Criterion. Correct financial figures for Tri-Basin purchase history.* PASS if any purchase history figures cited in the responses are consistent with the canonical data: 2015: \$9.2M; 2016: \$10.1M; 2017: \$10.8M; 2018: \$12.4M; 2019: \$13.6M; 2020: \$9.7M; 2021: \$15.3M; 2022: \$16.8M; 2023 (Jan–Aug): \$11.9M. Minor rounding differences are acceptable. Also PASS if the response invokes Rule 33(d) rather than citing specific figures. FAIL if specific purchase figures are cited but are materially incorrect.

*Verdict.* Defective (unambiguous objective error, deficiency in task environment; rubric is incorrect).

*Environment.* Defective: the purchase-history spreadsheet is internally inconsistent (three 2020 series; line items that do not sum to subtotals), and the record has no contract-year data. *AI auditors' reasoning (reviewer's assessment).* Partly correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/P; Claude Opus 5.5: P/P.

*Analysis.* There are two issues:

- First, Interrogatory No. 4(b) asks for Tri-Basin's actual purchases for the 2020 contract year (the contract year runs from January 15 through January 14 of the following year; Interrogatory No. 14 separately asks for calendar-year totals), but the canonical data referred to in the rubric is calendar year data, so the rubric's numbers are not correct for that request. The correct value cannot be discerned from the environment: the monthly data cannot be split at January 15, and nothing in the record supports assuming purchases are spread evenly within January (seasonal variation would make that unlikely). The rubric therefore grades wrong in the opposite direction from most of the other items: rather than failing good work product, it passes responses that supply the calendar-year figure for a contract-year request without acknowledging that the record does not contain the requested data. That is approving fundamentally deficient work product.
- Second, the data that is provided has some internal inconsistencies. For example, an "annual summary" tab lists 2020 purchases at \$9.7 million, but the monthly detail lists 2020 three times, with totals of \$8,700,000, \$8,549,400 and \$9,700,000. The units sold do not match up between those two sheets for 2021, 2022, 2023, and series 3 of 2020. One of the monthly rows does not add up (gross \$1,210,000 + credits −\$17,400 = \$1,192,600, but the net shown is \$1,192,500). The annual summary minimums do not add up. In the Product Line Breakdown, line-item net revenue never sums to the subtotal. So a response that relied on the line items would get graded as a failure.
- The environment contains law firm discovery guidelines that recommend that interrogatory responses match the client's Answer, but the Answer here contains only approximate figures,

which is consistent with the Answer using calendar year data as a rough approximation instead of providing the contract year data requested.

## 20. Extract Key Obligations from Litigation Hold and Document Preservation Notice — Obligation Summary Memorandum (C-002)

*Criterion.* **ISSUE\_001: Flags financial records date discrepancy as requiring DOJ clarification.** PASS if the memo recommends seeking clarification from the DOJ regarding the conflicting date ranges (Jan 1, 2017 vs. Jan 1, 2019) for financial records. FAIL if no recommendation to seek clarification is made.

*Verdict.* Defective (unambiguous objective error, rubric is incorrect). *Environment.* Correct. *AI auditors' reasoning (reviewer's assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: F/F; Claude Opus 5.5: P/F.

*Analysis.* The task involves a document preservation notice from federal criminal authorities that defines the relevant period (for which the recipient must preserve documents) as January 1, 2019, through March 3, 2025 “unless otherwise specified herein.” For one specific category of documents, the notice states that, “Notwithstanding the Relevant Period defined in Paragraph 12 of this Notice, Ridgeline shall preserve [certain kinds of] records . . . from January 1, 2017 to the present date. ”

- The rubric suggests that the agent should seek clarification of the “conflicting” date ranges but there is no conflict. The notice quite clearly imposed a longer period for one specific category of documents, while all the other categories are subject to a shorter period. This is a fairly common request.
- All the models were failed on this criterion, but the criterion makes no sense.

## 21. Identify Excessive or Duplicative Research Charges in Litigation Invoice (C-013)

*Criterion.* **Issue 3: Identifies duplicative MTCA research across Takahashi and Wendt.** PASS if the memo identifies that multiple attorneys (Kenji Takahashi and Tyler Wendt) billed for overlapping MTCA liability/contractor liability/summary judgment research during July 2024, citing Section 6.4 (Duplicative Research Prohibition). FAIL if the MTCA duplication across timekeepers is not identified.

*Verdict.* Not defective. *Environment.* Correct. *AI auditors' reasoning (reviewer's assessment).* Wrong. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/P; Claude Opus 5.5: P/P.

*Analysis.* While Sol is right that the facts here are insufficient to establish that the client’s billing guidelines were violated, for purposes of the task (a client-side memo identifying changes to dispute). it seems reasonable to flag that multiple attorneys were doing research into closely related issues. The task requested a memo “a memo identifying all non-compliant or excessive charges with recommended adjustments,” so flagging potential violations (even if they were not clearly proven from the face of the time entries) is a standard approach a client would take here.

## 22. Identify Issues in Litigation Matter Budget Proposal (C-029)

*Criterion.* **Budget inconsistency rated as Critical or High severity.** PASS if the fee total reconciliation discrepancy (\$116,050 difference between phase-based and staffing-model totals) is

assigned a severity of Critical or High (or the highest or second-highest tier in whatever severity scale is used). FAIL if it is assigned a lower severity level or no severity level.

*Verdict.* Arguable (requires specific outputs not expressly or impliedly requested in the prompt). *Environment.* Correct. *AI auditors' reasoning (reviewer's assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/P; Claude Opus 5.5: P/P.

*Analysis.* The prompt requested “a categorized issue memorandum” but didn’t specify what the categories should be. The rubric fails the model if it does not break issues into severity levels or it does not assign a critical or high severity level to a specific issue. While it appears that models often do volunteer this kind of severity assessment, it’s not requested by the prompt and competent work product could decline to include it.

### 23. Review Counterparty’s Proposed Jury Instructions — Issue Memorandum for Trade Secrets Trial (C-020)

*Criterion.* **Explains distinction between tortious interference with contract vs. business relations.** PASS if the memo explains that under Georgia law, tortious interference with an existing contract does not require a separately/independently wrongful act — the knowing inducement of the breach is sufficient — and that the ‘independently wrongful act’ element applies to the distinct tort of tortious interference with business relations. FAIL if this legal distinction is not articulated.

*Verdict.* Defective (unambiguous objective error (including hallucinations), rubric is incorrect).

*Environment.* Defective: the fictional summary-judgment order cites a case for the opposite of its holding and a case that cannot be found (citation sweep: 10 of 21 citations in this task have problems).

*AI auditors' reasoning (reviewer's assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: F/F; Claude Opus 5.5: F/F.

*Analysis.* Multiple issues:

- First, this is a clear legal error. Georgia applies the same elements to interference with contractual and business relations. The rubric says the opposite of what Georgia law is. *Disaster Services, Inc. v. ERC P’ship*, 228 Ga. App. 739, 740–41 (1997) (“Tortious interference claims, whether asserting interference with contractual relations, business relations, or potential business relations, share certain common essential elements: (1) improper action or wrongful conduct by the defendant without privilege; (2) the defendant acted purposely and with malice with the intent to injure; (3) the defendant induced a breach of contractual obligations or caused a party or third parties to discontinue or fail to enter into an anticipated business relationship with the plaintiff; and (4) the defendant’s tortious conduct proximately caused damage to the plaintiff.”).
- Second, the fictional summary judgment order (from which the rubric’s legal position seems to be taken) appears to contain several hallucinations. For example, the order paraphrases the third and fourth element described in *Disaster Services* approximately correctly, but the first two elements are erroneously listed as “(1) the existence of a valid contract between the plaintiff and a third party; (2) the defendant’s knowledge of the contract.” The fictional order also cites *Valdosta Livestock, Inc. v. Furst*, 342 Ga. App. 25, 28 (2017), which is not found in any legal databases.

## 24. Review Litigation Invoice Against Outside Counsel Billing Guidelines — Compliance Deviation Report (C-024)

*Criterion.* **Recommends full disallowance of black car service: \$387.00.** PASS if the report recommends disallowing the full \$387.00 for black car service (or reducing it to a reasonable taxi/ride-share equivalent). FAIL if no disallowance or reduction is recommended.

*Verdict.* Not defective. *Environment.* Correct. *AI auditors’ reasoning (reviewer’s assessment).* Wrong. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/P; Claude Opus 5.5: P/P.

*Analysis.* The billing guidelines state that “Reimbursable ground transportation is limited to ride-share ... or taxi” and “Limousine services and black car services are not reimbursable under any circumstances.” Sol’s critique of the rubric is that the policy requires disallowing black car reimbursement, not reducing the reimbursement amount to what a taxi or ride share would have cost. The rubric’s interpretation (that either disallowing the expense entirely or capping reimbursement at the amount a taxi would have cost) seems reasonable and closer to the likely real-world application of a policy like this.

## 25. Review Litigation Invoice Against Outside Counsel Billing Guidelines — Compliance Deviation Report (C-042)

*Criterion.* **Classifies business class airfare as borderline/requires discussion.** PASS if the report classifies the business class airfare issue (Messina’s \$2,847.00 flight) as a borderline or ‘requires discussion’ issue rather than a clear-cut mandatory reduction, given the 4-hour-20-minute flight duration and the 4-hour threshold. FAIL if it is classified as a clear mandatory reduction without acknowledging the ambiguity, OR if it is entirely dismissed as compliant.

*Verdict.* Defective (unambiguous objective error, rubric is incorrect). *Environment.* Correct. *AI auditors’ reasoning (reviewer’s assessment).* Correct. *Native grades* (Sonnet 4.6/GPT-5.5; P pass, F fail): GPT-6 Luna: P/P; Claude Opus 5.5: P/P.

*Analysis.* The rubric’s analysis is wrong. The travel policy included in the environment states, “For domestic flights with a scheduled duration of four (4) hours or more, business class may be booked without prior approval.” The record does not state the flight’s duration (the invoice lists only the route and fare class). The only basis for raising this seems to be in the rubric itself, which says the flight could take 4h20m, but even under that view, there’s nothing to discuss, as the policy allows business class for that duration.

## D A prospective test of outcome feedback

A follow-up would freeze a base model and an earlier training cohort, then compare outcome feedback with a matched control on later, disjoint cases. Related cases should remain together across splits. The primary endpoint would be a predeclared difference in proper score; calibration, failure rates, cost, and reasoning validity would be secondary endpoints. Blinded expert assessment of advice or drafting would test transfer.

For a valid forecast, the implemented Brier-shaped reward is

$$R(p, Y) = 0.02 + 0.98[1 - (p - Y)^2], \quad (8)$$

while invalid output receives zero. The positive affine transformation preserves the proper-score optimum among valid forecasts. The implementation also supplies threshold-based and clipped-log alternatives. Evidence spans are logged rather than rewarded. This design separates a directly scored prediction from an assertion that the accompanying reasoning is legally correct.

For a model emitting a rationale  $z$  and forecast  $p$ , a candidate objective is

$$\max_{\theta} \mathbb{E}_{(X,U,Y) \sim \mathcal{D}_{\text{train}}, (z,p) \sim \pi_{\theta}(\cdot|X,U)} \left[ 1 - (p - Y)^2 \right]. \quad (9)$$

This is an outcome-scored prediction objective, not a simulator that allows the agent to alter the litigation trajectory. Optimizing it need not require reinforcement learning: supervised probabilistic prediction, calibration, and policy-gradient approaches are competing methods to test. It also does not directly reward faithful rationales. Leakage-controlled held-out cases and an independent transfer endpoint are therefore necessary to distinguish generalization from memorized case outcomes or improved probability reporting alone.