

LegalForecastBench: Forecasting Judicial Decisions as a Test of Legal Reasoning Ability

John J. Hughes, III
Independent researcher
<https://www.johnjhughes3.com/research/>

Working paper — September 30, 2026

Abstract

Legal AI evaluations increasingly assess professional work products with expert-written rubrics. Judicial forecasting offers a complementary target: an uncertain legal judgment whose realized outcome can subsequently be observed. We introduce LegalForecastBench, a benchmark of claim–defendant-level dismissal probabilities from pre-decision records in U.S. federal litigation. The beta contains 91 cases decided between June 30 and August 7, 2026, with 387 scored units. In a panel of fourteen document-enabled model configurations, GPT-6 Sol has the lowest observed micro-Brier loss (0.11763) and the highest unit accuracy (84.50%). All fourteen improve on a retrospectively fitted constant-probability reference; thirteen improve on majority-class accuracy. Confidence differs substantially across configurations, and several strong models make consequential high-confidence errors. In a separate shared-summary experiment, two Luna configurations substantially outperform Jev; Jev underpredicts dismissal in the aggregate. We distinguish these evaluation findings from the hypothesis that outcome feedback can improve legal judgment, describe LegalForecastGym as a setting for testing that hypothesis, and examine task–rubric alignment through a bounded case study of Harvey’s Legal Agent Benchmark. The contribution is a legally specified probabilistic evaluation task and empirical baseline, with explicit limits concerning label adjudication, selection, contamination, repeated-run uncertainty, and transfer beyond forecasting.

1 Introduction

Reinforcement learning (RL) techniques have led to significant recent gains in the reasoning abilities of frontier AI models, but the largest and most significant gains appear to have come in verifiable domains. RL training environments specify what a model can observe and do, and which outcomes receive reward. In *reinforcement learning with verifiable rewards* (RLVR), correctness checks (like mathematical answer matching, executable tests, or formal proof checking) supply a scalable training signal that does not require a human to assess every model response [1, 2]. Earlier self-play systems demonstrated the power of outcome feedback in games with explicit rules [3].

The efficacy of these techniques in verifiable domains has been exceptional. During reinforcement learning, DeepSeek-R1-Zero’s average pass@1 on AIME 2024 rose from 15.6% to 77.9%; the final DeepSeek-R1 system, which used a broader training pipeline, reached the 96.3rd percentile on Codeforces [1]. AlphaProof coupled reinforcement learning with Lean’s formal proof checker and proved three problems from the 2024 International Mathematical Olympiad, including its hardest problem, solved by only five of 609 human contestants [4]. And OpenAI’s September

2026 announcement that it had solved the Navier–Stokes Millennium Prize Problem reportedly relied on an internal model developed through large-scale reinforcement learning, followed by Lean formalization and verification. [5, 6].

Despite its economic importance, the legal domain has not seen similar capability gains in longer-horizon tasks that are not as susceptible to verifiable rewards. Models have increasingly saturated some legal benchmarks that focus on bounded tasks with objectively scored answers, reducing their usefulness for distinguishing frontier models. For example, the Multistate Bar Examination, administered to prospective lawyers, contains multiple-choice questions that test the application of well-settled legal principles to stylized facts against an answer key [13]. LegalBench, a benchmark consisting of 162 tasks published in August 2023, similarly focused on tasks with unambiguous answers [12]. But frontier models have moved well beyond these benchmarks. Models have achieved perfect scores on an MBE practice-question benchmark. [66]. And Vals AI describes LegalBench as “largely saturated.” [67].

Researchers have tried at least two approaches to benchmark and train models on higher-level legal tasks: preference evaluation (comparing which of two outputs practicing attorneys prefer) and rubric-based reinforcement learning (grading model outputs against specified criteria to produce a reward). [7]. The rubrics as rewards approach has been studied in medicine and science [8], and leading vertical AI companies such as the Harvey AI Corporation report using the technique to train specialized legal models. [9, 10]. Rather than using human-written demonstrations for supervised training [11], the rubric defines what the system is rewarded for doing. Separate LLMs then judge outputs against the rubric to produce a reward signal.

In practice, however, the efficacy of these techniques in the legal domain is unproven, and some reinforcement learning environments appear to be artificial in ways that are likely to limit their usefulness for making AI models better at real-world tasks. In the course of working on this research, I discovered a number of errors in Harvey’s Legal Agent Benchmark. I used frontier models (GPT-6-Sol and Claude Opus 5.5, each using their first-party coding harness) to review all the litigation task rubrics for similar errors; both models agreed that hundreds of rubrics contained errors. In some cases, the answer key was inaccurate (such that, if the benchmark tasks were used for training, they could train the models into giving incorrect answers). In other cases, the task specifications contradicted information provided in the environment, or the requested work product diverged from the task environment. These construct validity issues could mean that the score does not adequately represent the professional capability being evaluated (or, if used for training, that the reward is misspecified.) [28]. Past research into reward-model overoptimization shows that optimizing a proxy does not necessarily improve the underlying objective [29, 30].

LegalForecastBench is an experiment in exploiting real-world litigation data as an evaluation metric (and potentially, later, as a training signal) for legal reasoning. The key insight motivating this approach is that legal tasks should be understood as lying on a spectrum of verifiability and specificity. At one end of the spectrum, some legal propositions are so objectively settled that they can give rise to simple and clear rules like those that LegalBench uses, for example, the rule that a car must stop before entering the intersection if a traffic light is red. Professor Hart describes such bright-line rules as having “a core of settled meaning” [14]. But even seemingly simple rules have “a penumbra of debatable cases in which words are neither obviously applicable nor obviously ruled out.” [14, 15]. And some legal questions, such as the meaning of constitutional provisions, may be heavily laden with value judgments or policy preferences and thus lack a clear and objective answer.

For lawyers and the legal system, litigation is the core mechanism for resolving uncertainty about

the law and reducing disputes over meaning into binding law. For example, if a contract dispute arises between two sophisticated parties, each side’s lawyer may advance arguments for why their preferred reading of the contract is the correct one. But if the dispute is submitted to a court, either in an action for breach or pursuant to the Declaratory Judgment Act, the court’s final judgment resolves that issue once and for all and constitutes binding law amongst those parties. [16]. Even an erroneous judgment is still binding law. [17]. Litigation thus turns a disputed question into an observable, legally operative disposition. From the perspective of training AI models, a difficult-to-verify legal question becomes a verifiable disposition every time a court enters judgment on a matter.

That makes predicting litigation outcomes an essential skill for lawyers in a variety of high-level judgment tasks. Oliver Wendell Holmes Jr., one of the most influential jurists in American history and a major precursor to American Legal Realism, famously wrote: "The prophecies of what the courts will do in fact, and nothing more pretentious, are what I mean by the law." [20]. While there are many components of high-level legal advice, decomposing many high-level legal judgment tasks into their constituent elements reveals that predicting litigation outcomes is often a key sub-task even in cases that never go to court. For example, lawyers frequently advise on commercial disputes. Prior work establishes that, even for disputes that never to go to a judge, the legal system provides a framework that constraints the parties’ respective bargaining range. [21]. In many disputes, each side’s best alternative to a negotiated agreement is to take the dispute to court, so the likely litigated outcome often sets a floor (or ceiling) on what each side is willing to accept (or offer) out of court. Even when a lawyer is merely giving general advice or preparing a contract between two parties who do not expect litigation, a good lawyer anticipates and prepares for potential disputes. [21, 22]. The value of prediction therefore extends beyond forecasting cases that ultimately reach judgment.

For the initial run of LegalForecastBench, I used AI models to collect documents from 91 real cases in federal court that had written decisions on motions to dismiss. Models, with my oversight, identified the moving defendants and the claims they challenged, and labeled the judge’s ruling based on whether it granted entirely a motion to dismiss a specific claim asserted against a defendant (or group of similarly situated defendants). AI models received the same briefing and docket that was before the judge (including the identity of the real parties and the name of the assigned judge, to the extent that information was on the docket prior to the judge’s disposition) and were asked to predict the judge’s ruling as to each claim-defendant unit. The headline metric is micro-Brier loss, a score for probabilistic forecasts. [18, 19]. GPT-6 Sol had the best observed scores, but this small, single-retained-run comparison is intended primarily to establish the framework for using real-world litigation data as an evaluation metric (and potentially a training signal). We separately tested a discriminative, non-autoregressive decision model (TypeSafe AI’s Jev model), using LLM-generated summaries of the court records (in order to fit the task within Jev’s small context window). Jev was not competitive with frontier models on this specific task.

Although this benchmark is merely evaluative, we are separately developing a training dataset intended to directly test the hypothesis that training models to predict litigation outcomes might give them skills that generalize to other legal judgment tasks. Just as training LLMs to predict the next token helped LLMs to understand latent facts about the world, in principle, teaching frontier models to predict litigation outcomes might help them understand the principles that drive real-world legal outcomes. Prior work shows that outcome-based forecasting can be a useful training signal, at least in helping models to forecast other events outside the training distribution. [23–25]. Whether improving models’ ability to predict litigation outcomes transfers to other legal judgment

tasks remains an empirical question.

2 Related work

The approach here is consistent with previous research that (i) used older AI techniques to build smaller, custom models for predicting the outcome of litigation (especially U.S. Supreme Court decisions), and (ii) more recent efforts to use frontier models to predict the outcomes of events defined in prediction markets. Some research in China has also used frontier models to predict criminal sentences.

Legal outcome prediction. Using AI models and other techniques to predict litigation outcomes has a long lineage. Ruger and colleagues used hand-coded metadata and statistical techniques to predict the 2002 Term of the U.S. Supreme Court [33]. Later work incorporated more sophisticated machine-learning and natural language processing techniques. Aletras and colleagues in 2016 used a text-based approach to predict the outcomes of European Court of Human Rights decisions [35]. Katz, Bommarito, and Blackman in 2017 used a random-forest classifier over structured Supreme Court data to predict Justice votes and case outcomes. [34]. CAIL2018 used a dataset of more than 2.6 million Chinese criminal cases to predict charges and sentences, with text-classification models providing initial baselines. [36]. As the transformer architecture became dominant, later legal prediction systems increasingly used models like BERT and T5. CLC-UKET, released in 2024, compared fine-tuned BERT and T5 models to zero- and few-shot GPT-3.5 and GPT-5 to predict U.K. Employment Tribunal cases. [38]. The field has thus progressed from statistical models over hand-coded metadata through conventional MLP and supervised machine learning to transformer-based classifiers and, more recently, large language models. Much of this literature differs from the approach studied here, because the models operate on structured metadata or facts and claims extracted from judicial opinions. The approach here simply asks frontier models to reason about the actual pre-decision litigation record, essentially the documents filed in court before the judge who decided the motion, without any access to statistical data or models showing prior decisions.

Legal benchmarks and evaluation. One of the most impressive academic benchmarks is LegalBench, a collaborative benchmark to which 40 authors contributed 162 tasks that aimed to assess legal domain expertise. Although some of the legal issues tested could give rise to disputes in real litigation (for example, there are tests of whether a given statement is hearsay or not), the benchmark deliberately uses relatively clean examples that can be reduced to a binary answer. [12]. The hearsay tasks, for example, explicitly articulate what proposition a given statement is being used to prove or disprove, and then ask whether that use is hearsay; in real-world litigation, what proposition the statement is being offered to prove may itself be disputed so providing a clean prompt on that point narrows the task to a relatively mechanical application of an explicit rule.

Most recent benchmarks test higher-level judgment tasks that lack an objective ground truth. JudgmentBench, released in May 2026, and Harvey’s LAB targeted open-ended professional legal work, like drafting, analysis, and advice, for which there are potentially multiple defensible outputs and no independently verifiable ground truth. [7, 39]. These benchmarks take different approaches to reward specification. JudgmentBench explicitly compared rubric-based scoring to comparative judgment (pairwise preference judgments collected from practicing attorneys), and found that comparative judgments recover the intended quality ordering substantially better than rubrics. LAB

relies on a rubric-based approach. There are potential downsides to both approaches. Preference-based evaluation is a form of RLHF that can amplify human biases and errors into the rewards signal [31, 32]; for example, reviewers could reward briefs that are longer, sycophantic, overly confident, or more consistent with the reviewer’s own judicial philosophy or political preferences, rather than rewarding briefs that are most effective at persuading a judge. As discussed herein, my preliminary audit of the LAB rubrics also found evidence of a significant number of errors and misspecified tasks, with the tasks apparently having been created by leading data companies that are selling significant amounts of training data to AI labs.

Verification, rewards, and forecasting. Frontier models are trained on several different forms of feedback, as described in prior work. InstructGPT in 2022 provided an influential demonstration of reinforcement learning from human feedback [11]. In 2023, Lightman and colleagues showed that, on mathematical reasoning tasks, process supervision (assigning feedback to individual reasoning steps) outperformed outcome supervision, which gives rewards based on whether the final answer is correct [40].

More recent work has used externally resolved real-world events to supply an analogous signal outside domains like mathematics. For example, Autocast, released in 2022, assembled forecasting questions from tournaments and a date-indexed news corpus to allow language models to make predictions based on the information a human would have had at the time. [41]. Halawi and colleagues in 2024 used a retrieval-augmented forecasting system that could search for information, and generate and aggregate multiple forecasts; the system approached the performance of competitive human forecasters [42]. ForecastBench, introduced in 2024, extended this approach into a dynamic benchmark that continually collects questions about unresolved future events and scores models once the event outcome is known [43].

Outcome-Based Reinforcement Learning to Predict the Future showed in 2025 that resolved binary forecasting questions could serve as reinforcement-learning rewards, training a 14-billion-parameter model on more than 100,000 events using adaptations of GRPO and ReMax [23]. OpenForecaster similarly explored training language models directly on resolved forecasting questions [24]. LegalForecastBench is an effort to apply a similar approach to litigation outcomes, at least as an evaluation metric for now.

Measurement validity. The correct scoring rules are those that encourage honest probabilities in expectation [18, 19]. Statistical evaluation also must respect dependence among questions sharing a record [44]. OpenAI’s published work on SWE-bench Verified also emphasizes the importance of avoiding specification error (which is relevant to the findings about rubric-based benchmarks, such as the preliminary audit of LAB). After published that benchmark in August 2024, OpenAI noted concerns in February 2026 that the benchmark no longer measures the coding capabilities of frontier models because the tests reject correct solutions and models were increasingly contaminated by having seen the open-source repositories often used for training purposes [45, 46]. These specification problems apply both to rubric-based approaches and to our own outcome labels.

3 Dataset and task

3.1 Prediction target and information set

Let $c \in \{1, \dots, C\}$ index cases and $u \in \mathcal{U}_c$ index scored units within case c . Each unit identifies a challenged claim or a separately resolveable pleaded subclaim against a moving defendant (or a

group of similarly situated defendants, for whom the outcome is predicted together since the legal and factual arguments for dismissal are materially identical as to all of them). We used models to identify appropriate units based on the arguments articulated in each motion; an affirmative defense, an element of a legal cause of action, a requested remedy, or a distinct factual theory was not automatically treated as a separate unit. Rather, units generally tracked claims asserted in the complaint (in some cases with a single claim asserted against multiple defendants being treated as different units, if different defendants advanced different grounds for dismissal). The model receives the pre-decision record that the judge received X_c and the unit specification U_{cu} and returns

$$p_{cu} \in [0, 1], \quad p_{cu} \approx \Pr(Y_{cu} = 1 \mid X_c, U_{cu}), \quad (1)$$

where $Y_{cu} = 1$ means that the judge dismissed that claim-defendant unit in its entirety, in the target motion’s first written disposition. That event is not a final victory in the lawsuit, since a plaintiff might be given an opportunity to file an amended complaint with additional allegations supporting a dismissed claim, other claims might survive even if some are dismissed, and the plaintiff normally has a right to appeal dismissal and seek a de novo adjudication of the issue from the circuit court of appeals.

A full dismissal with leave to amend still has label 1 (i.e., a dismissal is coded the same way, even if it is without prejudice and the plaintiff may seek leave to amend the complaint and try again). A partial dismissal of some claims or theories that leaves the frozen unit alive in a material respect has label 0, meaning the claim-defendant unit must be entirely dismissed. Later amendment, settlement, reconsideration, and appellate reversal do not change the target; the model is specifically asked to predict only the trial judge’s ruling on the specific motion to dismiss. Dispositions that cannot be reliably mapped to a unit were excluded from primary scoring. These choices sacrifice some procedural detail for a stable, objectively interpretable forecast event. Even with these limitations, label construction is a substantive legal task in its own right as the labeling models had to analyze the filed motions to determine what claims were being challenged and ascertain from the judge’s ruling how those were resolved. Two labeling models had to agree on the units, and had the option to escalate ambiguous cases to me for review; in practice, the labeling process was iterative in this preliminary run, as we discovered a number of edge cases that needed to be resolved. The process started with 100 cases but seven were eliminated during the labeling process because the order did not decide the frozen claims, one was eliminated because it was a stipulated dismissal, and one was a plaintiffs’ Rule 12(c) motion for judgment on the pleadings rather than a defendant’s motion to dismiss under Rule 12(b).

First written disposition of the frozen unit	Primary label
Entire claim dismissed, with or without leave to amend	1: fully dismissed
Only one nonseparable theory or remedy dismissed; claim remains	0: not fully dismissed
Claim dismissed against one defendant but survives against another	Separate labels for separately frozen units
Disposition does not reliably resolve the frozen target	Excluded from primary scoring

Table 1: Illustrative label semantics, not additional dataset observations. The documented protocol defines unit boundaries before consulting the outcome.

3.2 Cohort construction and labeling

The analyzed set consists of 91 U.S. federal cases with written decisions dated June 30–August 7, 2026. The cases were labeled as contained 409 claim-defendant units that could be forecasted. Of




Figure 1: Information separation in the forecasting task. Models inspect pre-decision material and frozen targets. The subsequent judicial disposition is used in the labeling and scoring path, not in the model-visible packet. In the retrospective beta, this is an information boundary rather than a claim that forecasts were made before the decisions occurred.

those, 387 are scored; 22 units therefore do not enter the headline metrics. The unscored units involved situations where the plaintiff voluntarily withdrew the frozen claim, the order did not cleanly grant or deny the frozen claim, or (for four units) we determined after freezing the unit that the frozen defendant was the wrong target or the motion did not address the frozen claim. Of the scored claim-defendant units, 285 were fully dismissed and 102 were not. The shared-summary preparation inventory contains 482 selected documents. The sample was not random; our pipeline favored cases where some of the necessary documents could be obtained from online databases, although we typically had to purchase a number of documents from the federal PACER system to ensure that the task included all the substantive filings on the motion that were before the judge.

The case construction and labeling pipeline used models as well as human legal oversight from me. We did take pains to maintain separation between *unitization* and *outcome coding*. In particular, we wished to avoid the possibility that the prediction units could leak information about the judge’s ultimate outcome to the models that were performing prediction. AI models first were asked to identify the claim-defendant units that a motion sought to dismiss, based only on the pre-decision materials (the motion and the docket). A later stage was given access to the claim-defendant units and the judge’s order disposing of the motion, to interpret the disposition for each frozen target. The labeling pipeline we used requires structured outcome agreement and supporting disposition excerpts for automatic acceptance; disagreements between models, ambiguity, and structural defects in the cases were routed to me for lawyer review and a decision. I resolved disputed units and corrected labels as necessary. For purposes of the initial release, this process was iterative and we did not track a blinded human-audit error estimate.

3.3 Temporal design and contamination

The models predicting outcomes were not given access to the target decision or the outcome. The construction pipeline also screens docket and filing text to avoid outcome leakage. Models had access to document text tools in a sandbox with no Internet access. Model API calls were made with provider-side web search tools disabled where applicable.

The decisions used in the benchmark were handed down after the reported knowledge or training cutoffs for twelve of the fourteen tested models. Kimi and Muse did not publish training cutoffs,

but based on the model release dates, it appears unlikely that those models were exposed to the decisions being tested in training.

The acquisition and labeling pipeline is designed to create new cohorts in the future, when new decisions become available. This enables fresh comparisons even after models might have been trained on the decisions at issue here. However, the downside is that refreshing a benchmark changes its case mix and difficulty, scores from different versions cannot be interpreted as a time series of model improvement.

4 Evaluation protocol and scoring

4.1 Execution

The main comparison uses a common document-access interface. Models can read documents using tools and produce a probability for each unit, together with a short rationale. Provider-specific reasoning settings are retained in the run registries. Appendix A lists the configuration settings. The study population is the fourteen configurations in Table 2.

4.2 Proper scores and descriptive references

For $N = \sum_c |\mathcal{U}_c| = 387$ and $C = 91$, the primary metric is the micro-Brier loss

$$B_{\text{micro}} = \frac{1}{N} \sum_{c=1}^C \sum_{u \in \mathcal{U}_c} (p_{cu} - Y_{cu})^2. \quad (2)$$

The equal-case score changes the estimand by giving each case the same weight:

$$B_{\text{case}} = \frac{1}{C} \sum_{c=1}^C \frac{1}{|\mathcal{U}_c|} \sum_{u \in \mathcal{U}_c} (p_{cu} - Y_{cu})^2. \quad (3)$$

Lower is better for both. Micro-Brier treats each scored claim–defendant unit equally, so cases that involve a motion to dismiss many separate claims receive more weight than cases where only one claim and defendant was subject to a motion to dismiss. Equal-case Brier weights each prediction unit in proportion to the number of units from that case.

The Brier rule is strictly proper for the binary target. If the true conditional event probability is q , then

$$\mathbb{E}[(p - Y)^2 \mid X, U] = (p - q)^2 + q(1 - q), \quad (4)$$

which is minimized by reporting $p = q$. The second term is irreducible uncertainty for that information set. A set of observable judicial events therefore create a verifiable score across the entire probability distribution, even though a single realized outcome standing alone does not reveal whether the individual probability a model assigned to that outcome was correct ex ante [18, 19].

We report accuracy at threshold $p \geq 0.5$ as a secondary descriptive metric; in other words, accuracy percentage is simply the percentage of cases where the model’s assigned 50

The observed dismissal frequency is $\hat{\pi} = 285/387 = 0.73643$. Always predicting dismissal has accuracy 73.64%, so a constant forecast $p = \hat{\pi}$ has micro-Brier $\hat{\pi}(1 - \hat{\pi}) = 0.19410$. The constant is fitted retrospectively on this cohort, so it is a descriptive reference, not a baseline derived from earlier cases. The uninformative forecast $p = 0.5$ has Brier 0.25.

Log loss is a complementary diagnostic:

$$L = -\frac{1}{N} \sum_{c,u} [Y_{cu} \log p_{cu} + (1 - Y_{cu}) \log(1 - p_{cu})]. \quad (5)$$

For a wrong forecast assigning 0.99 to the predicted class, Brier loss is 0.9801 and log loss is approximately 4.605; at confidence 0.90 the corresponding losses are 0.81 and 2.303. Thus log loss imposes a large penalty on extreme errors, specifically, instances where the model assigned a very high probability to the outcome that did not occur.

4.3 Uncertainty and scope of inference

Units from the same case share facts, and distinct claims may suffer from closely related legal defects. We therefore measure within-case dependence using an observed-scale intraclass correlation coefficient (ICC), rather than treating claim–defendant units as independent observations [68].

Across 387 scored units nested in 91 cases, the observed binary outcome ICC was 0.76 (95 percent whole-case bootstrap CI 0.61–0.87). Within-case agreement was 998 of 1,066 pairs (93.6 percent), compared with 44,623 of 73,625 pairs (60.6 percent) across cases. GPT-6 Sol’s Brier-loss ICC was 0.42 (95 percent CI 0.24–0.58). The estimator is an unbalanced method-of-moments one-way ANOVA ICC; the interval uses 20,000 whole-case resamples, assuming cases are independent for interval estimation.

The observed dependence supports resampling whole cases while retaining paired model predictions when quantifying uncertainty in score differences [69]. The fourteen-configuration table is descriptive and does not claim a significance ordering. The shared-summary experiment has its own three-condition comparison and its own multiplicity correction. Neither analysis establishes independence across related cases or multidistrict-litigation families, and neither measures variation across repeated model runs [44].

5 Results

5.1 Full-document comparison

Table 2 reports the fourteen-configuration panel. GPT-6 Sol achieves the best observed micro-Brier (0.11763), equal-case Brier (0.12949), and accuracy (327/387, 84.50%). Relative to the retrospectively fitted constant, its micro-Brier is lower by about 39.4%. All fourteen configurations beat that Brier reference; thirteen beat majority-class accuracy. Muse’s 279/387 correct forecasts (72.09%) illustrate why accuracy and probability quality need not give the same comparison.

More expensive or newer configurations do not uniformly score better. Astra’s micro-Brier is worse than those of GPT-5.6 Sol and Luna; Astra improves on Luna’s equal-case Brier but not its accuracy. Grok 4.7 has worse observed scores than Grok 4.6. GPT-6.1 Sol is worse than GPT-6 Sol on micro-Brier (0.14948 versus 0.11763). Claude Sonnet 5.5 is worse than Claude Sonnet 5 on micro-Brier (0.17298 versus 0.15848), though its unit accuracy is slightly higher (76.23% versus 75.71%). These are cohort-specific observations from retained successful runs, not evidence of a general capability regression or equivalence across model families.

Configuration	Micro Brier	Case Brier	Correct	Acc. (%)	Cost (\$)
GPT-6 Sol	0.11763	0.12949	327/387	84.50	11.01
Claude Fable 5.1	0.12877	0.13653	318/387	82.17	83.18
Grok 4.6	0.13368	0.13480	312/387	80.62	15.72
Claude Opus 5	0.13987	0.13548	310/387	80.10	45.89
GPT-5.6 Sol	0.14009	0.13319	313/387	80.88	24.91
Kimi K3	0.14130	0.15367	303/387	78.29	16.87
GPT-5.6 Luna	0.14322	0.15084	313/387	80.88	1.48
GPT-6.1 Sol	0.14948	0.13885	304/387	78.55	10.68
GPT-6 Astra	0.15221	0.14275	305/387	78.81	59.74
Gemini 3.8 Flash	0.15331	0.13900	310/387	80.10	35.61
Grok 4.7	0.15833	0.16454	298/387	77.00	33.00
Claude Sonnet 5	0.15848	0.15964	293/387	75.71	24.68
Claude Sonnet 5.5	0.17298	0.17183	295/387	76.23	6.99
Muse Spark 1.3	0.18121	0.17845	279/387	72.09	8.78
Cohort-rate reference	0.19410	—	285/387	73.64	—
Constant $p = 0.5$	0.25000	0.25000	—	—	—

Table 2: Full-document results, ordered by micro-Brier loss. Each model row covers 91 cases and 387 scored units. The last column reports the retained successful-workload cost estimate under heterogeneous accounting assumptions detailed in the text; it is not a controlled price comparison. The cohort-rate reference is retrospectively fitted, and its accuracy column uses always-dismiss classification. No fourteen-model significance ordering is claimed.

5.2 Sharp forecasts and confident errors

Define confidence $h_{cu} = \max(p_{cu}, 1 - p_{cu})$ and the high-confidence subset $H_m = \{(c, u) : h_{cu} \geq 0.9\}$ for configuration m . Comparing the mean confidence in H_m with classification accuracy in H_m gives a useful conditional calibration diagnostic. It is not a full reliability curve: the subset differs by model, and its units remain clustered by case.

GPT-5.6 Sol, GPT-5.6 Luna, and Astra make high-confidence forecasts on 58.9%, 47.0%, and 53.5% of units respectively, compared with 20.9% for Fable. Their mean confidence in those subsets is approximately 95–97%, while realized accuracy is approximately 89–90%. Astra misses 21 of 207 such predictions. However, greater sharpness is not unique to OpenAI: Gemini makes 276 high-confidence forecasts (71.3% of units) and misses 35. GPT-6 Sol provides a useful counterexample within OpenAI, missing 8 of 169 high-confidence forecasts, for 95.3% subset accuracy. GPT-6.1 Sol is sharper and less accurate in that subset, missing 18 of 199. Claude Sonnet 5.5 makes 63 high-confidence forecasts and misses 5.

Fable has no errors among 81 high-confidence units, while Kimi and Grok 4.6 each miss one. These observations do not establish population error rates of zero or one percent. A model may achieve excellent accuracy on a smaller, easier high-confidence subset. The practical finding is that confidence behavior varies substantially and can materially affect proper scores; model price or flagship status is not a sufficient calibration guarantee.

5.3 Inference cost

The retained normalized successful-workload figures range from \$1.48 for GPT-5.6 Luna to \$83.18 for Fable. These are not total experiment budgets or uniform invoice charges. Fable’s figure is a reconstruction of hypothetical cached execution; its historical uncached successful workload was




Figure 2: Micro- and equal-case Brier scores for the same fourteen configurations. The connecting segments show weighting sensitivity, not uncertainty intervals. The dashed constant-probability reference applies only to micro-Brier.

\$189.70. Luna’s billing-derived figure includes eight extra requests. Several other figures normalize provider discounts or impute missing cache information.

The retained corrected estimate for Grok 4.7 is \$33.00 at standard rates, alongside \$11.52 in provider-reported successful-call charges. Cache-adjusted pricing is not fully observed. GPT-6 Sol’s corresponding standard estimate is \$11.01, versus a saved pricing-based workload estimate of \$5.50. Failed attempts and unresolved charges are generally outside these totals. On these stated estimates Luna and GPT-6 Sol form the observed cost–score frontier; its precise location depends on the accounting assumptions. GPT-6.1 Sol’s recorded successful-call cost is \$10.68, and Claude Sonnet 5.5’s is \$6.99. Both are the totals saved on the run receipts.

6 A shared-summary experiment with Jev

Jev is described by its developer as a probabilistic “System One” model, exposing native Boolean probabilities rather than generating an agentic legal analysis [47]. Its context and interaction constraints differ from those of the document-enabled configurations. We therefore treat its evaluation as a separate experiment, not a row in the full-document model ranking.

GPT-5.6 Luna summarized each selected document using instructions to preserve material facts, arguments, authorities, and procedural posture without predicting the result. The summarizer received neither outcome labels nor other models’ forecasts. Jev and two Luna comparator configurations then received the same frozen summary cache and unit questions in a single request per case, without document tools or follow-up turns. The conditions comprise Jev through the historical `typesafe-ai/jev` Gateway alias, GPT-5.6 Luna with high reasoning, and GPT-6 Luna without




Figure 3: Conditional confidence diagnostic on predictions with $\max(p, 1 - p) \geq 0.9$. Each row uses the subset selected by that model, rather than a common set of units. Wrong/selected counts: GPT-6 Sol: 8/169; Fable 5.1: 0/81; Grok 4.6: 1/112; Opus 5: 4/121; GPT-5.6 Sol: 26/228; Kimi K3: 1/92; GPT-5.6 Luna: 20/182; GPT-6.1 Sol: 18/199; GPT-6 Astra: 21/207; Gemini 3.8 Flash: 35/276; Grok 4.7: 4/98; Sonnet 5: 5/60; Sonnet 5.5: 5/63; Muse Spark 1.3: 15/135. This is not a full reliability diagram or an estimate of statistical significance.

reasoning. The two Luna conditions differ in model generation as well as reasoning setting, so their contrast does not isolate a reasoning effect.

Jev’s micro-Brier of 0.36754 is worse than the 0.25 loss of always reporting $p = 0.5$. Its mean predicted dismissal probability is approximately 0.305 against an observed frequency of 0.736. This directly shows severe underprediction in the aggregate, rather than inferring calibration solely from a large Brier loss. Its 32.6% threshold accuracy is also below majority-class prediction.

GPT-5.6 Luna on the identical summaries achieves micro-Brier 0.15489 and accuracy 78.6%. The paired case-cluster bootstrap gives a Luna-minus-Jev difference of -0.21265 , with a 99.44% Bonferroni-adjusted interval (95% familywise coverage) of $[-0.32644, -0.06433]$ across the three model pairs and three metrics. The result remains exploratory because potential cross-case dependence was not assessed. The two Luna conditions are not distinguishable under this correction.

Condition	Micro Brier	Case Brier	Acc. (%)
GPT-5.6 Luna full documents	0.14322	0.15084	80.88
GPT-5.6 Luna summary + high reasoning	0.15489	0.16119	78.55
GPT-6 Luna summary + no reasoning	0.17152	0.18694	74.16
Jev summary + one-shot	0.36754	0.34416	32.56

Table 3: Shared-summary results and the full-document Luna reference. All rows use the same 91 cases and 387 scored units. Only the last three rows share a one-request summary interface; the first row is a descriptive full-document reference.




Figure 4: Probability quality with a shared summary cache. Both Luna conditions exploit predictive signal in the summaries, while Jev scores worse than the constant $p = 0.5$ forecast. The full-document reference changes both information access and interaction. Connecting segments indicate weighting sensitivity, not uncertainty intervals.

The shared-summary control shows that the summaries contain useful predictive signal that this Jev configuration does not exploit. It does not show that every material detail survived compression. Full-document Luna performs better than summary-only Luna, but that contrast changes information access and interaction together. Nor does the experiment establish that probabilistic or non-generative architectures are unsuitable for litigation. Domain exposure, task adaptation, the event formulation, and interface behavior remain possible explanations; Jev’s litigation-specific training coverage is unknown.

Summary preparation cost an estimated \$10.32. Successful inference estimates are approximately \$0.044 for Jev, \$1.62 for GPT-5.6 Luna, and \$0.103 for GPT-6 Luna. These registry-based estimates exclude retries and are not invoice totals. Omitting the common preparation cost would misstate the economics of the complete pipeline.

7 Prediction, adjudication, and legal strategy

7.1 An observable disposition is not an infallible legal answer

Litigation can resolve the parties’ rights even when the underlying interpretation remains contestable. Holmes’s contract example makes the distinction concrete: the parties anticipate different dates for a promised lecture, yet are bound by the court’s construction requiring performance within a reasonable time [20]. Finality between parties is also different from precedent governing later cases. A final, unappealed judgment on the merits can retain preclusive effect despite legal error [17]; a ruling in one case need not settle every future application of the same rule. LegalForecastBench uses an even narrower event: the first written disposition of a specified motion. Later amendment, settlement, or appellate reversal does not alter whether that event occurred.

Politically salient constitutional disputes illustrate the difference between an authoritative decision and continuing disagreement about its correctness. Before June 24, 2022, *Roe* and *Casey* supplied a federal constitutional protection against certain state restrictions on abortion, including the undue-burden framework for pre-viability restrictions. In *Dobbs v. Jackson Women’s Health Organization*, the Supreme Court overruled those precedents and held that the federal Constitution does not confer a right to abortion [48]. The decision changed the governing federal constitutional rule without ending disagreement about it or resolving every issue under state constitutions and other law. For

forecasting, a change of this kind also matters because past outcomes may cease to represent the legal regime governing future disputes.

7.2 Why forecasts affect settlement

The bargaining-in-the-shadow-of-the-law account connects expected adjudication to private agreement [21]. A simple risk-neutral model makes the role of probability explicit [49]. Suppose a plaintiff would receive a fixed award $J > 0$ if successful and zero otherwise. Let q_P and q_D be the plaintiff’s and defendant’s respective estimates of the probability of that success, and let $k_P, k_D \geq 0$ be their remaining litigation costs. Assuming each side bears its own costs, ignoring settlement costs, and restricting attention to monetary payoffs, the plaintiff’s expected trial payoff is $q_P J - k_P$, while the defendant’s expected trial cost is $q_D J + k_D$. A payment s from defendant to plaintiff is weakly acceptable to both when

$$q_P J - k_P \leq s \leq q_D J + k_D. \quad (6)$$

The interval is nonempty when

$$(q_P - q_D)J \leq k_P + k_D. \quad (7)$$

This is an illustrative reservation-value calculation in the economic analysis of litigation, not an estimated model of this dataset [50, 51]. Agreement about the probability of success creates a settlement surplus equal to the avoided costs. The settlement interval disappears when the plaintiff’s estimate of success exceeds the defendant’s estimate of the plaintiff’s success by enough that $(q_P - q_D)J > k_P + k_D$. Even when the interval exists, private information or strategic bargaining can prevent agreement [51]. A party should reject a particular settlement when its expected litigation payoff is higher under these assumptions. That does not imply that every other possible settlement is unacceptable. Nonmonetary goals, risk preferences, fee shifting, and uncertain damages require a richer model. A motion-to-dismiss probability informs this analysis but is not itself a probability of ultimate recovery or an estimate of settlement value.

7.3 Prediction as a component of legal advice

High-level strategic advice involves several tasks. A lawyer must (i) understand the client’s objectives; (ii) analyze the documents and other facts; (iii) analyze the applicable law; (iv) predict how a court might resolve the dispute; (v) assess counterparties’ incentives and objectives; (vi) identify potentially acceptable resolutions; and (vii) develop a strategy for obtaining the best realistically achievable outcome for the client. This is our decomposition of the advice task, consistent with the MacCrate skills framework and the New York State Judicial Institute’s account of counseling and negotiation [22]. Prediction informs these other activities, but does not determine the client’s goals or exhaust a lawyer’s professional responsibilities.

Prediction also bears on the selection and organization of arguments. Each additional point consumes briefing space and judicial attention. In *Jones v. Barnes*, the Supreme Court recognized the value of selecting stronger appellate arguments and the risk that adding weaker ones can dilute them [52]. This supports testing whether a model can prioritize consequential issues, rather than merely identify every technically available objection. A separate evaluation could ask experts to assess issue selection and the connection between the model’s rationale and the court’s stated grounds; forecast scores alone do not establish those abilities.

Professional experience should not be equated with demonstrated calibration. Goodman-Delahunty and colleagues’ prospective study of lawyers’ case-outcome predictions found overconfidence and

no calibration advantage from more years in practice [53]. The study used lawyers’ own minimum case goals and reported outcomes; years of practice did not measure specialization or experience with similar disputes. The practical importance of forecasting therefore motivates measuring both human and model performance. Our study does not include a matched human baseline or a direct comparison between models’ research skills and strategic judgment.

8 Outcome feedback and LegalForecastGym

LegalForecastGym implements a separate, single-step forecasting environment intended for post-training experiments. It renders a static litigation packet and receives a structured action containing a dismissal probability and cited evidence spans. Training loaders admit only explicitly designated training collections and reject locked test matters. Episode assembly keeps the label out of the observation. The current phase does not implement a multi-turn litigation simulator or train an agent to take procedural actions.

For a valid forecast, the implemented Brier-shaped reward is

$$R(p, Y) = 0.02 + 0.98[1 - (p - Y)^2], \quad (8)$$

while invalid output receives zero. The positive affine transformation preserves the proper-score optimum among valid forecasts. The implementation also supplies threshold-based and clipped-log alternatives. Evidence spans are logged rather than rewarded. This design separates a directly scored prediction from an assertion that the accompanying reasoning is legally correct.

The checked Gym artifacts contain no completed parameter-update experiment or held-out before/after training comparison. They do contain a live sampling comparison on 50 synthetic development episodes, used to assess output format and select an open-weight sampling model. That run is infrastructure evidence, not evidence of legal forecasting improvement. A separately recorded Gym export of 100 matters and 425 scoreable units is also distinct from the 91-case, 387-unit benchmark analyzed here; those populations are not pooled.

The intended learning study asks whether earlier outcome feedback improves predictions on later, disjoint matters and whether any gains extend beyond forecasting. Parameter learning, in-context feedback, and changes to the surrounding agent system are different interventions. Each needs a matched control and a held-out endpoint. Forecasting literature makes such learning plausible, and OpenForecaster reports calibration transfer to other question-answering benchmarks; those findings do not establish improved litigation strategy [23, 24]. Appendix D specifies a concrete prospective design.

9 Task validity and rubric-based legal evaluation

9.1 Why the reward definition matters

A June 2026 Harvey industry article characterizes litigation as slower to adopt agents than some volume-driven workflows; this is a vendor observation, not a measured cross-practice adoption study [54]. Harvey launched LAB on May 6, 2026 as an open-source collection of legal-agent tasks with matter files, requested deliverables, and expert-written criteria [39]. The launch describes 1,200 tasks spanning 24 practice areas and more than 75,000 criteria. Its initial-results report announced a partnership with Artificial Analysis, which subsequently introduced a 120-task LAB-AA variant

[55, 56]. This establishes a prominent industry evaluation effort, not a field-wide consensus that LAB or its publisher uniquely defines legal capability.

LAB-style criteria also enter training. Harvey and Applied Compute report full-parameter reinforcement learning of a legal agent using rubric-based rewards [9]. Harvey’s August 2026 Tenet research preview describes asynchronous RL of a Kimi K3-based model, with rewards combining criterion coverage, issue-level success, and perfect-completion bonuses. It reports gains on LAB and other legal benchmarks [10]. These are first-party training reports, not independent replications, but they make task validity consequential: a misspecified reward can shape learned behavior as well as a leaderboard.

Rubrics are useful when they encode meaningful requirements and are reliably graded. The risk is a mismatch among the requested work, the record supplied, the legal answer prescribed by a criterion, and the grader’s application of that criterion. These are distinct failure modes. JudgmentBench’s comparison of expert rubric scores and pairwise preferences on 30 BigLaw Bench tasks provides related empirical motivation, although its constructed quality levels and task selection limit generalization [7].

The evaluation question is one of construct validity: whether the measurements adequately represent the capability the benchmark claims to assess [28]. If the same criteria supply training rewards, a defective criterion also raises a question of reward misspecification. Reward hacking is a possible learned response to such a specification, and reward-model overoptimization is a possible divergence between an improving proxy score and the underlying target. These diagnoses require different evidence [29, 30]. Gao and colleagues study this divergence using a synthetic gold reward model; our audit examines task and criterion validity and does not measure a training trajectory. Outcome scores also have a limited target: they reward accurate forecasts, without directly rewarding good advice or faithful explanations.

9.2 Preliminary audit and a targeted task review

Our retained AI-assisted review covers the pinned litigation-and-dispute-resolution subset: 52 tasks and 2,858 criteria. GPT-6 Sol flagged 616 criteria; Claude Opus 5.5’s final reconciliation flagged 564. Their union contains 835 flagged criteria, leaving 2,023 unflagged by either final report. Of 345 overlapping flags, both reviewers called 83 problematic and 135 arguable; 127 received different levels of concern. The union represents 29.2% of the 2,858 criteria; the other 70.8% were not flagged, which does not establish that they are correct. These are review outputs, not a verified error rate. Opus’s reconciliation saw Sol’s findings, so agreement in the final reports is not independent confirmation [26].

The author’s targeted review of **analyze-counterparty-motion-to-dismiss** motivates a narrower case study (Appendix C). Among its 34 criteria, the AI reviewers flagged 16. Source inspection supports concrete concerns: a categorical integration-clause proposition despite express non-reliance language; a required characterization of accrual as factually disputed despite dated allegations requiring careful analysis; and output requirements, including a minimum number of recommendations, that the user-facing instruction does not specify [27]. These examples do not justify treating every flagged criterion as wrong, or treating every solver failure as a grading defect.

The task also illustrates why a realism critique must be precise. The complaint quotes a Texas forum-selection clause while filing in Georgia, but it also advances venue and scope arguments. That is a contestable litigation position, not proof that real litigants never behave this way. Similarly, the inspected motion contains a caption, table of contents, and table of authorities. Programmatic

document metadata establishes templating, not LLM authorship. The strongest criticism concerns how the rubric evaluates the supplied legal problem, rather than inferring professional invalidity from the documents’ appearance.

9.3 Implications for evaluation and training

All-pass aggregation can magnify one defective criterion: a useful work product fails if any required item fails. This does not make all-pass evaluation inherently inappropriate; some legal requirements are indispensable. It does mean that criteria should distinguish dispositive analysis, necessary completeness, and optional presentation. A benchmark should permit a well-supported rejection of a planted argument when the record or law makes that argument weak.

OpenAI’s 2026 analysis of SWE-bench Verified makes an analogous concern concrete: in a selected audit of 138 difficult tasks, it reported problems involving test requirements and task specification [46]. That finding supports examining task–test alignment; it provides no defect estimate for LAB. Neither the preliminary LAB audit nor the present forecasting results establish that rubric-based RL is ineffective, that synthetic tasks generally harm performance, or that training on litigation tasks caused degradation. A stronger experiment would compare matched environments built from synthetic and authentic records, independently adjudicate their criteria, and evaluate transfer to held-out professional tasks.

There is also evidence that rubric-based training can improve performance under a different evaluation format. Gunjal and colleagues report gains on GPQA-Diamond, a multiple-choice science benchmark, relative to direct Likert-based judge rewards [8]. This result does not establish transfer to legal tasks, but it cautions against assuming that rubric rewards merely teach compliance with the training rubric.

Real litigation is valuable because it supplies naturally occurring arguments, procedural complexity, and an outcome that the benchmark designer did not choose. It still requires expert construction. Authentic records can support both outcome scoring and carefully validated rubrics. The research opportunity is to test which supervision signals teach useful judgment, rather than assuming that either the court’s outcome or an expert checklist is a complete measure of legal competence.

Industry use and the limits of that evidence. Harvey’s product announcements report LAB all-pass results for Opus 4.8 (10.4%, May 28), Fable 5 (13.3%, June 9), and Sonnet 5 (5.8%, June 30) [57–59]. These examples show repeated use in model evaluation and product positioning. They are Harvey’s evaluations of third-party models, not evidence that each model developer independently adopted LAB as its definitive legal benchmark. The LAB-AA partnership broadens distribution, but neither prominence nor a commercial partnership establishes evaluation validity.

10 Limitations and generalization

Selection into adjudication. The target population is a set of accessible cases with written motion decisions, not all disputes, all filed cases, or all contemplated motions. Parties settle, abandon claims, amend pleadings, or choose not to sue. Cases reaching a published disposition may differ in uncertainty, stakes, representation, and bargaining conditions. Calibration conditional on reaching this dataset does not establish calibration for counterfactual outcomes in the broader dispute population [50]. At the same time, selection does not prove that learned skills cannot transfer: that is an empirical question requiring an independently defined evaluation population.

Prediction is not legal correctness. Judicial outcomes reflect doctrine and record facts, but also procedure, judicial behavior, and institutional context. Judge identity remains visible. We have not established how much performance comes from close legal analysis rather than judge, court, party, or case-type priors. A metadata-only comparison, judge masking, and a human forecasting baseline are important future tests. Neither fluent rationales nor a correct forecast demonstrate valid legal reasoning. Training to imitate outcomes can reproduce historical disparities or errors; professional advice also has obligations that prediction scores do not capture.

Labels and completeness. Unitization can omit, merge, or split legally distinct targets; dispositions can be ambiguous or procedural. Human oversight helps but does not quantify residual error. The absence of a complete blinded label audit and acquisition/exclusion census limits claims about both label accuracy and representativeness. Excluding 22 unscored units preserves an interpretable target but may also change difficulty. These limitations deserve the same scrutiny as disputed rubrics in another benchmark.

Statistical and execution limits. One successful prediction per unit does not characterize run-to-run variance, operational reliability, or expected total cost. The cohort is small and imbalanced. Case clustering does not resolve unverified dependence among related cases. Reported scores compare model–harness–provider configurations, not isolated model weights.

Temporal and transfer limits. A provider-reported cutoff and disabled browsing reduce specific leakage risks without proving contamination absent. New benchmark versions change the case distribution. Generalization to other jurisdictions, procedural stages, settlement decisions, drafting, or strategic advice is not established by this study. Outcome feedback may improve a narrow probability estimate without improving explanations or real client decisions.

11 Conclusion

LegalForecastBench operationalizes an uncertain but consequential legal judgment as a probabilistic forecast of a later judicial event. In the beta, all fourteen document-enabled configurations outperform a cohort-fitted constant in micro-Brier, while confidence behavior and model ordering differ in informative ways. GPT-6 Sol has the strongest observed scores, and the shared-summary comparison exposes a large gap between Jev and Luna that cannot be explained by the summaries containing no signal.

The broader opportunity is to make legal evaluation more empirically grounded. Real outcomes provide an external scoring reference, but useful training signals still require careful task definition, information boundaries, and legal labeling. The next research question is whether feedback on such forecasts improves performance on later cases and transfers to independently assessed legal work. The present evidence motivates that question; it does not answer it.

Reproducibility, access, and disclosure

This manuscript reports the fourteen-configuration comparison and checks its numerical claims against the retained repository artifacts as of September 30, 2026. The accompanying aggregate data and figure script regenerate the plots without court-document bytes. Code and result artifacts are maintained in the [LegalForecastBench repository](#); access to restricted source documents and

private adjudication material is a separate audit process. Availability of code is not equivalent to unrestricted source-data replication.

The author designed the benchmark and supplied domain judgments. AI assistants supported this manuscript’s drafting, literature retrieval, and numerical and source checks. The LAB audit discussed below is explicitly an AI-assisted preliminary analysis; the author’s targeted task review is distinguished from a completed independent audit. No model-weight reinforcement-learning experiment is claimed unless explicitly identified as a result of an external cited study.

References

- [1] Daya Guo, Dejian Yang, Haowei Zhang et al. DeepSeek-R1 incentivizes reasoning in LLMs through reinforcement learning. *Nature* 645, 633–638 (2025); *arXiv:2501.12948v2*, 2025. <https://www.nature.com/articles/s41586-025-09422-z>.
- [2] Zhiyuan Zeng, Hamish Ivison, Yiping Wang et al. RLVE: Scaling Up Reinforcement Learning for Language Models with Adaptive Verifiable Environments. *ICML 2026, Proceedings of Machine Learning Research* 306:153406–153428, 2026. <https://proceedings.mlr.press/v306/zeng26j.html>.
- [3] David Silver, Thomas Hubert, Julian Schrittwieser et al. Mastering Chess and Shogi by Self-Play with a General Reinforcement Learning Algorithm. *arXiv:1712.01815 (2017)*, preprint; precursor to the distinct 2018 *Science* article, 2017. <https://arxiv.org/abs/1712.01815>.
- [4] Thomas Hubert, Rishi Mehta, Laurent Sartran et al. Olympiad-level formal mathematical reasoning with reinforcement learning. *Nature* 651, 607–613; version of record 13 March 2026, 2026. <https://www.nature.com/articles/s41586-025-09833-y>.
- [5] OpenAI. Finite Time Blowup for Navier–Stokes. *Technical manuscript*, September 2026. <https://cdn.openai.com/pdf/32d9f210-8b73-45e0-91bc-82a30aef8a9a/navier-stokes.pdf>.
- [6] OpenAI. On the Navier–Stokes Millennium Prize Problem. *Research announcement*, September 8, 2026; updated September 10, 2026. <https://openai.com/index/navier-stokes-solution/>.
- [7] Russell Yang, Ruishi Chen, Pierce Kelaita et al. JudgmentBench: Comparing Rubric and Preference Evaluation for Quality Assessment. *arXiv:2605.25240v3*, preprint, 2026. <https://arxiv.org/abs/2605.25240>.
- [8] Anisha Gunjal, Anthony Wang, Elaine Lau, Vaskar Nath, Yunzhong He, Bing Liu, and Sean Hendryx. Rubrics as Rewards: Reinforcement Learning Beyond Verifiable Domains. In *International Conference on Learning Representations*, 2026. First posted as *arXiv:2507.17746* in 2025. https://proceedings.iclr.cc/paper_files/paper/2026/hash/cfd7bee7a651ee9af525098ef67a9e45-Abstract-Conference.html.
- [9] Niko Grupen, Julio Pereyra, Gabe Pereyra et al. Training a Legal Agent with Applied Compute. *Primary online source*, 2026. <https://www.harvey.ai/blog/training-a-legal-agent-with-applied-compute>.
- [10] Calvin Qi, Vasudha Rengarajan, Julio Pereyra et al. Harvey Tenet Research Preview. *August 20, 2026 company research report*, 2026. <https://www.harvey.ai/blog/post-training-update-harvey-tenet>.
- [11] Long Ouyang, Jeffrey Wu, Xu Jiang et al. Training language models to follow instructions with human feedback. *NeurIPS 2022*, 2022. <https://arxiv.org/abs/2203.02155>.
- [12] Neel Guha, Julian Nyarko, Daniel E. Ho et al. LegalBench: A Collaboratively Built Benchmark for Measuring Legal Reasoning in Large Language Models. *NeurIPS 2023, Datasets and Benchmarks*, 2023. <https://arxiv.org/abs/2308.11462>.
- [13] National Conference of Bar Examiners. About the Multistate Bar Examination. Official examination description, n.d.; accessed September 29, 2026. <https://www.ncbex.org/exams/mbe/about-mbe>.

- [14] H. L. A. Hart. Positivism and the Separation of Law and Morals. *Harvard Law Review*, 71(4):593–629, 1958. Especially pp. 607–608. <https://www.jstor.org/stable/1338225>.
- [15] Lon L. Fuller. Positivism and Fidelity to Law—A Reply to Professor Hart. *Harvard Law Review*, 71(4):630–672, 1958. Especially pp. 661–664. <https://www.jstor.org/stable/1338226>.
- [16] United States Congress. Declaratory Judgment Act, 28 U.S.C. § 2201(a). Text consulted September 29, 2026. <https://www.law.cornell.edu/uscode/text/28/2201>.
- [17] Supreme Court of the United States. Federated Department Stores, Inc. v. Moitie, 452 U.S. 394, 398–399 (1981). <https://www.govinfo.gov/content/pkg/USREPORTS-452/pdf/USREPORTS-452-394.pdf>.
- [18] Glenn W. Brier. Verification of Forecasts Expressed in Terms of Probability. *Monthly Weather Review* 78(1) (1950), 1–3, 1950. https://journals.ametsoc.org/view/journals/mwre/78/1/1520-0493_1950_078_0001_vofeit_2_0_co_2.xml.
- [19] Tilmann Gneiting, Adrian E. Raftery. Strictly Proper Scoring Rules, Prediction, and Estimation. *Journal of the American Statistical Association* 102 (2007), 359–378, 2007. <https://doi.org/10.1198/016214506000001437>.
- [20] Oliver Wendell Holmes, Jr. The Path of the Law. *Harvard Law Review*, 10(8):457–478, 1897. Especially pp. 457, 460–461. <https://www.jstor.org/stable/1322028>.
- [21] Robert H. Mnookin and Lewis Kornhauser. Bargaining in the Shadow of the Law: The Case of Divorce. *Yale Law Journal*, 88(5):950–997, 1979. <https://doi.org/10.2307/795824>.
- [22] Paul C. Saunders. Legal Skills and Professional Values: A Handbook. *Journal of the New York State Judicial Institute on Professionalism in the Law*, 7(1), Fall 2019. Reproducing and discussing the MacCrate skills framework. <https://www.nycourts.gov/LegacyPDFS/ip/JIPL/pdf/JIPL%20-%20FINAL%20-%20Skills%20and%20Professional%20Values.pdf>.
- [23] Benjamin Turtel, Danny Franklin, Kris Skotheim et al. Outcome-based Reinforcement Learning to Predict the Future. *Transactions on Machine Learning Research (November 2025)*; *arXiv:2505.17989v4*, 2025. <https://arxiv.org/abs/2505.17989>.
- [24] Nikhil Chandak, Shashwat Goel, Ameya Prabhu et al. Scaling Open-Ended Reasoning To Predict the Future. *arXiv:2512.25070v2 (January 2026)*; *preprint*, 2026. <https://arxiv.org/abs/2512.25070>.
- [25] Junjie Chen, Haitao Li, Qilei Zhang et al. Multi-Source Retrieval and Reasoning for Legal Sentencing Prediction. *arXiv:2602.04690v3*; *arXiv record comments 'EMNLP 2026 main' (publisher proceedings record not independently checked)*, 2026. <https://arxiv.org/abs/2602.04690>.
- [26] GPT-6 Sol audit workers, Claude Opus 5.5 audit workers, Coordinating agents; individual provenance not fully recorded. Harvey LAB litigation and dispute resolution audit: GPT-6 Sol and Claude Opus 5.5. *Local research artifact, pinned September 29, 2026*, 2026. <https://github.com/johnhughes3/LegalForecastBench/tree/58371a0e13494e591710e2ae3da8e428eebc359/docs/harvey-lab-audit/litigation-dispute-resolution>.

- [27] Harvey AI task authors; upstream authorship not individually listed. Pinned LAB task.json for analyze-counterparty-motion-to-dismiss. *Primary online source*, 2026. <https://github.com/harveyai/harvey-labs/blob/1dd81403b2fbb60596f7aea3fcecfaad7bf73143/tasks/litigation-dispute-resolution/analyze-counterparty-motion-to-dismiss/task.json>.
- [28] Abigail Z. Jacobs and Hanna Wallach. Measurement and Fairness. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency*, pp. 375–385, 2021. <https://doi.org/10.1145/3442188.3445901>.
- [29] Dario Amodei, Chris Olah, Jacob Steinhardt, Paul Christiano, John Schulman, and Dan Mané. Concrete Problems in AI Safety. *arXiv:1606.06565v2*, 2016. <https://arxiv.org/abs/1606.06565>.
- [30] Leo Gao, John Schulman, and Jacob Hilton. Scaling Laws for Reward Model Overoptimization. In *Proceedings of the 40th International Conference on Machine Learning*, PMLR 202:10835–10866, 2023. <https://proceedings.mlr.press/v202/gao23h.html>.
- [31] Mrinank Sharma, Meg Tong, Tomasz Korbak et al. Towards Understanding Sycophancy in Language Models. *ICLR 2024; first posted as arXiv:2310.13548 in 2023*, 2024. <https://arxiv.org/abs/2310.13548>.
- [32] Itai Shapira, Gerdus Benade, and Ariel D. Procaccia. How RLHF Amplifies Sycophancy. *ICML 2026; arXiv:2602.01002*, 2026. <https://arxiv.org/abs/2602.01002>.
- [33] Theodore W. Ruger, Pauline T. Kim, Andrew D. Martin, Kevin M. Quinn. The Supreme Court Forecasting Project: Legal and Political Science Approaches to Predicting Supreme Court Decision-Making. *Columbia Law Review* 104 (2004), 1150–1210, 2004. [Author-hosted article PDF](#).
- [34] Daniel Martin Katz, Michael J. Bommarito II, Josh Blackman. A General Approach for Predicting the Behavior of the Supreme Court of the United States. *PLOS ONE* 12(4): e0174698 (2017), 2017. <https://journals.plos.org/plosone/article?id=10.1371%2Fjournal.pone.0174698>.
- [35] Nikolaos Aletras, Dimitrios Tsarapatsanis, Daniel Preotiuc-Pietro, Vasileios Lampos. Predicting Judicial Decisions of the European Court of Human Rights: A Natural Language Processing Perspective. *PeerJ Computer Science* 2:e93 (2016), 2016. <https://peerj.com/articles/cs-93/>.
- [36] Chaojun Xiao, Haoxi Zhong, Zhipeng Guo et al. CAIL2018: A Large-Scale Legal Dataset for Judgment Prediction. *arXiv:1807.02478 (2018); CAIL challenge dataset paper*, 2018. <https://arxiv.org/abs/1807.02478>.
- [37] Masha Medvedeva, Pauline McBride. Legal Judgment Prediction: If You Are Going to Do It, Do It Right. *Proceedings of the Natural Legal Language Processing Workshop (ACL)*, 2023, pp. 73–84, 2023. <https://aclanthology.org/2023.nllp-1.9/>.
- [38] Huiyuan Xie, Felix Steffek, Joana Ribeiro de Faria et al. The CLC-UKET Dataset: Benchmarking Case Outcome Prediction for the UK Employment Tribunal. *Proceedings of the Natural Legal Language Processing Workshop (ACL)*, 2024, pp. 81–96, 2024. <https://aclanthology.org/2024.nllp-1.7/>.
- [39] Niko Grupen, Gabe Pereyra, Julio Pereyra. Introducing Harvey’s Legal Agent Benchmark. *Primary online source*, 2026. <https://www.harvey.ai/blog/introducing-harveys-legal-agent-benchmark>.

- [40] Hunter Lightman, Vineet Kosaraju, Yura Burda et al. Let’s Verify Step by Step. *ICLR 2024*; *arXiv:2305.20050*, 2024. <https://arxiv.org/abs/2305.20050>.
- [41] Andy Zou, Tristan Xiao, Ryan Jia et al. Forecasting Future World Events with Neural Networks. *NeurIPS 2022, Datasets and Benchmarks*, 2022. <https://arxiv.org/abs/2206.15474>.
- [42] Danny Halawi, Fred Zhang, Chen Yueh-Han, Jacob Steinhardt. Approaching Human-Level Forecasting with Language Models. *ICML 2024, Proceedings of Machine Learning Research 235*, 2024. <https://arxiv.org/abs/2402.18563>.
- [43] Ezra Karger, Houtan Bastani, Chen Yueh-Han et al. ForecastBench: A Dynamic Benchmark of AI Forecasting Capabilities. *ICLR2025*, 2025. <https://arxiv.org/abs/2409.19839>.
- [44] Evan Miller. Adding Error Bars to Evals: A Statistical Approach to Language Model Evaluations. *arXiv preprint*, 2024. <https://arxiv.org/abs/2411.00640>.
- [45] Carlos E. Jimenez, John Yang, Alexander Wettig et al. SWE-bench: Can Language Models Resolve Real-World GitHub Issues?. *ICLR 2024*, 2024. <https://arxiv.org/abs/2310.06770>.
- [46] OpenAI. Why SWE-bench Verified no longer measures frontier coding capabilities. *Primary online source*, 2026. <https://openai.com/index/why-we-no-longer-evaluate-swe-bench-verified/>.
- [47] TypeSafe AI. Models: Jev and the System One Interface. *Product documentation*, accessed September 29, 2026. <https://docs.typesafe.ai/models>.
- [48] Supreme Court of the United States. Dobbs v. Jackson Women’s Health Organization, 597 U.S. 215 (2022). Decided June 24, 2022. https://www.supremecourt.gov/opinions/21pdf/19-1392_6j37.pdf.
- [49] Steven Shavell. Economic Analysis of Litigation and the Legal Process. *NBER Working Paper 9697*, 2003. Chapter 17 of the draft *Foundations of Economic Analysis of Law*. <https://www.nber.org/papers/w9697>.
- [50] George L. Priest, Benjamin Klein. The Selection of Disputes for Litigation. *Journal of Legal Studies* 13 (1984), 1–55, 1984. <https://www.journals.uchicago.edu/doi/10.1086/467732>.
- [51] Lucian A. Bebchuk. Litigation and Settlement Under Imperfect Information. *RAND Journal of Economics* 15(3) (1984), 404–415, 1984. <https://hls.harvard.edu/bibliography/litigation-and-settlement-under-imperfect-information/>.
- [52] Supreme Court of the United States. Jones v. Barnes, 463 U.S. 745, 751–753 (1983). <https://tile.loc.gov/storage-services/service/ll/usrep/usrep463/usrep463745/usrep463745.pdf>.
- [53] Jane Goodman-Delahunty, Pär Anders Granhag, Maria Hartwig, and Elizabeth F. Loftus. Insightful or Wishful: Lawyers’ Ability to Predict Case Outcomes. *Psychology, Public Policy, and Law*, 16(2):133–157, 2010. <https://doi.org/10.1037/a0019060>.
- [54] Harvey Team. How AI Agents are Changing the Way Lawyers do Legal Work. *Primary online source*, 2026. <https://www.harvey.ai/blog/ai-agents-for-legal-work>.
- [55] Niko Grupen, Gabe Pereyra, Julio Pereyra. Legal Agent Benchmark: Initial Results. *Primary online source*, 2026. <https://www.harvey.ai/blog/legal-agent-benchmark-initial-results>.

- [56] Artificial Analysis. Harvey LAB-AA: New Benchmark Evaluating AI on Legal Tasks. *July 7, 2026, benchmark announcement*, 2026. <https://artificialanalysis.ai/articles/harvey-lab-aa>.
- [57] Harvey Team. Opus 4.8, Now Live in Harvey. *Primary online source*, 2026. <https://www.harvey.ai/blog/opus-4-8-now-live-in-harvey>.
- [58] Harvey Team. Fable 5, Now Available in Harvey. *Primary online source*, 2026. <https://www.harvey.ai/blog/fable-5-now-available-in-harvey>.
- [59] Harvey Team. Claude Sonnet 5, Now Live in Harvey. *Primary online source*, 2026. <https://www.harvey.ai/blog/sonnet-5-in-harvey>.
- [60] United States Supreme Court. Atlantic Marine Construction Co. v. United States District Court. *571 U.S.49,59–66 (2013)*, 2013. <https://www.law.cornell.edu/supremecourt/text/12-929>.
- [61] Texas Legislature. Texas Civil Practice and Remedies Code, Chapter 16. *Texas Civil Practice and Remedies Code § 16.070*, accessed September 29, 2026, 2026. <https://statutes.capitol.texas.gov/Docs/CP/htm/CP.16.htm#16.070>.
- [62] Supreme Court of the United States. Jones v. Bock. *549 U.S.199,215*, 2007. <https://tile.loc.gov/storage-services/service/ll/usrep/usrep549/usrep549199/usrep549199.pdf>.
- [63] Supreme Court of Texas. Italian Cowboy Partners, Ltd. v. Prudential Insurance Co. of America. *341 S.W.3d 323, 334–37 (Tex. 2011)*, 2011. <https://www.txcourts.gov/media/819896/OpinionsFY2011.pdf>.
- [64] Clay Mathematics Institute. Navier–Stokes Announcement. September 11, 2026. <https://www.claymath.org/news/navier-stokes-announcement/>.
- [65] Matthew Stubenberg, Chloe Berridge, Thomas Smith, and Joshua Casey. How AI Stacks Up Against the Multistate Bar Exam. *University of Hawai’i Richardson School of Law Research Paper No. 5291811; SSRN working paper*, May 12, 2025. <https://doi.org/10.2139/ssrn.5291811>.
- [66] Matthew Stubenberg. AI model achieves 100% on MBE bar exam questions. *Researcher’s online benchmark update, LinkedIn*, 2026; accessed September 30, 2026. https://www.linkedin.com/posts/matthew-stubenberg_well-it-finally-happened-an-ai-model-correctly-activity-7482410775534448641-urAX.
- [67] Vals AI. LegalBench Leaderboard and Methodology. *Benchmark-provider results and commentary*, updated September 27, 2026; accessed September 30, 2026. https://www.vals.ai/benchmarks/legal_bench.
- [68] Tak K. Mak. Analysing Intraclass Correlation for Dichotomous Variables. *Journal of the Royal Statistical Society, Series C (Applied Statistics)*, 37(3):344–352, 1988. <https://doi.org/10.2307/2347309>.
- [69] A. Colin Cameron and Douglas L. Miller. A Practitioner’s Guide to Cluster-Robust Inference. *Journal of Human Resources*, 50(2):317–372, 2015. <https://doi.org/10.3368/jhr.50.2.317>.

A Model configuration and temporal evidence

The table records the maintained local evidence, rather than certifying provider training histories. “K” denotes a reported knowledge cutoff, “T” a reported training cutoff, and “?” no established cutoff. Month-only dates are compared as the first day of the reported month. The reported cutoffs for twelve configurations precede June 30, 2026. Kimi and Muse remain qualified. A release date is not used to fill an unknown cutoff.

Configuration	Reasoning	Reported cutoff	Relation to window
GPT-6 Sol	High	2026-04-20 (K)	Precedes window
Claude Fable 5.1	Adaptive/default	2026-06 (T)	Precedes window
Grok 4.6	High	2026-02-01 (K)	Precedes window
Claude Opus 5	Adaptive/default	2026-05 (T)	Precedes window
GPT-5.6 Sol	High	2026-02-16 (K)	Precedes window
Kimi K3	High	Unknown (?)	Unknown
GPT-5.6 Luna	High	2026-02-16 (K)	Precedes window
GPT-6.1 Sol	High	2026-04-30 (K)	Precedes window
GPT-6 Astra	High	2026-04-30 (K)	Precedes window
Gemini 3.8 Flash	High thinking	2026-03 (K)	Precedes window
Grok 4.7	High	2026-05 (K)	Precedes window
Claude Sonnet 5	Adaptive/default	2026-01 (T)	Precedes window
Claude Sonnet 5.5	High	2026-06 (T)	Precedes window
Muse Spark 1.3	High	Unknown (?)	Unknown

Table 4: Configuration and cutoff evidence retained by the benchmark. “Precedes window” is a date comparison under the maintained rule, not an independently verified training-history claim.

The common document interface standardizes access, not every inference setting. Claude Fable 5.1, Claude Opus 5, and Claude Sonnet 5 use adaptive thinking at provider defaults. Claude Sonnet 5.5 pins high effort. The other configurations use the high reasoning or thinking setting recorded in their registries. Provider-default sampling and serving updates may still introduce variation. Full-document predictions and the one-request shared-summary conditions are kept separate throughout the analysis.

B Result provenance

The primary cohort identifier is `cycle-1-91-2026-09-08-luna-r5`. Costs in the main table are the totals recorded on the run receipts from provider-reported token usage and the frozen rates.

The checked source revisions are [LegalForecastBench 58371a0e](#), [LegalForecastCorpus b392ddaf](#), and [LegalForecastGym 4cc0f528](#). The LAB source snapshot is [1dd81403](#). These identify inspected versions; they do not certify every execution against the current source tree.

C A bounded LAB motion-to-dismiss case study

This case study uses the task and supplied documents at the pinned upstream revision recorded by the local audit. The requested output is an opposition issues memorandum. The fictional dispute concerns a software-services agreement, alleged pre-contractual representations, implementation

defects, and a motion asserting jurisdictional, venue, and pleading defenses. We distinguish direct record observations from legal interpretations and unverified AI flags.

Forum selection and litigation strategy. Complaint paragraph 48 sets out a clause selecting state or federal courts in Travis County, Texas, while the action is filed in the Northern District of Georgia. Paragraphs 18–19 assert statutory venue in Georgia, and paragraph 123 argues that the statutory deceptive-practices claim falls outside the clause. Those arguments may be weak, but the complaint is not wholly silent about the forum issue. A rigorous task should ask whether transfer is likely, which claims the clause covers, and whether a procedural objection changes the client’s position. Under *Atlantic Marine*, a valid clause selecting another federal forum is generally enforced through section 1404(a), while a state or foreign forum implicates forum non conveniens [60]. Correcting an opponent’s procedural vehicle is not the same as establishing that the chosen Georgia forum will retain the dispute. A criterion that emphasizes the former without measuring the latter can reward issue identification without strategic assessment.

Accrual and pleaded dates. Criteria C-013–C-015 require an accrual dispute and alternative discovery dates. The complaint alleges discovery of integration problems in August 2022, latency problems by November 2022, an acknowledgment in January 2023, and an unsuccessful patch on April 15, 2023. Paragraph 83 alleges that the plaintiff could not know the full extent of the breach until the patch failed, claims reliance on a promised fix, and states a March 22, 2024 filing date. These allegations require analyzing knowledge sufficient for accrual, distinct breaches, tolling or estoppel, and enforceability of the one-year contract term. Texas law itself restricts contractual limitations periods, subject to statutory exceptions [61].

The appropriate critique is therefore not that the dates mechanically prove the entire claim time-barred. Nor is it enough to announce that knowledge is always a fact issue immune from dismissal. A plaintiff’s allegations can establish an affirmative defense on the face of the pleading, while a genuine factual or legal issue may prevent that result [62]. The rubric should require applying the governing rule to the actual chronology and contract. It should allow the model to distinguish stronger from weaker opposition arguments instead of requiring the conclusion that every accrual question is unresolved.

Integration versus non-reliance. Criterion C-012 requires the model to characterize the merger-clause defense as legally incorrect. MSA section 14.7 also contains express non-reliance language, which is analytically distinct from an ordinary integration clause. Texas authority distinguishes these contractual concepts and examines the language and surrounding circumstances [63]. The source record therefore does not justify a categorical rubric requiring the ordinary-merger-clause answer without that analysis. This observation identifies an overbroad criterion; it does not establish that the particular disclaimer is enforceable or that the fraud claim must fail.

Hidden output requirements. C-028 requires concrete opposition recommendations for at least eight issues, although the prompt supplies no minimum count. C-022–C-025 prescribe high-severity treatment of specified issues. These criteria permit equivalent severity language; it would be inaccurate to say they necessarily require literal labels. C-026 prescribes organization tied to the motion’s framework and expressly permits omitting Rule 12(b)(1) when no genuine issue exists. Whether those requirements improve usefulness to counsel is a measurement question. They should not be confused with independently established legal errors.

Document construction and limits of the review. The supplied documents have programmatic-generation metadata and include a stale table-of-contents instruction in the agreement. The motion nevertheless has recognizable filing components. We have not established a court-rule violation warranting striking it, proven LLM authorship, or estimated how often such artifacts occur across LAB. Synthetic design can be useful for controlled testing; the concern is whether the rewarded conduct transfers to actual matters. The local audit also records genuine solver mistakes and unresolved legal disagreements. A benchmark-wide conclusion would require a representative expert review of source documents, criteria, model outputs, and grading decisions.

D A prospective test of outcome feedback

A follow-up would freeze a base model and an earlier training cohort, then compare outcome feedback with a matched control on later, disjoint cases. Related cases should remain together across splits. The primary endpoint would be a predeclared difference in proper score; calibration, failure rates, cost, and reasoning validity would be secondary endpoints. Blinded expert assessment of advice or drafting would test transfer.

For a model emitting a rationale z and forecast p , a candidate objective is

$$\max_{\theta} \mathbb{E}_{(X,U,Y) \sim \mathcal{D}_{\text{train}}, (z,p) \sim \pi_{\theta}(\cdot|X,U)} \left[1 - (p - Y)^2 \right]. \quad (9)$$

This is an outcome-scored prediction objective, not a simulator that allows the agent to alter the litigation trajectory. Optimizing it need not require reinforcement learning: supervised probabilistic prediction, calibration, and policy-gradient approaches are competing methods to test. It also does not directly reward faithful rationales. Leakage-controlled held-out cases and an independent transfer endpoint are therefore necessary to distinguish generalization from memorized case outcomes or improved probability reporting alone.